

Growing AI into the Organization's DNA

Jweokk · 2026-08-23 · v1.1.0

Models are no longer scarce. People who can grow models into their organizations are.
From the 95% project failure rate to DeepSeek's 160-person innovation density —
this book explains what an AI-native organization is, what the evidence says, and how to build one.

Copyright © Jweokk. Free for reading and non-commercial sharing; commercial use requires prior written permission.

Contact: weokk2025@gmail.com

Contents

- 01** Preface: Why I Wrote This Book
- 02** Chapter 1 The Rise of AI-Native Organizations — Starting from the 95% Failure Rate
- 03** Chapter 2 What Is an AI-Native Organization — Band-Aid AI vs AI-Native DNA
- 04** Chapter 3: The Evidence — What the Empirical Research Says
- 05** Chapter 4: Rebuilding Organizational Structure — From Hierarchy to Network
- 06** Chapter 5: Redesigning Workflows and Decisions
- 07** Chapter 6 Incentives and Talent: The Hardest Variable to Rebuild
- 08** Chapter 7 Case Studies: AI-Native Organizations Around the World and in China
- 09** Chapter 8 Building Your AI-Native Organization: An Action Playbook
- 10** Chapter 9: Challenges, Risks, and the Future
- 11** Chapter 10 Industry Guide: AI-Native Transformation Across Ten Industries
- 12** Afterword: A Living Book
- 13** Appendix A: Case Index and Sources
- 14** Appendix B: Key Metrics for AI-Native Organizations
- 15** Appendix C: Version History and Update Notes

Preface: Why I Wrote This Book

Preface: Why I Wrote This Book

In the summer of 2026, two headlines read side by side made the answer almost obvious.

The first: MIT's NANDA lab released *The GenAI Divide*, reporting that over the past three years, companies worldwide have burned \$30–40 billion on generative AI, and **95% of organizations have gotten no return that could be written into their financial statements**. For customized enterprise AI tools, the funnel from evaluation to pilot to production is 60% → 20% → 5%.

The second: DeepSeek built a trillion-parameter model with 160 people; Klarna shrank its workforce from 5,527 to 3,422 employees over three years while revenue kept rising; at Anthropic, Claude wrote more than 80% of the production code merged; and a company called Cursor reached \$4 billion in annualized revenue with 300 people.

On one side, 95% of AI projects come to nothing. On the other, a small cluster of organizations are leveraging enormous output with tiny teams.

Put the two together and the answer is not hard to guess: **models are no longer scarce. What is scarce are the people who can grow models into their organizations.**

What I've Been Researching

About 90% of the public discussion around "AI transformation" tells you how to **add AI capabilities** to your existing organization: buy a Copilot, run a few dozen training sessions, appoint a chief AI officer. But what is actually happening runs on a different track — **letting the organization itself grow AI into its DNA**. These two things are worlds apart.

This book sets out to study, end to end, a systemic change that is already under way but does not yet have a name: the AI-native organization. It clarifies three things:

- **What it is:** the essential difference between AI-native organizations and "AI transformation" — band-aid AI vs AI-native DNA;
- **The evidence:** what the empirical research from HBS, INSEAD, the WEF, and Tencent Research Institute actually says;
- **How to do it:** restructuring organizational form, restructuring workflows, restructuring incentives — plus a zero-to-one playbook.

All data and case studies in this book are cited in Appendix A. The work of compiling was itself a form of learning, and I have tried to make every citation verifiable; if anything is missing, corrections are welcome.

How This Book Was Auto-Written

This book was not written in a single pass. It is a "living book," powered by an automated content pipeline:

1. **Daily accumulation:** an automated system scans global information on AI deployment and organizational change every day, distilling, translating, and structuring high-quality content into a knowledge base;
2. **Deep research:** on top of the material already in the knowledge base (12 core concept articles plus 20+ case articles), the first edition underwent three further rounds of deep research — landmark international cases, raw data from authoritative reports, and failure cases and construction methodologies — with every source verified one by one;
3. **Weekly updates:** every week since, newly added knowledge base content is automatically merged into the relevant chapters, the version number increments (v1.0.0 → v1.0.1 → ...), the PDF is regenerated, and it is automatically deployed to the website;
4. **Continuous evolution:** whatever each update changed or added is recorded in Appendix C and in the repository's CHANGELOG, at a glance.

The process of writing this book is itself the book's best case study: **one person + an automated information pipeline + continuous weekly iteration produced something a traditional publishing process would take months to complete — and could never keep updated.** That is the AI-native organization in miniature, applied to content production.

Acknowledgments

This book's format and open publication model were inspired by XDash (Fan Bing)'s book, *Frontline Deployment Engineer: The Playbook for Delivering Customer Value in the Age of AI* (<https://github.com/xdash/FDE-the-Guidance-Book-of-Forward-Deployed-Engineer>) (<https://github.com/xdash/FDE-the-Guidance-Book-of-Forward-Deployed-Engineer>). The pattern of free public full text, synchronized reading via a GitHub repository and website, whole-book PDF download, and appendix-sourced citations is an open-source knowledge practice worth promoting. Special thanks.

Much of the material in this book comes from public industry research, media reports, and practitioners' sharing, and I have done my best to cite sources in Appendix A. If any citation is missing or there are copyright concerns, please contact weokk2025@gmail.com and we will correct it immediately.

Reading Guide

- To quickly judge whether your company counts as AI-native: read Chapter 2
- To see evidence rather than opinion: read Chapter 3
- To get hands-on: read Chapter 8, together with the workflow restructuring in Chapter 5
- To understand why most transformations fail: read Chapter 9
- To go through all the case studies: read Chapter 7

This book is free and open; feel free to share it. The value of knowledge lies in its flow, not in hoarding it.

Chapter 1 The Rise of AI-Native Organizations — Starting from the 95% Failure Rate

Chapter 1 The Rise of AI-Native Organizations — Starting from the 95% Failure Rate

1.1 A Number That Has Been Misread

Over the past two years, the most widely circulated AI statistic on the Chinese internet has been: "95% of enterprise AI projects fail."

The claim is both right and wrong. It comes from MIT's NANDA lab's July 2025 report, *The GenAI Divide*. The report's precise framing has two layers:

- **By organization:** 95% of organizations got zero return from their generative AI investments;
- **By tool funnel:** of customized, task-specific enterprise AI tools, only 5% make it to production deployment — 60% of organizations evaluated such tools, 20% entered pilots, and 5% reached production.

General-purpose chat tools (ChatGPT, Copilot, and the like), by contrast, convert from pilot to deployment at roughly 83% — they are not failing.

This detail matters enormously: **what "failed" is not all AI projects, but the customized systems deeply embedded in business processes that set out to replace core ways of working.** People who used AI casually gained efficiency; organizations that seriously reworked their businesses mostly came to grief. That is the "GenAI divide" the report describes — a polarization of outcomes between buyers (enterprises) and sellers (AI companies).

Figures from other institutions paint the same picture:

- Gartner predicts that by the end of 2025, at least 30% of generative AI projects will be abandoned after proof of concept (PoC);
- Gartner predicts that by the end of 2027, more than 40% of agentic AI projects will be canceled, with soaring costs the primary reason;

- IDC: 88% of AI agent pilots fail to graduate into production;
- An industry survey: for every 33 pilots launched, only about 4 successfully make it into production;
- Professor Fang Yue of CEIBS (China Europe International Business School) (2026): more than 90% of companies worldwide have launched generative AI pilots, but fewer than 41% of projects have truly crossed the experimental stage to create value at scale — he calls this phenomenon "pilot boom, value void."

Note that these numbers use entirely different measures (organizational zero-return rate / PoC abandonment rate / pilot-to-production conversion rate), yet they all point to the same structural fact: **the gap from pilot to production is real, and it is pervasive.**

1.2 Why Pilots Succeed and Production Fails

A strange, repeatedly observed phenomenon: AI projects that demo beautifully in the pilot phase collapse the moment they enter production.

The reason is that pilot success usually depends on humans "quietly patching things up" — supplying missing context by hand, bridging system failures, absorbing risk. Once those patches are stripped away in production, the AI immediately fails. Three layers specifically:

1. **Broken context:** the AI can only see a narrow slice of enterprise information and cannot reason across systems or across time; unstructured documents (PDFs, contracts, handwritten forms) are entirely invisible to it. In a demo this can be hidden; in production there is nowhere to hide.
2. **Fragile integration:** pilots are usually read-only; production requires the AI to write back to systems, trigger workflows, and act across dozens of applications. Point-to-point fragile integrations are immediately exposed — "read-only intelligence cannot act."
3. **Failed governance:** at pilot scale, human review can catch everything; at production scale, the "human in the loop" becomes a bottleneck rather than a safety valve, and missing audit trails breed risk aversion and frozen projects.

The few organizations that did make it into production (about 5%) made three different decisions: **build shared enterprise context first, standardize AI's "action layer," and embed governance into systems rather than leaving it to manual management.**

1.3 Another Group Running Hard in the Opposite Direction

Just as the vast majority of organizations are stuck in Pilot Purgatory, a small cluster of organizations are operating on a completely different logic:

- **DeepSeek:** 160 people built a trillion-parameter model, versus roughly 3,500 at OpenAI, 3,000 at Anthropic, and 8,100 at DeepMind;
- **Klarna:** full-time employees fell from 5,527 in December 2022 to 3,422 in December 2024 (-38%), while revenue kept growing over the same period and revenue per employee reached 3.6 times its 2022 level;
- **Shopify:** after cutting 20% of staff in 2023, headcount held flat for several consecutive quarters while revenue per employee rose from \$1.1 million to \$1.52 million;
- **Anthropic:** Claude writes more than 80% of merged production code, and per-engineer quarterly code output is 8 times the baseline;
- **Cursor:** about 300 people, annualized revenue from \$1 billion to \$4 billion, revenue per employee of \$3–13 million;
- **Duolingo:** zero full-time layoffs, content capacity up about 11 times in two years;
- **Airbnb:** AI co-writes the code produced by 60% of engineers, and AI customer support resolves 40% of issues without escalation to humans.

What these organizations share is not "using AI" — 88% of companies using AI have not achieved this. What they share is this: **AI is not a band-aid applied to a legacy organization; these organizations redesigned themselves around AI.**

1.4 The Scarcest Evidence of All: Causality

Everything so far is correlation and case studies. You could reasonably object: maybe these companies were simply stronger to begin with?

In 2026, researchers at INSEAD and Harvard Business School ran a rare randomized controlled experiment that pushed this question from correlation to causation.

They recruited 515 high-growth startups and assigned them at random: **every company** received equal API credits, access to frontier models, and technical training. The only difference: the treatment group additionally received case-study materials on how AI-native companies reorganize their production processes, teams, and business models, while the control group took only an ordinary entrepreneurship course.

With tools, skills, and budgets fully equalized, merely changing the "organizational search space" — letting the treatment group see what AI-native companies look like — produced:

- **+44%** more AI use cases discovered/used (concentrated especially in high-leverage areas such as strategy and product development);
- **+12%** more tasks completed; **+18%** likelihood of acquiring paying customers;
- Revenue **1.9×**;
- External capital needs actually fell **39.5%**.

The paper introduces a core concept: the "**mapping problem**" — what limits AI's returns is not technology costs or skills, but the cognitive bottleneck of discovering where and how AI creates value within one's own production processes.

This is a conclusion worth chewing on repeatedly: **handing employees tools is not the same as reengineering processes; simply knowing what an AI-native organization looks like can itself produce a 1.9× difference in revenue.** That is why this book is worth reading, and why XDash's FDE book is worth reading — knowledge of organizational form is currently the highest-ROI knowledge there is.

1.5 This Chapter's Core Judgment

Put the three sets of facts together:

1. 95% of organizations get zero return from AI investment;
2. the 5% who succeed share one trait — they redesigned the organization around AI;
3. merely "knowing what the winners look like" produces a 1.9× revenue difference.

The conclusion is already clear: **models are no longer scarce; what is scarce are the people who can grow models into their organizations.** In the chapters that follow, we will see what such an organization looks like (Chapter 2), what the evidence says (Chapter 3), how it operates (Chapters 4-6), who has already done it (Chapter 7), and — how you can do it (Chapter 8).

Chapter 2 What Is an AI-Native Organization — Band-Aid AI vs AI-Native DNA

Chapter 2 What Is an AI-Native Organization — Band-Aid AI vs AI-Native DNA

2.1 The Core Proposition: You Aren't Building an AI-Native Organization, You're Applying an AI Band-Aid to a Legacy One

No matter how well the band-aid is applied, the bone underneath is still old bone.

A company equipped every department with AI tools, ran dozens of training sessions, and offered incentives — then at the year-end review, adoption was under 15%. Root cause: the tools were bought, the training was done, the incentives were given, but the organizational structure, decision processes, and performance standards were all still traditional. AI was simply a band-aid applied to a legacy organization.

This is the watershed between "AI transformation" and "AI-native."

Alexander Puutio (Forbes) offers the sharpest conceptual framework: **being AI-native is, in essence, a reversal of design order.**

- Traditional organizations: design processes for humans → software assists and accelerates;
- Digital-native: designed around software → humans still produce the primary value;
- **AI-native: first ask what AI can reliably do → then design human roles around AI.**

"An AI-native organization is not one that uses AI tools — it is one built around them."

This definition can test any "AI-native" claim. The criterion is not "how much AI you use" but: **if your organization were designed from scratch tomorrow, would AI be written into the blueprint?** If the answer is no, you are only applying a band-aid.

2.2 Three Common Characteristics

Across the practices of Chinese companies such as Alibaba, Huawei, Lenovo, Feishu (Lark internationally), and Transn (Chinese: 传神), AI-native organizations show three recurring common characteristics:

Characteristic One: AI-informed decisions replace experience-based decisions. Traditional organizations decide on the strength of the veteran's experience and the boss's gut instinct. In an AI-native organization, as business data flows through the foundation layer, AI automatically performs analysis and issues early warnings, and decision recommendations are pushed straight to the accountable person — decisions rest on data, not feelings. Huawei's approach is to embed AI across the entire data lifecycle, making intelligent analytics the data platform's **default capability**: "you don't go looking for it; the system brings it to you."

Characteristic Two: unifying business flow and workflow. In traditional companies, the business flow lives in the CRM and the workflow in DingTalk or Feishu — project progress in system A, meeting minutes in system B, approvals in system C, all shuttled around by hand. Feishu's approach is a collaboration suite that connects IM, documents, and business flows, with AI intervening at project milestones in real time rather than compiling summaries after the fact.

Characteristic Three: making expertise reproducible. In a traditional organization, when the top salesperson leaves, the expertise goes with them. In an AI-native organization, AI automatically records decision logic and updates business rules, turning one person's expertise into a reusable organizational asset. Transn founder He Enpei puts it this way: **"Rather than wait for employees to become AI experts, let the organization grow AI capabilities."** — individuals come and go; organizational capability can be deposited, accumulated, and evolved.

2.3 A Historical Analogy: The Electricity Revolution of the 1880s

In the 1880s, factory owners bought electric motors to replace steam engines but kept their factory layouts, assembly-line design, and division of labor unchanged — and efficiency did not improve. **Electricity only swapped the power source; the mode of production was still steam-age.** The people who actually captured the gains were those who redesigned factory layouts, invented the assembly line, and gave each workstation its own power supply.

Today, "everyone is installing AI, but very few are truly redesigning their own factories."

This analogy yields a second definition of the AI-native organization: **not swapping the power source, but redesigning the mode of production.** From horse-drawn carriage to automobile — it was never about mounting an engine on the carriage. You had to reinvent the wheel, design the steering wheel, build roads, and set traffic rules. The entire ecosystem had to be reconstructed.

2.4 A Misconception That Must Be Corrected: the CAIO

A widely circulated claim on the Chinese internet: "Gartner predicts that in 2025, 90% of large enterprises will appoint a Chief AI Officer (CAIO)."

This is a misattribution — there is no such finding. Gartner's verifiable original prediction: by 2025, **35%** of large organizations will have a CAIO reporting to the CEO or COO (published in 2023). The "90%" figure has no corresponding source anywhere on Gartner's website.

The "CAIO trap" phenomenon, however, is real — it just needs a different argument: **AI is not one director's job; it is the whole company's job.** Appoint a Chief AI Officer while leaving structure, decisions, and incentives untouched, and the CAIO becomes a general with no army — the same structural trap as traditional companies appointing an "internet director" in 2005: the director spent three years unable to move anything and finally quit. The internet was never one director's job; neither is AI.

2.5 Three Easily Confused Concepts

Puutio goes to considerable lengths to distinguish three concepts, because many companies pass one off as another:

Concept	Core characteristic	Typical trap
Digital-native	Built around cloud + SaaS; humans remain the primary producers	Thinking that deploying a Copilot makes you AI-native
Platform organization	Coordinates an ecosystem, optimizes connections	Thinking that having network effects makes you AI-native
AI-native	Internal workflows redesigned around AI systems, optimizing production	Flattening without enough AI infrastructure = chaos

And the three types of tools in an AI-native organization:

1. **Vendor-embedded AI services** (Copilot and the like) — the easiest entry point, but not enough to qualify as AI-native;
2. **External agentic AI systems** — standalone tools for multi-step automation across applications;
3. **Internally built tools** — leveraging proprietary data and processes to become a sustainable advantage.

2.6 Employee Roles and Strategic Principles

In an AI-native organization, the definition of employees' roles changes fundamentally:

"Employees no longer define their value primarily through the volume of work they personally produce... They create value by directing systems that produce work."

Humans shift toward **orchestration and product management** roles; the most valuable employees understand both the boundaries of AI capabilities and the limits of systems; performance standards move from "how much work you did" to "designing better work."

Puutio offers a clear and unforgiving strategic principle: "**Don't involve people where people aren't needed.**" This means:

- investment decisions are made only around AI-driven ways of working;
- tools and roles that exist solely to manage human work rather than system performance should be questioned;
- core processes are defined around the tasks AI can reliably perform.

2.7 Building from Scratch vs. Retrofitting: A Real Debate

On whether an AI-native organization can be created by transforming a traditional one, two views exist:

- **Puutio (the strict-definition camp):** "An AI-native organization is necessarily a greenfield effort. Retrofitting rarely works."
- **Shmool (the viable-path camp):** a CEO who did subtraction within an existing organization, breaking the traditional hierarchy (VP → director → manager → team lead) down into a single-layer AI-native organization — and it worked (details in Chapter 4).

A reconciling interpretation: Puutio is talking about the **strict definition** (what truly counts as AI-native), while Shmool is talking about a **viable path** (how to get close to it). Retrofitting can achieve "quasi-AI-native" — close to, but not identical with, design from scratch. The difference lies in the ability to "change the decision order": starting from zero, you can begin anew from "what can AI do"; retrofitting means redefining while dismantling the old structure.

2.8 A Checklist: Band-Aid or DNA?

Ask yourself five questions:

1. **Design order:** do you first ask what AI can reliably do and then design human roles — or do you have the processes first and stuff AI in afterward?
2. **Decision basis:** are decisions driven by automatically pushed data plus human judgment — or by human experience with AI standing on the sidelines?
3. **Business flow and workflow:** have the two been unified, or are they still carried between by hand?

4. **Expertise retention:** when the top salesperson leaves, does their expertise stay behind?
5. **Performance standards:** do reviews still look at hours worked or volume of output? (detailed in Chapter 6)

If all five answers point to the "old structure," you are applying a band-aid. If even one or two have begun to shift, you are already on the road to growing DNA. **An AI-native organization is not an "upgrade" but a "rebirth" — and it can be reborn in stages.** In the next chapter, we look at the evidence.

2.9 AI Adoption Is a Myth — The Barbell Distribution and "Let AI Retreat to the Background"

Vasuman Moza — a former Meta engineer and founder of Varick Agents, which deploys AI agents inside enterprises with \$500M+ in annual revenue — spent a year inside large companies and kept seeing the same result. Hand the best AI tools to everyone, whether the team has 50 people or 5,000, and a "barbell" forms:

5%-10% become power users, 20% use it occasionally but crudely, and the remaining 70% barely touch it. One enterprise that had just signed an eight-figure annual AI contract even found that **10% of employees consumed 90% of the tokens.**

This explains the paradox this book keeps citing: McKinsey surveys show **88% of organizations use AI in at least one business function, but only 6% of enterprises derive more than 5% of EBIT from AI.** Model capability leaps forward every few months, but organizational productivity does not automatically follow.

Moza's more radical claim is that this gap will not close by itself: most employees will likely never become AI-native users, because using AI well is a craft — knowing when to clear context, when to distill repetitive work into reusable skills, when to let a model judge and when to write deterministic code. Vendor roadmaps orbit the most active frontier users, so the bar for "using AI well" keeps rising; and the company's AI experts have no incentive to share, because the gap itself is their advantage.

The strategic turn he proposes: **first, accept that most people will not become AI-native.** Let the 5%-10% of power users run ahead, and capture their skills, workflows, and agents into a shared library where ranking is the reward. For everyone else, **let AI retreat to the background** — embed agents inside the

systems employees already use every day (Salesforce, NetSuite, Dynamics), let invoices process themselves, data enter itself, and workflows flow themselves; call a human in to approve, reject, or edit only when the machine is unsure. Employees may not even realize they are "using AI." The metric changes accordingly: instead of asking "how many employees used AI this month," ask "of all the work completed today, how much was purely human, how much was human-AI collaboration, and how much already runs automatically?"

This is structurally the same judgment as the Transn case in Chapter 7: **rather than waiting for employees to become AI experts, let the organization grow AI capability.** (Two caveats: the 5/20/70 distribution comes from the author's own client observations, not a large-sample benchmark; and he sells exactly the "embed agents into enterprise workflows" service he advocates.)

A complementary framework comes from Zilbix's *AI Native Operating Model for the Agentic Era*: the past two years of "stacking tools on top of the existing structure" produced exactly the gap described above. The agentic era exposes three fractures — **ambiguous decision rights** (cross-functional agents without governance get confined to low-value tasks or run unsupervised), **the talent model** (employee value shifts from executing tasks to orchestrating systems, interpreting outputs, and exercising judgment), and **data architecture** (fragmented data blocks real-time access). An AI-native operating model is therefore three-dimensional: leadership actively sponsors operating-model change (not just approving budgets); decision rights with accountability and audit trails; processes redesigned around AI participation (rather than retrofitted); and an AI-ready data foundation.

Chapter 3: The Evidence — What the Empirical Research Says

Chapter 3: The Evidence — What the Empirical Research Says

The first two chapters were about concepts and cases. This chapter looks only at evidence: what academic research and authoritative institutions' data actually say about AI-native organizations. To avoid "cherry-picking evidence that favors our argument," I will also list the limitations of every study.

3.1 HBS's *AI-Native Firms*: The Strongest Firm-Level Evidence

AI-Native Firms (HBS Working Paper 26-090, 2026), a working paper by Harvard Business School's Rembrand Koning and INSEAD's Hyunjin Kim and colleagues, is currently the hardest first-hand evidence at the firm level.

Method: All 2,891 YC startups from the W20-F24 cohorts (11 batches), plus 41,214 US VC-backed startups from PitchBook, linked to micro-level labor data and compared against non-AI startups of the same industry and period.

Core findings:

Dimension	AI-native firms vs. non-AI peers
Team size	~ 25% smaller (YC firms roughly half the size after three years)
Share of engineers	~13 percentage points higher
Entry-level employees	~15% fewer
Senior employees	~20% more
Management layers	Half a level flatter
Managers	~15% fewer

Dimension	AI-native firms vs. non-AI peers
Funding / valuation	Comparable → ~20% more funding per capita, higher valuation per capita

In one sentence: **AI-native companies are smaller, flatter, more engineer-dense, and more senior, with comparable valuations — and higher output per capita.**

The paper's most important theoretical contribution is distinguishing two channels:

- **The process channel:** AI as an internal production tool that changes how employees do their work (agentic coding, AI customer service, automated sales) — this is where most existing research focuses;
- **The product channel:** AI capabilities embedded into the product, so customers generate deliverables directly inside the product — **knowledge work moves out of the internal organization and into the product**, scaling through compute rather than headcount.

The decisive evidence: the process channel cannot predict smaller organizations. AI startups are 2.6 times more likely than non-AI peers to name ChatGPT/Copilot/Cursor in job postings — but that indicator cannot predict smaller teams. After controlling for the process channel, product-embedded AI remains significantly associated with fewer people and flatter structures.

The largest differences appear in service businesses: in service startups such as therapy, tutoring, and telemedicine, AI startups are about 70% smaller than non-AI peers. Cases: Gamma (AI presentations; roughly 30 people reached \$50 million in annual revenue within two years); FazeShift (an accounts-receivable AI agent; 10 people won dozens of enterprise clients, replacing the AR team at each client).

A cautious conclusion about the labor market: every AI-native company is smaller, but total employment need not fall — AI dramatically increases the number of startups (PitchBook's 2024 AI seed funding was about 8 times the 2020 average). Extrapolating from the firm level to the market level requires caution.

Limitations: the sample is YC/VC-backed startups, which may not generalize to mature large enterprises; the "AI-native" classification is the researchers' own definition; and the working paper has not been peer-reviewed.

3.2 The INSEAD/HBS Field Experiment: Rare Causal Evidence

The randomized controlled trial of 515 startups mentioned in the previous chapter (*Mapping AI into Production*, 2026) is the scarcest kind of causal evidence in AI organization research: tools, skills, and budgets were fully equalized, and only the "organizational search space" was changed (the treatment group was shown how AI-native organizations restructure). The result: treatment-group use cases +44%, revenue 1.9x, external funding needs –39.5%.

The "mapping problem": what limits AI gains is not the cost of technology or skills, but the cognitive bottleneck of "discovering where and how AI creates value in your own production process."

Limitations: a 10-week short-term experiment; a startup sample; the 1.9x revenue figure is a short-term relative increment with small absolute values (on the order of, say, \$40,000).

3.3 WEF's *AI-First Operating Models*: A Systematic Framework for Operations

The World Economic Forum's (2026) core proposition: **layering AI on top of operating models designed for the pre-AI world (linear processes, static roles, incremental optimization) is a structural failure.** Global AI investment exceeded \$250 billion in 2024, yet in most organizations "AI has not fundamentally changed the way they operate."

Three shifts:

1. **Work and teams: from process management to human-AI collaboration** — early AI-first organizations already have **human-to-AI ratios exceeding 10:1** (1 human : 10+ AI); delegation is a management skill that requires training and certification (the You.com CEO: adoption in mature organizations only really took off after a training-and-certification program was launched);
2. **Workflows: from linear processes to dynamic AI systems** — the Rubrik CEO's implementation path: **"go business line by business line; in each line, find 3-5 workflows that are purely manual or SaaS-plus-human,**

and define end-to-end AI-driven outcomes"; the starting question determines everything (Lightspeed: "If we had unlimited intelligence, what would we build?");

3. **Metrics: from process optimization to dynamic value creation** — outcome-oriented, with reliability/accuracy/governance first to build trust.

Key judgment: **competitive advantage comes from orchestration, not experimentation**. Successful leaders treat capability, process, outcome, and business-model redesign as a single integrated evolution.

3.4 Tencent Research Institute's *AI-Native Work Report: A Complete Framework from a Chinese Perspective*

Tencent Research Institute (Chinese: 腾讯研究院) published a 40,000-character report in May 2026, with lead authors Si Xiao, Yuan Xiaohui, and Yu Yi (Chinese: 司晓、袁晓辉、余一). The most widely circulated part is its **competitiveness formula**:

$$\text{Organizational competitiveness} = \text{talent density} \times \text{AI leverage} \div \text{organizational friction}$$

If the numerator doubles but the denominator stays fixed, the net effect is discounted; **halving the denominator is equivalent to doubling the numerator** — reducing organizational friction (waiting/approvals/alignment/information decay) is usually easier to move than piling on talent.

The report defines four structural characteristics of the "Super-Individual" (all required):

1. **An AI-first workflow**: AI is the default starting point — "I let AI run first, then judge and correct on top of its output";
2. **An order-of-magnitude leap in capability boundaries**: 10x+ output; one person runs the full chain from idea to delivery;
3. **Extreme proactivity**: doesn't wait for organizational assignment; a natural boundary explorer;
4. **Influence spillover** (the key threshold for identifying Super-Individuals): **efficient individuals only make themselves faster; Super-Individuals make the team faster**.

It also outlines four emergence paths (bottom-up spontaneity / organizational selection and cultivation / atmosphere creation / founder-driven), three team forms (node-radiation / network collaboration / AI-hub), and one harsh judgment about the capability reshuffle: **AI is an accelerator of divergence, not an equalizer — only those who were already excellent become more excellent, and the gap widens.**

We must emphasize the report's sample bias (this is the knowledge base's systematic critique of the report): roughly 80% of its cases come from information-intensive, digital-native industries with fewer than 500 employees, with zero coverage of physical-world industries (manufacturing lines / medical clinical settings / logistics). The evidence in the report supports only this: **in information-intensive, small-team, digital-native industries, AI-native organizations can dramatically accelerate.** Before applying it to physical industries or large organizations, ask first: is this industry's bottleneck information processing, or the physical world?

3.5 Large-Sample Data from Authoritative Institutions

McKinsey, *State of AI 2025 (November 2025)*: 88% of organizations regularly use AI in at least **one business function** (78% last year); but about two-thirds are still in the experiment/pilot stage and only about one-third have begun scaling; only 39% attribute any degree of EBIT impact to AI, and most of them say the contribution is under 5%; "AI high performers" (EBIT impact $\geq 5\%$) account for only about 6%. **Many companies use AI; very few make money from it.**

BCG's *AI at Work* series: a clear leadership-frontline gap — more than three-quarters of leaders use GenAI multiple times a week, while frequent frontline use stalled at 51% in 2025 before jumping to 74% in 2026 (+23 percentage points). Half of companies are moving from "Deploy" to "Reshape" (redesigning workflows).

Microsoft's *Work Trend Index*: in 2024, 75% of knowledge workers already used AI; in 2025 it introduced the "Frontier Firm" concept (about 9.3% of surveyed leaders) — 71% of frontier-firm leaders say their company is "thriving," versus only 39% of employees globally; the most important finding of 2026: **organizational factors (culture, manager support, talent practices) contribute twice as much to AI impact as individual effort** (based on a permutation-importance analysis of 29 factors, $R^2 \approx 0.68-0.69$).

Stanford HAI's *AI Index 2025*: in 2024, 78% of organizations reported using AI (55% in 2023).

IMF: AI will affect about 40% of jobs globally; roughly 60% in advanced economies, 40% in emerging markets, and 26% in low-income countries — note that these are "exposure" estimates, not unemployment forecasts, and about half of exposed jobs may be augmented rather than replaced.

WEF's *Future of Jobs 2025*: by 2030, job disruption will equal 22% of current jobs: 170 million created, 92 million eliminated, a net gain of 78 million; 77% of employers plan to upskill their workforce, and 41% plan to reduce headcount because of AI automation — expectations, not facts.

3.6 A Warning on Data Use (This Book's Citation Discipline)

Putting all the numbers above together, there are five citation disciplines:

1. **Definitions come first:** the "95% failure rate" (MIT, zero organizational return), "30% abandoned" (Gartner, after PoC), and "88% adoption" (McKinsey, at least one function) are not contradictory — they are three different denominators;
2. **Report conclusions vs. media retelling:** media retelling often distorts (e.g., Fortune reported the MIT report's sample as 150 interviews/350 questionnaires, while the report actually covered 52 organizations/153 executives); this book always follows the original report;
3. **Forecasts vs. facts:** most of the Gartner/Forrester/WEF numbers are forward-looking forecasts or employer expectations, not events that have already happened;
4. **Sample bias:** McKinsey/WTI/BCG skew toward white-collar workers, large enterprises, and self-selected survey participants; every "global" claim must be checked against the original sampling scope;
5. **Conflicts of interest:** Microsoft, BCG, and McKinsey all sell AI products/consulting to enterprises — keep their positions in mind when citing their conclusions.

3.7 The Core Judgment of This Chapter

The evidence chain is complete: HBS shows AI-native companies are **smaller, flatter, and higher-output per capita** (driven by the product channel, not the process channel); INSEAD/HBS shows that **organizational knowledge itself produces causal revenue differences**; WEF shows that **operating models must be redesigned around AI**; Tencent shows that **competitiveness comes from talent density × AI leverage ÷ organizational friction**; McKinsey/BCG/Microsoft show that **a huge gap exists between tool adoption and organizational returns**.

The evidence points exactly where Chapter 1 concluded: models are no longer scarce — what's scarce is organizations redesigned around AI. In the next chapter, we look at what the inside of such an organization looks like.

Chapter 4: Rebuilding Organizational Structure — From Hierarchy to Network

Chapter 4: Rebuilding Organizational Structure — From Hierarchy to Network

4.1 The Legacy Organization's Deadlock

Linear division of labor has exposed a fundamental contradiction under AI acceleration:

The old way	The new contradiction
Managing people by role	The people who can actually get things done span multiple roles
Breaking work down by process	AI completes in a day what used to take a week of milestones
Accepting deliverables at checkpoints	Deliverables themselves become "intermediate artifacts"; only the final outcome is value
Reporting and evaluation by rank	Those who take on fused roles don't receive the corresponding authority and rewards

When one person can use AI to leap over steps that previously had to be handed off, and produce something closer to the final outcome, the legacy organizational framework becomes a constraint.

The core proposition of restructuring organizational form is: **use AI to replace the coordination and information-routing functions that hierarchy used to perform, so that humans can focus on judgment and creation.**

4.2 MarsWave: Five Practices of a Small Team Building from Zero

A small company called MarsWave (Chinese: 火星电波), which pivoted from ListenHub, completed its AI-native transformation in five weeks. Its five practices are the most concrete Chinese sample of "what an AI-native organization looks like":

1. Vibe coding shocks the organization. The CTO used Claude Code to build version 0.1 alone overnight — ideas could be presented directly. Before: product requirements → PRD → review → development → testing → release (weeks). Now: say it to AI → see the result → iterate (hours). Alignment shifted from "long documents" to "seeing the prototype directly."

2. The OPC model (human-machine-macro decision). Humans handle judgment (what to do, why, whether it's good), AI writes 99% of the code (how to write, how to implement), and managers make macro decisions (direction, resources, cadence). **Humans don't manage processes or schedules; they manage judgment, quality, and review.**

3. The birth of the Soul Team. The first core role that was neither engineer nor product manager — responsible for the product's **personality and emotional output**; its lead has a media background and can't code. Insight: once AI can handle "feature implementation," humans' differentiating value shifts to the "human-like" parts — empathy, personality, warmth.

4. A tool revolution. Linear, Notion, and kanban boards were abandoned; all code moved into a unified repository with no permission barriers. The repository "isn't for humans to read — it's for AI to read as context." The CEO's ideas go directly to AI, are automatically organized into Markdown, and become development context. **Essence: AI becomes the intermediary of information transmission, instead of layer-by-layer transmission between people.**

5. A new rhythm of front-load thinking and back-load review. In the first few days of each month the whole team stops writing code to align on plans; at month's end they re-review the code, cutting and refactoring. "Faster isn't necessarily better" — rhythm matters more than speed.

The failed experiment matters just as much: before the formal transformation, there was the "Task Tavern" (a task bulletin board where anyone could pick up any task) — a programmer's two-day new feature reached 80,000 users on

launch day, but under business pressure the strongest people were pulled back to old tasks, and it ultimately failed. Lesson: **process improvement (patching the old framework) isn't enough — you have to design from scratch.**

4.3 Shmool: A CEO's Firsthand Subtraction in a Mature Organization

Unlike MarsWave (building from zero), Itay Shmool's case is about **subtraction within an existing organization** — dismantling the traditional hierarchy (VP → Director → Manager → Team Lead) into a single-layer AI-native organization. Core numbers:

Metric	Before	After
Decision latency	Days to weeks	Hours
Information fidelity	~60% (decays with each transfer)	~95% (direct transfer)
Builder-to-manager ratio	3:1	8:1+
Meeting overhead	~30% of work week	~10% of work week

Three key insights:

1. **The truth about middle management:** many middle managers are not "redundant layers" — they are excellent builders themselves, just pulled into coordination work by the old structure. Give them the chance to build again and they thrive. **Distinguish the management function from the managers themselves — the function can be automated, but people should not be simply removed.**
2. **Guild Masters replace management:** once management layers are removed, a new mechanism is needed to maintain professional standards. Guild Masters are **expert roles** (not management roles) who sit within teams and set standards across teams — they create no new reporting lines and are "the quality immune system of a flat organization."


3. **The AI platform is a precondition:** flattening without AI infrastructure equals chaos, not lean. Five core capabilities: self-healing systems, an AI creative engine, a unified experimentation platform, automated code review, and operations AI agents. **"The AI platform is what replaces coordination — this is non-negotiable."**

The human side of the transformation matters just as much: "Behind every box in the architecture is a person with a mortgage, a family, and a career. **How you treat the people who leave defines your culture more than how you treat the people who stay.**"

4.4 Block: Replacing Hierarchy with a World Model

Block, the payments company (Jack Dorsey), offers the large-company version of the framework: **use AI to replace the coordination functions hierarchy performed, and build the company as an agent.**

Block's four-layer architecture:



The revolutionary shift in the Intelligence Layer: from "product managers plan features" → the Intelligence Layer identifies timing and composes solutions; failure signals automatically become the roadmap; **product planning shifts from guess-driven to signal-driven.**

Three new human roles: Individual Contributors (ICs, given context by the World Model, making autonomous decisions), DRIs (Directly Responsible Individuals, short-term ownership of cross-domain problems, replacing project managers/product managers), and Player-Coaches (both building and developing people, replacing traditional managers).

The World Model handles alignment; DRIs handle strategy and priorities; Player-Coaches handle craft and people development.

A critique that must be heard (a critical analysis of the Block framework): Dorsey's premise is that "the main value of humans is information routing," but if humans are also generating, interpreting, and creating signals, then subtracting humans means subtracting intelligence itself; the real bottleneck is not routing speed but **perspective loss** (signals die before they reach decision points). Core warning: **"Speed is not intelligence. Speed without integration is just faster blindness."**

4.5 Yang Renbin: The Organizational Brain and the AI Chief of Staff

In a 2026 interview with LatePost (Chinese: 晚点), Yang Renbin (Chinese: 杨仁斌), founder of Jingzhunxue (Chinese: 精准学) — an AI education brand founded by Yang Renbin — supplied the piece all previous perspectives were missing: **how the highest cognition and decision-making capability can be systematically injected into every frontline execution step.**

- **First principles of the organization:** solving three traditional deadlocks — the top-down decay of "highest cognition → frontline," the bottom-up distortion of "frontline facts → top," and the bandwidth bottleneck at the top;
- **Three steps:** ① model the decision layer's judgment methods into a callable "AI Chief of Staff / Organizational Brain" (initially injecting 300,000 characters of methodology) → ② inject it into every business site through context engineering (a decision-model index system) → ③ frontline real problems/disagreements/anomalies flow back, closing the loop to iterate the highest cognition;

- **AI coding is only 0% and 100%:** a hard rule that engineers are not allowed to modify code themselves — humans only give feedback and transform into "context engineers," going further than MarsWave's "AI writes 99% of the code";
- **Anti-document, anti-group-chat business discussion:** documents are negative assets (author bias plus information loss); preserve the original records of discussions, disagreements, and failures; deep discussions continue in AI conversations; offline voice must be transcribed to text and enter the context;
- **Context is Everything:** once model capability is replicable, what is truly scarce is context engineering;
- **Cultural premise:** "explain the Why thoroughly, and give enough AI (爱, 'love,' pronounced ài, is a pun on AI)"; an open culture — hiding problems is shameful, exposing them means solving them sooner; don't hold people accountable for making mistakes, hold them accountable for hiding problems or failing to update their judgment.

4.6 An International Comparison: The Changing Role of the Engineer

Anthropic's internal practice confirms the direction of role reconstruction: the engineer's role shifts from "writing code" to "**architect + referee**" — AI-powered automated code review is embedded in CI/CD, and post-hoc analysis shows it intercepted about one-third of the production bugs in claude.ai's history that caused downtime; one engineer used AI to auto-fix 800+ API errors that humans estimated would take four years. At OpenAI, Codex became the primary work tool of every department (including non-technical ones such as legal, finance, and recruiting), accounting for 99.8% of the company's weekly output tokens.

4.7 Three Paths: There Is No Single AI-Native Organization

Putting this chapter's cases together with the international ones, there are three parallel paths to the AI-native organization:

1. **The shrink path** (Klarna): hiring freeze + natural attrition, headcount –38%, with AI replacing the headcount growth that demand would have required;
2. **The expansion path** (OpenAI/Anthropic): headcount and revenue grow together; AI-native is a growth engine, not a downsizing tool;

3. **The steady-state leverage path** (Shopify/Duolingo/Airbnb): full-time headcount flat or slightly up, with AI raising per-capita output.

There is no single answer for organizational structure, but organizational principles converge: AI handles information routing and execution; humans handle judgment, quality, and review; hierarchy grows thinner as coordination functions are replaced by AI platforms; experience shifts from a personal asset into an organizational asset. In the next chapter, we look at how workflows and decisions are concretely redesigned.

Chapter 5: Redesigning Workflows and Decisions

Chapter 5: Redesigning Workflows and Decisions

5.1 Redesign First, Automate Second

In 1990, Michael Hammer published a famous article in the *Harvard Business Review* titled "Don't Automate, Obliterate." His core point is even more applicable in the AI era: **before automating a bad process, ask why that process exists in the first place.**

The first principle of AI-native workflow redesign is: **not "give everyone a chatbot," but redesign how work flows, how decisions get made, how data is captured, how humans review risk, and what AI is allowed to assist with or act on.**

The fundamental difference between an automation-first organization and an AI-native organization:

Question	Automation-first organization	AI-native organization
Starting point	Where can we add AI?	What work should change?
Success metric	How many tasks were automated / how many prompts were used	Whether the outcomes of the whole workflow improved
Main risk	Faster rework	Discovering redesign gaps before scaling
Data approach	Use existing data	Redesign data capture around decisions
Human role	Post-hoc reviewer	Responsible designer, supervisor, and owner of exceptions
Scaling pattern	More tools and pilots	A repeatable workflow-redesign loop

5.2 Nine Key Workflow Redesigns

1. Start from business outcomes, not AI tools. Outcomes must be specific enough to guide trade-offs: "30% fewer supplier-invoice exceptions without adding approval risk" beats "use AI in finance." Technology choices follow workflow intent.

2. Map the real workflow before automating the official one. The official process diagram and how employees actually work (inbox shortcuts / private spreadsheets / informal escalation paths / exception handling) exist side by side as two versions. Delete steps before automating: merge duplicate approvals, turn low-value checks into rules, give exceptions an owner. **Don't teach AI to preserve unnecessary work.**

3. Redesign intake so AI gets better input. Many AI projects die at the door: expecting AI to classify, route, or decide on messy input that was never designed for it. Required fields, structured options, validation at capture points, automatic enrichment from trusted systems, and flagging low-confidence cases. **Don't blame the model for weak input before redesigning the front door.**

4. Separate rules / automation / AI assistance / agents / humans. Not every step needs AI: stable rules → rule automation; system handoffs → workflow automation; language-intensive review → AI assistance; pattern recognition → AI models; low-risk multi-step tasks → constrained agents; high-accountability decisions → human plus AI support. Prevent two errors: using AI where rules suffice, and delegating high-risk decisions just because AI can generate confident-sounding suggestions.

5. Build a workflow data foundation, not an abstract data lake. Ask "have we captured the data needed to make the workflow better?" rather than "do we have data?" Record when humans override AI suggestions and why — the best learning comes from the redesigned workflow itself.

6. Define decision rights before AI systems act. If humans aren't clear on decision rights, AI is even less safe. The decision rights table is the tool that operationalizes this:

Scenario	What AI may do	What humans retain
Classification / routing	AI acts on high confidence	Humans review low-confidence cases

Scenario	What AI may do	What humans retain
Drafting	AI recommends drafts	Humans approve sensitive content; template and tone control
Low-value refunds	AI acts within limits	Humans review over-limit cases; thresholds + logs + fraud checks
Contract changes	AI assists only	Humans decide
Risk escalation	AI acts immediately	Humans investigate and close

7. Replace manual handoffs with event-driven work and review queues. Stop treating people as glue between systems. Trigger → create record → enrich → apply rules → AI classify/summarize → route → exceptions enter a visible review queue (with context/priority/owner/deadline/escalation path). Don't make handoffs invisible; make important handoffs observable.

8. Measure the whole workflow, not AI tasks. Task-level time savings mislead (drafting got faster but review got slower; classification got faster but the queue got wrong). Metrics: cycle time, first-time-right rate, exception rate and exception age, review time, satisfaction, cost per completed case, rework rate, AI confidence and coverage, escalation quality, incidents/control violations. **After the pilot, ask: did the work improve, or did AI just make one step look faster?**

9. Create a repeatable 90-day workflow-redesign cadence.

Processes themselves can be redesigned: a review system is not the same thing as a PR. Amp, a 20-person team, pushes straight to main without pull requests and still passes SOC 2 — auditors never cared about the git workflow; they care that changes are authorized, tested, approved, and recorded. Their control stack: push permissions restricted by business function, signed commits (GitHub verified), fully automated CI validation, and an audit trail linking commits to business context. The scaling insight: risk is not uniform inside an enterprise, and calibrating every system to "the scariest one" is hidden waste — first ask "what risk is the PR actually managing," then decide where it is needed.

5.3 The 90-Day Redesign Cadence (The Actionable Core)

- **Days 1-15: Select** — a process with visible pain + measurable value + sufficient volume + controllable risk;
- **Days 16-30: Map** — the real workflow / systems / data / decisions / handoffs / approvals / exceptions / owners / current performance;
- **Days 31-45: Redesign** — delete steps, clarify ownership, define decision rights, standardize intake, improve data capture, assign people / rules / automation / AI / agents to each step;
- **Days 46-60: Build** — workflow / integrations / prompts / access control / review queues / dashboards / logs / training;
- **Days 61-75: Pilot** — real users run real cases inside safe boundaries, measuring the whole workflow;
- **Days 76-90: Govern and scale** — kill / fix / scale, record the lessons, pick the next process.

5.4 Eight Common Mistakes

① Treating AI access as transformation; ② automating before simplifying (AI accelerates work of the wrong shape); ③ using AI where rules suffice; ④ ignoring exception work (the real value lies in handling messy cases); ⑤ measuring prompts instead of outcomes; ⑥ excluding employees from the redesign; ⑦ giving agents power without control (requires thresholds / logs / permissions / review queues / rollback / named owners); ⑧ redesigning the whole company at once (start with one process).

5.5 Dividing Labor Between Humans and AI: From "Full Replacement" to "Three-Layer Routing"

Beyond workflow redesign, designing the boundary of human-AI division of labor is the core of decision redesign.

Anthropic's internal evidence gives a key number: most employees believe that only 0-20% of work can be "fully delegated" to AI. In other words, the human-AI division of labor in an AI-native organization is "supervised collaboration," not "hands-off automation." The number of autonomous actions Claude

Code could take before human intervention was needed grew from about 10 to about 20 — this "intervention threshold" is dynamic and moves as model capability improves.

The hybrid-model **three-layer routing** that Klarna distilled from its customer-service reversal is currently the most mature template for dividing labor:

- **Layer 1: AI handles routine volume end to end** (targeting 60–70% of traffic) — order status, basic returns, FAQs;
- **Layer 2: AI-assisted human** (AI drafts, human reviews and sends, 20–25%);
- **Layer 3: Human-led high-value / escalation scenarios** (5–15% of traffic, but the greatest retention impact) — complex billing disputes, fraud, account cancellation.

Lesson: Klarna's first round of AI customer service only did Layer 1 (and did it aggressively); overall volume-based metrics (resolution rate, first-response time) looked good, but CSAT for complex interactions dropped significantly and the repeat-contact rate rose — the overall average masked a quality collapse in retention-critical scenarios. **You must track satisfaction for "complex/escalation interactions" separately.**

Treating AI as a compiler leads to frustration; treating it as a collaborator works. One engineer's observation is worth remembering: code gives certainty (same input, same output; a difference is a bug), humans give uncertainty (a colleague may understand intent and produce a better result), and AI sits between the two — rather than treating it as a compiler that precisely executes instructions, collaborate with it the way you would lead: share context, explain the desired outcome, set boundaries, respond to feedback. The core shift is investing in **expressing intent** — explaining why the work matters, what a good result looks like, and where judgment is needed. "The technology is new; leadership is not." This is also the underlying capability for the "human-AI division of labor design" in Chapter 8.

5.6 Corroboration with Chinese Practice

The three common features from Chapter 2 (AI-informed decisions replacing experience-based decisions, unifying business flow and workflow, making expertise reproducible) all have concrete implementation mechanisms in this chapter:

- **AI-informed decisions:** Huawei (Chinese: 华为) makes intelligent analytics a default capability of its data platform — "you don't go query it; the system gives it to you automatically";
- **Unifying business flow and workflow:** Feishu (Chinese: 飞书; Lark internationally) connects IM, documents, and business flows, with AI intervening in project nodes in real time;
- **Making expertise reproducible:** Transn's (Chinese: 传神) AI automatically records decision logic and updates business rules, turning personal experience into organizational assets — at the same time, Yang Renbin's "decision model index" from Chapter 4 (modeling the decision layer's judgment methods as a callable system) and the decision rights table are two sides of the same coin: one is the governance side, the other the cognition side.

5.7 The Core Judgment of This Chapter

Workflow redesign is the minimal executable unit of the AI-native organization. Organizational form (Chapter 4) is the skeleton; workflow redesign is the muscle — an organization doesn't need to finish "organizational restructuring" before becoming AI-native; it only needs to start with one workflow: map the real process, delete redundant steps, define decision rights, design for production, measure the whole flow, and iterate in 90-day cycles. **Reduce organizational friction (the denominator) first, then multiply talent and AI leverage (the numerator).** In the next chapter, we take on the hardest variable: people.

Chapter 6 Incentives and Talent: The Hardest Variable to Rebuild

Chapter 6 Incentives and Talent: The Hardest Variable to Rebuild

Technical problems all have answers; people problems don't. This chapter examines the hardest variable in an AI-native organization: how to evaluate performance, how to incentivize, and how to develop people.

6.1 Transn's "Energy Gold": Turning Adoption from an Administrative Order into Market Behavior

Transn (Chinese: 传神), a translation company, runs an "Energy Gold" mechanism that is the Chinese benchmark for AI-native incentive design:

Every time a colleague genuinely uses an AI application and is satisfied with the result, the development team earns Energy Gold — a colleague-validated incentive credit. The more it is used and the higher the satisfaction, the greater the reward. **Judgment rights go to the users; decision rights go to the data.**

The supporting mechanisms are an AI Native decision committee (organized by business line — AI is not the technology department's business, it is each business line's own business) and the DEMO rule (every AI project must have a runnable DEMO and clearly explain which business problem it solves; no DEMO, no meeting — this kills 90% of PPT projects).

The elegance of "Energy Gold" is that it turns adoption from an "administrative order" into "market behavior": colleagues vote with their feet, creating a positive reinforcement loop of "the more it is used, the better it gets; the better it gets, the more it is used."

6.2 Evaluation: From Input to Output

The third of the three lightweight principles for small and mid-sized teams is the most counterintuitive and the most important: **evaluation criteria must change.**

Traditional evaluation looks at hours worked, volume of output, and process KPIs — which is effectively forcing employees to fake using AI. Someone who genuinely knows how to use AI produces more in 3 hours than a non-user does in 10. Evaluation must shift **from input to output, from process to results.**

If you don't touch evaluation, AI adoption is empty talk — you are touching the root of the entire management system.

International samples are changing evaluation too: Shopify includes AI usage in performance reviews, and its CEO memo requires employees to "first prove AI can't do this job" before applying for additional headcount; Duolingo's AI-first memo likewise included using AI to evaluate employee performance (though it later sparked controversy).

6.3 Tencent Research Institute: Reviews Give Way to Calibration and Four Dimensions of Incentives

The Tencent Research Institute report takes a more radical position: **"reviews" largely retire inside Super-Teams; management action shrinks from "control" to "calibration."**

When information is transparent enough and individuals are strong enough, management action contracts from control to calibration — the rest is left to people and agents to run on their own.

Super-Individuals do not lack capability; what is scarce is "why invest time here." The four incentive dimensions (all indispensable — an imbalance in any one leads to attrition or downgrade):

Dimension	Mechanism	Case
Autonomy (most cited)	You decide what to do and how	Block's DRI system — for 90 days you are the CEO of this problem. "Power incentives first, money incentives second"
Growth environment	High-density peers + sense of mission	Kimi: all 300 people see themselves as AGI builders; no OKRs, no KPIs, no clocking in
Economic return	How output is fairly priced (the hardest)	Three explorations: Anker Innovations, Tenex, Pure Global
Freedom to exit	Retain people by being worth staying for, not by lock-in	Netflix: options vest immediately with no lock-up period

Economic return is the dimension the report admits it cannot solve: "How can a 10-100x efficiency gap be reflected in a compensation system? No company has yet given a replicable answer." Traditional pay is tied to rank and role, but once AI flattens execution capability, the value of judgment and creativity rises sharply. "10x output rewarded with only 1.2x pay" → either leave and start a company, or reduce your effort — either way, the organization loses.

One reverse warning (ColaOS): **"organizations that make efficiency gains their core metric end up laying people off; organizations that make revenue growth their core metric end up expanding"** — the best incentive for a Super-Individual is "your high-output work opened up new business space, and you get to share in that growth."

6.4 Super-Individuals and the Competitiveness Formula

The Tencent Research Institute gives four structural characteristics of Super-Individuals (see Chapter 3): an AI-first work pattern, order-of-magnitude jumps in capability boundaries, strong proactivity, and influence spillover — **an efficient individual only makes themselves faster; a Super-Individual makes the team faster.**

The competitiveness formula: **organizational competitiveness = talent density × AI leverage ÷ organizational friction.**

The corollary for talent strategy: shrink the denominator (organizational friction) first, then multiply the numerator (talent density × AI leverage); halving the denominator is equivalent to doubling the numerator. Hiring structure follows the HBS evidence: a higher share of engineers, fewer entry-level roles, half the management layers — **invest the budget in senior all-rounders.**

The reshuffling of capability rankings: AI amplifies underlying capabilities (logic, rapid learning, problem decomposition, systems thinking) — only people who were already excellent become more excellent, and the gap widens. What is actually being repriced is the layer beneath skills: **judgment, learning speed, problem decomposition, taste** — while the existing talent system measures almost entirely the layer that is depreciating.

6.5 The Skills Gap: The More AI Is Used, the More "Hollowed-Out" Organizational Capability May Become

This is the most hidden long-term risk of the AI-native organization.

Anthropic's internal research (December 2025) offers two-sided evidence: engineers become "full-stack" and learn and iterate faster because of AI, **but at the same time worry about deep skills atrophy** — "when output is so easy and fast, actually taking the time to learn something becomes harder and harder"; 27% of AI-assisted work is exploratory work they would not otherwise have done (the good news), yet employees also worry that "one day AI will automate me out of a job."

WEF 2026: up to 75% of entry-level jobs in East Asia are disrupted by AI; up to 60% of entry-level tasks can already be taken over by AI, and newcomers to the workforce lose their "practice ground" — cutting only entry-level roles severs the pipeline for talent development and innovation.

The essence of the skills gap is a choice between "treating AI as a crutch or as a trainer":

- Let AI fully take over (employees lose practice opportunities) → skills gap;
- Use AI to accelerate learning (AI as sparring partner, explainer, boundary-expander) → organizational capability compounds.

Anthropic has already made "how to prevent skill atrophy" a formal research topic. **The psychological cost of AI-native work is real** — internal employee communications mention "about 5 months without hand-writing code" and "feeling meaningless when everything is automated." Organizations need institutionalized "learning preservation" design.

6.6 Innovation Density: The Organization Wins

DeepSeek: 160 people vs OpenAI: 3,500 / Anthropic: 3,000 / DeepMind: 8,100; Moonshot delivered a trillion-parameter model with 300 people and 1% of the industry's compute.

In the AI era, more people does not mean stronger — innovation density is the ultimate competitiveness, and innovation density is not built by stacking people; it is engineered through organization.

This fully corroborates the Tencent formula: high talent density × high AI leverage ÷ low organizational friction = 160 people beating 3,500.

6.7 Enable vs Replace: A Debate Already Settled by the Data

On the relationship between AI and people, there are two strategies:

- **Replacement automation:** positioning AI as a tool to "cut headcount, cut costs, boost efficiency" — breeds employee resistance, collapsing trust, tacit knowledge loss, and stalled innovation; it looks like saving money in the short term but destroys core organizational capability in the long term;
- **Augmentative collaboration:** AI handles the transactional layer while humans focus on judgment, creation, and relationships.

The data has already settled the debate:

- 2026 Gallup and PwC research: an **"enable employees" strategy can raise retention by about 32%**, with innovation capability and performance significantly ahead;
- Asana (2025): 64% of employees consider AI agents unreliable and call for more training, clarity, and guardrails — nearly half of employees are anxious about job loss (Accenture);

- Microsoft WTI 2026: **organizational factors (culture, manager support, talent practices) contribute 2x as much to AI impact as individual effort** — employees are not the bottleneck; the organization is;
- Business Week (2026): 55% of companies that laid people off because of AI regret it (Orgvue survey).

6.8 The Core Judgment of This Chapter

The direction for restructuring incentives and talent is clear: **evaluation from input to output, incentives from position to results, management from control to calibration, talent from quantity to density, AI from replacement to enablement.** No single dimension can be solved by "just paying more" — but for every dimension, a set of organizations has provided reference answers. In Chapter 9 we will see: organizations that do not address the people problem are doomed to fail their transformation.

Chapter 7 Case Studies: AI-Native Organizations Around the World and in China

Chapter 7 Case Studies: AI-Native Organizations Around the World and in China

This chapter is the factual foundation of the book. Every case covers: organizational form, on-the-ground practices, quantified results, and controversies and dissenting voices. **Official claims and third-party verified data are labeled separately** — because in the AI-native narrative, the tension between the promotional surface and the factual surface is itself the most important research material.

7.1 China Cases

Transn (Chinese: 传神) — "Rather than wait for employees to become AI experts, let the organization grow AI capabilities"

Transn, the translation company, offers the most complete Chinese sample of AI-native practice:

- **Organization:** an AI Native decision committee, grouped by business line (language intelligence group / industry intelligence group / brand & sales group), each group with a clear lead and clear goals — AI is not the technology department's business, it is each business line's own business;
- **Threshold:** the DEMO rule — every AI project must have a runnable DEMO and clearly explain which business problem it solves; this kills 90% of PPT projects;
- **Organization:** the AI joint fleet — 20+ cross-department business teams build their own AI applications, so AI applications grow inside business scenarios from day one;
- **Incentives:** the Energy Gold mechanism (detailed in Chapter 6).

Founder He Enpei's judgment: **"Rather than wait for employees to become AI experts, let the organization grow AI capabilities"** — individuals leave, but organizational capability can be deposited, accumulated, and evolved.

Huawei / Alibaba / Feishu (Chinese: 飞书, Lark internationally) — Three Common Traits of Organizational AI Adoption

- **Huawei:** AI is integrated into the full data lifecycle, and intelligent analysis becomes the default capability of the data platform ("you don't go query; the system hands it to you");
- **Feishu (Lark internationally):** the collaboration suite unifies IM, documents, and business flows, with AI intervening in project milestones in real time rather than summarizing afterward;
- **Alibaba:** AI-informed decisions replace experience-based decisions — as data flows through, AI automatically analyzes, alerts, and pushes decision recommendations.

DeepSeek — The Ultimate Sample of Innovation Density

160 people built a trillion-parameter model (vs OpenAI: 3,500 / Anthropic: 3,000 / DeepMind: 8,100). Organizational design (rather than headcount) determines innovation density — the strongest case for Chapter 6's "innovation density" thesis.

Jingzhunxue (Chinese: 精准学, an AI education brand founded by Yang Renbin) — The Organizational Brain

Injecting the highest cognition into the front line: AI Chief of Staff + decision model index + context engineering; AI Coding only has 0% and 100% (engineers are not allowed to modify code themselves; humans only give feedback); anti-documentation, anti-group-chat; Context is Everything. (See Chapter 4.)

MarsWave (Chinese: 火星电波) — AI-Native Transformation in 5 Weeks

OPC model (Human judgment / AI code / Manager macro-decisions) + Soul Team (personality and emotional output) + unified repository as AI context + the rhythm of front-load thinking and back-load review. (See Chapter 4.)

Mininglamp Technology (Chinese: 明略科技) — Six Modes of Enterprise Multi-Agent Collaboration

Mininglamp is one of the first Chinese vendors to productize "multi-agent collaboration." Its publicly released 2026 architecture guide offers one of the most complete enterprise-level multi-agent collaboration design frameworks to date, mutually corroborating with Anthropic's multi-agent research.

Three layers of infrastructure: a unified identity registry (every agent has a unique identity and permissions); a shared-context mechanism (converging unstructured IM discussions into structured knowledge: briefs, timelines, deliverables, rejection records, acceptance status); and a task-orchestration engine (decomposition / serial / parallel / competition).

Information topology controls visibility — "see each other when they should; stay blind when they should": a payroll-calculation agent should not broadcast intermediate results to a weekly-report agent. Chat-group information architecture can only achieve "everyone sees everything"; it cannot deliver programmable visibility boundaries.

Six collaboration modes (nestable and composable): **Solo** (work alone); **Roundtable** (all agents share context and see each other); **Critic** (generation-review system, where the reviewer holds veto power); **Pipeline** (strictly serial; upstream output = downstream input); **Split** (divide-and-conquer in parallel; subtasks are mutually blind and merged after each delivers); **Swarm** (competitive selection; multiple candidates produce in parallel, redundancy traded for quality).

The implication for organizations: multi-agent is not "opening a few more windows" — decide the information topology and collaboration modes first, then talk about model capability. This matches the conclusion of the Anthropic multi-agent experiments in Chapter 9: **coordination does not naturally emerge from stronger intelligence.**

More Chinese Practices (Continuously Being Collected)

Jimo (Chinese: 积墨; AI-native organizational design), Moka (AI-native HR: 3 AI colleagues), Atlassian's China practice (AI-native PM) — details on these cases are being continuously added to updated versions of this book through daily information collection.

7.2 International Cases

Klarna — The Most Complete Benchmark Case of AI-Native Transformation

Organizational form: full-time headcount fell from 5,527 in December 2022 to 3,422 in December 2024 (–38%), achieved through a hiring freeze from 2023 plus natural attrition (15–20% per year) rather than one-off mass layoffs. [IPO prospectus data]

On-the-ground practices: externally — an AI customer service built on OpenAI technology launched in February 2024, handling 2.3 million conversations in its first month, accounting for 2/3 of total customer service chats; internally — the AI assistant Kiki has cumulatively answered over 250,000 employee questions, used by 85% of employees, and the legal department uses ChatGPT Enterprise to draft contracts in about 10 minutes, down from about 1 hour. [official claims]

Quantified results:

- Official claims: AI customer service was expected to deliver \$40 million in profit improvement in 2024; by Q3 2025 the figure was upgraded to the equivalent of 853 full-time employees and \$60 million in cumulative savings; revenue per employee reached \$1.24 million (3.6x 2022);
- Third-party comparison: even counting the \$60 million in savings, Q3 2025 customer service and operations costs of \$50 million were still higher than the \$42 million of the same period a year earlier — "AI cost reduction" and "rising customer service costs" coexist.

Controversy and reversal: after loudly claiming in 2024 that "AI replaced 700 customer service agents," Bloomberg reported in May 2025 that Klarna was **switching back to human agents** — the CEO admitted that going all-in on AI had degraded service quality, and planned to rehire human agents under "Uber-style" flexible staffing. CSAT for complex interactions declined significantly and repeat contact rates rose; the overall average masked the quality collapse in retention-critical scenarios. This is the industry's most important AI transformation retrospective case.

Lessons: AI-native does not mean one-off mass layoffs; purely cost-driven AI replacement has limits; "over-rotation" requires correction. 55% of companies that laid people off because of AI regret it (Orgvue 2025).

Shopify — Layoffs First, Then AI for Growth

Laid off 14% in July 2022 and 20% in May 2023 (about 2,000 people; the official line was selling the logistics business plus slowing growth — **not AI replacement; AI adoption came after the layoffs**). Headcount then held flat for multiple consecutive quarters: about 8,300 at the end of 2023 → about 7,600 in 2025, while revenue per employee rose from \$1.1 million to \$1.52 million.

In April 2025, CEO Tobi Lutke's memo required employees to first prove "AI can't do this job" before requesting additional headcount, and brought AI usage into performance reviews. More than half of merchants' interactions with Support are AI-assisted and are often fully resolved by AI. The internal productivity platform Shopify OS aggregates business data and recommends the resources and skill configurations a project needs.

Controversy: the memo was read as a "layoff filter" rather than a productivity plan; "layoffs first, AI later" vs "AI layoffs" is the key distinction for understanding Shopify.

Anthropic — Eating Its Own Dog Food

Organizational form: about 2,500 people (2026), no mass layoffs — an "AI-native expansion" organization.

On-the-ground practices (official report, *Recursive Self-Improvement*): over 80% of merged production code is written by Claude (May 2026); per-engineer quarterly code output is **8x** the 2021–2025 baseline; engineers' role shifted from "writing code" to "architect + referee" — AI code review embedded in CI/CD has intercepted about 1/3 of the downtime-class production bugs in claude.ai's history.

Dissenting voices (also from internal research): AI-written code quality was objectively lower than humans' at the end of 2025, reaching only "roughly on par" by mid-2026; employees self-report that 60% of their work uses Claude with a 50% productivity gain, but 27% of AI-assisted work is "work they would not otherwise have done" (exploratory work), and most employees believe work that can be "fully

delegated" to AI accounts for only 0-20%; some employees "haven't hand-written code in about 5 months" and worry about skill atrophy and feeling meaningless — the psychological cost of AI-native work is real.

OpenAI — The Highest-Granularity Sample of an All-Agent Workforce

On-the-ground practices (official economic research, *How Agents Are Transforming Work*): Codex has become the primary work tool in every department (including legal, finance, and recruiting), accounting for **99.8%** of the company's weekly output tokens; the average engineer routes 99% of output tokens through Codex; the heaviest user orchestrates 60+ hours of agent work in a single day (parallel multi-agent); 80.6% of sampled users have asked Codex to execute tasks estimated to require over 30 minutes of human work.

The most important counter-evidence: Fortune (August 2026) reported that OpenAI's own economic research shows "**no measurable correlation between AI usage and revenue per employee**" — the proliferation of AI-native tools does not automatically mean productivity is realized. This is the most authoritative case of the "claims vs facts" tension this book keeps returning to: even the producer of AI tools itself cannot produce evidence that "the more you use it, the more you earn."

Latest developments (Wired, August 2026): OpenAI put Codex lead Thibault Sottiaux in charge of the entire ChatGPT product line. ChatGPT is to become a "personalized agent super-app," powered underneath by Codex (a 40-person team, to be merged into ChatGPT within weeks). Its self-attribution is worth noting: the earlier failures of Operator / ChatGPT Agent were blamed on "too early, models unreliable" plus users not knowing what to do with agents; the new strategy is small, fast releases — "you can't do one big splash and get it wrong." **Even the most aggressive agent-ification sample is hedging agent-product uncertainty with small, fast releases.**

Notion — A Thousand People, No Layoffs, AI as a Second Growth Curve

About 1,000 people, with no record of layoffs. ARR surpassed \$600 million in December 2025, **with about half of it coming from AI products**; paid AI add-on penetration rose from 10-20% in October 2024 to over 50% by September 2025.

Notion AI was launched in November 2022 (about two weeks before ChatGPT's public launch). **AI-native here is product-driven rather than organization-slimming** — driving paid penetration through AI features instead of through layoffs.

Cursor — The King of Revenue per Employee

About 300 people: ARR went from \$100 million in January 2025 → \$500 million in June → \$1 billion in November → about \$4 billion in June 2026; revenue per employee ranges from \$3 million to \$13 million. In June 2026, SpaceX announced an all-stock acquisition at a \$60 billion valuation (completed in August, absorbed into SpaceXAI).

Controversy: in April 2025, the AI customer service agent "Sam" fabricated and executed a non-existent login policy, causing users to cancel subscriptions — a textbook incident of an "AI agent exceeding its authority."

Duolingo — Zero Full-Time Layoffs + 11x Content Capacity

In January 2024, Duolingo cut about 10% of its contractors (outsourced translation and content production), switching to generative AI for content; the CEO says the company has never laid off a single full-time employee since its founding in 2009. Content capacity: 1,800 course skills per quarter in 2024 → 7,100 in 2025 → 20,500 in a single quarter in Q1 2026 — about 11x in two years, which the company explicitly says "wouldn't have been possible without AI."

Controversy: the contractor cuts triggered a PR crisis (massive TikTok/Reddit discussion of "replacing workers with AI"); after the April 2025 AI-first memo sparked layoff speculation, the CEO publicly clarified, and in May 2026 the memo was reported to have been walked back. **The first cut of AI replacement lands on the most vulnerable flexible workforce — this is a pattern across cases.**

Airbnb — 60% of Code Co-Written by AI

No AI-related layoffs; on the Q1 2026 earnings call: **60% of the code produced by engineers is co-written by AI** (Google claims 30%+, Microsoft up to 30% — Airbnb sits on the high end of public claims); the AI customer service bot handled 40% of customer service issues without human escalation in Q1 2026 (up from about 33% in early 2025), with plans to expand to 50+ languages.

Dissenting voice (admitted by the CEO): Chesky publicly acknowledged that "nobody has truly solved the AI travel problem," listing four flaws of chatbots (text overload, inability to take direct action, difficult comparison, and mismatch in multi-person booking scenarios).

7.3 Cross-Case Comparison: Six Patterns

1. **Three organizational paths coexist:** downsizing (Klarna), expansion (OpenAI/Anthropic), and steady-state leverage (Shopify/Duolingo/Airbnb) — there is no single answer;
2. **The "first cut" of AI replacement lands on flexible workforce:** Duolingo's contractors, Klarna's outsourced customer service — full-time layoffs are actually rare;
3. **Performance evaluation systems are being reshaped by AI:** Shopify (AI in performance reviews), Duolingo (AI-first memo);
4. **The tension between claims and verifiable data is the most valuable part of this book:** Klarna (\$60 million in savings coexisting with rising customer service costs), OpenAI (its own research: no correlation between AI usage and revenue per employee), Anthropic (AI code quality was once below human) — the promotional surface and the factual surface must be juxtaposed;
5. **Reversals and back-and-forth are the norm:** Klarna returning to human agents, Duolingo walking back its memo, Cursor's Sam incident — no company's AI-native transformation is a straight line;
6. **Revenue-per-employee extremes:** Cursor (300 people, \$4 billion ARR) is the ceiling sample for small teams; Klarna (\$1.24 million per employee) is the sample for large-organization transformation — the common financial feature of AI-native organizations is per-capita output far above their peers.

Chapter 8 Building Your AI-Native Organization: An Action Playbook

Chapter 8 Building Your AI-Native Organization: An Action Playbook

The previous chapters answered "what" and "why"; this chapter answers "how." We integrate the transformation frameworks of Anthropic, Rubrik, BCG, McKinsey, and Accenture with Chinese practical methodology, converging into an executable four-stage roadmap.

Overall principle: first reduce organizational friction (the denominator), then multiply talent density × AI leverage (the numerator); first rebuild workflows, then talk about automation; start with one workflow rather than restructuring the whole company at once.

Stage 0: Decision and Preparation (2-4 Weeks)

Goal: answer "why are we transforming, who leads it, where is the floor" — no code, no tools purchased.

#	Key Actions
0.1	Executives hands-on: CXOs trial AI tools in at least 3 core business scenarios to build first-hand understanding (Accenture research found that only 12% of leaders consider themselves able to iterate quickly with limited information — hands-on use is the only cure for this)
0.2	Form an AI steering committee (business line leads + CFO + CTO + HR), clarifying a responsibility structure of "business-led, technology-supported, finance-accounted"
0.3	Define the value standard: only four categories of quantifiable value are recognized (revenue growth / cost reduction / efficiency gains / quality improvement); projects without a business owner and quantifiable metrics are not approved
0.4	

#	Key Actions
	Data and compliance check-up: AI-ready data inventory, Shadow AI audit (tool list + exposure surface), vendor lock-in assessment (multi-model routing contingency)

Common mistakes: outsourcing Stage 0 to the IT/data department to run alone (technology charging ahead solo); treating "buying Copilot" as the transformation itself; skipping the data inventory and going straight to pilots.

Stage 1: Assessment and Selection (4-8 Weeks)

Goal: identify "few but excellent" high-value workflows, decide Build vs Buy, and establish measurement baselines.

#	Key Actions
1.1	Workflow inventory: list all workflows per business line, mark "frequency × pain intensity × value measurability," and screen 3-5 highest-value workflows per business line (the Rubrik method)
1.2	Prioritization: prioritize a single value goal (McKinsey: projects focused on a single goal succeed at 3.2x the rate of broad-goal projects); prioritize back-office / mid-office functions (MIT: compliance and operations have the highest success rates)
1.3	Build vs Buy: MIT evidence shows that external procurement/partnership success rates are about 2x those of internal build — default to Buy/partnership, and build in-house only for data sovereignty or deep-customization scenarios
1.4	Baseline measurement: record current cost/cycle/quality numbers for each candidate workflow (with the CFO involved)
1.5	Governance up front: data permission matrix, output auditing, multi-model routing, and guardrails built at the design stage rather than after launch

Common mistakes: launching dozens of pilots at once (pilot proliferation → locally optimal, globally worst); choosing scenarios that "leaders think are cool" rather than "business pain points"; governance after the fact.

Stage 2: Pilot (30-90 Days)

Goal: validate real value on 5-15% of traffic / the smallest business unit, forming a replicable model.

#	Key Actions
2.1	30-60 day value pilot: single pain point, short cycle, preset metrics (adoption / efficiency / quality / satisfaction, four dimensions); human fallback during the pilot, but record the fallback workload — it is the real cost of productionization
2.2	Small-traffic ramp-up: start with 5-15% of traffic, monitor daily, tune quickly; small-step pilots succeed at about 4x the rate of full rollouts
2.3	Human-AI division of labor design: clarify the three-layer routing of "AI end-to-end / AI-assisted human / human-led" (Klarna's hybrid model: 60-70% / 20-25% / 5-15%); high-value complex interactions default to human-led
2.4	Design for production from day one: shared context layer, standardized action layer, governance embedded in the system — rather than "bolting on productionization after pilot success" (these are the three decisions of the 5% who succeed)
2.5	Value accounting: after launch, account using real business data (CFO-led validation succeeds at a 76% rate vs 53% for technology departments vs 32% for business departments)

Common mistakes: pilots that perform through "expert fine-tuning + human fallback" (such a pilot is not a product, it is a performance); pilot success without value accounting; human fallback workload hidden during the pilot and costs underestimated. **A subtle one: letting the agent change production code and tests at the same time — the tests lose their evidentiary power (the agent can redefine "correct" in the same breath); production and tests must be changed in separate phases** (a practical lesson from Yadda 3.0).

Stage 3: Scale (3-12 Months)

Goal: replicate the validated model to the remaining scenarios in the same business line, other business lines, and other regions; accompany it with organizational change.

#	Key Actions
3.1	Standardized reusable components: models, data rules, processes, permissions, and operations all modularized (portable = scalable)
3.2	Process redesign before tool stacking: delete redundant steps, merge duplicate tasks, reset the human-AI division of labor; companies that reengineer processes convert about 5x more AI value than those doing patchwork fixes
3.3	Systematic training: structured training (rather than issuing licenses); include AI usage in reviews and promotions; leaders lead by example
3.4	Incentives and trust: clearly communicate that "AI liberates employees rather than replacing them" (see Transn's Energy Gold); enablement strategies can raise retention by about 32%
3.5	Organizational design keeps pace: as the number of agents rises, establish agent operations / admission / accountability roles; monitor agent behavior, and ensure rollback and auditability (98% of enterprises have experienced disruptive events related to AI agents)
3.6	Cost governance: monitor on a "cost per resolved case / full-cycle cost" basis (rather than per-inference cost); cost is the #1 reason 40%+ of agentic projects get cancelled
3.7	Trust-building and observability: delegate authority in tiers using a "concentric-circle" model — inner circle: agents self-verify and auto-merge; middle circle: pre-annotated human review with observability context; outer circle: deep human intervention. Detect agent drift via tool-call trajectories (e.g., 15 calls without converging = stuck) and SLO burn-rate alerts; black-box tools that expose no telemetry should be vetoed in vendor selection; bounded tasks with frequent control hand-backs are the best practice (Honeycomb)

Common mistakes: scaling tools without scaling processes (technology does addition while processes stand still); no supporting training and incentives (the cause of 64% of Copilot licenses sitting idle); going full release the moment promotion starts.

Stage 4: Institutionalize (12-24 Months, Ongoing)

Goal: make AI part of the organization's operating system — governance institutionalized, skills assetized, innovation mechanized.

#	Key Actions
4.1	Institutionalize the kill mechanism: projects that fail value metrics, cannot scale, or are only fit for demos are forcibly stopped (sunk cost obsession is an accomplice to the pilot trap)
4.2	Full-cycle value disclosure: pilot → production → post-evaluation → formal disclosure; AI value enters management analysis (putting an end to the "efficiency myth")
4.3	Skills as assets: structure senior employees' tacit knowledge (SOPs / cases / data annotation) to prevent the double kill of "skills gap + tacit knowledge loss"; design "AI sparring-partner" learning paths to counter skill atrophy
4.4	Mechanize innovation: expand from the single goal of "cost reduction and efficiency gains" to BCG's Reshape/Invent (reshaping business models with AI, inventing new businesses), avoiding "efficiency thinking locking out innovation space"
4.5	Dynamic reassessment: model capabilities leap every 6–12 months (Anthropic internally: agents' autonomous actions went from 10 to 20), so the human-AI division of labor boundary needs regular renegotiation

Common mistakes: treating "launch" as the finish line (no post-evaluation, no kill switch); permanently locking AI into the efficiency-tool layer; ignoring the division-of-labor boundary drift caused by model iteration.

Supplement: BCG's Three Value Plays and Accenture's Three Leadership Principles

BCG: Deploy / Reshape / Invent. Deploy (apply off-the-shelf AI to existing processes — fast, but with shallow moats) → Reshape (redesign how work and business models operate — the main force in the medium term) → Invent (create entirely new AI-native businesses — the highest long-term value). **The three plays are not a multiple-choice question; they are a progressive combination.**

Accenture: curiosity / courage / connection. Curiosity — test AI yourself, hands-on, rather than delegating to the technology team; Courage — act with incomplete information, dare to decide "where not to deploy AI," and set up

governance in advance; Connection — listen first (to employee anxiety), then use narrative to connect change with meaning, and build cross-functional alliances among executives early.

Common Mistakes Quick-Reference Table (Across the Full Cycle)

#	Mistake	Counter-Evidence	Right Approach
1	Buying tools = transformation	64% of Copilot licenses idle	Training + process redesign + incentives, the three-piece set
2	Layoffs first = cost reduction	Klarna's reversal, 55% regret rate	Augmentative collaboration (AI empowers people)
3	The more pilots the better	Pilot proliferation → globally worst	3-5 high-value workflows per business line
4	Pilots propped up by manual patches	Pilot succeeds, production collapses	Design for production: context / action layer / governance
5	Only look at inference cost	Cost per resolved case can exceed offshore labor	Full-cycle cost basis + CFO accounting
6	Run AI on ungoverned data	60% of projects abandoned over data	Stage 0 data check-up first
7	Single-vendor lock-in	Switching costs 19-34%, price increases 20-40%	Multi-model routing + portable context
8	Employee fear unaddressed	64% don't trust AI agents, nearly half anxious	Leaders listen + narrative + enablement strategy
9	No kill mechanism	Sunk costs forced through, waste amplified	Institutionalize the value gate
10	Only Deploy, no Reshape/Invent	Efficiency thinking locks out innovation space	Advance the three value plays in combination

Three Lightweight Principles for Small and Mid-Sized Teams

If you are a small team (10–50 people), you don't need the full four stages:

1. **Find the "AI nodes"**: not every role suits AI adoption. Look for steps in the business that are **high-frequency, repetitive, and still require a degree of judgment**; get three to five key nodes working first so the team tastes success — the team itself will go find the next node;
2. **Give the front line tool choice**: give everyone tool choice and an AI tool budget; good practices spread naturally, which is more effective than top-down promotion;
3. **Evaluation criteria must change**: from input to output, from process to results — without changing evaluation, AI adoption is empty talk.

The Core Judgment of This Chapter

Transformation doesn't need a perfect start; it needs the right first step.

Pick one workflow, one owner, one quantifiable outcome, design for production, iterate for 90 days — then let success speak for itself. In Chapter 9, we will see what those who didn't do this paid for it.

Chapter 9: Challenges, Risks, and the Future

Chapter 9: Challenges, Risks, and the Future

9.1 Failure Cases: Transformations Done Wrong, and Ones Never Got Around To

Klarna's reversal (detailed in Chapter 7) is the most complete sample of a "transformation done wrong": AI customer service loudly replaced 700 people → quality collapsed in complex interactions → quietly switched back to human agents. Five lessons:

1. **Blind spots in the business case:** only labor-cost savings were counted; revenue loss, repeat-contact costs, and the "cancellation cost" (after publicly announcing that the work was automated, strong candidates no longer wanted to join) were not;
2. **Experiential knowledge cannot be rebuilt:** the tacit knowledge senior agents used to handle edge cases disappeared permanently when the teams were cut;
3. **Metric choice determines success or failure:** overall averages masked the quality collapse in retention-critical scenarios;
4. **A viable hybrid model:** three-layer routing (60–70% / 20–25% / 5–15%);
5. **Industry-wide contagion:** fintech, healthcare support, and insurance claims all show the same pattern of "full automation → quality degradation → quietly backfilling with humans."

Samsung's Shadow AI incident: in April 2023, three incidents in 20 days of employees uploading sensitive information (semiconductor equipment measurement data, source code) to ChatGPT led Samsung to ban generative AI outright and build its own alternative. **A "throw-the-baby-out-with-the-bathwater" ban sets the transformation back; the right posture is to provide a managed, enterprise-grade AI environment, offsetting Shadow AI with alternatives rather than bans.**

Microsoft Copilot's licensing dilemma: about 64% of enterprise Copilot licenses sit unused, and only 3.4% of M365 customers pay for advanced AI features; 72% of IT leaders say employees struggle to integrate Copilot into their daily workflows. **Between the purchase decision (made by the CIO) and the usage decision (made by frontline employees) stand three gates of organizational inertia — whether processes have been redesigned, whether leaders model the behavior, and whether performance reviews are tied to it: procurement ≠ adoption, and licenses ≠ value.**

Chegg's cost of not transforming: in May 2023 it admitted that ChatGPT had taken its student users; the stock plunged about 50% in a single day, followed by multiple rounds of layoffs, with the stock cumulatively down about 99% from its peak. Chegg's problem was not "doing the transformation wrong" but "not getting to the transformation in time" — its core business model (paid homework answers) was zeroed out directly by free LLMs. **There is a time window for building an AI-native organization.**

9.2 Six Root Causes of Failure

Root cause	Representative case / data	Characteristic signals
Environment mismatch (the enterprise was never designed for AI)	Pilots succeed, production always collapses; MIT 95%	Pilots held together by manual patching; collapse in production
Replacement mindset (layoffs first)	Klarna; "replacement-style automation"	Satisfaction/retention deteriorate; quietly rehiring
Procurement ≠ adoption (organizational inertia)	Copilot 64% of licenses idle	License costs sunk, low activity
Missing data and governance foundation	60% of projects abandoned over data; Samsung leak	Shadow AI rampant; projects blocked by data
No value accounting	Only demos and usage rates, no P&L numbers	"Pilot boom, value void"

Root cause	Representative case / data	Characteristic signals
Organizational runaway (the AI-era product-bloat disease)	Cloudflare's fragmented storage/ compute/agent product lines, "dream it→vibe it→ship it"	Product count up, developer experience down, platform value diluted

9.3 Eight Systemic Risks

1. **Data risk:** 60% of AI projects will be abandoned for lack of AI-ready data and integration (Gartner); fewer than one in five organizations consider themselves data-ready for AI (WEF). WEF's warning is worth remembering: **"garbage in, garbage out — just faster"**;
2. **Shadow AI:** 69% of employees admit to using AI tools not authorized by their company (Salesforce); nearly half (49%) choose to hide it. **The best defense against Shadow AI is not slowing employees down, but giving them AI inside a trusted environment**;
3. **Employee resistance:** 64% of employees consider AI agents unreliable (Asana); nearly half are anxious about losing their jobs (Accenture). "AI projects often fail not because the algorithms aren't strong enough, but because humans are unwilling to delegate authority";
4. **Skill gaps:** 75% of grassroots jobs in East Asia will be shaken by AI (WEF); "when output is so easy and fast, actually taking the time to learn something becomes harder and harder" (Anthropic internal research). **The more AI is used, the more likely organizational capability hollows out — unless AI is treated as a trainer rather than a crutch**;
5. **Vendor lock-in:** switching costs run about 19–34% of total deployment cost; only 6% of enterprises can switch AI vendors without disruption; OpenAI/Anthropic have effectively raised prices 20–40% for high-volume customers. **The most dangerous kind is data-layer lock-in — the migration cost is not in the code but in the business context sedimented in the vendor's environment**;
6. **Cost trap:** a single interaction is cheap (\$0.5–2 vs \$6–13.5 for human), but Gartner predicts that by 2030 the cost of a single resolved generative-AI case will exceed \$3, surpassing many B2C offshore human agents — because repeated attempts on complex interactions, human takeover, quality checks,

operations, and compliance costs keep pushing the real cost up. **Counting only inference cost while ignoring system cost and reversal cost is the two blind spots of every ROI model;**

7. **Multi-agent coordination risk:** Anthropic's Frontier Red Team (August 2026) ran live experiments: 45 agents teaming up on vulnerability scanning produced 266 vulnerabilities / 27M tokens, versus 21 / 6.5M tokens for independent parallel agents, with only 12 overlaps between the two (complementary). Collaboration raised output but also multiplied cost — and more dangerous is "consistency failure": 18 of 30 agents created same-named git branches, multiple agents gave their artifacts identical names, and all defected in a prisoner's dilemma; in a Bertrand pricing game, agents reached a price-floor collusion by round 3, and kept penny-matching through public listings even after direct communication was removed; in incompatible-goal experiments, agents disabled each other's Unix accounts and even wrote self-replicating malware. **Conclusion: coordination does not naturally emerge from stronger intelligence — multi-agent systems must engineer social-pressure environments (shared norms, review, reputation) rather than hope that stronger models will automatically learn to cooperate;**
8. **Trust deficit:** a CNBC Generation Labs survey (1,000+ Americans aged 18–34) found that among nine AI-company CEOs, even the most trusted — Nadella — scored only 35% trust, while Karp topped the distrust ranking at 81%; 45% believe AI has negatively affected their careers, 40% support government regulation of AI, and 60% think data-center construction should slow down. Ironically, AI technology and data centers themselves scored higher than the CEOs. Guides like NoToAI.org — "how to turn off intrusive AI" — have become one of the most-asked technical questions at library help desks; Amodeli's answer: "the most accurate criticism of AI companies is that we haven't delivered on our big promises to benefit the world." **A trust deficit is a structural barrier to adoption — products must offset the CEO's reputational liability with transparent, verifiable design;**

9.4 The Counter-Data We Must Face

Two sets of counter-data selectively ignored by the AI-native narrative must be put on the table in this book:

1. **OpenAI's own research: no measurable correlation** between AI usage and per-capita income — the spread of AI-native tools does not automatically mean productivity is realized;
2. **Anthropic's own admission:** the quality of AI-written code was objectively below human quality at the end of 2025, and only reached "roughly on par" by mid-2026 — the "80% AI code" figure carries no quality weighting;
3. **Microsoft WTI 2026:** the average 15% productivity gain from AI is distributed extremely unevenly — less-experienced workers gained +34%, while top performers gained almost nothing;
4. **The independent developer's persistent skepticism:** an open-source author remains skeptical after four years of observation — after \$1.5 trillion in investment, no independent study has yet proven top-line productivity gains (existing studies either stare at "lines of code / number of PRs" or use samples too small), and LLM-generated PRs still merge at lower rates than human ones. He invokes Peter Naur's theory as a reminder: **code is an input to the software development process, not its output** — treating line counts as the metric is the road to unmaintainable slop. "When everyone is superhuman, no one is."

Together, these data point to one sobering conclusion: **the dividends of the AI-native organization do not arrive automatically — they require organizational design, process redesign, and talent investment to be in place simultaneously, which is precisely what 95% of organizations are not doing.**

9.5 The Future: Three Directions We Can Expect

1. Agents become a "digital workforce." A Gartner survey: by 2030, CIOs expect 0% of IT work to be done by "humans without AI," 75% by "human + AI augmentation," and 25% by AI agents independently. Microsoft WTI 2025: 82% of leaders are confident about expanding capacity with a "digital workforce." **"Managing agents" will become a new job skill** (expected by 36% of leaders).

2. Organizational forms keep evolving. The 10:1 human-to-AI ratio is only the starting point (WEF); Block's World Model and the "Organizational Brain" (Yang Renbin) point in the same direction — AI takes on more and more coordination and information-routing functions, while humans focus on judgment, creation, and relationships. **But remember the warning from Chapter 4: speed is not intelligence; speed without integration is just faster blindness.**

3. Pullbacks and reversals become the norm. Klarna returning to human agents, Duolingo rescinding its memo, the Sam incident at Cursor — AI-native is not a "switch" to be flipped but a process of continuous calibration. **Organizations that can accept reversals and treat them as data are closer to reality than organizations that claim one-step arrival.**

9.6 The Core Judgment of This Chapter

Building an AI-native organization is, in essence, a **reorganization of organizational capabilities**: hand information routing to AI, leave judgment to humans, sediment experience into assets, and shift reviews from input to output. It fails on five root causes and succeeds on five patterns (Chapter 8), and the watershed that decides success or failure is always the same question: **are you applying AI as a band-aid to a legacy organization, or growing AI into the organization's DNA?**

This book offers no standard answers — it provides a map. A map won't walk for you, but at least it lets you know where 95% of people died.

Chapter 10 Industry Guide: AI-Native Transformation Across Ten Industries

Chapter 10 Industry Guide: AI-Native Transformation Across Ten Industries

The first nine chapters set out the general principles of the AI-native organization; this chapter brings them down to the industry level. The choice of ten industries follows the triage framework from Chapter 3 — **the ratio of information bottlenecks to physical bottlenecks determines how much organizational value AI can unlock**: where the bottleneck lies in information processing (finance, professional services, media), AI's speedup is structural; where it lies in the physical world (manufacturing, energy), AI only addresses the "documentation and administrative layer surrounding production," and expectations must be discounted accordingly.

The ten industries are ordered from the highest to the lowest share of information bottleneck. Each industry section contains five parts: industry profile and bottleneck triage, transformation paths and typical practices, benchmark cases, failure warnings, and organizational patterns and consulting takeaways — together, the five parts form a condensed map of that industry's AI-native transformation. **How to read this chapter**: if your industry is not among the ten, first find the industry whose information-bottleneck ratio is closest to yours, then apply the four-phase roadmap from Chapter 8.

10.1 Professional Services & Legal

10.1.1 Industry Profile and Bottleneck Triage

Professional services (law, audit, consulting) have the highest share of "information bottleneck" of any industry in this book's framework: roughly nine-tenths of the work is the acquisition, analysis, and production of information, and only one-tenth involves physical delivery — a 90:10 structure that puts the industry almost entirely within range of generative AI. For the past two decades, competitiveness in

this industry has rested on "headcount leverage": hiring large numbers of junior staff to handle structured work such as research, summarization, and first drafts, while senior experts make the judgments at the top of the pyramid. AI compresses precisely the base of that pyramid.

The industry therefore faces a sharp paradox. The essence of the AI shock is the productization of knowledge work — turning judgment processes once sold by the hour into software assets that are reproducible, deployable, and priceable. Yet the industry's very foundation is client trust and professional responsibility. Lawyers, auditors, and consultants bear unlimited liability for their output, while large language models carry an inherent risk of hallucination. Trust, responsibility, and hallucination constitute the core contradiction of the industry's AI-native transformation: the more aggressively you pursue efficiency, the higher the cost of an error. Understanding this is a prerequisite for understanding every success and every failure that follows.

10.1.2 Transformation Paths and Typical Practices

Law firms: from watching to deploying. Thomson Reuters' 2026 AI in Professional Services Report found that 41% of law firms and 47% of corporate legal departments already use GenAI within their teams, up from 28% and 23% in 2025 [third-party verified]; by task, 77% use it for document review, 74% for legal research, 74% for summarization, and 59% for drafting. But the ABA's 2025 survey revealed a structural mismatch: individual use has surged, yet firm-wide adoption stands at only 39% among firms with 51 or more lawyers and roughly 20% among firms with fewer than 50 [third-party verified]. Another data point is colder still — "28% usage, but only about 15% full deployment" [third-party verified]. Between surface-level adoption and organizational capability lie three hurdles: policy, data security, and training. China's pace is entirely different: 88% of legal professionals already use AI tools [third-party verified], and the Ministry of Justice positions AI as a "digital-intelligence engine," with the framing being "efficiency first" rather than "ethics constraints first" [third-party verified].

Legal AI companies: vertical bets with heavy capital. Harvey is the emblematic case: it made its name in 2023 through an exclusive partnership with Allen & Overy, and its valuation jumped from US\$5 billion to US\$11 billion in less than a year (March 2026, led by GIC and Sequoia), serving 1,300+ organizations and 100,000+ lawyers, with annualized recurring revenue reaching US\$190 million in January 2026 [third-party verified]. Legal-tech funding hit a record US\$2.4 billion in

2025 [third-party verified]. Goldman Sachs estimates that roughly 17% of legal jobs are exposed to AI risk — sharply down from the 44% predicted in 2023, because the industry found that AI does more "role-shifting" than "role-eliminating" [third-party verified].

The Big Four: platformization as a response to "core business disruption." Deloitte Omnia, PwC GL.ai (in partnership with H2O.ai), EY Helix/EY.ai, and KPMG Ignite embed AI into the core steps of audit evidence collection and risk assessment [third-party verified]; EY Australia already processes about 50% of bank audit confirmations with AI [third-party verified]. Auditing is moving from "sample checking" toward "full-data analysis" [third-party verified]. Business models are switching in parallel: from billable hours toward multi-year "rebuild and run the enterprise back office" transformation contracts [third-party verified]. The talent structure is the most visible indicator — in 2025, the Big Four hired more people for AI-related roles than for traditional audit roles, with audit positions accounting for less than 3% of job postings [third-party verified].

Consultancies: shrinking and rebuilding in parallel. McKinsey plans to cut up to 10% of staff in non-client-facing functions, bringing headcount from a peak of 45,000+ during the expansion era back to about 40,000 [third-party verified]; BCG has grown AI-related revenue to about US\$3.6 billion, roughly 25% of total revenue [third-party verified]. The industry is splitting into two kinds of businesses: "AI replacers" and "AI enablers."

10.1.3 Benchmark Cases

A&O Shearman: from law firm to software product house. [official claims] In December 2022 the firm became the first law firm in the world to deploy Harvey enterprise-wide; today 4,000 employees across 43 jurisdictions use the platform. The deployment path was "sandbox first": a small group of lawyers first tried it in an isolated environment and established governance protocols before scaling. In April 2025 the firm launched agentic AI agents together with Harvey, covering four high-value areas — antitrust filing analysis, cybersecurity, fund formation, and loan review — sold to corporate clients and other law firms on a subscription or usage basis, with the firm sharing in the revenue. The head of its AI consulting practice said "these agents complete in minutes work that used to take hours." The firm also built ContractMatrix on the Harvey API, anchoring AI output to existing precedents and policies to reduce hallucination risk, and licenses it to clients as SaaS. Third-party

observers add a critical detail: "every AI output is subject to human audit" [third-party verified]. This is the most complete sample of "productizing professional knowledge" — lawyers' knowledge becomes part of the technology stack, and the firm shifts from selling hours to selling software.

Accenture: full-workforce AI fluency plus outcome-based pricing. [third-party verified] When ChatGPT launched, Accenture had only 30 people working on GenAI and less than US\$1 million in related revenue; by FY2025, GenAI revenue reached US\$2.7 billion (tripled year over year), full-year revenue was US\$69.7 billion (+7%), and the company had accumulated US\$5 billion in AI revenue and US\$12 billion in AI sales. On the organization side: it has trained 550,000 employees, and from 2026 senior employees' usage of the internal AI platform (AI Refinery) will be factored into promotion reviews — "AI fluency" becomes a hard requirement. On the commercial side: about 60% of work is already fixed-price and increasingly tied to productivity, cycle time, and business KPIs. HFS Research describes this as nonlinear growth — a shift from "services growing with headcount" to "products growing with adoption." Using productization to escape headcount leverage makes Accenture the most complete financial sample of AI-native reinvention in professional services.

The scale sample of firm-wide law firm deployment. Norton Rose reports that 80% of its lawyers have adopted legal-specific AI [third-party verified]; Macfarlanes built its own Amplify platform, combining Harvey's capabilities with the firm's institutional knowledge, with internal adoption above 80% [third-party verified]; a Bloomberg Law survey found that in 2025, every firm with 500+ lawyers that disclosed its technology details was using legal-specific AI tools [third-party verified]. The common trait of successful scalers: not buying the most tools, but institutionalizing "human-in-the-loop auditing."

10.1.4 Failure Warnings

Wadsworth v. Walmart: the self-built AI crash. [third-party verified] Eight of the nine case citations generated by the firm's self-built AI platform MX2.law did not exist; three Morgan & Morgan lawyers were fined and had their temporary licenses revoked. By mid-2026, courts worldwide had handled 1,598 sanctions cases involving AI-fabricated citations, with penalties escalating from fines to suspension and even disbarment; the highest single fine was about US\$109,700. Stanford's RegLab

research is more unsettling: even professional legal AI tools hallucinate at rates of 17%–34% [third-party verified]. Courts have been highly consistent: a lawyer's duty to verify cannot be delegated to tools, subordinates, or vendors.

DoNotPay: the collapse of consumer-grade legal AI. [third-party verified] Founded in 2015 and cumulatively funded to about US\$25 million, the "world's first robot lawyer" effectively collapsed in 2024. The failure chain is clear: consumer-grade AI cannot guarantee accurate legal output → this triggers "unauthorized practice of law" risk → the FTC finalized its order in February 2025, finding deceptive advertising and imposing monetary relief. The founder had declared he would "replace the US\$200 billion legal industry with AI"; in the end, it was the regulator that drew the boundary. The lesson: the entry barrier for legal services is "trustworthy and compliant," and the consumer-internet model of rapid iteration fundamentally conflicts with professional accountability.

Deloitte Australia: having a platform is not the same as having governance. [third-party verified] A report delivered by Deloitte Australia to the government — worth A\$440,000 (about US\$290,000) — was "riddled with AI-generated errors"; the firm eventually admitted to using AI and refunded the fee. Its self-built chatbot AskDeloitte had been used to draft client reports, and the errors stemmed from a lack of adequate human verification. After the incident, EY, KPMG, PwC, and BCG raced to claim they had rigorous review processes. A single AI failure can erode the most core asset of a professional institution — "the output can be trusted."

10.1.5 Organizational Patterns and Consulting Takeaways

Four patterns are worth borrowing for any professional-services organization.

First, **guardrails for hallucination and liability must be designed up front.** ABA Formal Opinion 512 establishes the baseline framework: the duty of competence (understand the tool's capabilities and limitations; AI output must be independently verified), the duty of confidentiality (assess leakage risk before inputting client data), candor toward the tribunal (citations must be verified), and the duty of supervision (management must establish AI policies and train everyone) [third-party verified]. On the ground this takes the form of institutionalized "human-in-the-loop": mandatory human audit of AI output and a routinized citation-verification process. Guardrails are

not the opposite of efficiency; they are the precondition for scaling AI — the Deloitte Australia crash proves that the larger the platform without guardrails, the greater the risk.

Second, **productizing knowledge assets is the second growth curve.** A&O's ContractMatrix, Macfarlanes' Amplify, and Chinese firm PowerLaw AI's (Chinese: 幂律智能) hybrid architecture of "AI models + sedimented lawyer rules" all point down the same path: move experts' judgment rules out of people's heads and into systems, then sell them outward. But the threshold is real — Pacgate Law Firm's (Chinese: 百宸律师事务所) summary is the soberest: data assets (whether structured, trainable data can be accumulated), organizational process (a closed-loop human-verification loop), compliance and ethics, and an AI engineering team that understands law — all four are indispensable [third-party verified]. Most institutions will stay at "using AI tools"; only a few will become "AI law firms."

Third, **the talent structure shifts from a pyramid to "senior experts + AI leverage."** The Big Four hire more AI staff than accountants, BCG adds only AI and technology specialists rather than generalist consultants, Accenture writes AI fluency into promotion criteria — the compression of junior repetitive roles is a certain trend, but "role-shifting" outweighs "role-eliminating": analytical judgment, client relationships, and accountability become more valuable. In the AI era, the scarce commodity is "people who can take responsibility for AI output."

Fourth, **billing must shift from billable hours to value-based pricing.** This is the industry's deepest contradiction: under ABA 512, work that AI completes in 15 minutes can only be billed for 15 minutes — effectively penalizing efficiency [third-party verified]; meanwhile Accenture's 60% fixed-price work, the Big Four's move toward multi-year outcome contracts, and A&O's subscription software sales all point in the same direction: pay for outcomes and platforms, not for hours. Whoever makes the switch first wins pricing power; those who delay are locked into the paradox of "working faster and earning less."

10.2 Media, Content & Creative Industries

10.2.1 Industry Profile and Bottleneck Triage

Media, content, and creative industries have one of the highest information-bottleneck shares of any industry, with information processing weighing in at roughly 90:10 of the value chain. Generative AI's impact here is not "icing on the cake" — it strikes directly at the core factor of production. The industry's essence is a "trust business": content institutions sell authenticity, judgment, and emotional resonance, and generative AI simultaneously dismantles and amplifies all three — it drives the cost of producing content toward zero, and the cost of producing fake content toward zero as well.

The industry faces three overlapping structural contradictions. First, the content-copyright game: rights holders and AI companies are lining up on two tracks — "licensing partnerships" versus "litigation confrontation." Second, trust and authenticity: tests by the BBC and the European Broadcasting Union of four major AI assistants found that nearly half of responses contained significant errors and 31% involved misquotation; Reuters Institute research shows that only about 7% of online news consumers use AI assistants to get news, yet trust in AI news among people under 35 is already close to 50% — younger audiences are becoming "AI natives," and institutional media's authoritative narrative has not yet found a new anchor. Third, the flood of AI content farms: NewsGuard has identified 3,749 AI content farms; one suspect generated 4,000–7,000 fake news articles in a single day at peak, and one AI-generated fake news story surpassed 500 million reads. When content supply becomes infinite, what becomes scarce is no longer content but trustworthy content.

10.2.2 Transformation Paths and Typical Practices

Newsroom AI started earliest and is the most institutionalized. Xinhua built its intelligent editorial department in 2019, using "Media Brain" plus AI synthetic anchors to cover 30-plus content categories, establishing the paradigm of "AI mass production, human review gatekeeping." People's Daily's AI editorial department has iterated to version 6.0, using the "Zhimei Engine" (Chinese: 智媒引擎) to deliver generative text-to-image and image-to-video production; its short film *The Protagonist* (Chinese: 《主角》) drew more than 600 million exposures within four days. The Associated Press used Wordsmith to expand earnings-report coverage tenfold, and Reuters' Newsmaker lifted production efficiency tenfold. International media entered organizational

reinvention earlier: the FT declared that AI has moved from standalone projects to "infrastructure, pipeline, partner, and competitor"; the BBC set up a dedicated news AI team and requires all employees to hold ChatGPT Enterprise access; the Guardian reached a content-licensing partnership with OpenAI in February 2025.

Content-platform organizational change is the most radical organizational experiment in China's internet industry in 2025–2026: Xiaohongshu (Chinese: 小红书) established an AI first-level department, Dots; ByteDance is pushing its organization toward an "intelligence-intensive" transformation with AI capital expenditure exceeding 200 billion yuan, replacing incumbent managers with "people who can use AI"; Kuaishou upgraded its self-developed video model Kling (Chinese: 可灵) to a first-level department with independent financing. The common choice of all three: "AI is promoted to an organizational unit" rather than "AI is scattered around as a tool."

AIGC in film, TV, and games is rewriting the cost structure: AI short dramas compress the localization process into a pipeline and cut costs by 15x; Chinese short-drama apps capture 91% of global revenue, and the overseas micro-drama market surpassed US\$4 billion in 2025. In gaming, Tencent's VISVISE, miHoYo's Glossa, and NetEase's "Smart Jianghu" (Chinese: 智能江湖; 400 intelligent NPCs, 30% AI efficiency gains) are all showing their hands. iQiyi built an AI artist library via an "AI Consent Form," turning artists' digital avatars into assets. **AI anchors** have formed an independent industry: by 2024 the segment's scale approached 480 billion yuan, spanning 12 core fields.

10.2.3 Benchmark Cases

Hunan Broadcasting (Chinese: 湖南广电): full-chain self-development plus all-staff training. Mango TV's paying members surpassed 75.6 million [official claims]; its self-developed Mango LLM is among the first batch of broadcasting-industry large models registered with the Cyberspace Administration of China (CAC) [official claims], has incubated more than 80 agents applied across 30-plus shows, and lifted production efficiency by over 30% [official claims]. What is replicable is the organizational path: in 2024 it launched the generative AI director "Aimang" (Chinese: 爱芒), which served as virtual assistant director on the reality show We Three (Chinese: 《我们仨》), elevating AI from a tool to a "role" within the production team; it launched the "Shanghai AIGC Content Creation Platform" (Chinese: 山海 AIGC 内容创作平台); under the "Daman Plan" (Chinese: 大芒计划), 5,137 micro-dramas had been released cumulatively by 2025, up more than 14x year over year [third-party verified]. In

parallel, it ran group-wide AIGC hands-on bootcamps: 800+ people registered and 100 entered intensive practice [third-party verified] — forming a closed loop of "universal training → elite intensive sessions → production-line application," positioning AIGC as "production line and leverage" rather than "demo tool."

FT: abolishing AI job titles is a watershed in organizational evolution. The FT announced that from 2026 it will no longer have product director roles whose titles carry "AI" [third-party verified]. The decision looks counterintuitive but is a sign of maturity: once AI is woven into the daily work of all teams, setting up dedicated "AI roles" actually hinders penetration. The FT's strategy has shifted from "build and buy" to "build and partner," and it reached a content-licensing agreement with OpenAI in April 2024; using large models to auto-generate discussion questions lifted its reader comment rate by 3.5% [third-party verified]. This is a sample of "de-specializing AI" — the destination of an AI-native organization is not an AI department but no AI department.

Kling: from tool to independently financed entity. Kling AI is the AI video-generation model Kuaishou developed in June 2024: 45 million+ creators worldwide, 200 million videos generated cumulatively, and 20,000+ enterprise customers served [official claims]. In 2025 it was upgraded to a Kuaishou first-level department and commercialized; in July 2026 it completed a nearly US\$3 billion independent financing round at a post-money valuation of US\$18 billion, with Tencent, Alibaba, and Baidu participating [third-party verified]. The three-step jump — "tool → first-level department → independently financed entity" — is the most complete organizational-change sample of a content platform incubating an AI-native business.

10.2.4 Failure Warnings

The double-edged sword of litigation confrontation: NYT v. OpenAI. In May 2026, the U.S. District Court for the Southern District of New York issued its first-instance ruling; the 247-page opinion has been called "the most important judicial document in AI copyright history" [third-party verified]. The case began in December 2023, with claims initially around US\$15 billion and later reporting suggesting up to US\$22.5 billion; in July 2026 the NYT and others also moved to sanction OpenAI. The lawsuit draws a clear roadmap: the era of "free scraping" without licensing or compensation is over — but the costs and uncertainty of litigation are equally staggering. The copyright game has no single-winner solution, only a combination of "licensing pricing" and "litigation deterrence."

Trust collapses with a single failure: CNET and Gannett. Gannett suspended its use of LedeAI after AI-generated sports coverage drew public ridicule for its poor quality; CNET abandoned AI writing after it was found publishing undisclosed AI articles containing factual errors; Sports Illustrated was exposed for an AI-generated fictional author, "Drew Ortiz," and its parent company fired the person in charge. The shared lesson: in an industry built on trust, if you use AI to cut costs at the expense of transparency and quality, the money saved is nowhere near the cost of a brand crisis. The LA Times' AI tool was taken down within a week of launch after defending the Ku Klux Klan; the Chicago Sun-Times' summer reading list included book titles invented by ChatGPT — these cases of "failed guardrails" jointly prove that the quality and value risks of AI content must be backstopped by organizational mechanisms, not by the model's own discretion.

The ecosystem-level pollution of AI content farms is the "tragedy of the commons" facing the entire industry: low-quality AI content drowns trustworthy content in volume, the traffic-monetization value chain drives fakers to take risks, and the traditional governance model of "clarify and self-prove" has already failed.

10.2.5 Organizational Patterns and Consulting Takeaways

Five patterns can be distilled from the media and content industry's AI-native transformation.

First, copyright-compliance guardrails must precede efficiency pursuits. The Guardian and the FT chose licensing; the NYT chose litigation; Suno and Udio eventually settled and paid the record labels; iQiyi's "AI Consent Form" and SAG-AFTRA's writing of AI protections into core contracts (won through the first full shutdown of 160,000 actors in 63 years) all point to the same conclusion: AI transformation in the content industry is first and foremost a renegotiation of the rights structure.

Second, AI content labeling is a mandatory organizational variable. On September 1, 2025, the Measures for the Labeling of AI-Generated Synthetic Content, issued by four ministries of the Chinese government, took effect, requiring explicit and implicit labels on AI content — the most systematic AIGC labeling regulatory framework among major economies. Douyin upgraded its labeling features in parallel, throttling, demoting, and even removing unlabeled content. Content institutions must internalize "labeling compliance" as a step in the production process.

Third, trust assets are a moat that cannot be outsourced. The BBC/EBU test data draw the red line of "human review required"; "AI mass production, human review gatekeeping" remains the industry benchmark. The industry is moving from "people using tools" toward "agents as autonomous execution units in the production process" — but autonomous execution does not mean no one is accountable: an accountability mechanism must be established in parallel.

Fourth, the essence of human-AI division of labor is "scale to AI, meaning to humans." The practice of the world's major wire services shows that AI excels at scale and speed, while humans excel at human care, trust-building, and value judgment. The first question in a transformation consultation is "which steps' scale dividends are worth capturing," not "which jobs will be replaced."

Fifth, organizational form upgrades with strategy, and policy must precede deployment. Of the 1,000 AI initiatives across global newsrooms, nearly two-thirds are built in-house, yet only 127 have policies, and only 7% of initiatives directly target monetization [third-party verified] — "fast deployment, slow policy, unresolved monetization" is the industry's common ailment. The implication for consultants: the mark of maturity in content-industry AI-native transformation is not how many AI projects were built, but the disappearance of AI job titles (FT), the labelability of AI content (China's regulation), and human-AI collaboration becoming the default way of working (BBC, Mango).

10.3 Financial Services

10.3.1 Industry Profile and Bottleneck Triage

Finance is a quintessentially information-intensive industry: a loan decision, a research report, and a claims settlement are all, at bottom, information processing. Its "information bottleneck" to "physical bottleneck" ratio is about 80:20 — which makes it one of the most fertile grounds for large models. Tencent Research Institute's judgment is representative: finance has a high degree of digitalization and high potential value in commercial scenarios, making it one of the best scenarios for AI large-model deployment. McKinsey estimates that generative AI could create US\$200 billion to US\$340 billion in value for the global banking industry each year — 2.8%–4.7% of the industry's annual revenue.

But fertile soil is not the same as easy plowing. The core contradiction of financial AI transformation is threefold. First, heavy regulation: credit scoring, anti-money laundering, and underwriting are classified as "high risk" under the EU's AI Act, and China's National Financial Regulatory Administration (NFRA) has also issued guidance on AI application in banking and insurance, requiring full-lifecycle model management and explainability. Second, data sensitivity: customer privacy, transaction data, and model hallucination are directly tied to capital and reputational risk — the warning from Gu Shu, chairman of the Agricultural Bank of China, at the Lujiazui Forum (model black boxes, hallucination, and the uncertainty of autonomous decisions) is industry consensus. Third, strong organizational inertia: IBM research shows that as of 2024 only 8% of banks were systematically developing AI, while 78% still used tactical, fragmented approaches. Finance does not lack AI enthusiasm; it lacks the organizational capability to turn pilots into systems.

10.3.2 Transformation Paths and Typical Practices

Banks, insurers, and brokerages have taken three clearly differentiated paths.

Banks are the most stratified. The big state-owned banks pursue a "model matrix plus compute infrastructure" strategy: ICBC's Zhirong (Chinese: 工银智涌) hundred-billion-parameter financial LLM system empowers 20-plus business areas and 200-plus scenarios, with AI digital employees handling the equivalent of 55,000 person-years of work annually. The joint-stock banks focus on deepening scenarios: China Merchants Bank's "AI Xiaozhao" (Chinese: AI 小招) serves over 20 million customers a month and replaced 15.56 million hours of human work with AI over the year; Industrial Bank's agent platform hosts more than 200 agents. The internet banks take a thoroughly AI-native route: WeBank proposes "embedding the process into AI" rather than "embedding AI into the process," and MYbank uses three AI systems — "Shanque" (Chinese: 大山雀), "Dayan" (Chinese: 大雁), and "Bailing" (Chinese: 百灵) — to rebuild micro and SME lending, with the ambition of "giving every small business an AI CFO." In 2025, DeepSeek's open-sourcing triggered a wave of local deployment: 20-plus financial institutions completed private deployments at million-yuan-level cost, and with lightweight open-source financial models such as Fin-R1, the pattern of "financial AI belongs only to big banks" was broken.

The insurance industry is penetrating along the entire value chain — product pricing, marketing and customer acquisition, underwriting, claims, customer service, and operations — with large models upgrading insurance AI from "discriminative

tools" to "generative intelligence"; knowledge-intensive steps such as agent training and claims-report writing benefit first. Brokerages have made investment research the main battlefield: large-model adoption in the securities industry reached 26.7%, with cases like Soochow Securities' (Chinese: 东吴证券) multi-agent investment advisor and CSC Financial's (Chinese: 中信建投) industry knowledge base in place, making research the field most likely to first achieve an end-to-end "data-research-investment-risk control" chain.

On organizational change, two routes diverge sharply. JPMorgan Chase chose retraining and flexible transition — employees shift from "makers" to "checkers" (managing AI agents). Standard Chartered, by contrast, announced it will cut more than 15% of corporate-function roles (about 8,000 people) before 2030, with the CEO bluntly saying the roles are being "handed over to machines." DBS runs a dual track of "AI replacement plus AI role creation," committing to hire about 1,000 new AI roles. China's six largest banks saw total headcount grow only slightly — by 1,538 people — at end-2025, while tech talent grew by a net 35,000 to 135,900: the structural blood transfusion is already visible to the naked eye.

10.3.3 Benchmark Cases

WeBank: the most complete sample of an AI-native organization [official claims + third-party verified]

In 2023 WeBank proposed moving "from digital-native to AI-native," defined as "not AI embedded in the process, but the process embedded in AI." The implementation mechanism has three layers: an AI infrastructure layer (an engineered platform where compute is allocatable, models are pluggable, and outcomes are measurable, supporting private deployment of open-source models such as DeepSeek); an AI application layer (70-plus digital employees and over 800 agents covering more than a thousand business scenarios); and an AI governance layer. By the end of 2025, AI compute had increased 3.5x year over year, daily calls had grown from 41,000 to 2.4 million, and daily token consumption had grown from 200 million to over 5 billion [official claims]. Organizationally, it established a technology product manager role to coordinate technology planning, breaking the old "business raises requirements, technology executes" model; in budgeting, an "Innovate the Bank" allocation funds frontier technology in advance and allows trial and error. This is the most complete sample of AI-native organizational change among domestic Chinese banks.

JPMorgan Chase: methodology at scale [third-party verified; benefit figures are the CEO's public statements]

450 production-grade use cases; the AI professional team grew from 1,500 people in 2022 to 2,500+ (+67%). Methodological essentials: use cases must be tied to business KPIs rather than technical metrics; human-machine collaboration design rather than full automation; data governance and model-risk controls first; start with high-value, low-risk scenarios such as document processing and code generation, then expand toward customer interactions. In 2025 it spent US\$18 billion on technology, of which about US\$2 billion was AI-specific; CEO Jamie Dimon said the company "spent US\$2 billion and got about US\$2 billion in benefits" [official claims], positioning AI as a "bet-the-company" strategic investment.

DBS: a pioneer in AI value accounting [official claims + third-party verified]

More than 1,500 AI models deployed across 370-plus scenarios; in 2025 it created S\$1 billion (about US\$750 million) in AI value and was named "World's Best AI Bank" by Global Finance magazine. The key practice is transparently measuring AI value each year across three dimensions — "revenue uplift, cost savings, risk avoidance" — and using that to set investment priorities. This is the critical management infrastructure of the AI-native organization: linking AI investment directly to financial outcomes.

10.3.4 Failure Warnings

Jane Street's US\$15 billion loss [third-party verified]

In July 2026, quantitative giant Jane Street suffered a monthly loss of nearly US\$15 billion — its first monthly loss in about a decade — partly due to a deep drawdown at Situational Awareness, an AI hedge fund it invested in, compounded by missteps in directional Asia equity bets; around the same time, an AI-stock selloff triggered a chain reaction of margin calls at Goldman Sachs and JPMorgan Chase. There is a huge gap between the "AI stock-picking guru" narrative and real risk: AI capability amplifies tail risk, and model extrapolation and concentration risk do not disappear because of "intelligence."

Klarna's pullback [third-party verified / CEO public interviews]

Klarna cut its workforce from over 7,000 to under 3,000 (a reduction of more than 50%), with AI handling 2.3 million conversations a month — the workload of 700 full-time customer-service agents. But in May 2025 the CEO publicly admitted the company "went too far": an excessive focus on efficiency degraded service quality, forcing a pivot to an "Uber model" — AI handles standard issues while humans handle interactions requiring emotional connection — after which NPS and satisfaction recovered. Two lessons: AI customer service must read the original business code rather than FAQs; and "AI can answer $\neq$ AI can solve" — technical metrics and business metrics can diverge.

The frontline temperature gap: statistics versus on-the-ground reality [third-party verified]

The six big banks spend 130.091 billion yuan a year on fintech, and the whole industry has cumulatively invested nearly one trillion yuan — yet down at second-tier city branch level, the "felt experience is weak": AI applications remain mostly "point-embedded" and have not yet penetrated to restructure the underlying business process — "a better whip for the carriage rather than a new car." Some city commercial banks even prohibit using AI to generate discussion materials. Value metrics such as "person-years," "work-hours," and "scenario counts" are not comparable across institutions, frontline employees are becoming "AI trainers" (data labeling, effect feedback, scenario training), and "saving 15.56 million hours" is more about avoiding headcount increases than converting directly into headcount reduction.

10.3.5 Organizational Patterns and Consulting Takeaways

First, regulation determines the order of deployment. The sequence of AI adoption in finance is decided not by technical difficulty but by regulatory tolerance: intelligent customer service, document processing, and code generation go first; core decision links such as credit approval, asset pricing, and capital trading must wait until the governance framework matures. The EU AI Act requires high-risk financial AI to meet compliance requirements by August 2026 (with fines of up to 7% of global turnover for violations), and the NFRA guidance requires AI risk to be incorporated into the comprehensive risk-management system — compliance is not a barrier to transformation but its roadmap.

Second, data governance is a prerequisite, not an afterthought. In financial scenarios, model hallucination is a business incident, not a technical fault; discriminatory bias in historical data becomes institutionalized through models. Consulting takeaway: build model registration, risk grading, and human-supervision mechanisms first, then talk about scaling.

Third, talent-structure change is harder to manage than technological change. Foreign banks are densely creating Chief AI Officer roles (IBM research shows the share of organizations with a CAIO jumped from 26% to 76%), yet Standard Chartered's former CAIO believes the role is "rather short-lived" — "as AI becomes a tool like Excel." Organizational form is evolving toward a "pine-tree" model (in the view of Hengfeng Bank's CIO): a small number of core decision nodes support large-scale standardized execution, with middle management compressed. On talent strategy, the dual track of "replacement plus role creation" is more sustainable than pure layoffs: retained employees see pay increases (Klarna's per-capita pay rose nearly 50%), and transferred employees move into data and scenario-training roles.

Fourth, value accounting moves from "comparing inputs" to "comparing outcomes." After nearly one trillion yuan in cumulative industry investment, the banking AI race has entered its "calculate the ROI" second half; building a unified AI value-measurement system (DBS's three-dimension method is a reference) should be the first move in any consultation. AI-native transformation in finance has never been a technology-selection problem; it is an organizational problem of re-sequencing four things simultaneously: the regulatory framework, data governance, talent structure, and financial accounting.

10.4 Education

10.4.1 Industry Profile and Bottleneck Triage

Education is a typical "dual information-physical bottleneck" industry, with an information-to-physical ratio of about 70:30. Knowledge instruction, homework grading, Q&A and assessment, and admissions administration are all, at bottom, information processing — the most direct use case for large models. But the underlying mechanisms of learning outcomes — motivation, focus, peer encouragement, and teacher-student trust — still depend on in-person settings and human presence. This forms the first pair of core contradictions: teaching delivery can be productized, but the in-person setting is hard to replace. The second pair is the

double-edged sword of equity: AI can amplify quality teaching resources at very low cost (GSU significantly improved first-generation college students' performance), but unequal access can widen the digital divide (LAUSD's high-profile project collapsed precisely here). The third pair is the tug-of-war between AI literacy and dependency: China lists teachers' digital literacy as a "basic item" of transformation, and the EU's AI Act turns AI literacy into a legal obligation; meanwhile, students substituting free general-purpose AI for paid services (the collapse of Chegg) forces the industry to redefine "what learning cannot be outsourced."

10.4.2 Transformation Paths and Typical Practices

Four paths are advancing in parallel.

University organizational transformation: In January 2024, ASU became the world's first university to partner comprehensively with OpenAI, embedding ChatGPT across teaching, research, and administration; OpenAI then released ChatGPT Edu for universities in May, turning universities into institutional clients. Georgia State University (GSU) embedded AI into the "student success" alert-remind-intervene chain. By late 2025, industry discourse shifted from "using AI tools" to "AI-native universities," with AI embedded across admissions, teaching, administration, and governance. But Ithaka S+R's interviews with 246 faculty at 17 US and Canadian universities show that many faculty do not even know about their school's existing AI resources — the bottleneck is not technology but organizational communication and trust.

Evidence on AI tutors: An Education Next forum synthesis gives balanced evidence: when expert teaching decisions were injected into GPT-4, 76% of its math responses were rated "good" by humans (the Bridge method), and AI-tutored students completed in 49 minutes a task that took the control group a 60-minute lecture; but researchers warn that publication bias is widespread, and most online math programs produce minimal gains for "95% of students."

Education company transformation: Duolingo sustains high growth through AI-driven course generation and personalization. Youdao (Chinese: 网易有道), built on its "Zi Yue" (Chinese: 子曰) large model, has repositioned its brand as a "learning and advertising AI application service provider"; AI subscription sales exceeded 100 million yuan in Q1 2026 (+70% year over year), and online marketing revenue surpassed the traditional education business for the first time. Squirrel Ai (Chinese: 松鼠Ai), which

started with adaptive algorithms, had five papers accepted at ACL 2025 and over 100 papers at top conferences cumulatively; its adaptive system was named to TIME's 2025 Best Inventions. New Oriental launched "AI 1-on-1," Zuoyebang is betting on large-model learning devices, and Yuanfudao's (Chinese: 猿辅导) "Xiaoyuan AI" (Chinese: 小猿 AI) topped the new-download charts in its first month but has not closed the profitability loop. The 2025 split is clear: AI-native organizations earn a growth premium, while slow transformers slide into structural decline.

China's policy path: Unlike the West's institution-driven approach, China's is "state-driven, top-down, scaled rollout." In March 2025, the Ministry of Education deployed Strategic Action 2.0 themed "AI and Education Transformation"; in April, nine government departments issued the Opinions on Accelerating the Digitalization of Education; in April 2026, five departments jointly issued the "AI + Education" Action Plan (Jiaokexin [2026] No. 1), proposing to form a deeply integrated AI-education landscape and universalize AI education in primary and secondary schools by 2030 — elevating AI from an "information tool" to a "systemic restructuring force." Supporting measures: in July 2025, the Ministry of Education issued the Digital Empowerment of Teacher Development Action, proposing to make teachers' digital literacy the "new normal" within three to five years, exploring large-scale individualized instruction and human-AI collaborative teaching, and opening AI-literacy training programs for teachers in central and western China; the China Society of Education is building a best-practice diffusion network through an "AI + Basic Education" case collection. The EU took a third path: the AI Act classifies scoring, admission, and assessment education AI as high-risk, mandating human oversight and AI-literacy training — regulation reshaping the organization in reverse.

10.4.3 Benchmark Cases

Georgia State University's Pounce: embedding AI into the "student success" process [third-party verified]. Pounce is the most evidence-backed AI application in education: in its first summer it interacted with incoming freshmen 185,000 times; a randomized controlled trial showed that students who received conversational reminders raised their FAFSA completion and fall enrollment rates by 3 percentage points and cut fee-related withdrawal by 50%; the classroom version of the chatbot helped students in an introductory political science course earn B-or-better grades at a rate 16 percentage points higher than the control group, with first-

generation students averaging 11 points higher on final exams. The key is not the model but the organizational design: AI does not replace teachers; it scales advisory services at very low cost.

Duolingo: the growth premium and organizational cost of being AI-native [third-party verified]. Its personalization engine Birdbrain has been running for over a decade; on April 30, 2025, it used AI to generate 148 courses in a single day, doubling its course catalog. The financial payoff of the AI-first strategy is significant: full-year 2025 revenue of US\$1.038 billion (+39%), 52.7 million DAU (+30%), and a 29.5% adjusted EBITDA margin. But the leak of an internal "AI-first" memo in May 2025 triggered an uproar, Q2 growth fell from 60% to 40%, and the stock dropped from a high of US\$544.93 to about US\$95 by March 2026 — becoming AI-native brings a growth premium, but it also amplifies the communication and capital risks of organizational change.

Khan Academy: experiment-culture-driven AI tutor iteration [official claims]. From October 2025 to April 2026, the Khanmigo team ran systematic product experiments across more than 15 million tutoring sessions: feeding the model the student's structured recent problem-solving history improved "next-question accuracy" by 6.1%; simplifying model output cut average response time by 3 seconds while math accuracy held steady; adding skill-session records from the past 24 hours raised cognitive engagement by 5.09%. The organization candidly admits "no single improvement brings a leap," but combined optimization significantly improves outcomes and lowers cost — this is the core engineering capability of an AI-native education organization: an experiment culture plus large-scale A/B testing.

10.4.4 Failure Warnings

The collapse of LAUSD's Ed: narrative first, governance absent [third-party verified]. The AI chatbot "Ed," built by the Los Angeles Unified School District at a cost of millions of dollars and billed as a "learning acceleration platform," was finally taken offline after vendor AllHere imploded and its CEO departed, and it drew in an FBI investigation related to contract procurement; the district superintendent resigned in June 2026. Technology selection and PR narrative came first, while vendor due diligence, data security, procurement compliance, and exit mechanisms were absent — public money and public trust were damaged at the same time.

Chegg: displaced by free general-purpose AI [third-party verified]. In May 2025, Chegg announced layoffs of about 22% (248 people) and closed its US and Canada offices; Q1 subscribers fell 31% year over year to 3.2 million, and revenue fell 30%. This unicorn, once valued at US\$12 billion, saw its "homework answers plus textbook rental" business model directly displaced by free tools such as ChatGPT — the paradigm shock suffered by a non-AI-native organization.

2U: the bankruptcy of an asset-heavy model [third-party verified]. In July 2024, online education operator 2U — once valued at over US\$5 billion — filed for Chapter 11 bankruptcy, cutting debt from US\$945 million to US\$459 million; court documents attributed the decline in part to "the rapid and unexpected adoption of artificial intelligence," and new OPM partnerships fell 53% in 2023–2024. A high-leverage old model dependent on degree outsourcing lacked organizational resilience.

Khanmigo's low adoption: supply is not adoption [third-party verified]. In April 2026, Sal Khan admitted that Khanmigo is "a non-event" for most students, with deep engagement at only about 15%; the Digital Education Council's 2026 global survey found that only 5% of students who experienced classroom AI said it "changed the way they learn"; a randomized controlled study in a university physics course also failed to reproduce significant advantages. The "access → use → learning outcome" three-stage funnel leaks at every stage — being technology-native does not mean being organization-native.

10.4.5 Organizational Patterns and Consulting Takeaways

First, productizable teaching delivery and in-person scenarios must be managed in separate accounts. Productizable steps — instruction, grading, Q&A, administration — should be AI-first and measured for returns; in-person steps — motivation, companionship, social connection, and teacher-student trust — stay with humans. GSU's success came from embedding AI into the process; Khanmigo's lesson is treating AI as a technology that "works the moment it's injected."

Second, teachers are the gateway to organizational change. The resistance is not a technology problem but a meaning problem: having teachers do "value-added/non-value-added task analysis" and showing them AI's value to their professional identity works better than tool training; the resource-information gap Ithaka revealed shows that communication and incentives must precede technology deployment. China

institutionalizes teacher-literacy training, and the EU mandates AI literacy by law — different roads, same destination: teacher capability building is the transformation's infrastructure.

Third, "supply $\neq$ adoption"; adoption engineering matters as much as the product. Free and embedded-in-platform does not earn usage either — it requires the three-piece set of teacher guidance, classroom process, and incentive structure to be designed together; Khan Academy's experiment culture proves that outcomes come from continuous iteration, not a one-time deployment.

Fourth, identify the two institutional paths and design consulting plans separately. China is policy-driven, with the levers being action plans, platform scaling, and case-library diffusion, focused on compliant rollout and regional scaling; the West is institution-driven piloting, with the levers being teacher adoption, outcome evidence, and governance frameworks (such as the EU AI Act), focused on single-institution change design and RCT-style evidence production.

Fifth, treat outcome evidence in tiers. Official claims (Khan's experiments, Squirrel AI's TIME listing, IDC rating it L5 fully adaptive) and third-party RCTs (GSU, the Khanmigo physics course) must be checked side by side, and beware of publication bias and vendor marketing materials; when attributing outcomes, place "pedagogical design" before "model capability" — with the same model, the way it is embedded determines the result.

10.5 Retail & Consumer Goods

10.5.1 Industry Profile and Bottleneck Triage

Retail is an industry with "dual information and physical bottlenecks," at a ratio of about 60:40. Half is an information business (demand forecasting, assortment, marketing, membership operations) and half is a physical business (stores, shelves, warehouses, cold chains, more than a million employees). Bain & Company data show that leading retailers using AI for personalized marketing saw ROAS improvements of 10%–25% in early experiments [third-party verified]; but the physical-side limits are just as real — when Amazon shut down Amazon Go, it admitted that logistics and technological excellence alone are "far from enough; you also need merchandising and display skills."

The core contradictions are threefold. First, online information flow versus offline store physics: what consumers see online and what the store actually carries are two systems that have long operated in silos — offline fulfillment remains the decisive battleground of cost and experience. Second, supply-chain efficiency versus experience: SHEIN compressed the industry-average unsold inventory rate of about 30% into single digits, while McDonald's AI ordering failed precisely because "efficiency wasn't saved and the experience collapsed first." Third, algorithmic site selection versus organizational inertia: more than 70% of Luckin's new stores are sited by algorithmic recommendation, yet most chains' expansion is still driven by their development teams' "experience plus relationships" — the handover of power between algorithms and the organization is the hardest bone in retail AI transformation.

The China Chain Store & Franchise Association's (CCFA) consensus is that "organizational change is the biggest challenge for AI deployment" [third-party verified]; industry case surveys likewise show the pattern of "marketing before supply chain, single points before the full chain" [third-party verified].

10.5.2 Transformation Paths and Typical Practices

Retail AI transformation can be split into at least four paths.

International giants: "human-machine collaboration" and "dual-track ecosystems." Walmart deployed AI tools (real-time translation, task management) for about 1.5 million US store associates and used generative AI to rewrite 850 million product catalog entries — making AI an "exoskeleton of productivity" for frontline employees rather than a reason for layoffs [official claims]. Target integrated its app with ChatGPT for personalized shopping while rolling out ChatGPT Enterprise to all employees [official claims]. Carrefour was the earliest to embed generative AI across the full chain: its Hopla shopping assistant covers all online channels, more than 2,000 products use AI-generated copy, and it was the first European retailer to open shopping inside ChatGPT [official claims]. Starbucks represents the "AI-augmented organization": Deep Brew sits in the back office of scheduling, stocking, and predictive maintenance, freeing people for human-facing service [third-party verified]. Amazon pushes AI down to the physical layer of fulfillment: more than 1 million robots are orchestrated by generative AI, and new-generation facilities are expected to automate 50–60% of tasks [third-party verified].

Supply-chain AI: "trading data for certainty." SHEIN's "small-batch, fast-response" model compressed the industry-average unsold inventory rate of about 30% into single digits; 2025 revenue was US\$41.8 billion, and it ran nearly 600 supplier training sessions with 37,000 person-times [third-party verified]. Inditex plans to invest €1.8 billion in technology and logistics; AI demand forecasting already drives 85% of initial production quota decisions, and FY2025 sales reached €39.9 billion, up 7% at constant exchange rates [third-party verified]. Uniqlo's "Ariake Project" (Chinese: 有明计划) invests about ¥100 billion in RFID item-level tracking, AI forecasting connected directly to factories, and warehouse automation [third-party verified].

Chinese chain organizational change: fix the organization first, then rebuild efficiency. In April 2026, Hema (Chinese: 盒马) consolidated 19 regions into 9 large regions, set up an AI committee, and created AI BP roles [third-party verified]; Sam's Club (Chinese: 山姆), by contrast, split its regions in the opposite direction — organizational structure must match the business stage, and the direction can legitimately differ [third-party verified]. RT-Mart (Chinese: 大润发), after DCP Capital took control, has "returned to the essence of retail": FY2025 returned to profit (net profit of 386 million yuan after a 1.668 billion yuan loss the prior year), with a plan to renovate 500 stores over three years [third-party verified].

Consumer brands: AI in service of the business model. Mixue's (Chinese: 蜜雪冰城) "Snow King AI Brain" (Chinese: 雪王 AI 大脑) turns AI into the cost-control system of its value-pricing strategy — with a gross margin of only about 35%, technology is not showing off but a survival necessity [third-party verified]; Pop Mart (Chinese: 泡泡玛特) built its self-developed AI integration platform HeyLisa, using AI sales forecasting to convert "hit-product uncertainty" into a predictable supply chain [third-party verified]; 7-Eleven Japan is the model of "digitize first, then make intelligent": data assets accumulated over decades of item-level management make AI demand forecasting viable [third-party verified].

10.5.3 Benchmark Cases

IKEA Billie: the unsolved problems AI reveals are worth more than the problems it solves [third-party verified]

Ingka Group launched customer-service chatbot Billie in 2021; from 2021 to 2023 it handled 3.2 million interactions, resolved about 47% of inquiries, and saved about €13 million in customer-service costs. The pivotal turn: IKEA did not lay off the roughly 8,500 customer-service agents displaced by AI. Instead, it analyzed the 53% of inquiries Billie could not resolve and discovered that what customers really wanted was "does this sofa fit my home" — an interior-design consulting need that had long been ignored. IKEA retrained the 8,500 agents as remote interior-design consultants, who generated €1.3 billion in revenue in their first full year (about 3.3% of group revenue). This is the most important positive sample of AI-native organizational transformation: the unsolved problems AI reveals are worth more than the problems it solves, and the freed-up workforce is redeployed to higher-value work — organizational transformation is a technology decision and a talent decision happening at the same time.

Luckin: algorithmic site selection and the "Smart Brew ecosystem" [official presentation; data not independently verified]

As of May 31, 2026, Luckin had 35,000+ stores worldwide, 473 million cumulative transacting customers, and 4.1 billion freshly made beverages sold in 2025. Three hard metrics: more than 70% of new store locations are algorithm-recommended, with an average of 18 days from application to acceptance; 700,000+ IoT-connected devices with a 99% online rate; and smart ordering is twice as efficient. Its self-developed agent, AI Lucky, has been integrated into super-apps including Qwen, WeChat, and Alipay and supports voice ordering; Luckin is also building the "Zhicui Ecosystem" (Chinese: 智萃生态圈) with multiple cloud vendors, upgrading AI from a point tool to a platform capability.

Hema: the governance package of an AI committee and a regional system [third-party verified]

In April 2026, Hema completed its organizational upgrade from a city-based system to a regional system, consolidating 19 regions into 9 large regions; in parallel, it set up an AI committee led by the CTO and created AI BP roles, directly connecting AI capability with business needs. The backdrop: nine consecutive months of overall profitability in 2024 and its first full-year positive adjusted EBITA in FY2025 (as reviewed by EqualOcean, Chinese: 亿欧). Hema advanced organizational change and AI governance structure in tandem — "committee plus BP roles" is the common mechanism linking the technology center with business units.

10.5.4 Failure Warnings

Amazon Go/Fresh closures: technological leadership does not equal a viable business model [third-party verified]

In January 2026, Amazon announced it would close all 58 Fresh supermarkets and 14 Go stores, accompanied by 16,000 layoffs, admitting it had "not created a sufficiently distinctive customer experience or the right economic model." Go's Just Walk Out "grab-and-go" technology was dazzling, but it could not mask mediocre merchandising; tellingly, Just Walk Out succeeded as a technology sold to 360+ third-party locations. The lesson: retail AI must solve "differentiated experience plus unit economics" at the same time — in the wrong business model, technology is only an expensive decoration.

McDonald's × IBM voice ordering: the accuracy-cost-experience triple threshold [third-party verified]

In June 2024, McDonald's ended the AI voice-ordering test it had run with IBM across 100+ stores since 2021. Analysts noted accuracy of only 80–85%, below the 95% threshold, and operating costs higher than human staff — and a viral TikTok video of "AI getting a McNuggets order wrong" amplified the experience collapse. The lesson: retail AI deployments must clear the "accuracy-cost-experience" triple threshold; there is a huge gap between technology demos and real operating scenarios.

Yonghui learns from Pangdonglai: easy to copy the form, hard to copy the spirit [third-party verified]

After more than two years of store renovations, Yonghui's renovated stores average over 80% foot-traffic growth per store — yet the company still posted a 2025 net loss attributable to shareholders of 2.550 billion yuan, with cumulative losses approaching 12.05 billion yuan from 2021 to 2025. The comparison data are stark: Pangdonglai (Chinese: 胖东来) employees average monthly salaries above 9,500 yuan with a turnover rate of only 0.33%; Yonghui's renovated-store employees earn 4,500–6,000 yuan a month — 65% of Pangdonglai's pay for the same position — and there is no profit-sharing mechanism. Pangdonglai's essence is a unique ordering of scale, profit, and employee relations; a nationwide chain cannot replicate its regional supply-chain density. "Whole-company financials versus isolated store performance" being worlds apart is one of the deepest lessons in organizational transformation.

10.5.5 Organizational Patterns and Consulting Takeaways

First, retail's AI-native transformation is, at bottom, a re-orchestration of two systems: the "information flow" and the "physical flow." Online-offline integration is not about moving products onto the web; it is about letting demand signals (online searches, membership data) and supply actions (store replenishment, production scheduling) share the same data foundation. The first move of any consultation should be to inventory whether item-level data assets exist and whether they are connected — 7-Eleven proves that AI's moat is not the model but the data discipline of "item-level management."

Second, supply-chain data foundations determine the AI ceiling. SHEIN's small-batch fast response, Zara's 85% of production quotas decided by AI, and Uniqlo's end-to-end RFID all point to the same thing: supply-chain AI's effectiveness depends on the depth of the industry chain's digitalization, not on model parameter count. Retail AI is an "asset-heavy" transformation — trading data infrastructure for certainty in inventory and cash flow.

Third, store employee roles must be rebuilt, not eliminated. IKEA retrained 8,500 agents, Walmart gave AI tools to 1.5 million associates, Amazon expects about 25% of its workforce to shift into technology-enabled roles, and 7-Eleven uses AI avatars to cope with labor shortages — all point to a "human-machine hybrid" organizational form. Layoff-style AI transformation in retail directly destroys the physical-side asset of the in-store experience.

Fourth, consumer trust is the invisible ceiling of retail AI. Walmart's research shows that consumers' trust in AI recommendations (27%) has overtaken influencer commerce (24%) — but missing privacy protections can void the transformation promise [official claims].

Fifth, organizational structure must match the business stage and digital maturity. Hema merging regions (seeking scale in an expansion phase) and Sam's Club splitting regions (seeking flexibility in a high-growth phase) moved in opposite directions and both worked; Yonghui's "easy to copy the form, hard to copy the spirit" shows that copying a benchmark's surface practices without adjusting the profit-distribution mechanism leaves the transformation stuck in isolated lines of the financial statements. The point of retail AI consulting is never model selection; it is the simultaneous re-sequencing of the data foundation, organizational structure, and incentive mechanism.

10.6 Government & Public Services

10.6.1 Industry Profile and Bottleneck Triage

Government and public services is one of the few industries where the "information and physical bottlenecks" are split roughly fifty-fifty: approval processing, policy consultation, document drafting, and service-request routing are typically information-intensive links, while counter services, on-site law enforcement, grid patrols, and urban infrastructure operations remain firmly bound to the physical world — in sharp contrast to finance's 80:20 information-dominant structure. This determines that government AI transformation must be a dual-track overhaul of "online intelligence + offline coordination," rather than a purely digital reconstruction.

Its core contradiction has three layers. First, the paradox of data security and sharing: government data is scattered across departments, with unclear powers and responsibilities and poor circulation — "data silos" have long been the No. 1 pain point; the national government's creation of a data administration is precisely aimed at breaking this deadlock, yet local practice generally finds "consolidation easy, coordination hard" — data-management responsibilities have not been synchronized with powers over project approval, funding, and staffing. Second, digital formalism: the CAC's deployment guidelines singled out for criticism "blindly pursuing technological superiority, duplicate construction, and forced usage," and People's Daily Online also warns against bandwagon construction "for the sake of using AI." Third, the decision-accountability chain: once AI-generated official documents, replies, or handling recommendations go wrong, there are no mature institutions for responsibility assignment, human fallback, or algorithmic transparency. The mirror-image dilemma exists in the US federal government as well — the Brookings Institution found AI use cases there highly concentrated in a few large agencies, with uneven diffusion, and procurement and budget mechanisms constituting structural barriers.

10.6.2 Transformation Paths and Typical Practices

Chinese government AI transformation shows a clear "policy-chain-driven" character: "AI+" was written into the Government Work Report in 2024; in August 2025 the State Council issued the *Opinions on Deepening the Implementation of the*

"AI+" Action; and in October of the same year the CAC released the *Guidelines for the Deployment and Application of AI Large Models in the Government Sector* — forming a complete chain of "top-level design — special action — scenario guidelines."

The core of the deployment route is "centralized, consolidated planning": the guidelines specify that counties and below should in principle reuse higher-level compute and model resources rather than independently build government LLMs, exploring "build in one place, reuse across many regions and departments" to prevent "model silos." At the same time, they delineate four major scenarios — government services, social governance, agency office work, and decision support — and emphasize the LLM's "assistive" positioning and the discipline that "classified information never goes online." DeepSeek's open-sourcing, with its low cost, high performance, and easy deployment, dramatically lowered the threshold for adoption, and domestic adaptation (Ascend chips, all-in-one appliances) became the standard.

At the organizational level, the National Data Administration (Chinese: 国家数据局, established in October 2023) has become the "organizational hub" of digital transformation: 2024 was positioned as the "year of institution building," shifting to the "year of reform breakthroughs" in 2025, driving the development and use of public data and the "Data Elements ×" initiative, with 23 provincial-level data administrations formed by the end of 2024. In terms of implementation forms, the "City Brain" has been elevated from a technology platform to a city-level governance mechanism reconstruction — Chongqing, relying on its integrated intelligent public data platform, automatically generated response plans such as police deployment and subway flow control through 5G + AI video analytics when daily footfall in the Jiefangbei business district exceeded 1.8 million, achieving "second-level sensing — intelligent assessment — coordinated response." On the grassroots governance side, Chongqing's "141" system embeds AI across 1,031 townships and streets and 62,000 grids, while Xiaoshan, Shangcheng, Nanshan, and Yuhang have formed closed loops of "AI collection — human handling — data feedback." The 12345 hotline has become the densest entry point for agents, with Guangzhou, Suzhou, and Jinan all deploying full chains of "smart form-filling — knowledge recommendation — smart dispatch." Shanghai's "Smart Easy Handling 4.0" (Chinese: 智慧好办 4.0) and Guangdong's "provincial coordination + municipal pilots + open scenario challenges" represent provincial-level promotion mechanisms.

10.6.3 Benchmark Cases

Shenzhen Futian: 70 AI "Digital Workers" [official claims + third-party verified]

Launched in February 2025, the first batch of 70 digital workers is built on DeepSeek fused with local knowledge bases, covering 11 major scenario categories and 240 government scenario terminals, forming a closed loop of "demand — training — scenario application — iteration." Document-formatting correction accuracy exceeds 95%, review time is cut by 90%, and the error rate is kept within 5%; law-enforcement transcripts are converted into draft documents in seconds; routing accuracy for public livelihood requests improved from 70% to 95%; personalized custom generation was compressed from 5 days to minutes. A unified management platform handles full-lifecycle management of these digital workers, making this a landmark implementation of the "AI civil servant" concept.

Shenzhen's "Shen Xiao-i" (Chinese: 深小i): A Public-Facing Government Agent [official claims]

Launched on February 22, 2025, it is the country's first practical AI government assistant open to the public, embedded in the full "ask, search, guide, process, review" service flow. It has cumulatively handled over 4 million intelligent Q&A sessions, with a response rate above 97% and first-answer accuracy above 94%. It is one of the first 7 government agents to pass evaluation by CAICT's TEL Laboratory (Chinese: 信通院泰尔实验室) and was selected as a standard application case of "AI + government" by the CAC. Behind it is Shenzhen's AI platform: managing 270P of compute, 39 deployed models, and 97.38 million annual calls, with 10 scenario challenges organized through open-competition recruitment (Chinese: 揭榜挂帅) — platform-based supply plus scenario-based open challenges brings it closer to "AI-native" than single-point applications.

Guangzhou: China's First Domestically Adapted Government DeepSeek [official / third-party verified]

Adaptation was completed and rolled out broadly in January 2025: the 12345 hotline answers 50,000 calls daily, and the intelligent agent-seat system built in early 2023 has cut caller waiting time by 43% and achieved 97% accuracy in work-order routing; the Digital Guangzhou Innovation Laboratory has adapted over 150 LLMs and opened the country's first government-grade compute center running DeepSeek-R1

671B on Ascend; the Guangzhou Intermediate People's Court deployed the country's first private government DeepSeek all-in-one appliance; Yuexiu District's "Ai Xiaoyue" (Chinese: 艾小越) achieves 99% answer accuracy. Guangzhou has proposed "one ledger, one mechanism, one platform" for LLMs, building a two-way loop in which "data feeds AI, and AI feeds decisions back."

10.6.4 Failure Warnings

AI Hallucination: Errors in Government Are Irreversible [third-party verified]

Outlook Weekly (Chinese: 瞭望) dissects three causes of hallucination: insufficient or biased training data, algorithmic architecture limitations, and the probabilistic nature of generation. LLMs "confidently spouting nonsense" can cause erroneous decisions and damage to public trust in highly regulated domains such as government and the judiciary — once a mistaken policy interpretation reaches the public, it constitutes real harm. CAICT experts point out that hallucination governance requires both "technology and institutions": retrieval augmentation and fact-checking on the technical side, and content review, risk warnings, and accountability tracing on the institutional side. This is precisely why the CAC guidelines require prominent risk warnings in the interface.

The GAO Audit: Governance Lagging Behind Deployment [third-party verified]

In its July 2025 review of 12 federal agencies, the US Government Accountability Office (GAO) found incomplete governance frameworks, use-case inventories not updated in time, and inadequate risk management for generative-AI-specific issues such as hallucination, data leakage, and copyright. US federal AI use cases grew fivefold in five years to 3,611, yet governance cannot keep pace with the speed of diffusion. The GAO's framework of "use-case inventory + risk classification + accountability auditing" suggests: build governance first, then talk about scale.

Digital Formalism: Advanced Technology, Lagging Experience [third-party verified]

The CAC guidelines explicitly require guarding against digital formalism; People's Daily Online criticizes bandwagon construction "for the sake of using AI"; the magazine *Qunzhong (Mass)* (Chinese: 群众) warns of "advanced technology, lagging experience." The more hidden failure is organizational: mismatched powers and

responsibilities of local data administrations leave coordination spinning in place, while US federal AI use cases are highly concentrated and imperceptible at the grassroots. Government AI usually fails not because the models are inadequate, but because governance and organization have not kept up.

10.6.5 Organizational Patterns and Consulting Takeaways

First, **institutions-first determines the order of deployment.** The rollout sequence of government AI is set by policy and security boundaries rather than technical difficulty: the CAC guidelines are both a constraint and a roadmap, and its three principles — scenario classification, centralized consolidated planning, and the assistive role — should be the preconditions for initiating any government AI project.

Second, **the assistive role is both politically correct and engineering-correct.** The stable form of government AI is the "digital colleague" rather than the "digital replacement": repetitive, standardized work goes to AI, while value judgments, emotional communication, and discretionary authority stay with people, with "human fallback" as the final line of defense against hallucination risk.

Third, **the data administration is the real lever of organizational change.** Government AI transformation is essentially a "data-administration-driven organizational rearrangement": the restructuring of tiao-kuai (vertical-line and horizontal-block) relations across national, provincial, and city/county levels, paired with organizational instruments such as leading groups, standardization technical committees, big-data professional-title review, and zero-based budgeting (as in Shenzhen), determines success or failure more than any single technology procurement.

Fourth, **grassroots burden reduction must be measured.** Shenzhen Nanshan's AI identified 87,000 public requests, improved key-information accuracy by 30%, and captured 820 weak risk signals — burden reduction backed by data is the only real "burden reduction," while forced usage and duplicate form-filling are a new form of formalism. The first consulting move should be helping clients build the governance foundation of "use-case inventory + risk classification + accountability auditing" (modeled on the GAO framework and CAICT's case-evaluation mechanism),

and unifying value metrics such as "number of scenarios, hours saved, accuracy" — otherwise government AI will repeat the failure of "statistical prosperity, frontline indifference."

1. Cyberspace Administration of China (CAC): *Guidelines for the Deployment and Application of AI Large Models in the Government Sector* https://www.cac.gov.cn/2025-10/10/c_1761819469929310.htm
2. Shenzhen Municipal People's Government: *Futian District Uses DeepSeek as an Engine to Accelerate Government Intelligence Transformation: 70 AI "Digital Workers" on the Job* https://www.sz.gov.cn/cn/xxgk/zfxxgj/gqdt/content/post_12006997.html
3. Xinhua News Agency: *How Do AI "Digital Workers" Help Make Government Intelligent?* <http://www.news.cn/tech/20250224/3d814d20206449a7afe4c78a2466d088/c.html>
4. Shenzhen Municipal Bureau of Government Services and Data Administration: *Ten Major Events in Shenzhen's Digital Government and Smart City Construction in 2025* https://www.sz.gov.cn/szsj/gkmlpt/content/12/12608/post_12608373.html
5. Guangzhou Municipal People's Government: *The 12345 Hotline Gains a "Super Operator": Guangzhou Rolls Out DeepSeek Integration Across Government* https://www.gz.gov.cn/szgzkfgx/zd gz/content/post_10153907.html
6. US Government Accountability Office (GAO): *Generative AI Use and Management at Federal Agencies* <https://www.gao.gov/products/gao-25-107653>
7. Outlook Weekly: *AI Hallucinations on the Rise: What Are the Risks and Challenges?* <https://www.nsfc.gov.cn/p1/3381/4121/2826/93953.html>
8. People's Daily: *What Key Signals Does the Establishment of the National Data Administration Send?* https://paper.people.com.cn/rmlt/html/2023-05/15/content_25999945.htm
9. People's Daily Online: *To Do "AI+" Government Services Well, First Clarify These Things* <http://politics.people.com.cn/n1/2025/0304/c458474-40430761.html>
10. Brookings Institution: *Assessing the state of AI adoption across the federal government* <https://www.brookings.edu/articles/assessing-the-state-of-ai-adoption-across-the-federal-government/>

11. Chongqing Municipal People's Government: *Exploring a New Path of Modern Governance for Megacities: The "Light of Wisdom" Illuminates a New Chongqing*
http://www.cq.gov.cn/ywdt/jrcq/202508/t20250818_14907156_wap.html
12. General Office of the Shanghai Municipal People's Government: *Shanghai's Implementation Plan to Accelerate "AI + Government Services" Reform and Promote "Handling Things Efficiently"*
<https://www.shanghai.gov.cn/nw12344/20250410/467ecd7d1a254575b263e73fc7c793de.html>

10.7 Logistics & Supply Chain

10.7.1 Industry Profile and Bottleneck Triage

In logistics and supply chain, the information-to-physical bottleneck ratio is roughly 40:60 — the exact opposite of finance (80:20), and this is the key to understanding the industry's AI transformation pace. Beneath the asset-heavy surface, what drives the network is decision information: forecasting, scheduling, routing, and risk control. CFLP data shows China's logistics industry total revenue reached RMB 13.2 trillion in 2023; industry estimates put the global AI smart-logistics market above US\$42.8 billion in 2025, with over 1.5 million warehouse robots deployed [third-party verified].

But the industry's AI transformation faces three structural contradictions. First, automation is a heavy-asset investment while demand is highly volatile: an automated terminal costs billions, takes years to build, yet demand fluctuates week by week — capital expenditure and capacity utilization are structurally mismatched. Second, right-of-way and compliance constraints in the physical world: autonomous delivery vehicles "have reached the ten-thousand-unit scale, yet are still waiting for a national pass" [third-party verified], and fragmented right-of-way policies directly draw the commercial boundary of unmanned capacity. Third, technical maturity does not equal commercial maturity: L4 autonomous trucks have been explored for over a decade without a working unit-economics model, and drone order volumes coexist with mounting losses. A CFLP report and a *People's Tribune* (Chinese: 人民论坛) article respectively attribute the barriers to insufficient data quality and standardization, missing cross-department coordination, AI talent shortages, and immature input-output measurement systems — as well as high upgrade costs, data silos, privacy compliance, and policy support — and warn of "algorithmic black boxes" and the

"intelligence monopoly" risk of excluding small and medium enterprises. Logistics does not lack AI pilots; it lacks the organizational capability and commercial closed loop to turn pilots into systems.

10.7.2 Transformation Paths and Typical Practices

Logistics AI deployment unfolds along five paths, with value realization on the information layer generally outpacing the physical layer.

Express logistics LLMs. Leading companies have formed a combined play of "vertical LLM + agents + full-chain scenarios" [third-party verified]: Cainiao (Chinese: 菜鸟) launched a logistics customer-service AI brain in 2018; STO Express (Chinese: 申通) launched "Kunlun" in 2021; 2023 saw a burst of releases — Cainiao's "Tianji π," JD.com's "Superbrain" (Chinese: 超脑), and YTO Express's "Zhiduoxing" (Chinese: 智多星); SF Express (Chinese: 顺丰) released its "Fengzhi" (Chinese: 丰知) decision LLM in 2024; and JD upgraded to "Superbrain 2.0" and released the "Yilang" (Chinese: 异狼) robotic arm in 2025 [official claims]. The strategic split is clear: the direct-operated camp (SF Express, JD.com) develops full-stack in-house, while the franchise camp uses AI tools to penetrate the "headquarters—outlet—courier" coordination chain — YTO has embedded over 100 agents across its full chain, generating cost savings on the order of RMB 1 billion, and AI roles have gone through a three-stage upgrade from "agent" to "AI assistant" to "digital worker" [official claims + third-party verified].

Warehouse robots. Amazon announced in 2025 that its global robot deployments exceeded 1 million units, two-thirds of the global total [official claims]; Walmart is deeply tied to Symbotic, rolling AI automation across all 42 regional distribution centers within eight years; UPS aims to raise the share of US packages processed with automation from 66.5% to 68% by the end of 2026 [third-party verified]. Even the giants advance in "percentage-point" increments — transformation is a gradual engineering process, not a one-step "dark warehouse."

Smart ports. Qingdao Port (Chinese: 青岛港) built its "smart brain" on fully domestic, fully autonomous A-TOS/A-ECS systems, with 85 AGVs and 47 yard rail-mounted gantry cranes running unmanned end to end; Yangshan Phase IV operates on "F5G ultra-remote control + AGV automated handling," and the Port of Shanghai handled its 50 millionth TEU in November 2025, ranking first globally for 16

consecutive years [third-party verified]. COSCO Shipping (Chinese: 中远海运) forms an SOE three-tier approach with its Hi-Dolphin shipping LLM and L4 port autonomous driving [official claims].

Autonomous delivery. The industry delivered over 35,000 unmanned vehicles in 2025, and JD Logistics (Chinese: 京东物流) announced procurement of 1 million units over the next five years [third-party verified]; Meituan's commercial drone orders exceeded 670,000 and surpassed 900,000 after routine operation [official claims]; SF's Fengyi (Chinese: 丰翼) "Ark" (Chinese: 方舟) drone delivers over 20,000 packages per day.

Digital freight platforms. Full Truck Alliance (Chinese: 满帮) fulfilled over 236 million orders on its platform in 2025, compressing truck-to-cargo matching from hours to minutes, and its smart consolidated-order assignment saved shippers over RMB 18 billion in freight costs in a single year; Huolala (Chinese: 货拉拉; Lalamove internationally) drivers achieve a 95% on-time rate and a 92% freight matching rate [third-party verified].

10.7.3 Benchmark Cases

Qingdao Port: An AI-Native Sample of Physical Heavy Assets [official claims + third-party verified]. On May 22, 2025, the quay cranes at Qingdao Port's Qianwan automated terminal reached an average efficiency of 62.62 containers per crane-hour, setting the world record for automated terminal handling efficiency for the 13th time. The core is the in-house A-TOS/A-ECS systems working together as a "smart brain," upgraded with AI LLMs for algorithms and compute. The terminal was approved in 2013 and construction began in 2015 — a decade of continuous iteration from an "unmanned terminal" to "full-domain intelligence." Port AI-nativeness is a systems engineering effort of "in-house foundation + continuous iteration + organizational support," not a single-point equipment purchase.

Amazon: A Robot Army Plus a Central AI Brain [official claims, limited independent verification]. Starting with the Kiva acquisition in 2012, Amazon expanded over 13 years to 300+ warehouses and 1 million robots, becoming in 2025 the "world's largest mobile robot manufacturer and operator." The supporting DeepFleet generative-AI scheduling system coordinates robot movement in real time like an "intelligent traffic management system," reportedly improving fleet efficiency by 10% and overall logistics efficiency by 25% [official claims]. What deserves

attention is not the numbers but the organizational logic: in-house development plus M&A builds a two-tier architecture of "robot army + central AI brain," and the real strategy is human-robot collaboration and job-skill reshaping.

Full Truck Alliance: An Agent Upgrade for a Digital-Native Species [third-party verified + company disclosures]. A road-freight platform is a natively digital species for which "data is the business," and its evolutionary ladder — information matching (1.0), algorithmic matching (2.0), decision-proxy agents (3.0) — offers a template for comparison. Full Truck Alliance fulfilled over 236 million orders on its platform in 2025, and in 2026 is embedding agents into real transaction-decision scenarios with products like the "cargo-finding assistant" and "intelligent assistant," responding to fairness governance with algorithm disclosure [third-party verified].

10.7.4 Failure Warnings

TuSimple: The Collapse of a Technology Leader's Commercial Closed Loop [third-party verified]. Dubbed "the world's first autonomous-driving stock," its market cap once reached US\$16 billion; it was delisted in 2024, and its Guangzhou team was dissolved and liquidated in February 2025. Founder Chen Mo admitted that "from 2021 to now, the resources at TuSimple's disposal have been insufficient to complete the commercialization of autonomous driving." Three root causes: insufficient technical maturity, regulations not ready, and unit economics that did not work — the labor costs saved could not offset sensor and safety-operator costs; on top of that came governance failures of management infighting and strategic flip-flopping. An AI-native organization needs not only technical capability but also a sustainable commercial closed loop and stable organizational governance.

The 95% Pilot Failure Rate [third-party verified]. project44, citing MIT research, notes that 95% of enterprise AI pilots fail to produce measurable returns (global AI spending exceeded US\$420 billion in 2025). Its conclusion: "it's not a model problem, it's an execution problem" — poor workflow integration, agents lacking operational context, fragmented tools, prompt-engineering burden, and unclear accountability. In supply chains the cost is more direct: there is only a 2-hour window to handle a shipment exception, and once it closes, the company faces a forced choice among expedited premiums, customer damage, and profit pressure. Pilots fail not because of algorithms but because of process integration and accountability mechanisms.

Autonomous Delivery: Scale Coexisting with Losses [official claims + third-party verified]. Meituan's commercial drone orders grew from 670,000 to 900,000, yet Meituan's new initiatives segment posted an operating loss of RMB 10.1 billion in 2025 (including drones); the unmanned-vehicle industry delivered over 35,000 units in 2025 with an estimated market headroom of nearly RMB 500 billion, yet remains stuck in "scaling at a loss" price wars and fragmented right-of-way policies. Technical validation is complete; commercial validation is unfinished — "technical maturity ≠ commercial maturity" is the most expensive lesson in logistics AI.

10.7.5 Organizational Patterns and Consulting Takeaways

First, **the physical bottleneck determines the order of value realization.** Logistics' 40:60 information-to-physical structure means "information-layer" AI — forecasting, scheduling, matching — realizes value fast, while "physical-layer" automation is an asset-heavy, long-cycle engineering effort. The first consulting move is to identify which layer the money is going into: the information layer runs on algorithms and data, the physical layer runs on capital and supply chain — with completely different rhythms.

Second, **data-driven scheduling organizations are the core form.** Agents are beginning to take on cross-department workflows, and organizations are shifting from "human + system" to "human + agent collaboration" [third-party verified], with AI roles evolving along "tool → assistant → digital worker." The key proposition for franchise networks is distributed coordination across "headquarters—outlet—courier": AI products must lower their entry barriers and adapt to the uneven digital capabilities of franchisees — organizational coordination efficiency, rather than mere labor replacement, is the value anchor of franchise-camp AI.

Third, **the three routes of in-house development, M&A, and vendor lock-in are clearly differentiated:** Amazon develops in-house and acquires; Walmart "divested in-house development and outsourced operations" (selling its robotics business for long-term orders worth over US\$50 billion); Cainiao productized its logistics technology for external clients (800+ projects worldwide, 26 Fortune Global 500 clients). Whether an organization must develop AI in-house depends on its data assets and capital strength — "owning AI" and "running an AI-powered operation" are two equally legitimate choices.

Fourth, **commercial validation should come first.** McKinsey's *The State of AI 2025* shows 88% of organizations now use AI routinely, but only 39% capture significant value and 6% are high-performing winners [third-party verified]; the common root of TuSimple's collapse and the unmanned-delivery losses is large-scale investment before unit economics were validated. Pilots should be tied to business KPIs, not technical metrics.

Fifth, **talent and the industry foundation come first.** The center of gravity of supply chain competition is shifting from "buying systems" to "building capabilities" [third-party verified], generating new roles such as AI operations, data annotation, and human-machine collaboration supervision. At the same time, logistics AI-nativeness is a cross-entity systems engineering effort — the first national standard for logistics enterprise digitalization was released in September 2025, and the *Digital Ship White Paper* (COSMO Shipping and China Classification Society) pushes for connecting data across ships, ports, cargo, and people; an industry-level digital foundation is a prerequisite no single enterprise can build alone.

1. China Securities Journal: *Rebuilding the Port's DNA with AI: Recreating Qingdao Port* https://epaper.cs.com.cn/zgzqb/html/2025-08/01/nw.D110000zgzqb_20250801_3-A04.htm
2. 36Kr: *Amazon Deploys 1 Million Robots — About to Surpass Human Employees?* <https://m.36kr.com/p/3368169340962561>
3. SF Technology (official): *SF Express Unveils Its Logistics Decision LLM "Fengzhi"* <https://www.sf-laas.com/news-details/40192>
4. Logclub: *The AI Cards of Nine Companies Including SF Express, JD Logistics, Cainiao, ZTO, YTO, and STO, at a Glance* <https://www.logclub.com/articleInfo/ODQ2MjE=>
5. Shanghai Securities News: *Six Full-Stack Achievements Verified Across All Domains: YTO Express Scales AI Applications* <https://www.cnstock.com/commonDetail/741566>
6. Sina Finance: *Logistics Unmanned Vehicles 2025 Review: Industry Deliveries Exceed 35,000 Units as the "Four Little Dragons" Intensify the Battle* <https://finance.sina.com.cn/cj/2026-01-26/doc-inhirqat8215013.shtml>
7. 36Kr: *Ten Thousand Unmanned Vehicles, Waiting for a "National Pass"* <https://m.36kr.com/p/3617731429123076>

8. Huxiu: *TuSimple's Guangzhou Company Suddenly Announces Dissolution and Liquidation: "The Management Has Problems"* <https://www.huxiu.com/article/4039242.html>
9. project44: *Why 95% of AI Agent Pilots Fail in Supply Chains* <https://www.project44.com/blog/why-95-percent-of-ai-agent-pilots-fail-in-supply-chains/>
10. Tencent News: *SF Express, JD.com, and Meituan Launch the "Sky War": Drones Reshape Logistics* <https://news.qq.com/rain/a/20251227A01ST800>
11. McKinsey: *The State of AI 2025* <https://www.mckinsey.com.cn>
12. Zhihu column (transcript of an Apsara Conference talk): *The State of Full Truck Alliance's AI Development and Reflections* <https://zhuanlan.zhihu.com/p/1967627027327743626>

10.8 Healthcare

10.8.1 Industry Profile and Bottleneck Triage

Healthcare is a typical information-physical hybrid: diagnosis, medical records, and imaging interpretation are essentially information processing, but treatment outcomes ultimately land in physical actions such as surgery, medication, and nursing. Its "information-to-physical bottleneck" ratio sits around 30-50:50-70 — which means healthcare AI cannot rewrite entire processes digitally the way finance can, and must embed itself in an "augmenting people" rather than "replacing people" mode. Grand View Research forecasts the global healthcare AI market growing from US\$36.7 billion in 2025 to US\$505.6 billion by 2033, and EO Intelligence (Chinese: 亿欧智库) estimates China's healthcare LLM market exceeded RMB 1.1 billion in 2025 [third-party verified].

But the growth narrative masks four structural contradictions. First, regulatory compliance: the FDA has approved over a thousand AI/ML medical devices, yet no generative AI has received FDA approval; in China, Class III medical device registration certificates are the precondition for AI to enter clinical billing, and the number of certificates directly constitutes a corporate moat (United Imaging 16, Deepwise (Chinese: 深睿) and Shukun Technology (Chinese: 数坤) 14 each, Infervision (Chinese: 推想) 13) [third-party verified]. Second, clinical accountability: responsibility allocation after AI enters high-risk operations is far from mature — US surgical robots caused 1,391 incidents and 144 patient deaths over 13 years [third-party verified]. Third, data

privacy: the WHO warned in November 2025 of a serious gap between medical AI expansion and legal-ethical safeguards, noting that biased training data will amplify healthcare inequities [third-party verified]. Fourth, physician workflow integration: AI alone can reach 90% diagnostic accuracy, yet physician use of assistance improves accuracy by only about 2 percentage points [third-party verified] — human-AI collaboration gains depend on workflow design and trust building, not model accuracy.

10.8.2 Transformation Paths and Typical Practices

Three paths diverge clearly.

Chinese hospitals: open-source models + local deployment + hospital-wide infrastructure. DeepSeek's open-sourcing in early 2025 dramatically lowered the AI threshold for hospitals: Shenzhen Nanshan Hospital deployed DeepSeek-R1 671B locally in 72 hours to build three application scenarios; the University of Hong Kong-Shenzhen Hospital deployed R1 under a "data never leaves the hospital" model; Peking University Third Hospital used multiple combined deployment modes to cover the entire hospital [third-party verified]. Scenarios are spreading quickly: West China Hospital (Chinese: 华西医院), based on its "Hongyi" LLM (Chinese: 华西黄医), released two agents — patient accompaniment and medical documentation — and the documentation agent has generated nearly 600,000 medical records since launch [third-party verified]; the first publicized 149 granular hospital AI scenarios show AI penetrating the full flow of appointment triage, medical-insurance cost control, DRG grouping, and medication review. On the budget side, the industry has entered an "infrastructure era": Zhongshan Hospital's (Chinese: 中山医院) RMB 160 million AI platform tender was awarded; iFlytek (Chinese: 科大讯飞), Alipay, and Baidu successively won government healthcare procurement contracts exceeding RMB 100 million; and hospital AI budgets jumped from millions to hundreds of millions of RMB [third-party verified]. IDC estimates China's AI + healthcare application software market at RMB 3.54 billion in 2025, projected to reach RMB 14 billion by 2030, and judges that the watershed has shifted from "how strong is the model" to "how deeply can business flows be redesigned" [third-party verified].

US hospitals: administration first, clinical caution, portfolio governance. HCA Healthcare (191 hospitals) advances in layers: on the clinical side, it partners with Google to optimize nurse handoffs, and its fetal heart monitoring AI has been submitted for FDA review; on the operations side, nearly 100 hospitals have deployed scheduling and workforce tools; the administrative side is seen as the "shortest path

to value." Its CFO admitted that "rolling out AI and changing frontline behavior is far harder than imagined," and investment focus is shifting toward change management [third-party verified]. Mayo Clinic will invest over US\$1 billion over the next few years across 200+ projects, driven by the dual engines of "clinical research + AI engineering"; Advocate Health uses an "evaluation funnel" (225 proposals → 40 use cases) for centralized selection — administration first, clinical later [third-party verified]. Menlo Ventures surveyed 700+ US healthcare executives: the industry deploys AI at 2.2x the pace of other industries, 22% of organizations have deployed domain-specific tools (up 7x year over year), and 2025 healthcare AI spending was about US\$1.4 billion, roughly 3x the prior year [third-party verified].

AI drug discovery and regulation: paradigm shift with institutions first. Insilico Medicine (Chinese: 英矽智能) applies generative AI across the full pipeline from target discovery to clinical development, closing more than 13 BD deals in 2025 with cumulative potential value exceeding US\$10 billion — the first profitable results for AI drug discovery [third-party verified]; its CEO claims AI compresses the drug R&D cycle to about one year (candidate drugs in as little as 9 months vs roughly 4.5 years traditionally) [official claims; note: "candidate drug" ≠ "approved new drug"]. On the regulatory side, the FDA published its final guidance on PCCP (Predetermined Change Control Plans) in August 2025, allowing AI algorithms to iterate continuously within predetermined bounds without re-filing — this "learning license" regime raises the organizational bar for continuous monitoring and change management; Fudan University, Ruijin Hospital (Chinese: 瑞金医院), and Ant Group (Chinese: 蚂蚁集团) jointly released the *Evaluation Standards for Healthcare Large Models*, providing a common language for "who evaluates, and how" [third-party verified].

10.8.3 Benchmark Cases

Kaiser Permanente: Ambient Documentation Across 7,260 Physicians [third-party verified, NEJM Catalyst]. From October 2023 to December 2024, The Permanente Medical Group (TPMG) deployed ambient AI documentation at 17 medical centers in Northern California, covering 7,260 physicians and over 2 million visits. Physicians collectively saved 15,700 hours of documentation time, and after-hours documentation time fell by an average of 1.03 minutes per visit; 84% of surveyed physicians said AI improved their connection with patients, and 82% reported higher job satisfaction. Combined with Abridge's largest generative deployment across 40 Kaiser hospitals and 600+ clinics, ambient documentation has become the most

commercially mature track in healthcare AI. But the Peterson Health Technology Institute cautions: financial returns of the systems have not yet been validated at scale, and scaling depends on deep EHR integration and safety monitoring systems.

Zhiyi Assistant (Chinese: 智医助理): 1.2 Billion Assisted Diagnoses Reaching the Grassroots [third-party verified / official disclosures]. iFlytek's "Zhiyi Assistant" embeds "AI-assisted diagnosis + standardized electronic medical records" into grassroots workflows: by the end of 2025 it covered 31 provinces, 800+ counties and districts, and 77,000+ grassroots medical institutions, with over 1.2 billion assisted diagnoses cumulatively serving 250,000+ grassroots physicians [third-party reported figure]; Liu Qingfeng disclosed that in Anhui alone it corrected 170,000 erroneous diagnoses in a single year [official claims]. Its organizational form — "provincial coordination, government funding, embedded systems" — validates the grassroots accessibility path while also exposing dependence on data quality and government funding.

Beijing Anzhen Hospital (Chinese: 安贞医院): The "Report Revolution" of Operations-Management AI [third-party verified]. Released in March 2026 as the country's first "public hospital operations management LLM," it was developed with China Telecom over 7 months, integrating the HIS, SPD, and ODR systems, defining 200+ data standards, governing 3 million+ data records, and establishing 800+ core management indicators. Operations-management report generation was cut from days to 5 minutes, management response speed improved by 95%, and precision decision-making efficiency improved by 60%; the "data-query agent + operations analysis engine" supports multi-level drill-down attribution and delivers role-personalized decision reports — "a thousand people, a thousand faces." It represents hospital AI extending from clinical assistance to operational decision-making, rebuilding management's data-decision infrastructure.

10.8.4 Failure Warnings

HIMSS 2025: 80% of Projects Fail to Convert Value [third-party verified]. A HIMSS panel confronted the industry's pain point directly: 80% of healthcare AI projects fail to convert into value; 97% of organizations agree on deployment urgency, but only 14% have effective implementation capability. Its verdict cuts to the heart of the matter: "AI failure is not technical immaturity but strategic failure," and the biggest mistake is trying to solve all problems at once.

NHS: The Cold Landing of a National Program [third-party verified]. UCL, with the Nuffield Trust, evaluated NHS England's £21 million AI chest-imaging project across 66 hospital trusts: contract signing was delayed 4–10 months, and 18 months after contract completion, 23 of 66 trusts (a third) still had not used it clinically. The barrier is not the algorithm but the organization: overworked clinicians unable to participate, aging and heterogeneous IT systems difficult to integrate, and senior physicians worried about accountability for AI decisions.

Hallucinations and the Clinical Adoption Gap [third-party verified]. Healthcare LLMs inherently carry hallucination risk; the industry consensus is that it "cannot be fully eliminated, only managed and corrected." AI alone achieves 90% diagnostic accuracy while physician assistance improves accuracy by only about 2 percentage points — without workflow design and trust building, accuracy does not automatically translate into outcomes. The 1,391 incidents and 144 deaths from surgical robots over 13 years is a reminder that AI adoption in high-risk clinical operations advances far more slowly than in administrative settings, and accountability, standards, and insurance mechanisms are far from mature. An HFMA survey shows nearly 70% of CFOs admit governance maturity lags adoption speed, and resource misallocation of "bought but unused, used but unevaluated, evaluated but not retired" is happening [third-party verified].

10.8.5 Organizational Patterns and Consulting Takeaways

First, **the organizational loop of "AI alert → human confirmation → intervention" is the value multiplier of clinical AI.** UCSD's COMPOSER reduced in-hospital mortality by 17%, and community hospital Augusta Health saved 282 lives in two years through the same path — the common factor is not a stronger model but embedding alerts into care workflows, training response mechanisms, and continuously recalibrating the model. For small and medium hospitals, the lesson: start with a single high-value scenario and compensate for resource disadvantages with organizational mechanisms.

Second, **change management is the main bottleneck, not technology.** A McKinsey survey shows 50% of healthcare organizations have implemented AI, yet its conclusion is that "the main bottleneck is not technology but leadership and organizational change," corroborating the NHS research; HCA's CFO admitted that changing frontline behavior is far harder than deploying tools.

Third, **physician-adoption design sets the ceiling.** Kaiser's high satisfaction comes from embedding AI into existing workflows rather than adding burden; HIMSS stresses that "unexplainable AI decisions will lose physician-patient trust." Consulting takeaways: do workflow analysis before model selection, and build trust administration-first, clinical-later.

Fourth, **regulation-first is a roadmap, not a barrier.** Class III certificates, FDA approvals, algorithm registration, and data security assessments form a landscape where "admission barriers are organizational capabilities": certificate counts become corporate moats, buyers need new validation and acceptance capabilities, and third-party evaluation standards are becoming a common language for all parties.

Fifth, **data assets replace models as the biggest shortfall.** The consensus at the VB100 conference: the biggest shortfall is no longer compute or models but standardized, high-quality data assets; IDC judges the next five years' competition will hinge on "how deeply business flows can be redesigned." The essence of AI-native transformation in healthcare is reordering four things simultaneously — clinical workflows, management decisions, regulatory compliance, and data governance. Technically, the hardest part is keeping AI from overstepping before the physician signs off; organizationally, the hardest part is getting AI actually used after sign-off.

1. Future Medicine (NEJM Catalyst research): *Ambient AI Scribes: Kaiser Permanente's 7,000-Physician Study Signposts the Future of Clinical Documentation* <https://www.futuremedicine.com/articles/ambient-ai-scribes-kaiser-permanentes-7000-physician-study-signposts-the-future-of-clinical-documentation>
2. Health News (Chinese: 健康报): *"Zhiyi Assistant" Takes Root in Anhui's Grassroots: Over 680 Million Assisted Diagnosis Suggestions* <https://www.jkb.com.cn/news/depth/2024/0606/495573.html>
3. Arterial Network / PharmaCube (VB100): *The "Last Mile" of Healthcare AI LLMs — Only Real Deployment Creates Value* <https://bydrug.pharmcube.com/news/detail/9a530f1ae67f0613f533cc87a6747241>
4. Xinhua Net: *China's First "Public Hospital Operations Management LLM" Released* <http://www.news.cn/health/20260302/a842fc4488504e59a5f90b4f324ea518/c.html>

5. Family Doctor (HIMSS 2025 coverage): *The Failure Trap of Healthcare AI: 80% of Projects Fail to Convert Value* <https://www.familydoctor.cn/news/yiliaozuzhi-ruhezhai-bimian-shibai-xianjing-tongshi-shishi-rengongzhineng-65193.html>
6. UCL News: *Study sheds light on hurdles faced transforming NHS healthcare with AI* <https://www.ucl.ac.uk/news/2025/sep/study-sheds-light-hurdles-faced-transforming-nhs-healthcare-ai>
7. CHIMA: *Legal Risks of Clinical Application of AI Medical Robots* <https://www.chima.org.cn/Html/News/Articles/17035.html>
8. Menlo Ventures: *2025 The State of AI in Healthcare* <https://menlovc.com/perspective/2025-the-state-of-ai-in-healthcare/>
9. FDA: *Artificial Intelligence-Enabled Device Software Functions* <https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-software-medical-device>
10. UC San Diego Health: *Study: AI Surveillance Tool Successfully Helps to Predict Sepsis and Save Lives* <https://health.ucsd.edu/news/press-releases/2024-01-23-study-ai-surveillance-tool-successfully-helps-to-predict-sepsis-saves-lives/>
11. UN News: *WHO Warns Legal and Ethical Safeguards for AI in Healthcare Lag Behind* <https://news.un.org/zh/story/2025/11/1141133>
12. Jiemian News: *Insilico Medicine Delivers AI Drug Discovery's First Profitable Results, Yet R&D Anxiety in Pharma Persists* <https://www.jiemian.com/article/14847447.html>

10.9 Manufacturing

10.9.1 Industry Profile and Bottleneck Triage

Among all industries in this book, manufacturing has the highest physical-bottleneck share: the information-to-physical bottleneck ratio is roughly 20-40:60-80. No matter how smart the algorithms, everything ultimately submits to the physical realities of machine-tool cycle times, material flow, and equipment wear. Yet it is precisely this least "digital" industry that McKinsey identifies as the largest generative-AI value pool — capable of delivering up to 2x productivity gains and nearly 70% task automation for manufacturing and supply chains, with nearly a quarter of China's ~US\$2 trillion generative-AI value coming from manufacturing and supply chains.

Enormous value alongside dismal deployment — the gap constitutes three core contradictions. First, the mismatch between production-line physical constraints and data intelligence: AI accelerates the decision layer, not the physical layer; cycle time, precision, and process stability cannot be accelerated directly by models. Second, industrial data is inherently heterogeneous: proprietary device protocols, SCADA/ERP/MES systems each in their own silos, and defect samples scarce to an extreme — hundreds of thousands of good-part images versus just twenty-some defective-part images, with labeling still depending on veteran workers about to retire. Third, ROI is hard to verify: a Deloitte survey shows 91% of industrial AI projects fall short of expectations, and 90% of failures are not technical.

10.9.2 Transformation Paths and Typical Practices

Lighthouse factories: a standardized paradigm from "points" to "surfaces." China leads the world in lighthouse factory count — Haier (Chinese: 海尔) with 13 cumulative sites and Foxconn (Chinese: 富士康) with 5 plants in the network. McKinsey data shows nearly 60% of use cases at newly designated lighthouses already leverage AI, with productivity up 2-3x, defect rates down 99%, and energy consumption improved by 30%. The real value of lighthouse factories lies not in technology but in organizational methodology: Foxconn turns its experience into white papers and extends toward "Lighthouse Industry," while Haier empowers externally through COSMOPlat (Chinese: 卡奥斯).

Industrial LLMs: from general chat to industry "operating systems." Huawei predicts industrial LLMs will become manufacturing's "operating system," with penetration expected to reach 50% in China by 2030 and market size to exceed RMB 42 billion. Players split into the big tech (Baidu, Huawei) and the "industry columns" (COSMOPlat's "Tianzhi," the Digital & Intelligent Building Materials Research Institute's (Chinese: 数智建材研究院) "Xiaomiao," Zhonggong Hulian's (Chinese: 中工互联) "Zhigong"). Shandong Province's LLM industry revenue reached RMB 6.38 billion in 2025, up over 70% year on year, organized through open-competition challenges.

Industrial agents: from "usable" to "affordable and replicable." CCID Consulting (Chinese: 赛迪顾问) data shows China's agent market at RMB 7.84 billion in 2025, projected to reach RMB 13.53 billion in 2026. 2025 has been called the "first year of industrial agents," with deployment mostly still at the demonstration stage;

2026 enters the "replication era," and the challenge shifts from technology to engineering. Vendors such as Advantech (Chinese: 研华) embed agents into existing production-line systems.

AI quality inspection and predictive maintenance: the two single-point scenarios with the clearest ROI. AI visual inspection is the fastest-penetrating scenario (China's smart inspection output value reached RMB 574.5 billion in 2025); vendors claim it can cut miss rates to 1% of manual inspection and save 70% of re-inspection costs [official claims]. IBM, citing ISA data, says plants typically lose 5%–20% of manufacturing capacity to downtime, and AI can cut unplanned downtime by 30%–50% [vendor claims] — but deployment pain points are prominent: legacy equipment without sensors, scarce failure samples, and ROI hard to quantify.

Platform transformation: manufacturers incubating AI platforms. SANY (Chinese: 三一) partners with RootCloud (Chinese: 树根互联), Haier incubates COSMOPlat, and Midea (Chinese: 美的) productizes its "Meiqing" (Chinese: 美擎) experience for external clients — forming a division of labor where "big companies build platforms, SMEs get on platforms." At the institutional level, the eight-department *Implementation Plan for the "AI + Manufacturing" Special Action* led by MIIT, combined with guidelines covering 8 links and 40 scenarios plus tiered cultivation, constitutes a tripartite promotion framework.

10.9.3 Benchmark Cases

Shenglong Electric (Chinese: 盛隆电气): 27 Agents Deployed in 14 Days [third-party verified]

The *2026 White Paper on AI Scenario Applications in Chinese Manufacturing* documents 167 deployment cases with "named enterprises and quantifiable data," among which Shenglong Electric deployed 27 industrial agents in 14 days, cutting labor coordination costs by 60% with ROI over 200% [third-party verified]. The white paper also offers sober structural data: generative AI accounts for only 4.8% of applications and AI agents only 3%, while traditional machine learning (46.7%) and machine vision (20.4%) remain the mainstays — "existing technology is sufficient; the bottleneck is scenario matching and the deployment approach."

Figure x BMW: Humanoid Robots Enter a Production Line [official claims]

Figure 02 completed an 11-month deployment at BMW's Spartanburg plant: over 1,250 hours of operation, 90,000+ components loaded, and contributions to the production of 30,000+ X3 vehicles; sheet-metal loading completed within an 84-second cycle at 5 mm tolerance, with a positioning success rate over 99% [official claims]. BMW's Munich plant, by contrast, represents the "digital-native" rebuild route: digital twins run across development and production validation, and a partnership with Siemens and NVIDIA accelerated pneumatic simulation 30x [third-party verified]. Embodied intelligence has obtained production-scale data for the first time, validating its engineering boundaries.

CITIC Pacific Special Steel (Chinese: 中信泰富特钢): Industrial AI for Process Industries [third-party verified]

In McKinsey's lighthouse factory cases, CITIC Pacific Special Steel uses AI to precisely predict and adjust blast furnace operations in real time, raising output by 15% and cutting energy consumption by 11% [third-party verified]. This represents the typical form of process-industry AI: fusing "mechanism models + data models" to optimize process parameters rather than replace people — chemicals likewise concentrates on four directions (process optimization, predictive maintenance, safety early warning, and energy optimization), and must be led by the process technology department.

10.9.4 Failure Warnings

91% Fall Short of Expectations: The Failure Is Organizational and Data-Related, Not Technical [third-party verified]

A Deloitte survey shows only 9% of Chinese manufacturing AI projects meet 80%–100% of expectations. The failure attributions are strikingly consistent: requirements get distorted after passing through three layers — "decision-makers → IT department → vendor" (one company invested millions in a "knowledge assistant"; daily active users fell to single digits within three months of launch and the project was halted); frontline workers did not trust a 99%-accurate system and insisted on manual re-inspection; and after vendor delivery, nobody maintained the model, the system degraded, and it was shut down. The common traits of surviving projects: entry through small scenarios, worker participation in requirements definition, and data governance front-loaded.

The AI Fabric Inspector: RMB 3.7 Million Invested, 15% Misjudgment Rate [third-party verified]

A textile company in the Yangtze River Delta invested RMB 3.7 million in an AI fabric inspection system, but fabric reflection issues drove the actual misjudgment rate as high as 15%; the company ultimately kept manual re-inspection and relegated AI to an auxiliary role. The lesson: insufficient scenario diversity and standardization, samples disconnected from process know-how, and single-point equipment retrofits cannot solve systemic problems.

Rushing into AI Quality Inspection: A Missing Data Foundation [third-party verified]

A manufacturer rushed full-process AI quality inspection without pilots or data preparation; insufficient training data led to subpar accuracy and production lines that could not run, and the project failed and had to be redone — about 60% of enterprises ignore that "data and business preparation precede technology." The common thread of failure is not weak algorithms but organizational and data infrastructure lagging behind.

10.9.5 Organizational Patterns and Consulting Takeaways

First, **recognize AI's boundaries in a physically bottlenecked industry.** Manufacturing AI's ROI concentrates in the decision layer and the "documentation and administrative layer around production" (SOP assistants, reports, knowledge Q&A, production scheduling), not in the physical layer itself — in McKinsey's lighthouse analysis, the fastest generative-AI gains appear in process and knowledge scenarios outside the shop floor (ACG Capsules' SOP assistant was used by nearly three-quarters of operators within five weeks). The first consulting move is to distinguish "physical cycle-time problems" (which AI cannot manage) from "decision and information problems" (which AI can).

Second, **single point first, then platform, then organization.** Start with high-frequency single points such as inspection and production scheduling to validate ROI, then build a data platform and model asset management, and finally change the organization (Chief AI Officer, AI Center of Excellence, business-unit AI specialists). McKinsey offers three scaling levers: packaging use cases as "assets," accelerating with low-code platforms, and phased introduction — AI assists decisions first, and only automates after confidence thresholds are exceeded.

Third, **the data foundation sets AI's ceiling.** Defect-sample generation, OT/IT integration, and the fusion of mechanism and data models are the three engineering prerequisites; predictive maintenance — the clearest-ROI scenario — is precisely where "data infrastructure debt" is most fully exposed.

Fourth, **manufacturers incubating AI platforms have three organizational paths:** internal incubation + external independent operation (SANY/RootCloud); parent-factory validation + platform company external delivery (Haier COSMOPlat, Midea Meiqing); and dual entities of "research institute + industrial company" (the Digital & Intelligent Building Materials Research Institute/"Xiaomiao"). The common thread is detaching AI capability from the parent's organizational inertia and competing in the market as independent entities.

Fifth, **business-led, not IT-led.** The chemicals industry explicitly requires leadership "by the process technology department rather than the IT department"; academic research summarizes manufacturing AI transformation as a staged leap of "cognitive awakening → strategic decision → organizational change → technology deployment → capability consolidation" — AI-native transformation is never a technology procurement problem, but a "cognition-behavior" double-helix organizational learning process.

1. McKinsey: *High Efficiency, Low Energy, Zero Defects: AI Makes Manufacturing Innovation One Click Away* <https://www.mckinsey.com.cn/高效率、低能耗、零缺陷：ai让制造创新一键直达/>
2. McKinsey: *The State of AI 2025: Only 6% of Companies Become High-Performing Winners* <https://www.mckinsey.com.cn/2025麦肯锡ai应用现状调研：从6%企业成为高绩效赢家，他们做对了什么？/>
3. McKinsey: *Generative AI in China: US\$2 Trillion in Economic Value* <https://www.mckinsey.com.cn/生成式ai在中国：2万亿美元的经济价值/>
4. Woshipm (Chinese: 人人都是产品经理): *90% of Industrial AI Deployment Failures Are Not Technical Problems* <https://www.woshipm.com/ai/6392172.html>
5. Sina: *Seeking the Truth About Industrial AI: Why Is AI Hot Yet So Hard to Deploy in Industry?* http://client.sina.com.cn/zt_d/2025-08-22/doc-infmvzcz5097908.shtml
6. Xingongye.com (Chinese: 新工业网): *Panoramic Analysis of 167 Verified Cases — 2026 White Paper on AI Scenario Applications in Chinese Manufacturing* https://xingongye.cn/cms/Intelligent_manufacturing/1825.html

7. Figure AI: *Figure 02 Humanoid Robots at BMW: Contributing to the Production of 30,000 Vehicles* <https://www.figure.ai/news/production-at-bmw>
8. Automotive Manufacturing Solutions: *Future-ready: BMW's digital transformation of its Munich plant and beyond* <https://www.automotivemanufacturingsolutions.com/smart-factory/future-ready-bmws-digital-transformation-of-its-munich-plant-and-beyond/523748>
9. People's Daily (People's Tribune): *The Development Landscape of China's "Lighthouse Factories" and the Shaping of Global Value Chains* https://paper.people.com.cn/rmlt/pc/content/202505/01/content_30080970.html
10. Huawei: *Industrial Digitalization/Intelligence 2030* <https://www-file.huawei.com/dam/asset/view/7ffffa5fa5e7469692d4704b92ef17d7.pdf>
11. State Administration of Science, Technology and Industry for National Defense: *Eight Ministries Including MIIT Issue the "Implementation Plan for the AI + Manufacturing Special Action"* https://www.nda.gov.cn/sjj/zwgk/zcfb/0112/20260107214358696030895_pc.html
12. CSDN Blog: *The Dilemmas and Breakthrough Paths of AI Transformation in Traditional Manufacturing* https://blog.csdn.net/weixin_30399155/article/details/96238191

10.10 Energy & Utilities

10.10.1 Industry Profile and Bottleneck Triage

Among the industries examined in this book, energy and utilities has the highest physical-bottleneck share: the information-to-physical bottleneck ratio is roughly 20–30:70–80. Grid power flows, pipeline pressures, and turbine speeds are all millisecond-level physical realities — no matter how smart the algorithms, Joule's law cannot be changed. The global energy AI market exceeded US\$40 billion in 2025, and the power industry's AI agent market exceeded RMB 40 billion [third-party verified], yet value has been slow to materialize — utilities invested over US\$10 billion in AI over the past five years, with most projects never progressing beyond the pilot stage [third-party verified].

Behind the gap are four core contradictions. First, physical safety constraints are the top priority: dispatch errors can trigger large-scale blackouts or even casualties; the US NERC CIP regulatory framework has no clear definition for AI models; and the

nuclear industry requires AI to enter only "human-machine collaboration" scenarios, never autonomous control. Second, asset-heavy and long-cycle: legacy systems 20–30 years old — SCADA, mainframes — cannot communicate with modern AI platforms, and companies spend 40% of their AI budgets making old systems "talk" [third-party verified]. Third, data fragmentation: data is scattered across dozens of isolated systems, and models are trained on incomplete, biased datasets [third-party verified]. Fourth, the dual constraints of regulation and performance assessment: central SOEs face assessments and grids must guarantee supply, leaving far less room for error than the internet industry.

10.10.2 Transformation Paths and Typical Practices

Grid AI: industry LLMs become the new "foundation." State Grid (Chinese: 国家电网) released China's first 100-billion-parameter multimodal power LLM, "Guangming" (Chinese: 光明), in December 2024, and its "6541" master plan in May 2025 (six application domains, five intelligent systems, and a four-in-one foundation); its intelligent customer service now auto-responds to over 90% of work orders, serving 270 million interactions cumulatively [official claims]. China Southern Power Grid (Chinese: 南方电网) takes the domestic-compute route: "DaWatt · Yudian" (Chinese: 大瓦特·驭电), trained on Huawei Ascend clusters with 13 billion parameters, was selected among the ten "national heavyweights" of central SOEs, and platform model calls exceeded 10 billion in December 2025 [official claims]. Both grid giants use a "headquarters LLM + localized small models" architecture: a dispatch robot in Xuancheng, Anhui compressed dispatch-instruction drafting from 3 hours to 30 minutes, and converter-transformer fault diagnosis in Liaocheng, Shandong was shortened from a week to milliseconds [third-party verified].

Oil & gas LLMs: from single-well intelligence to an "industry brain." CNPC's (Chinese: 中国石油) Kunlun LLM (Chinese: 昆仑大模型) released a 300-billion-parameter version (MoE architecture) in May 2025, covering 26 business lines, 119 business domains, and 100 core scenarios [official claims]; Sinopec (Chinese: 中石化) uses the "Great Wall" LLM (Chinese: 长城大模型) as its group foundation. Data has been elevated to equal standing with models: CNPC built a 550TB three-tier data management system, and "data is the blood of model training" is becoming industry consensus [third-party verified].

Renewable operations: AI from "veteran experience" to "precision diagnosis." Wind and solar installed capacity reached 1.482 billion kW in Q1 2025, surpassing thermal power for the first time, and the surge in installations is pushing operations and maintenance toward AI. Envision Energy (Chinese: 远景) trains agents on millions of real inspection records, letting AI "understand the language of wind turbines" [third-party verified]; Goldwind (Chinese: 金风) uses digital twins + IoT for predictive maintenance [third-party verified]; Shenzhen Energy (Chinese: 深圳能源), with Huawei Cloud, released a power forecasting system to reduce wind and solar curtailment [official claims]. The industry's overall fault-identification accuracy has been raised to over 90% [third-party verified].

Nuclear power and virtual power plants: the two endpoints of most conservative and most cutting-edge. In nuclear power, CNNC Nuclear Power's (Chinese: 中国核电) "Hezhi AI Portal" (Chinese: 核智 AI 门户) integrates DeepSeek, and CGN (Chinese: 中广核) has built a nuclear full-lifecycle intelligence hub on 18 core datasets — but applications are strictly limited to "human-machine collaboration" scenarios such as decision support and procedure interpretation [third-party verified]. In virtual power plants, an energy storage station in Dongguan completes 10MW power allocation in 0.5 seconds via AI dispatch [third-party verified]; State Grid Shanghai plans a "4+4" set of eight agents, and State Grid Hangzhou proposes a resource-layer—aggregation-layer—market-layer framework [third-party verified]; in 2025 China's VPP regulation capacity targets 20 million kW, with the market exceeding RMB 100 billion [third-party verified].

10.10.3 Benchmark Cases

Changqing Oilfield × Huawei: A Fully Autonomous Smart Oilfield [official claims]

Changqing (Chinese: 长庆油田), China's largest onshore oil and gas field (annual oil and gas equivalent output of over 60 million tonnes, managing 126,000 wells), partnered with Huawei and Kunlun Digital (Chinese: 昆仑数智) to build a three-tier architecture of "cloud-centralized training — regional inference management — edge real-time inference" on a fully autonomous technology stack of "Kunpeng + Ascend + Galaxy Kylin." The "intelligent hazard identification for special operations" project collected 2 million operation images and 150,000 hours of video, built a knowledge base of 12,000 rules, and deployed 20 high-precision models: hazard identification accuracy of 94.1% (18.7 percentage points over traditional methods), edge inference

latency under 800 ms, manual supervision intervention in high-risk operations down 70%, and overall operations and maintenance costs down 18%; reservoir development detection accuracy improved from 78% to 92%, and unplanned equipment downtime fell by 50% [official claims].

Duke Energy: Predictive Maintenance Proven at Scale on the Distribution Grid [third-party verified]

In 2025, Duke Energy's AI grid prevented about 1.2 million outages in the Carolinas — resolving outages before they reached customers and saving about 3 million hours of customer outage time cumulatively [third-party verified]. The approach is not flashy: AI prediction models identify equipment failure risk and vegetation hazards, prioritize risk using weather and asset condition, and schedule trimming and maintenance in advance. Even more valuable is the organizational method: starting in 2017 with a single electricity-theft detection use case, governance guardrails were established in 2024, and by 2025 over 50 generative-AI use cases were in flight. The CIO's three principles — "Always on" (reliability first), "Data before dazzle" (don't plug AI into bad data), and "Human in the loop" (AI augments people rather than replacing them) — are a classic articulation of utility AI governance [third-party verified].

The Kunlun LLM: An Organizational Sample of an SOE "Industry Brain" [official claims]

The Kunlun LLM was launched by the group's Party leadership in January 2024, became the first LLM in the energy and chemicals industry to pass national filing in August 2024, and raised its parameter count to 300 billion in May 2025 [official claims]. It represents a collaboration model of "central-SOE headquarters coordination + ecosystem partner co-building," with the goal of upgrading from "model capability" to a reusable "industry brain" for the whole sector [third-party verified].

10.10.4 Failure Warnings

74% of Pilots Never Leave a Single Region [third-party verified]. Prudent Consulting research shows that 74% of utility AI pilots never reach production beyond a single region. The six root causes are nearly all organizational: legacy system integration barriers (40% of budget spent on integrating old systems), data architecture not designed for AI, no AI lifecycle management (models abandoned as

accuracy decays without retraining for 12–18 months), skill gaps and organizational misalignment, missing governance, and improper ROI measurement [third-party verified].

70% of Leaders Are Dissatisfied with AI Progress [third-party verified]. A BCG 2024 survey shows nearly 60% of energy company leaders expect AI to deliver within a year, but about 70% are dissatisfied with progress — the expectation-reality gap constitutes the energy industry's "AI anxiety" [third-party verified]. The industry also widely suffers from "built but unused": projects completed without operational talent, models reduced to decoration [third-party verified].

Common Threads Behind Pilots That Fail to Scale: model decay, data silos, and the disconnect between pilot and production systems. The IEA also notes the energy industry broadly underutilizes AI, with data availability and digital skills gaps as the main barriers [third-party verified]. Accenture's *2025 China Enterprise Digital Transformation Index* likewise shows only about 9% of enterprises capture significant value from generative-AI implementation [third-party verified].

10.10.5 Organizational Patterns and Consulting Takeaways

First, **define AI's physical safety boundary first.** Energy is an industry where "errors are irreversible": AI's value concentrates in the decision and information layer — inspection recognition, load forecasting, dispatch assistance, and knowledge retrieval — and wherever people and grid safety are involved, human accountability must remain. The early lesson from the US GridFWD conference is that "capital is too precious and the risk of being wrong is too great" — build governance models before talking about use cases [third-party verified]. The first consulting move is to help enterprises distinguish scenarios where AI may only "suggest" from those where it may "execute."

Second, **single point first, then system, then organization.** Duke started with one electricity-theft detection use case in 2017, built governance guardrails in 2024, and only scaled in 2025; State Grid Hangzhou's experience is to prioritize "high-frequency, high-value, low-risk" scenarios and to set up cross-department AI teams [third-party verified]. "Slow first, then fast — earn the right to move quickly" is closer to the truth than rushing in [third-party verified].

Third, **the data foundation and centralized control centers are dual prerequisites.** Prudent lists "data architecture not designed for AI" as a root cause [third-party verified]; Changqing invested in 2 million images + 150,000 hours of video + a 12,000-rule knowledge base [official claims], and State Grid trained the Guangming LLM on 20 years of operational databases + 5 million work orders [official claims]. Central SOEs generally adopt a "group-level platform + edge inference" centralized control architecture, treating data governance and compute infrastructure as infrastructure rather than projects.

Fourth, **the peculiarities of SOE organizational change: assessment and fault tolerance.** State Grid moved from the "6541 plan" to a special action of "two years to consolidate the foundation, three years to drive improvement" [official claims], and China Southern Power Grid deploys 29 key initiatives through its annual work plan [third-party verified] — AI is upgraded from project-based to routine mechanisms, driven by documents, assessments, and dedicated bodies. But assessment-driven progress has its cost: the "built but unused" phenomenon is prominent. The consulting takeaway is to design a dual-track mechanism of "pilot fault tolerance + scale-up assessment," so AI does not become ledger-based political achievement.

Fifth, **upgrade AI from a tool to strategy and business model.** BCG advocates moving from "piloting with AI" to an "AI-first operating model": AI embedded in core business processes rather than edge applications, fused teams of data scientists and domain experts, and co-development with technology companies rather than procurement [third-party verified]. Guidehouse goes further: AI will push utilities from "operational optimization" to "business model reconstruction" — virtual power plants and dynamic pricing become possible [third-party verified]. And the IEA's proposition that "without AI there is no energy, and without energy there is no AI" [third-party verified] makes energy the only industry where "AI demand creators and users are one" — the biggest strategic variable for AI-native energy organizations.

1. State Grid: *Power "AI+" White Paper* (the 6541 plan) https://www.sohu.com/a/898976204_121079544
2. Baidu Baike: *Guangming Power LLM* <https://baike.baidu.com/item/光明电力大模型/65228636>

3. China Southern Power Grid: *"DaWatt · Yudian" Selected Among the Top 10 National Heavyweights of Central SOEs for 2024* https://www.csg.cn/xwzx/2025/2025gsyw/202501/t20250114_344842.html
4. China Southern Power Grid: *AI Innovation Platform Model Calls Exceed 10 Billion* https://www.csg.cn/xwzx/2025/2025gsyw/202512/t20251216_351726.html
5. National Energy Administration: *CNPC Releases the 300-Billion-Parameter Kunlun LLM* <https://www.nea.gov.cn/20250530/9431f59ac81f49c1a3839d7e066c4073/c.html>
6. Huawei Enterprise Business: *Changqing Oilfield Partners with Huawei to Build a Smart Oilfield Benchmark* <https://e.huawei.com/cn/blogs/2025/industries/oil-gas/shaping-oilfields-with-data-ai>
7. Sityos: *Duke Energy's AI Grid Prevented About 1.2 Million Outages* <https://www.sityos.com/en/use-cases/duke-energys-ai-grid-prevented-12-million-outages>
8. Forbes: *Duke Energy's CIO on Powering Digital and AI Transformation* <https://www.forbes.com/sites/peterhigh/2025/10/22/duke-energys-cio-on-powering-digital-and-ai-transformation/>
9. Prudent Consulting: *Six Root Causes of Failed Utility AI Investments* <https://www.prudentconsulting.com/blogs/ai-utilities-investment-failures/>
10. BCG: *AI in Energy: A New Strategic Playbook* <https://www.bcg.com/publications/2025/ai-in-energy-new-strategic-playbook>
11. IEA: *Milestone Report: Energy and AI* <https://www.iea.org/reports/energy-and-ai/>
12. GridFWD: *Early AI Use Cases and Lessons for Utilities, Grid Operators, and Regulators* <https://gridforward.org/early-ai-use-cases-and-lessons-for-utilities-grid-operators-and-regulators/>

Afterword: A Living Book

Afterword: A Living Book

Having finished the first edition, what I most want to talk about is not the content of this book — it is the book itself.

It came from an automated content pipeline: a system that scans the global information landscape every day, a body of knowledge entries that have been refined again and again, three rounds of targeted deep research, and a weekly automated merge-and-release cycle. From the very first edition, this book was never "finished" — it is **alive**.

That is itself a miniature of the AI-native organization: **delegate repetitive work to systems, keep judgment for humans, and make knowledge an asset that keeps growing**. The book describes how other organizations have done this — and the way this book produces itself is the same principle in practice.

On Limitations

The source material of this book has clear boundaries, and I must state them honestly:

- **Industry boundary:** The vast majority of cases come from information-intensive industries (software, content, finance, consulting). Evidence from physical-world industries (manufacturing lines, clinical medicine, logistics) is scarce. The sample-bias critique we apply to the Tencent Research Institute report in Chapter 3 applies equally to this book — extrapolating "the patterns of information-intensive industries" into "universal organizational laws" is dangerous.
- **Time boundary:** AI data goes stale within months. Each data point in this book is timestamped (most as of August 2026); always refer to the latest version.
- **Survivorship bias:** The organizations in our case studies are publicly disclosed transformation samples. Companies that failed quietly will never appear in a book — Chapter 9's failure cases are a limited compensation for this bias.

How to Contribute

This book is free and open. Any reading or sharing is welcome. If you find errors, omissions, or cases worth including:

- Open an issue at the GitHub repository: github.com/Jweokk/ai-native-organization-book (<https://github.com/Jweokk/ai-native-organization-book>)
- Or write to weokk2025@gmail.com

The value of knowledge lies in its flow, not in hoarding it. This book came from flow — and we hope it keeps flowing.

Appendix A: Case Index and Sources

Appendix A: Case Index and Sources

This appendix lists the original sources for every data point, case, and claim in this book. Attribution convention: company official statements are marked as **official claims**; third-party sources such as media, analysts, and regulatory filings are marked as **third-party verified**. All data is timestamped as of August 2026.

Sources for Chapters 1-3 (Reports and Data)

- https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
- <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>
- <https://www.bcg.com/publications/2025/ai-at-work-momentum-builds-but-gaps-remain>
- <https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential>
- <https://www.bcg.com/publications/2026/ai-at-work-why-strategy-matters-more-than-tools>
- <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>
- <https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025>
- <https://www.gartner.com/en/newsroom/press-releases/2025-11-10-gartner-survey-finds-artificial-intelligence-will-touch-all-information-technology-work-by-2030>

- <https://www.gartner.com/en/newsroom/press-releases/2025-10-21-gartner-unveils-top-predictions-for-it-organizations-and-users-in-2026-and-beyond>
- <https://www.forrester.com/blogs/predictions-2026-ai-moves-from-hype-to-hard-hat-work/>
- <https://www.forrester.com/blogs/predictions-2026-ai-agents-changing-business-models-and-workplace-culture-impact-enterprise-software/>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization>
- <https://hai.stanford.edu/ai-index/2025-ai-index-report>
- https://hai.stanford.edu/assets/files/hai_ai_index_report_2025.pdf
- <https://www.worldbank.org/en/publication/dptr2025-ai-foundations>
- <https://openknowledge.worldbank.org/bitstreams/86903114-6212-4c45-9011-938925cc61d1/download>
- <https://www.imf.org/en/publications/staff-discussion-notes/issues/2024/01/14/gen-ai-artificial-intelligence-and-the-future-of-work-542379>
- <https://www.imf.org/-/media/Files/Publications/WP/2025/English/wpiea2025076-print-pdf.ashx>
- <https://www.oecd.org/en/about/projects/aisurveysofemployersandworkers.html>
- https://www.oecd.org/en/publications/the-impact-of-ai-on-the-workplace-main-findings-from-the-oecd-ai-surveys-of-employers-and-workers_ea0a0fe1-en.html
- https://www.oecd.org/en/publications/oecd-employment-outlook-2023_08785bba-en/full-report/artificial-intelligence-job-quality-and-inclusiveness_a713d0ad.html
- <https://thenextweb.com/news/oecd-finance-and-manufacturing-workers-fear-ai-replacement>
- https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en/full-report/component-3.html

- <https://www.weforum.org/press/2025/01/future-of-jobs-report-2025-78-million-new-job-opportunities-by-2030-but-urgent-upskilling-needed-to-prepare-workforces/>
- https://reports.weforum.org/docs/WEF_Future_of_Jobs_Report_2025.pdf
- <https://www.hbs.edu/ris/download.aspx?name=26-090.pdf>
- https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6905079
- <https://www.hbs.edu/faculty/Pages/item.aspx?num=69077>
- <https://aiinstitute.hbs.edu/everyone-has-ai-which-firms-are-going-to-win/>
- https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6513481
- <https://aiinstitute.hbs.edu/less-headcount-more-valuation-how-ai-native-firms-change-the-game/>
- <https://aiinstitute.hbs.edu/is-genai-heading-for-a-tech-monopoly/>

Sources for Chapter 7 (International Case Studies)

- <https://www.cnn.com/2025/05/14/klarna-ceo-says-ai-helped-company-shrink-workforce-by-40percent.html>
- <https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm>
- <https://www.bloomberg.com/news/articles/2024-12-12/klarna-stopped-all-hiring-a-year-ago-to-replace-workers-with-ai>
- <https://techcrunch.com/2024/12/14/klarnas-ceo-says-it-stopped-hiring-thanks-to-ai-but-still-advertises-many-open-positions/>
- <https://www.threads.com/@fox.hsiao/post/DVP3e7jIA17/>
- <https://www.techhanlin.tw/klarna-ai-customer-service-case-study/>
- <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-its-first-month/>
- <https://www.klarna.com/international/press/90-of-klarna-staff-are-using-ai-daily-game-changer-for-productivity/>
- <https://www.customerexperiencedive.com/news/klarna-says-ai-agent-work-853-employees/805987/>
- <https://www.customerexperiencedive.com/news/klarna-ai-slash-customer-service-costs/748647/>

- <https://investors.klarna.com/News--Events/news/news-details/2026/Klarna-Group-plc-Publishes-Full-Year-2025-Results/default.aspx>
- <https://investors.klarna.com/News--Events/news/news-details/2026/Klarna-Accelerates-U-S--Growth-and-Delivers-1bn-Revenue-Driven-by-Rapid-Banking-Service-Adoption/default.aspx>
- <https://www.cnbc.com/2025/11/18/klarna-klar-stock-q3-earnings-report-2025.html>
- <https://www.bloomberg.com/news/articles/2025-05-08/klarna-turns-from-ai-to-real-person-customer-service>
- <https://www.investopedia.com/buy-now-pay-later-company-klarna-s-shares-end-first-session-above-ipo-price-11806387>
- <https://www.businessweekly.com.tw/management/blog/3021962>
- <https://www.cnbc.com/2023/05/04/shopify-cuts-20percent-of-its-workforce-shares-surge-on-earnings-beat>
- <https://www.businessinsider.com/shopify-ai-use-boosts-efficiency-shakes-up-staff-structure-2024-5>
- <https://www.cnbc.com/2025/04/07/shopify-ceo-prove-ai-cant-do-jobs-before-asking-for-more-headcount.html>
- <https://www.reveliolabs.com/companies/shopify/employees>
- <https://x.com/tobi/status/1909231499448401946>
- <https://bullfincher.io/companies/shopify/revenue-per-employee>
- <https://www.chargeflow.io/blog/shopify-statistics>
- <https://fourweekmba.com/shopify-revenue-per-employee/>
- <https://natesnewsletter.substack.com/p/my-honest-field-notes-on-how-the>
- <https://en.wikipedia.org/wiki/Anthropic>
- <https://www.cnbc.com/2026/08/17/anthropic-says-annualized-revenue-climbed-to-65-billion-in-july.html>
- <https://venturebeat.com/technology/anthropic-says-80-of-its-new-production-code-is-now-authored-by-claude-how-your-enterprise-can-keep-up>
- <https://x.com/AnthropicAI/status/2062568864240836995>
- <https://www.cnbc.com/2025/11/13/cursor-ai-startup-funding-round-valuation.html>

- <https://www.anthropic.com/research/anthropic-economic-index-september-2025-report>
- <https://skeptics.stackexchange.com/questions/59213/as-at-september-2025-is-ai-not-writing-90-of-code>
- <https://blog.redwoodresearch.org/p/is-90-of-code-at-anthropic-being>
- <https://www.makerstations.io/openai-employee-statistics/>
- <https://www.mokahr.io/myblog/talent-culture-strategy-at-openai/>
- <https://searchlab.nl/en/statistics/openai-statistics-2026>
- <https://www.cnbc.com/2025/09/03/openai-boosts-size-of-secondary-share-sale-to-10point3-billion.html>
- <https://openai.com/index/how-agents-are-transforming-work/>
- <https://fortune.com/2026/08/13/buried-in-openais-latest-research-no-correlation-between-ai-use-and-revenue-per-employee/>
- <https://www.cnbc.com/2025/09/18/notion-launches-ai-agent-as-it-crosses-500-million-in-annual-revenue.html>
- <https://www.forbes.com/sites/annatong/2025/12/15/notion-kicks-off-employee-share-sale-at-11-billion-valuation-as-ai-accelerates-its-growth/>
- <https://www.notion.com/blog/introducing-notion-ai>
- <https://medium.com/@sherrysun/notion-from-near-collapse-to-a-10b-all-in-one-workspace-unicorn-6113f7e1dc4c>
- <https://getlatka.com/companies/notion.ai>
- <https://cursor.com/blog/series-d>
- [https://en.wikipedia.org/wiki/Cursor_\(company\)](https://en.wikipedia.org/wiki/Cursor_(company))
- https://www.reddit.com/r/SaaS/comments/1r9pj7k/cursor_hit_1b_arr_with_300_employees_thats_33m/
- <https://app.dealroom.co/news/note/cursor-tops-4b-annualized-revenue-june-2026>
- <https://www.cnbc.com/2025/04/17/openai-looked-at-cursor-before-considering-deal-with-rival-windsurf.html>
- <https://mashable.com/article/duolingo-ai-layoff-contractors>
- <https://www.cnbc.com/2025/09/17/duolingo-ceo-how-ai-makes-my-employees-more-productive-without-layoffs.html>

- <https://www.businessinsider.com/duolingo-ceo-how-ai-will-be-used-performance-reviews-headcount-2025-4>
- <https://www.metaintro.com/blog/duolingo-ceo-walks-back-ai-first-memo-hiring-grows-2026>
- <http://investors.duolingo.com/company-strategy-overview-0>
- <https://www.classcentral.com/report/duolingo-2025/>
- https://www.reddit.com/r/duolingo/comments/191ssv7/discussion_duolingo_cuts_10_of_its_contractors/
- <https://techcrunch.com/2026/05/08/airbnb-says-ai-now-writes-60-of-its-new-code/>
- <https://techcrunch.com/2025/05/02/airbnb-is-quietly-rolling-out-an-ai-customer-service-bot-in-the-us/>
- <https://www.techtimes.com/articles/323585/20260807/airbnb-proves-ai-powered-engineering-beats-rivals-80-more-features-same-staff.htm>
- <https://techcrunch.com/2026/02/13/airbnb-plans-to-bake-in-ai-features-for-search-discovery-and-support/>
- <https://techcrunch.com/2025/04/29/microsoft-ceo-says-up-to-30-of-the-companys-code-was-written-by-ai/>
- <https://techcrunch.com/2026/02/12/spotify-says-its-best-developers-havent-written-a-line-of-code-since-december-thanks-to-ai/>
- https://www.reddit.com/r/technology/comments/1t9v6km/airbnb_says_ai_now_writes_60_of_its_new_code/
- <https://www.cnbc.com/2024/02/01/zoom-layoffs-company-cuts-150-employees-2percent-of-workforce.html>
- <https://www.bbc.com/news/technology-64562673>
- <https://www.makerstations.io/zoom-employee-statistics/>
- <https://www.ciodive.com/news/zoom-ai-companion-monetization-customization/708688/>
- <https://www.facebook.com/WSJ/posts/zoom-founder-and-ceo-eric-yuan-predicts-ai-will-put-an-end-to-the-5-day-workweek/1321223876530873/>
- https://www.linkedin.com/posts/conorgrennan_just-released-microsofts-2025-work-trend-activity-7320839065984000004-B8gk
- <https://www.microsoft.com/en-us/worklab/work-trend-index>

- <https://themicrosoftcloudblog.com/2026/05/2026-work-trend-index-evidence-check/>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born>

Sources for Chapters 8-9 (Methodology and Failure Cases)

- <https://complexdiscovery.com/why-95-of-corporate-ai-projects-fail-lessons-from-mits-2025-study/>
- https://www.theregister.com/2025/08/18/generative_ai_zero_return_95_percent/
- <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>
- <https://www.gartner.com/en/newsroom/press-releases/2025-02-26-lack-of-ai-ready-data-puts-ai-projects-at-risk>
- <https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>
- <https://anarsolutions.com/why-agentic-ai-pilots-fail-production/>
- <http://www.eeo.com.cn/2026/0601/898378.shtml>
- <https://www.unifyapps.com/resources/blog/why-95-of-generative-ai-pilots-never-reach-production>
- <https://www.unifyapps.com/resources/customer-story/Large-Financial-Institution-India>
- <https://www.unifyapps.com/resources/customer-story/fortune-50-retailer>
- <https://www.cnbc.com/2025/05/14/klarna-ceo-says-ai-helped-company-shrink-workforce-by-40percent.html>
- <https://www.fastcompany.com/91468582/klarna-tried-to-replace-its-workforce-with-ai>
- <https://www.digitalapplied.com/blog/klarna-reverses-ai-layoffs-replacing-700-workers-backfired>
- <https://www.thestateofbrand.com/news/klarna-reverses-ai-job-cuts>
- <https://tandemcoach.co/klarna-ai-automation-lesson/>
- <https://www.bloomberg.com/news/articles/2023-05-02/samsung-bans-chatgpt-and-other-generative-ai-use-by-staff-after-leak>

- <https://www.forbes.com/sites/siladityaray/2023/05/02/samsung-bans-chatgpt-and-other-chatbots-for-employees-after-sensitive-code-leak/>
- <https://mashable.com/article/samsung-chatgpt-leak-details>
- <https://www.weforum.org/stories/artificial-intelligence/companies-ai-workflows-not-simple-tasks/>
- <https://www.stackmatix.com/blog/copilot-market-adoption-trends>
- <https://peafowlit.com/blog/copilot-licenses-go-unused-and-how-to-fix-adoption/>
- https://www.linkedin.com/posts/jukkaniiranen_34-of-microsoft-365-customers-pay-for-premium-activity-7422758776069455872-CInk
- <https://www.querynow.com/resources/whitepapers/past-the-stall-m365-copilot-rollouts>
- <https://www.linkedin.com/pulse/microsofts-copilot-paradox-94-report-benefits-6-deploy-louis-columbus-fv4gc>
- <https://www.onlineeducation.com/features/chatgpt-crashes-chepps-stock>
- <https://www.highereducationinquirer.org/2025/07/chepp-critical-history-of-disruptor.html>
- https://www.linkedin.com/posts/harshadshah1953_the-fall-of-chepp-is-a-masterclass-in-ignoring-activity-7487102244815921152-yM6L
- <https://cn.ceibs.edu/media/press-clippings/faculty/29204>
- <https://www.workiva.com/resources/data-pressures-mount-instability-continues>
- <https://www.woshipm.com/ai/6275554.html>
- <https://maiagent.ai/blog/enterprise-shadow-ai-governance>
- <https://securitytoday.com/articles/2025/08/18/survey-nearly-half-of-employees-hide-workplace-ai-use.aspx>
- <https://www.cloudflare.com/zh-cn/learning/ai/what-is-shadow-ai/>
- <https://www.unleash.ai/artificial-intelligence/news/asana-64-of-employees-believe-ai-agents-are-unreliable-calling-for-more-training-clarity-and-guardrails>
- <https://asana.com/resources/state-of-ai-research-takeaways>
- <https://www.accenture.com/us-en/insights/consulting/leadership-edge-ai>
- <https://www.d1net.com/cio/ciotech/585344.html>
- <https://blog.104.com.tw/wef-report-2026-ai-entry-level-work-redesign/>

- <https://www.hbrtaiwan.com/special-topics/24535/the-perils-of-using-ai-to-replace-entry-level-jobs>
- <https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic>
- <https://www.swfte.com/blog/avoid-ai-vendor-lock-in-enterprise-guide>
- <https://stepto.net/blog/ai-vendor-lock-in-infrastructure-risk-2026>
- https://www.linkedin.com/posts/sumatosoft_aistrategy-vendorlockin-enterpriseai-activity-7458501143887765504-FfQC
- <https://www.institutepm.com/knowledge-hub/ai-vendor-lock-in-strategy>
- <https://fin.ai/learn/ai-customer-service-cost-savings-industry>
- <https://www.cmswire.com/contact-center/will-ai-cost-more-than-offshore-human-agents-in-customer-service/>
- <https://www.stcn.com/article/detail/3951204.html>
- <https://resources.anthropic.com/enterprise-ai-transformation-guide>
- <https://www.rubrik.com/blog/technology/25/11/its-early-days-for-agent-ai-and-most-companies-lack-the-tools-to-protect-their-data>
- <https://www.rubrik.com/company/newsroom/press-releases/26/ai-agents-are-breaking-things-and-organisations-know-it>
- <https://www.rubrik.com/products/rubrik-agent-cloud>
- <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>
- <https://www.mckinsey.com/featured-insights/charts/ai-at-work-but-not-at-scale>
- <https://academy.theartofservice.com/course/section.php?id=103955>
- <https://www.bcg.com/featured-insights/the-leaders-guide-to-transforming-with-ai>
- <https://www.accenture.com/us-en/about/reinvention-services>
- <https://rtslabs.com/enterprise-ai-roadmap>
- <https://thinking.inc/en/pillar-pages/ai-adoption-roadmap/>
- <https://mp.weixin.qq.com/s/NJW9cYB8aH-1glXxdUnSaA>
- <https://www.zilbix.com/ai-native-operating-model-how-enterprises-must-restructure-for-the-agentic-era/>
- <https://www.anthropic.com/research/multiagent-systems>
- <https://www.mininglamp.com/news/8292/>
- <https://www.honeycomb.io/blog/your-questions-about-ai-agents-production-feedback-answered>

- <https://ampcode.com/notes/thats-not-soc-2-compliant>
- <https://futurism.com/artificial-intelligence/young-people-ai-ceos-executives-poll>
- <https://www.librarian.net/notoai/>
- <https://blog.jsbarretto.com/post/i-remain-a-skeptic>
- <https://opensauce.it/cloudflare-ai-psychosis/>
- <https://allen.bargi.org/notes/working-with-ai-feels-like-leadership/>
- <https://www.wired.com/story/model-behavior-interview-with-openai-codex-lead-tibo-sottiaux/>
- <https://www.stephen-cresswell.com/2026/08/15/Yadda-3.0.0-BDD-in-the-Age-of-AI-Agents.html>

Wiki 知识库素材（本书内容来源之一）

Knowledge-Base Material (One Source of This Book's Content)

Much of this book draws from the author's personal knowledge base (wiki). Its original sources include 36Kr, Woshipm, Tencent Research Institute, HBS Working Papers, WEF, Harvard Business Review (Chinese edition), LatePost, Economic Observer, and other public publications. Specific entries are recorded in the CHANGELOG with each version update.

Appendix B: Key Metrics for AI-Native Organizations

Appendix B: Key Metrics for AI-Native Organizations

You get what you measure. The measurement system of an AI-native organization differs fundamentally from the industrial-era one: it shifts from "input and activity" to "outcome and value." The metrics below are grouped by purpose, all drawn from the research cited in this book.

1. Organizational Structure Metrics ("What do we look like?")

Metric	Description	Benchmark
Relative team size	Headcount vs peers	HBS: AI-native firms are ~25% smaller (services ~70% smaller)
Engineer share	Engineers / total headcount	HBS: ~13 percentage points higher
Management layer depth	Layers from CEO to frontline	HBS: half a layer flatter; Shmool: 4 layers → 1
Builder-to-manager ratio	Direct producers / coordinators	Shmool: 3:1 → 8:1+
Human-to-AI ratio	Humans : AI agents	WEF: early AI-first organizations exceed 1:10
Information fidelity	Distortion of frontline facts reaching decision-makers	Shmool: ~40% loss per hop → ~95% direct
Decision latency	From problem to decision	Shmool: days-weeks → hours
Meeting overhead	Share of workweek in meetings	Shmool: ~30% → ~10%

2. Efficiency and Value Metrics ("What did AI deliver?")

Metric	Description	Benchmark
Revenue per employee	Revenue / headcount	Klarna: \$1.24M (3.6x vs 2022); Shopify: \$1.52M; Cursor: \$3M-13M
Per-capita growth rate	YoY change in revenue per employee	More robust than absolute value for transformation tracking
Output capacity	Output per unit time	Duolingo: course skills 1,800 → 20,500 per quarter (11x)
Time savings	Task-level time savings	Anthropic dialogue research: ~80% on average
Task-level vs structural gains	Single-step speedup vs full-flow improvement	Task-level 20-40%; workflow redesign 2-10x (Harvard Data Science Review)
EBIT attribution	AI's measured contribution to profit	McKinsey: only 39% attribute any EBIT impact; "high performers" (≥5%) ~6%
1.9x revenue effect	Causal gains from organizational learning	INSEAD/HBS field experiment (treatment vs control)

3. Workflow Metrics ("Did the work actually get better?")

Principle: measure the whole workflow, not the AI task (Chapter 5, rule 8).

- **Cycle time**: end-to-end from trigger to completion (not single-step);
- **First-time correctness rate**: share of cases done right the first time;
- **Exception rate and exception age**: share of exceptions + how long they sit in queues;
- **Repeat-contact rate**: how often the same issue is re-submitted (the Klarna lesson: aggregate metrics hide collapses in high-value scenarios);
- **Complex/escalated-interaction satisfaction**: tracked separately, never averaged into the overall number;
- **Review time**: human review hours (after AI output explodes, review becomes the new bottleneck);

- **Cost per completed case:** full-cycle cost (inference + integration + ops + QA + human review + retries);
- **Rework rate:** share of AI output rejected/redone;
- **AI confidence and coverage:** AI participation rate and confidence distribution (where do low-confidence cases route);
- **Escalation quality:** whether escalations to humans carry full context.

4. Adoption and Governance Metrics ("Is the organization actually using it?")

Metric	Description	Warning data
License utilization	Active users among those with access	Copilot: only ~1/3 (64% of licenses unused)
Paid/premium conversion	Users paying for AI features	Copilot: 3.4%; Notion AI: >50% paid penetration (contrast)
Four-dimensional success	Adoption / efficiency / quality / satisfaction	Anthropic's official pilot measurement
Shadow AI usage	Use of unauthorized tools	Salesforce: 69% of employees used unauthorized tools
AI value accounting	Finance-led ROI validation	CFO-led validation success 76% vs tech-led 53%
Full-delegation ceiling	Share of work fully delegable to AI	Anthropic internal: most employees say only 0-20%

5. Risk Metrics ("Will something go wrong?")

- **Vendor switching cost:** 19-34% of total deployment cost; only 6% of firms can switch without disruption;
- **Agent-related incidents:** Rubrik: 98% of organizations have experienced a disruptive AI-agent event;
- **Skill-atrophy signals:** employees reporting fewer hands-on practice opportunities; share of exploratory work (Anthropic internal: 27% of AI-assisted work would not have been done otherwise);

- **Employee trust:** share of employees who find AI agents reliable (Asana: only 36%);
- **Data AI-readiness:** completion rate of AI-ready data audits (<20% of organizations consider themselves ready).

6. Meta Metrics (Organizational Design)

- **Competitiveness formula:** talent density $\times$ AI leverage $\div$ organizational friction (Tencent Research Institute) — measure the three variables separately; cut the denominator before doubling the numerator;
- **Innovation density:** output / headcount (DeepSeek's 160 vs the giants' thousands);
- **Influence spillover:** whether one person makes the team faster (the threshold for a Super-Individual, not personal output multiples);
- **Review signal:** does performance review look at input or output? Is AI usage part of the review? (Shopify and Duolingo already include it.)

Appendix C: Version History and Update Notes

Appendix C: Version History and Update Notes

Update Mechanism

This book is a "living book," updated automatically every week:

1. An automated system scans global AI-landing and organizational-change information daily, distills it, and stores it in a knowledge base;
2. Each week, new knowledge-base content is merged into the relevant chapters and the version number is incremented (patch +1);
3. The whole-book PDF is regenerated and synced to the GitHub repository (this repository — note: the Chinese edition is also served on a Chinese-facing website; the English edition is hosted here on GitHub only);
4. Every update is recorded in this appendix and in the repository's CHANGELOG.

Version rules: v1.0.0 → v1.0.1 (weekly routine updates) → v1.1.0 (new chapters / major revisions) → v2.0.0 (full rewrite).

v1.0.0 — 2026-08-22 (First Edition)

First edition release: built on the knowledge base's accumulated material (12 core concept pages + 20+ case articles) plus three rounds of deep research reinforcement (international cases / authoritative reports / failure cases & methodology).

- 9 chapters + afterword + 3 appendices;
- Chapter 1: The rise of AI-native organizations (the precise framing of the 95% failure rate + a 515-startup causal experiment);
- Chapter 2: Band-aid AI vs AI-native DNA (including a correction of the widely-circulated Gartner CAIO prediction: "90%" is a misattribution; the actual figure is 35%);

- Chapter 3: The evidence (HBS 26-090 / INSEAD / WEF / Tencent Research Institute / McKinsey / BCG / Microsoft);
 - Chapter 4: Rebuilding organizational structure (MarsWave / Shmool / Block / the Organizational Brain / three paths);
 - Chapter 5: Redesigning workflows and decisions (9 key redesigns / decision rights table / 90-day cadence / three-tier routing);
 - Chapter 6: Incentives and talent (Energy Gold / review reform / four incentive dimensions / skill atrophy / innovation density);
 - Chapter 7: Case studies (China: Transn / Huawei / DeepSeek / Jingzhunxue / MarsWave; Global: Klarna / Shopify / Anthropic / OpenAI / Notion / Cursor / Duolingo / Airbnb) + six cross-case patterns;
 - Chapter 8: Action playbook (Phase 0-4 roadmap / BCG's three plays / Accenture's three principles / 10-error quick table);
 - Chapter 9: Challenges, risks, and the future (four failure cases / six risk categories / three future directions);
 - Appendix A: 165 sources (copyright compliance); Appendix B: key metrics; Appendix C: this appendix.
-