

Growing AI into the Organization's DNA

Jweokk · 2026-09-06 · v2.0.2

Models are no longer scarce. People who can grow models into their organizations are.
From the 95% project failure rate to DeepSeek's 160-person innovation density —
this book explains what an AI-native organization is, what the evidence says, and how to build one.

Copyright © Jweokk. Free for reading and non-commercial sharing; commercial use requires prior written permission.

Contact: weokk2025@gmail.com

Contents

- 01** Preface: Why I Wrote This Book
- 02** Chapter 1 The Rise of AI-Native Organizations — Starting from the 95% Failure Rate
- 03** Chapter 2: What Is an AI-Native Organization — Band-Aid AI vs. AI-Native DNA
- 04** Chapter 3: The Evidence — What the Empirical Research Says
- 05** Chapter 4: Rebuilding Organizational Structure — From Hierarchy to Network
- 06** Chapter 5: Redesigning Workflows and Decisions
- 07** Chapter 6 Incentives and Talent: The Hardest Variable to Rebuild
- 08** Chapter 7: Case Studies — AI-Native Organizations in China and Around the World
- 09** Chapter 8 Building Your AI-Native Organization: An Action Playbook
- 10** Chapter 9: Challenges, Risks, and the Future
- 11** Chapter 10 Industry Guide: AI-Native Transformation Across Ten Industries
- 12** Afterword: A Living Book
- 13** Appendix A: Case Index and Sources
- 14** Appendix B: Key Metrics for AI-Native Organizations
- 15** Appendix C: Version History and Update Notes

Preface: Why I Wrote This Book

In the summer of 2026, two headlines read side by side made the answer almost obvious.

The first: MIT's NANDA lab released *The GenAI Divide*, reporting that over the past three years, companies worldwide have burned \$30–40 billion on generative AI, and **95% of organizations have gotten no return that could be written into their financial statements**. For customized enterprise AI tools, the funnel from evaluation to pilot to production is 60% → 20% → 5%.

The second: DeepSeek built a trillion-parameter model with 160 people; Klarna shrank its workforce from 5,527 to 3,422 employees over three years while revenue kept rising; at Anthropic, Claude wrote more than 80% of the production code merged; and a company called Cursor reached \$4 billion in annualized revenue with 300 people.

On one side, 95% of AI projects come to nothing. On the other, a small cluster of organizations are leveraging enormous output with tiny teams.

Put the two together and the answer is not hard to guess: **models are no longer scarce. What is scarce are the people who can grow models into their organizations.**

What I've Been Researching

About 90% of the public discussion around "AI transformation" tells you how to **add AI capabilities** to your existing organization: buy a Copilot, run a few dozen training sessions, appoint a chief AI officer. But what is actually happening runs on a different track — **letting the organization itself grow AI into its DNA**. These two things are worlds apart.

This book sets out to study, end to end, a systemic change that is already under way but does not yet have a name: the AI-native organization. It clarifies three things:

- **What it is:** the essential difference between AI-native organizations and "AI transformation" — band-aid AI vs AI-native DNA;

- **The evidence:** what the empirical research from HBS, INSEAD, the WEF, and Tencent Research Institute actually says;
- **How to do it:** restructuring organizational form, restructuring workflows, restructuring incentives — plus a zero-to-one playbook.

All data and case studies in this book are cited in Appendix A. The work of compiling was itself a form of learning, and I have tried to make every citation verifiable; if anything is missing, corrections are welcome.

How This Book Was Auto-Written

This book was not written in a single pass. It is a "living book," powered by an automated content pipeline:

1. **Daily accumulation:** an automated system scans global information on AI deployment and organizational change every day, distilling, translating, and structuring high-quality content into a knowledge base;
2. **Deep research:** on top of the material already in the knowledge base (12 core concept articles plus 20+ case articles), the first edition underwent three further rounds of deep research — landmark international cases, raw data from authoritative reports, and failure cases and construction methodologies — with every source verified one by one;
3. **Weekly updates:** every week since, newly added knowledge base content is automatically merged into the relevant chapters, the version number increments (v1.0.0 → v1.0.1 → ...), the PDF is regenerated, and it is automatically deployed to the website;
4. **Continuous evolution:** whatever each update changed or added is recorded in Appendix C and in the repository's CHANGELOG, at a glance.

The process of writing this book is itself the book's best case study: **one person + an automated information pipeline + continuous weekly iteration produced something a traditional publishing process would take months to complete — and could never keep updated.** That is the AI-native organization in miniature, applied to content production.

Acknowledgments

This book's format and open publication model were inspired by XDash (Fan Bing)'s book, *Frontline Deployment Engineer: The Playbook for Delivering Customer Value in the Age of AI* (<https://github.com/xdash/FDE-the-Guidance-Book-of-Forward-Deployed-Engineer>) (<https://github.com/xdash/FDE-the-Guidance-Book-of-Forward-Deployed-Engineer>). The pattern of free public full text, synchronized reading via a GitHub repository and website, whole-book PDF download, and appendix-sourced citations is an open-source knowledge practice worth promoting. Special thanks.

Much of the material in this book comes from public industry research, media reports, and practitioners' sharing, and I have done my best to cite sources in Appendix A. If any citation is missing or there are copyright concerns, please contact weokk2025@gmail.com and we will correct it immediately.

Reading Guide

- To quickly judge whether your company counts as AI-native: read Chapter 2
- To see evidence rather than opinion: read Chapter 3
- To get hands-on: read Chapter 8, together with the workflow restructuring in Chapter 5
- To understand why most transformations fail: read Chapter 9
- To go through all the case studies: read Chapter 7

This book is free and open; feel free to share it. The value of knowledge lies in its flow, not in hoarding it.

Chapter 1 The Rise of AI-Native Organizations — Starting from the 95% Failure Rate

1.1 A Number That Has Been Misread

Over the past two years, the most widely circulated AI statistic on the Chinese internet has been: "95% of enterprise AI projects fail."

The claim is both right and wrong. It comes from MIT's NANDA lab's July 2025 report, *The GenAI Divide*.^[1-02] The report's precise framing has two layers:

- **By organization:** 95% of organizations got zero return^[1-02] from their generative AI investments;
- **By tool funnel:** of customized, task-specific enterprise AI tools, only 5% make it to production deployment — 60% of organizations evaluated such tools, 20% entered pilots, and 5% reached production.

General-purpose chat tools (ChatGPT, Copilot, and the like), by contrast, convert from pilot to deployment at roughly 83%^[1-02] — they are not failing.

This detail matters enormously: **what "failed" is not all AI projects, but the customized systems deeply embedded in business processes that set out to replace core ways of working.** People who used AI casually gained efficiency; organizations that seriously reworked their businesses mostly came to grief. That is the "GenAI divide" the report describes — a polarization of outcomes between buyers (enterprises) and sellers (AI companies).

Figures from other institutions paint the same picture:

- Gartner predicts that by the end of 2025, at least 30% of generative AI projects will be abandoned after proof of concept (PoC)^[1-07];
- Gartner predicts that by the end of 2027, more than 40% of agentic AI projects will be canceled, with soaring costs the primary reason;
- IDC: 88% of AI agent pilots fail to graduate into production;
- An industry survey: for every 33 pilots launched, only about 4 successfully make it into production;

- Professor Fang Yue of CEIBS (China Europe International Business School) (2026): more than 90% of companies worldwide have launched generative AI pilots, but fewer than 41% of projects have truly crossed the experimental stage to create value at scale — he calls this phenomenon "pilot boom, value void."

Note that these numbers use entirely different measures (organizational zero-return rate / PoC abandonment rate / pilot-to-production conversion rate), yet they all point to the same structural fact: **the gap from pilot to production is real, and it is pervasive.**

1.2 Why Pilots Succeed and Production Fails

A strange, repeatedly observed phenomenon: AI projects that demo beautifully in the pilot phase collapse the moment they enter production.

The reason is that pilot success usually depends on humans "quietly patching things up" — supplying missing context by hand, bridging system failures, absorbing risk. Once those patches are stripped away in production, the AI immediately fails. Three layers specifically:

1. **Broken context:** the AI can only see a narrow slice of enterprise information and cannot reason across systems or across time; unstructured documents (PDFs, contracts, handwritten forms) are entirely invisible to it. In a demo this can be hidden; in production there is nowhere to hide.
2. **Fragile integration:** pilots are usually read-only; production requires the AI to write back to systems, trigger workflows, and act across dozens of applications. Point-to-point fragile integrations are immediately exposed — "read-only intelligence cannot act."
3. **Failed governance:** at pilot scale, human review can catch everything; at production scale, the "human in the loop" becomes a bottleneck rather than a safety valve, and missing audit trails breed risk aversion and frozen projects.

The few organizations that did make it into production (about 5%)^[1-02] made three different decisions: **build shared enterprise context first, standardize AI's "action layer," and embed governance into systems rather than leaving it to manual management.**

1.3 Another Group Running Hard in the Opposite Direction

Just as the vast majority of organizations are stuck in Pilot Purgatory, a small cluster of organizations are operating on a completely different logic:

- **DeepSeek:** 160 people built a trillion-parameter model, versus roughly 3,500 at OpenAI, 3,000 at Anthropic, and 8,100 at DeepMind;
- **Klarna:** full-time employees fell from 5,527 in December 2022 to 3,422 in December 2024 (-38%), while revenue kept growing over the same period and revenue per employee reached 3.6 times its 2022 level;
- **Shopify:** after cutting 20% of staff in 2023, headcount held flat for several consecutive quarters while revenue per employee rose from \$1.1 million to \$1.52 million;
- **Anthropic:** Claude writes more than 80% of merged production code, and per-engineer quarterly code output is 8 times the baseline;
- **Cursor:** about 300 people, annualized revenue from \$1 billion to \$4 billion, revenue per employee of \$3–13 million;
- **Duolingo:** zero full-time layoffs, content capacity up about 11 times in two years;
- **Airbnb:** AI co-writes the code produced by 60% of engineers, and AI customer support resolves 40% of issues without escalation to humans.

What these organizations share is not "using AI" — 88% of companies using AI^[1-03] have not achieved this. What they share is this: **AI is not a band-aid applied to a legacy organization; these organizations redesigned themselves around AI.**

1.4 The Scarcest Evidence of All: Causality

Everything so far is correlation and case studies. You could reasonably object: maybe these companies were simply stronger to begin with?

In 2026, researchers at INSEAD and Harvard Business School ran a rare randomized controlled experiment^[1-33] that pushed this question from correlation to causation.

They recruited 515 high-growth startups^[1-33] and assigned them at random: **every company** received equal API credits, access to frontier models, and technical training. The only difference: the treatment group additionally received case-study materials on how AI-native companies reorganize their production processes, teams, and business models, while the control group took only an ordinary entrepreneurship course.

With tools, skills, and budgets fully equalized, merely changing the "organizational search space" — letting the treatment group see what AI-native companies look like — produced:

- **+44%** more AI use cases discovered/used (concentrated especially in high-leverage areas such as strategy and product development);
- **+12%** more tasks completed; **+18%** likelihood of acquiring paying customers;
- Revenue **1.9×**^[1-33];
- External capital needs actually fell **39.5%**.

The paper introduces a core concept: the "**mapping problem**" — what limits AI's returns is not technology costs or skills, but the cognitive bottleneck of discovering where and how AI creates value within one's own production processes.

This is a conclusion worth chewing on repeatedly: **handing employees tools is not the same as reengineering processes; simply knowing what an AI-native organization looks like can itself produce a 1.9×**^[1-33] **difference in revenue.** That is why this book is worth reading, and why XDash's FDE book is worth reading — knowledge of organizational form is currently the highest-ROI knowledge there is.

1.5 This Chapter's Core Judgment

Put the three sets of facts together:

1. 95% of organizations get zero return from AI investment;
2. the 5% who succeed share one trait — they redesigned the organization around AI;
3. merely "knowing what the winners look like" produces a 1.9× revenue difference.

The conclusion is already clear: **models are no longer scarce; what is scarce are the people who can grow models into their organizations.** In the chapters that follow, we will see what such an organization looks like (Chapter 2), what the evidence says (Chapter 3), how it operates (Chapters 4-6), who has already done it (Chapter 7), and — how you can do it (Chapter 8).

Chapter 2: What Is an AI-Native Organization — Band-Aid AI vs. AI-Native DNA

2.1 The Core Proposition: You Are Not Building an AI-Native Organization — You Are Applying AI as a Band-Aid to a Legacy One

No matter how well you apply the band-aid, the bone underneath is still the old bone.

A company gives AI tools to every department, runs dozens of training sessions, and hands out incentives — and at the year-end review, adoption is below 15%. The tools were bought, the training was done, the incentives were paid, yet the organizational structure, decision processes, and evaluation criteria are all still traditional. AI is just a band-aid stuck onto a legacy organization.

This is the dividing line between "AI transformation" and "AI-native."

Alexander Puutio (Forbes) offers a sharp conceptual framework: **AI-native is, in essence, a reversal of the design order.**

- Traditional organization: design processes for humans → software assists and accelerates;
- Digital-native: design around software → humans still produce the primary value;
- **AI-native: first ask what AI can reliably do → then design human roles around AI.**

"An AI-native organization is not one that uses AI tools — it is one built around them."

This definition can test any "AI-native" claim. The criterion is not "how much AI is used," but: **if the organization were designed from scratch tomorrow, would AI be written into the blueprint?** If the answer is no, it is only a band-aid.

2.2 Three Common Traits

Drawing on the practices of Chinese companies including Alibaba, Huawei, Lenovo, Feishu, and Transn, AI-native organizations share three traits that recur again and again.

Trait 1: AI-informed decisions replace experience-based decisions. Traditional organizations decide based on Old Wang's experience and Director Zhang's gut. In an AI-native organization, as business data flows through the underlying systems, AI automatically performs analysis and early warning, and decision recommendations are pushed directly to the responsible person — decisions are based on data, not feelings. Huawei's approach is to embed AI across the entire data lifecycle, making intelligent analysis the default capability of its data platform: "You don't go looking for it; the system delivers it to you."

Trait 2: Business flow and workflow are unified. In traditional companies, the business flow lives in the CRM and the workflow lives in DingTalk or Feishu; project progress is in system A, meeting minutes in system B, approvals in system C — everything is carried across by hand. Feishu's approach is a collaboration suite that connects IM, documents, and business flow, with AI intervening in project milestones in real time rather than summarizing afterward.

Trait 3: Making expertise reproducible. In a traditional organization, when the top salesperson leaves, the expertise leaves with them. In an AI-native organization, AI automatically records decision logic and updates business rules, turning one person's expertise into a reusable organizational asset. Transn founder He Enpei puts it this way: "Rather than wait for employees to become AI experts, it is better to grow AI capability into the organization." Individuals come and go; organizational capability stays.

2.3 Historical Analogy: The Electricity Revolution of the 1880s

In the 1880s, factory owners bought electric motors to replace steam engines, but factory layouts, assembly line designs, and the division of labor did not change — and efficiency did not improve. Electricity only changed the power source; the mode of production was still steam-age. The real winners were the people who redesigned factory layouts, invented the assembly line, and gave every workstation its own power supply.

Today everyone is installing AI, but few are genuinely redesigning their own factories.

This analogy yields a second definition of the AI-native organization: **it is not a change of power source, but a redesign of the mode of production.** Going from horse-drawn carriage to automobile is not about bolting an engine onto a carriage — you have to reinvent the wheel, design the steering wheel, build the roads, and set the traffic rules. The entire ecosystem must be reconstructed.

2.4 A Misconception That Must Be Corrected: The CAIO

Chinese-language internet media widely repeat the claim: "Gartner predicts that by 2025, 90% of large enterprises will have a Chief AI Officer (CAIO)."

This is a misquote — there is no evidence for it. The verifiable original Gartner prediction (published in 2023) is that by 2025, 35% of large organizations will have a CAIO reporting to the CEO or COO. The figure "90%" has no corresponding source anywhere on Gartner's website.

The "CAIO trap," however, is real — the argument just needs to be corrected: **AI is not one director's job; it is the whole company's job.** If you create a Chief AI Officer position while leaving structure, decisions, and incentives untouched, the CAIO is a general with no army — the same structural trap as the traditional enterprises that created "Internet directors" in 2005: the director spent three years unable to move anything and finally quit. The Internet was never one director's job, and neither is AI.

2.5 Three Easily Confused Concepts

Puutio goes to great lengths to distinguish three concepts, because many companies use them to impersonate one another:

Concept	Core trait	Typical trap
Digital-native	Built around cloud + SaaS; humans remain the primary producers	Thinking that deploying Copilot makes you AI-native
Platform organization	Coordinates an ecosystem, optimizes connections	Thinking that having network effects makes you AI-native

Concept	Core trait	Typical trap
AI-native	Redesigns internal workflows around AI systems, optimizes production	Flattening without enough AI infrastructure = chaos

Plus the three categories of tools in an AI-native organization:

1. **Vendor-embedded AI services** (Copilot and the like) — the easiest entry point, but not enough to count as AI-native;
2. **External agentic AI systems** — standalone tools for cross-application, multi-step automation;
3. **Internally built tools** — built on proprietary data and processes, becoming a sustainable advantage.

2.6 Employee Roles and Strategic Principles

In an AI-native organization, the definition of employees' roles changes fundamentally:

"Employees no longer define their value primarily by the volume of work they personally produce... they create value by orchestrating systems that produce work."

Humans shift into orchestration and product-management roles; the most valuable employees understand both the boundaries of AI capabilities and the limitations of the systems; performance standards shift from "how much work did you do" to "designing better work."

Puutio states a clear and stern strategic principle: **"Don't involve people where people aren't needed."** This means:

- Investment decisions are made only around AI-driven ways of working;
- Tools and roles that exist only to manage human work rather than system performance should be questioned;
- Core processes are defined around tasks AI can execute reliably.

2.7 Building from Scratch vs. Retrofitting: A Real Debate

On whether "an AI-native organization can be transformed from a traditional one," there are two positions:

- **Puutio (strict-definition camp):** "An AI-native organization is necessarily a greenfield effort. Retrofitting rarely works."
- **Shmool (feasible-path camp):** A CEO who subtracted inside an existing organization, collapsing the traditional hierarchy (VP → director → manager → team lead) into a single-layer AI-native organization — successfully (details in Chapter 4).

A reconciliation of the two: Puutio is talking about the strict definition (what a true AI-native organization is), Shmool about a feasible path (how to get close to it). Retrofitting may achieve "quasi-AI-native" — close to, but not identical to, designing from scratch. The difference lies in the ability to "change the decision order": starting from zero allows you to re-ask "what can AI do" from the beginning; retrofitting means redefining while dismantling the old structure.

2.8 The Checklist: Are You Applying a Band-Aid or Growing the Gene

Ask yourself five questions:

1. **Design order:** Do you first ask "what can AI reliably do" and then design human roles, or do you have processes first and stuff AI in afterward?
2. **Decision basis:** Are decisions based on data pushed automatically plus human judgment, or on human experience with AI on the sidelines?
3. **Business flow and workflow:** Are the two systems unified, or are you still carrying data across by hand?
4. **Experience retention:** When the top salesperson leaves, does their expertise stay?
5. **Evaluation criteria:** Do reviews still look at hours worked and volume of output? (detailed in Chapter 6)

If all five answers point to "old structure," you are applying a band-aid. If even one or two have started to shift, you are already on the road to growing the gene. **An AI-native organization is not an "upgrade" but a "rebirth" — and it can be reborn in stages.** The next chapter looks at the evidence.

2.9 The Myth of AI Ubiquity — The Barbell Distribution and "Pushing AI into the Background"

Vasuman Moza (a former Meta engineer and founder of Varick Agents, which deploys AI agents for large enterprises with \$500M+ in annual revenue) spent a year embedded inside large companies and found that no matter whether a team has 50 people or 5,000, handing the best AI tools to every single person always ends in a "barbell":

5%-10% of people become power users, 20% use AI occasionally but crudely, and the remaining 70% barely use it at all. One large enterprise that had just signed an eight-figure-dollar annual AI contract even found that **10% of its employees consumed 90% of its tokens.**

This explains the paradox this book keeps returning to: McKinsey surveys show that 88% of organizations use AI in at least one business function, yet only 6%^[1-03] of enterprises earn more than 5% of their EBIT from AI. Model capability jumps every few months; organizational productivity does not automatically keep up.

Moza's more radical judgment is that this gap will not close by itself, and most employees will very likely never become AI-native users — because using AI well is a craft (knowing when to clear context, when to crystallize repetitive work into skills, when to hand judgment to the model, and when to write deterministic code). AI vendors' product roadmaps revolve around the most active frontier users, so the bar for "using AI well" keeps rising; the AI experts inside companies have no incentive to share — the gap itself is their advantage.

His strategic pivot: **first accept that most people will not become AI-native.** Let that 5%-10% of power users keep running ahead, and distill the skills, workflows, and agents they figure out into a shared library, where ranking is the reward. For the remaining majority, **push AI into the background** — embed agents into the systems employees use every day (Salesforce, NetSuite, Dynamics): invoices processed automatically, data entered automatically, processes flowing automatically, with humans called in to approve, reject, or modify only when the machine is unsure

— employees may not even realize they are "using AI." The metrics change accordingly: instead of asking "how many employees used AI this month," ask "of the work completed today, how much was done entirely by humans, how much by human-AI collaboration, and how much ran automatically."

This is structurally identical to Transn's judgment in Chapter 7: **rather than wait for employees to become AI experts, grow AI capability into the organization.** (Two caveats: the 5/20/70 distribution comes from the author's own client observations, not a large-sample benchmark; and embedding agents into enterprise workflows is precisely the service he sells.)

The same structural framing comes from Zilbix's *Enterprise Operating Models in the Agentic Era*: the past two years of "layering tools on top of existing structures" produced exactly the gap described above. The agentic era exposes three fractures — **unclear decision rights** (cross-functional agents without governance are either confined to low-value tasks or run unsupervised), **the talent model** (employee value shifts from executing tasks to orchestrating systems, interpreting outputs, and making judgments), and **the data architecture** (fragmented data blocks real-time access). An AI-native operating model is therefore three-dimensional: leadership actively sponsors the operational transformation (not merely approving budgets), decision rights with accountability and audit trails, and processes redesigned for AI participation (rather than retrofitted) — built on an AI-ready data foundation.

2.10 Buying AI Features vs. Building AI-Native Capability

A data point from HR technology puts this question on the table: more than 70% of Chinese enterprises claim to have deployed AI HR tools, but fewer than 15% have achieved an organizational capability leap^[8-72]. "Sticking an AI label on a legacy system" and "rebuilding the organization's memory hub and decision nerves" are two entirely different paths: the former is consumption — it compounds nothing; the latter is investment — the more it is used, the more it understands the enterprise.

AI-native capability can be measured with a three-layer architecture:

- **Organization-wide memory:** data assets shared across modules and scenarios — candidate tags, onboarding profiles, and talent reviews reused across the full chain, breaking down data silos;

- **Proactive initiative:** the system finds people instead of people finding the system — an agent judges "what should happen now" from real-time data (a candidate not interviewed for three days → reminder; two no-shows → lower intent score and recommend alternatives) rather than executing fixed scripts;
- **Continuous evolution:** every manual correction is written back into the underlying context asset layer, forming a data flywheel. The formula: **AI talent density × AI collaboration depth = organizational intelligence compounding.**

Why can't a legacy system leap forward by "adding an AI module"? Three barriers. First, **architecture debt:** recruiting, HR, and performance each have independent databases — a cross-database query needs 3 databases and 5 field mappings and takes more than 8 seconds, so AI can only be bolted on. Second, **interaction paradigm conflict:** menus-plus-forms coexisting with dialogue-plus-recommendations raises cognitive load instead of lowering it (one financial enterprise saw feature usage *drop* 18% after adding an AI assistant). Third, **the data-permission gap:** permissions are designed per module, so AI cannot read or write across modules — either rebuild the permission system (prohibitively expensive) or give AI low privileges (confining it to edge tasks like summarization and visualization).

How do you verify whether a system is truly "AI-native"? Four questions. Ask the AI, "Which department is Zhang from last year's campus recruiting in, what's his performance, and what internal position would he fit?" — if you must manually specify the data source, it's not integrated; a direct answer with sources is real integration. If a resume sits untouched for 48 hours, does the system react (push to the interviewer, remind, flag low intent)? After three months of continuous use, does recommendation accuracy improve with manual corrections (stuck at 40-50% = fixed algorithm; rising = learning)? Can you define enterprise rules in natural language ("tech roles must have a final interview with the CTO") instead of filing tickets and waiting for version releases?

The path scales by size: organizations of 200-500 people build their first data flywheel in recruiting and reuse the tags through onboarding and probation; 500-2000 build cross-scenario collaboration by connecting the recruiting-HR-talent chain on a shared data foundation; 2000+ sink enterprise context into an organizational capability platform that speaks "management dialect." The smaller you are, the more

you should start with a single point of breakthrough rather than rolling out every module at once. (The talent data cited here is vendor-reported and treated as official claims; see Appendix A.)

Chapter 3: The Evidence — What the Empirical Research Says

The first two chapters were about concepts and cases. This chapter looks only at evidence: what academic research and authoritative institutions' data actually say about AI-native organizations. To avoid "cherry-picking evidence that favors our argument," I will also list the limitations of every study.

3.1 HBS's *AI-Native Firms*: The Strongest Firm-Level Evidence

AI-Native Firms (HBS Working Paper 26-090, 2026)^[1-29], a working paper by Harvard Business School's Rembrand Koning and INSEAD's Hyunjin Kim and colleagues, is currently the hardest first-hand evidence at the firm level.

Method: All 2,891 YC startups from the W20-F24 cohorts (11 batches), plus 41,214 US VC-backed startups from PitchBook^[1-29], linked to micro-level labor data and compared against non-AI startups of the same industry and period.

Core findings:

Dimension	AI-native firms vs. non-AI peers
Team size	~ 25% smaller (YC firms roughly half the size after three years)
Share of engineers	~13 percentage points higher
Entry-level employees	~15% fewer
Senior employees	~20% more
Management layers	Half a level flatter
Managers	~15% fewer
Funding / valuation	Comparable → ~20% more funding per capita, higher valuation per capita

In one sentence: **AI-native companies are smaller, flatter, more engineer-dense, and more senior, with comparable valuations — and higher output per capita.**^[1-29]

The paper's most important theoretical contribution is distinguishing two channels:

- **The process channel:** AI as an internal production tool that changes how employees do their work (agentic coding, AI customer service, automated sales) — this is where most existing research focuses;
- **The product channel:** AI capabilities embedded into the product, so customers generate deliverables directly inside the product — **knowledge work moves out of the internal organization and into the product**, scaling through compute rather than headcount.

The decisive evidence: the process channel cannot predict smaller organizations. AI startups are 2.6 times^[1-29] more likely than non-AI peers to name ChatGPT/Copilot/Cursor in job postings — but that indicator cannot predict smaller teams. After controlling for the process channel, product-embedded AI remains significantly associated with fewer people and flatter structures.

The largest differences appear in service businesses: in service startups such as therapy, tutoring, and telemedicine, AI startups are about 70%^[1-29] smaller than non-AI peers. Cases: Gamma (AI presentations; roughly 30 people reached \$50 million in annual revenue within two years); FazeShift (an accounts-receivable AI agent; 10 people won dozens of enterprise clients, replacing the AR team at each client).

A cautious conclusion about the labor market: every AI-native company is smaller, but total employment need not fall — AI dramatically increases the number of startups (PitchBook's 2024 AI seed funding was about 8 times the 2020 average)^[1-29]. Extrapolating from the firm level to the market level requires caution.

Limitations: the sample is YC/VC-backed startups, which may not generalize to mature large enterprises; the "AI-native" classification is the researchers' own definition; and the working paper has not been peer-reviewed.

3.2 The INSEAD/HBS Field Experiment: Rare Causal Evidence

The randomized controlled trial of 515 startups mentioned in the previous chapter (*Mapping AI into Production*, 2026)^[1-33] is the scarcest kind of causal evidence in AI organization research: tools, skills, and budgets were fully equalized, and only the "organizational search space" was changed (the treatment group was shown how AI-native organizations restructure). The result: treatment-group use cases +44%, revenue 1.9x, external funding needs –39.5%.

The "mapping problem": what limits AI gains is not the cost of technology or skills, but the cognitive bottleneck of "discovering where and how AI creates value in your own production process."

Limitations: a 10-week short-term experiment; a startup sample; the 1.9x revenue figure is a short-term relative increment with small absolute values (on the order of, say, \$40,000).

3.3 WEF's *AI-First Operating Models*: A Systematic Framework for Operations

The World Economic Forum's (2026) core proposition: **layering AI on top of operating models designed for the pre-AI world (linear processes, static roles, incremental optimization) is a structural failure**. Global AI investment exceeded \$250 billion in 2024, yet in most organizations "AI has not fundamentally changed the way they operate."

Three shifts:

1. **Work and teams: from process management to human-AI collaboration** — early AI-first organizations already have **human-to-AI ratios exceeding 10:1** (1 human : 10+ AI); delegation is a management skill that requires training and certification (the You.com CEO: adoption in mature organizations only really took off after a training-and-certification program was launched);
2. **Workflows: from linear processes to dynamic AI systems** — the Rubrik CEO's implementation path: **"go business line by business line; in each line, find 3-5 workflows that are purely manual or SaaS-plus-human,**

and define end-to-end AI-driven outcomes"; the starting question determines everything (Lightspeed: "If we had unlimited intelligence, what would we build?");

3. **Metrics: from process optimization to dynamic value creation** — outcome-oriented, with reliability/accuracy/governance first to build trust.

Key judgment: **competitive advantage comes from orchestration, not experimentation**. Successful leaders treat capability, process, outcome, and business-model redesign as a single integrated evolution.

3.4 Tencent Research Institute's *AI-Native Work Report: A Complete Framework from a Chinese Perspective*

Tencent Research Institute (Chinese: 腾讯研究院) published a 40,000-character report in May 2026, with lead authors Si Xiao, Yuan Xiaohui, and Yu Yi (Chinese: 司晓、袁晓辉、余一). The most widely circulated part is its **competitiveness formula**:

$$\text{Organizational competitiveness} = \text{talent density} \times \text{AI leverage} \div \text{organizational friction}$$

If the numerator doubles but the denominator stays fixed, the net effect is discounted; **halving the denominator is equivalent to doubling the numerator** — reducing organizational friction (waiting/approvals/alignment/information decay) is usually easier to move than piling on talent.

The report defines four structural characteristics of the "Super-Individual" (all required):

1. **An AI-first workflow**: AI is the default starting point — "I let AI run first, then judge and correct on top of its output";
2. **An order-of-magnitude leap in capability boundaries**: 10x+ output; one person runs the full chain from idea to delivery;
3. **Extreme proactivity**: doesn't wait for organizational assignment; a natural boundary explorer;
4. **Influence spillover** (the key threshold for identifying Super-Individuals): **efficient individuals only make themselves faster; Super-Individuals make the team faster**.

It also outlines four emergence paths (bottom-up spontaneity / organizational selection and cultivation / atmosphere creation / founder-driven), three team forms (node-radiation / network collaboration / AI-hub), and one harsh judgment about the capability reshuffle: **AI is an accelerator of divergence, not an equalizer — only those who were already excellent become more excellent, and the gap widens.**

We must emphasize the report's sample bias (this is the knowledge base's systematic critique of the report): roughly 80% of its cases come from information-intensive, digital-native industries with fewer than 500 employees, with zero coverage of physical-world industries (manufacturing lines / medical clinical settings / logistics). The evidence in the report supports only this: **in information-intensive, small-team, digital-native industries, AI-native organizations can dramatically accelerate.** Before applying it to physical industries or large organizations, ask first: is this industry's bottleneck information processing, or the physical world?

3.5 Large-Sample Data from Authoritative Institutions

McKinsey, *State of AI 2025 (November 2025)*: 88% of organizations regularly use AI^[1-03] in at least **one business function** (78% last year); but about two-thirds are still in the experiment/pilot stage and only about one-third have begun scaling; only 39% attribute any degree of EBIT impact to AI, and most of them say the contribution is under 5%; "AI high performers" (EBIT impact $\geq 5\%$) account for only about 6%^[1-03]. **Many companies use AI; very few make money from it.**

BCG's *AI at Work* series: a clear leadership-frontline gap — more than three-quarters of leaders use GenAI multiple times a week, while frequent frontline use stalled at 51%^[1-04] in 2025 before jumping to 74%^[1-06] in 2026 (+23 percentage points). Half of companies are moving from "Deploy" to "Reshape" (redesigning workflows).

Microsoft's *Work Trend Index*: in 2024, 75% of knowledge workers already used AI^[1-13]; in 2025 it introduced the "Frontier Firm" concept^[1-14] (about 9.3% of surveyed leaders) — 71% of frontier-firm leaders say their company is "thriving," versus only 39% of employees globally; the most important finding of 2026: **organizational factors (culture, manager support, talent practices) contribute twice as much to AI impact as individual effort** (based on a permutation-importance analysis of 29 factors, $R^2 \approx 0.68-0.69$).

Stanford HAI's AI Index 2025: in 2024, 78% of organizations reported using AI^[1-16] (55% in 2023).

IMF: AI will affect about 40% of jobs globally; roughly 60% in advanced economies, 40% in emerging markets, and 26%^[1-20] in low-income countries — note that these are "exposure" estimates, not unemployment forecasts, and about half of exposed jobs may be augmented rather than replaced.

WEF's Future of Jobs 2025: by 2030, job disruption will equal 22% of current jobs: 170 million created, 92 million eliminated, a net gain of 78 million^[1-27]; 77% of employers plan to upskill their workforce, and 41% plan to reduce headcount because of AI automation — expectations, not facts.

JetBrains' 2026 Developer Ecosystem Survey (August 2026; 15,000+ professional developers): the adoption curve of AI coding agents is the hardest leading indicator on the engineering side — 90% of professional developers use AI coding agents weekly and 68% daily; Claude Code's workplace adoption rose from 18% in January to 39% (47% in the US), and 31% of developers use it as their only primary tool; GitHub Copilot fell from 29% to 21% and lost the top spot, while OpenAI Codex climbed from 3% to 16%^[1-36]. Tool rankings change, but the pattern does not: the adoption curve always arrives first where output is most measurable and feedback loops are shortest — the order in which AI penetrates an organization is the order of value measurability.

3.6 A Warning on Data Use (This Book's Citation Discipline)

Putting all the numbers above together, there are five citation disciplines:

1. **Definitions come first:** the "95% failure rate" (MIT, zero organizational return), "30% abandoned" (Gartner, after PoC), and "88% adoption" (McKinsey, at least one function) are not contradictory — they are three different denominators;
2. **Report conclusions vs. media retelling:** media retelling often distorts (e.g., Fortune reported the MIT report's sample as 150 interviews/350 questionnaires, while the report actually covered 52 organizations/153 executives)^[1-02]; this book always follows the original report;
3. **Forecasts vs. facts:** most of the Gartner/Forrester/WEF numbers are forward-looking forecasts or employer expectations, not events that have already happened;

4. **Sample bias:** McKinsey/WTI/BCG skew toward white-collar workers, large enterprises, and self-selected survey participants; every "global" claim must be checked against the original sampling scope;
5. **Conflicts of interest:** Microsoft, BCG, and McKinsey all sell AI products/consulting to enterprises — keep their positions in mind when citing their conclusions.

3.7 The Core Judgment of This Chapter

The evidence chain is complete: HBS shows AI-native companies are **smaller, flatter, and higher-output per capita**^[1-29] (driven by the product channel, not the process channel); INSEAD/HBS shows that **organizational knowledge itself produces causal revenue differences**^[1-33]; WEF shows that **operating models must be redesigned around AI**; Tencent shows that **competitiveness comes from talent density × AI leverage ÷ organizational friction**; McKinsey/BCG/Microsoft show that **a huge gap exists between tool adoption and organizational returns**.

The evidence points exactly where Chapter 1 concluded: models are no longer scarce — what's scarce is organizations redesigned around AI. In the next chapter, we look at what the inside of such an organization looks like.

Chapter 4: Rebuilding Organizational Structure — From Hierarchy to Network

4.1 The Legacy Organization's Deadlock

Linear division of labor has exposed a fundamental contradiction under AI acceleration:

The old way	The new contradiction
Managing people by role	The people who can actually get things done span multiple roles
Breaking work down by process	AI completes in a day what used to take a week of milestones
Accepting deliverables at checkpoints	Deliverables themselves become "intermediate artifacts"; only the final outcome is value
Reporting and evaluation by rank	Those who take on fused roles don't receive the corresponding authority and rewards

When one person can use AI to leap over steps that previously had to be handed off, and produce something closer to the final outcome, the legacy organizational framework becomes a constraint.

The core proposition of restructuring organizational form is: **use AI to replace the coordination and information-routing functions that hierarchy used to perform, so that humans can focus on judgment and creation.**

4.2 MarsWave: Five Practices of a Small Team Building from Zero

A small company called MarsWave (Chinese: 火星电波), which pivoted from ListenHub, completed its AI-native transformation in five weeks. Its five practices are the most concrete Chinese sample of "what an AI-native organization looks like":

1. Vibe coding shocks the organization. The CTO used Claude Code to build version 0.1 alone overnight — ideas could be presented directly. Before: product requirements → PRD → review → development → testing → release (weeks). Now: say it to AI → see the result → iterate (hours). Alignment shifted from "long documents" to "seeing the prototype directly."

2. The OPC model (human-machine-macro decision). Humans handle judgment (what to do, why, whether it's good), AI writes 99% of the code (how to write, how to implement), and managers make macro decisions (direction, resources, cadence). **Humans don't manage processes or schedules; they manage judgment, quality, and review.**

3. The birth of the Soul Team. The first core role that was neither engineer nor product manager — responsible for the product's **personality and emotional output**; its lead has a media background and can't code. Insight: once AI can handle "feature implementation," humans' differentiating value shifts to the "human-like" parts — empathy, personality, warmth.

4. A tool revolution. Linear, Notion, and kanban boards were abandoned; all code moved into a unified repository with no permission barriers. The repository "isn't for humans to read — it's for AI to read as context." The CEO's ideas go directly to AI, are automatically organized into Markdown, and become development context. **Essence: AI becomes the intermediary of information transmission, instead of layer-by-layer transmission between people.**

5. A new rhythm of front-load thinking and back-load review. In the first few days of each month the whole team stops writing code to align on plans; at month's end they re-review the code, cutting and refactoring. "Faster isn't necessarily better" — rhythm matters more than speed.

The failed experiment matters just as much: before the formal transformation, there was the "Task Tavern" (a task bulletin board where anyone could pick up any task) — a programmer's two-day new feature reached 80,000 users on launch day, but under business pressure the strongest people were pulled back to old tasks, and it ultimately failed. Lesson: **process improvement (patching the old framework) isn't enough — you have to design from scratch.**

4.3 Shmool: A CEO's Firsthand Subtraction in a Mature Organization

Unlike MarsWave (building from zero), Itay Shmool's case is about **subtraction within an existing organization** — dismantling the traditional hierarchy (VP → Director → Manager → Team Lead) into a single-layer AI-native organization. Core numbers:

Metric	Before	After
Decision latency	Days to weeks	Hours
Information fidelity	~60% (decays with each transfer)	~95% (direct transfer)
Builder-to-manager ratio	3:1	8:1+
Meeting overhead	~30% of work week	~10% of work week

Three key insights:


1. **The truth about middle management:** many middle managers are not "redundant layers" — they are excellent builders themselves, just pulled into coordination work by the old structure. Give them the chance to build again and they thrive. **Distinguish the management function from the managers themselves — the function can be automated, but people should not be simply removed.**
2. **Guild Masters replace management:** once management layers are removed, a new mechanism is needed to maintain professional standards. Guild Masters are **expert roles** (not management roles) who sit within teams and set standards across teams — they create no new reporting lines and are "the quality immune system of a flat organization."
3. **The AI platform is a precondition:** flattening without AI infrastructure equals chaos, not lean. Five core capabilities: self-healing systems, an AI creative engine, a unified experimentation platform, automated code review, and operations AI agents. **"The AI platform is what replaces coordination — this is non-negotiable."**

The human side of the transformation matters just as much: "Behind every box in the architecture is a person with a mortgage, a family, and a career. **How you treat the people who leave defines your culture more than how you treat the people who stay.**"

4.4 Block: Replacing Hierarchy with a World Model

Block, the payments company (Jack Dorsey), offers the large-company version of the framework: **use AI to replace the coordination functions hierarchy performed, and build the company as an agent.**

Block's four-layer architecture:



The revolutionary shift in the Intelligence Layer: from "product managers plan features" → the Intelligence Layer identifies timing and composes solutions; failure signals automatically become the roadmap; **product planning shifts from guess-driven to signal-driven.**

Three new human roles: Individual Contributors (ICs, given context by the World Model, making autonomous decisions), DRIs (Directly Responsible Individuals, short-term ownership of cross-domain problems, replacing project managers/product managers), and Player-Coaches (both building and developing people, replacing traditional managers).

The World Model handles alignment; DRIs handle strategy and priorities; Player-Coaches handle craft and people development.

A critique that must be heard (a critical analysis of the Block framework): Dorsey's premise is that "the main value of humans is information routing," but if humans are also generating, interpreting, and creating signals, then subtracting humans means subtracting intelligence itself; the real bottleneck is not routing speed but **perspective loss** (signals die before they reach decision points). Core warning: **"Speed is not intelligence. Speed without integration is just faster blindness."**

4.5 Yang Renbin: The Organizational Brain and the AI Chief of Staff

In a 2026 interview with LatePost (Chinese: 晚点), Yang Renbin (Chinese: 杨仁斌), founder of Jingzhunxue (Chinese: 精准学) — an AI education brand founded by Yang Renbin — supplied the piece all previous perspectives were missing: **how the highest cognition and decision-making capability can be systematically injected into every frontline execution step.**

- **First principles of the organization:** solving three traditional deadlocks — the top-down decay of "highest cognition → frontline," the bottom-up distortion of "frontline facts → top," and the bandwidth bottleneck at the top;
- **Three steps:** ① model the decision layer's judgment methods into a callable "AI Chief of Staff / Organizational Brain" (initially injecting 300,000 characters of methodology) → ② inject it into every business site through context engineering (a decision-model index system) → ③ frontline real problems/disagreements/anomalies flow back, closing the loop to iterate the highest cognition;
- **AI coding is only 0% and 100%:** a hard rule that engineers are not allowed to modify code themselves — humans only give feedback and transform into "context engineers," going further than MarsWave's "AI writes 99% of the code";
- **Anti-document, anti-group-chat business discussion:** documents are negative assets (author bias plus information loss); preserve the original records of discussions, disagreements, and failures; deep discussions continue in AI conversations; offline voice must be transcribed to text and enter the context;
- **Context is Everything:** once model capability is replicable, what is truly scarce is context engineering;

- **Cultural premise:** "explain the Why thoroughly, and give enough AI (爱, 'love,' pronounced ài, is a pun on AI)"; an open culture — hiding problems is shameful, exposing them means solving them sooner; don't hold people accountable for making mistakes, hold them accountable for hiding problems or failing to update their judgment.

4.6 An International Comparison: The Changing Role of the Engineer

Anthropic's internal practice confirms the direction of role reconstruction: the engineer's role shifts from "writing code" to "**architect + referee**" — AI-powered automated code review is embedded in CI/CD, and post-hoc analysis shows it intercepted about one-third of the production bugs in claude.ai's history that caused downtime^[7-28]; one engineer used AI to auto-fix 800+ API errors that humans estimated would take four years. At OpenAI, Codex became the primary work tool of every department (including non-technical ones such as legal, finance, and recruiting), accounting for 99.8% of the company's weekly output tokens^[7-38].

4.7 Three Paths: There Is No Single AI-Native Organization

Putting this chapter's cases together with the international ones, there are three parallel paths to the AI-native organization:

1. **The shrink path** (Klarna): hiring freeze + natural attrition, headcount –38%^[7-01], with AI replacing the headcount growth that demand would have required;
2. **The expansion path** (OpenAI/Anthropic): headcount and revenue grow together; AI-native is a growth engine, not a downsizing tool;
3. **The steady-state leverage path** (Shopify/Duolingo/Airbnb): full-time headcount flat or slightly up, with AI raising per-capita output.

There is no single answer for organizational structure, but organizational principles converge: AI handles information routing and execution; humans handle judgment, quality, and review; hierarchy grows thinner as coordination functions are replaced by AI platforms; experience shifts from a personal asset into an organizational asset. In the next chapter, we look at how workflows and decisions are concretely redesigned.

Chapter 5: Redesigning Workflows and Decisions

5.1 Redesign First, Automate Second

In 1990, Michael Hammer published a famous article in the *Harvard Business Review* titled "Don't Automate, Obliterate." His core point is even more applicable in the AI era: **before automating a bad process, ask why that process exists in the first place.**

The first principle of AI-native workflow redesign is: **not "give everyone a chatbot," but redesign how work flows, how decisions get made, how data is captured, how humans review risk, and what AI is allowed to assist with or act on.**

The fundamental difference between an automation-first organization and an AI-native organization:

Question	Automation-first organization	AI-native organization
Starting point	Where can we add AI?	What work should change?
Success metric	How many tasks were automated / how many prompts were used	Whether the outcomes of the whole workflow improved
Main risk	Faster rework	Discovering redesign gaps before scaling
Data approach	Use existing data	Redesign data capture around decisions
Human role	Post-hoc reviewer	Responsible designer, supervisor, and owner of exceptions
Scaling pattern	More tools and pilots	A repeatable workflow-redesign loop

5.2 Nine Key Workflow Redesigns

1. Start from business outcomes, not AI tools. Outcomes must be specific enough to guide trade-offs: "30% fewer supplier-invoice exceptions without adding approval risk" beats "use AI in finance." Technology choices follow workflow intent.

2. Map the real workflow before automating the official one. The official process diagram and how employees actually work (inbox shortcuts / private spreadsheets / informal escalation paths / exception handling) exist side by side as two versions. Delete steps before automating: merge duplicate approvals, turn low-value checks into rules, give exceptions an owner. **Don't teach AI to preserve unnecessary work.**

3. Redesign intake so AI gets better input. Many AI projects die at the door: expecting AI to classify, route, or decide on messy input that was never designed for it. Required fields, structured options, validation at capture points, automatic enrichment from trusted systems, and flagging low-confidence cases. **Don't blame the model for weak input before redesigning the front door.**

4. Separate rules / automation / AI assistance / agents / humans. Not every step needs AI: stable rules → rule automation; system handoffs → workflow automation; language-intensive review → AI assistance; pattern recognition → AI models; low-risk multi-step tasks → constrained agents; high-accountability decisions → human plus AI support. Prevent two errors: using AI where rules suffice, and delegating high-risk decisions just because AI can generate confident-sounding suggestions.

5. Build a workflow data foundation, not an abstract data lake. Ask "have we captured the data needed to make the workflow better?" rather than "do we have data?" Record when humans override AI suggestions and why — the best learning comes from the redesigned workflow itself.

6. Define decision rights before AI systems act. If humans aren't clear on decision rights, AI is even less safe. The decision rights table is the tool that operationalizes this:

Scenario	What AI may do	What humans retain
Classification / routing	AI acts on high confidence	Humans review low-confidence cases

Scenario	What AI may do	What humans retain
Drafting	AI recommends drafts	Humans approve sensitive content; template and tone control
Low-value refunds	AI acts within limits	Humans review over-limit cases; thresholds + logs + fraud checks
Contract changes	AI assists only	Humans decide
Risk escalation	AI acts immediately	Humans investigate and close

7. Replace manual handoffs with event-driven work and review queues. Stop treating people as glue between systems. Trigger → create record → enrich → apply rules → AI classify/summarize → route → exceptions enter a visible review queue (with context/priority/owner/deadline/escalation path). Don't make handoffs invisible; make important handoffs observable.

8. Measure the whole workflow, not AI tasks. Task-level time savings mislead (drafting got faster but review got slower; classification got faster but the queue got wrong). Metrics: cycle time, first-time-right rate, exception rate and exception age, review time, satisfaction, cost per completed case, rework rate, AI confidence and coverage, escalation quality, incidents/control violations. **After the pilot, ask: did the work improve, or did AI just make one step look faster?**

9. Create a repeatable 90-day workflow-redesign cadence.

Processes themselves can be redesigned: a review system is not the same thing as a PR. Amp, a 20-person team, pushes straight to main without pull requests and still passes SOC 2^[8-64] — auditors never cared about the git workflow; they care that changes are authorized, tested, approved, and recorded. Their control stack: push permissions restricted by business function, signed commits (GitHub verified), fully automated CI validation, and an audit trail linking commits to business context. The scaling insight: risk is not uniform inside an enterprise, and calibrating every system to "the scariest one" is hidden waste — first ask "what risk is the PR actually managing," then decide where it is needed.

5.3 The 90-Day Redesign Cadence (The Actionable Core)

- **Days 1-15: Select** — a process with visible pain + measurable value + sufficient volume + controllable risk;
- **Days 16-30: Map** — the real workflow / systems / data / decisions / handoffs / approvals / exceptions / owners / current performance;
- **Days 31-45: Redesign** — delete steps, clarify ownership, define decision rights, standardize intake, improve data capture, assign people / rules / automation / AI / agents to each step;
- **Days 46-60: Build** — workflow / integrations / prompts / access control / review queues / dashboards / logs / training;
- **Days 61-75: Pilot** — real users run real cases inside safe boundaries, measuring the whole workflow;
- **Days 76-90: Govern and scale** — kill / fix / scale, record the lessons, pick the next process.

5.4 Eight Common Mistakes

① Treating AI access as transformation; ② automating before simplifying (AI accelerates work of the wrong shape); ③ using AI where rules suffice; ④ ignoring exception work (the real value lies in handling messy cases); ⑤ measuring prompts instead of outcomes; ⑥ excluding employees from the redesign; ⑦ giving agents power without control (requires thresholds / logs / permissions / review queues / rollback / named owners); ⑧ redesigning the whole company at once (start with one process).

5.5 Dividing Labor Between Humans and AI: From "Full Replacement" to "Three-Layer Routing"

Beyond workflow redesign, designing the boundary of human-AI division of labor is the core of decision redesign.

Anthropic's internal evidence gives a key number: most employees believe that only 0-20% of work can be "fully delegated" to AI^[8-40]. In other words, the human-AI division of labor in an AI-native organization is "supervised collaboration," not "hands-off automation." The number of autonomous actions Claude

Code could take before human intervention was needed grew from about 10 to about 20 — this "intervention threshold" is dynamic and moves as model capability improves.

The hybrid-model **three-layer routing** that Klarna distilled from its customer-service reversal is currently the most mature template for dividing labor:

- **Layer 1: AI handles routine volume end to end** (targeting 60–70% of traffic) — order status, basic returns, FAQs;
- **Layer 2: AI-assisted human** (AI drafts, human reviews and sends, 20–25%);
- **Layer 3: Human-led high-value / escalation scenarios** (5–15% of traffic, but the greatest retention impact) — complex billing disputes, fraud, account cancellation.

Lesson: Klarna's first round of AI customer service only did Layer 1 (and did it aggressively); overall volume-based metrics (resolution rate, first-response time) looked good, but CSAT for complex interactions dropped significantly^[7-14] and the repeat-contact rate rose — the overall average masked a quality collapse in retention-critical scenarios. **You must track satisfaction for "complex/escalation interactions" separately.**

Treating AI as a compiler leads to frustration; treating it as a collaborator works. One engineer's observation is worth remembering: code gives certainty (same input, same output; a difference is a bug), humans give uncertainty (a colleague may understand intent and produce a better result), and AI sits between the two — rather than treating it as a compiler that precisely executes instructions, collaborate with it the way you would lead: share context, explain the desired outcome, set boundaries, respond to feedback. The core shift is investing in **expressing intent** — explaining why the work matters, what a good result looks like, and where judgment is needed. "The technology is new; leadership is not." This is also the underlying capability for the "human-AI division of labor design" in Chapter 8.

There is also an underappreciated ordering problem on the cost side: **choose the router first, then the model.** Most teams building agents pick the model first and the architecture second — backwards. Model choice should be the last decision, not the first. What actually determines cost and quality is the router: a small piece of code that decides which tier of model handles each request. Routing has three layers: a **skill classifier** turns the raw request into a concrete operation (intent recognition),

the **router** decides which tier executes it (reading only the classification label plus a few features: complexity, context size, historical success rate), and a **model selector** picks the cheapest model within the tier that clears a confidence threshold. The classifier is a language problem; the router is a scheduling problem — conflating them buries model choice inside the prompt and kills the ability to A/B different models on the same operation. Get routing right and 70-80% of traffic can run on local models or async models; Coinbase used "better defaults, routing, and caching" to halve AI spend^[8-80] while token usage grew exponentially. Ask one question first: how much work actually needs a real-time answer? Once you can queue work, the answer is often: very little.

5.6 Corroboration with Chinese Practice

The three common features from Chapter 2 (AI-informed decisions replacing experience-based decisions, unifying business flow and workflow, making expertise reproducible) all have concrete implementation mechanisms in this chapter:

- **AI-informed decisions:** Huawei (Chinese: 华为) makes intelligent analytics a default capability of its data platform — "you don't go query it; the system gives it to you automatically";
- **Unifying business flow and workflow:** Feishu (Chinese: 飞书; Lark internationally) connects IM, documents, and business flows, with AI intervening in project nodes in real time;
- **Making expertise reproducible:** Transn's (Chinese: 传神) AI automatically records decision logic and updates business rules, turning personal experience into organizational assets — at the same time, Yang Renbin's "decision model index" from Chapter 4 (modeling the decision layer's judgment methods as a callable system) and the decision rights table are two sides of the same coin: one is the governance side, the other the cognition side.
- **Governing AI coding:** Meituan (Chinese: 美团) reapplied its "eval agent" methodology to managing AI coding — first "human-human alignment" (one dictator beats ten democrats), then "human-AI alignment" (constraints written as AI Rules, with a human-AI agreement rate held at 90% or above; reverse the order and the AI Rules become empty words). AI code review uses higher-tier models to review lower-tier output, cross-vendor mutual review, plus a Pre-PR mechanism; technical debt is dissolved as "ride-along work" attached to

business requirements, with zero dedicated scheduling. This system keeps a codebase of 310,000 lines — about 90% AI-generated — maintainable^[8-83]. It also rewrites the definition of experience value: from "being able to see everything" to "knowing what matters."

5.7 The Core Judgment of This Chapter

Workflow redesign is the minimal executable unit of the AI-native organization. Organizational form (Chapter 4) is the skeleton; workflow redesign is the muscle — an organization doesn't need to finish "organizational restructuring" before becoming AI-native; it only needs to start with one workflow: map the real process, delete redundant steps, define decision rights, design for production, measure the whole flow, and iterate in 90-day cycles. **Reduce organizational friction (the denominator) first, then multiply talent and AI leverage (the numerator).** In the next chapter, we take on the hardest variable: people.

Chapter 6 Incentives and Talent: The Hardest Variable to Rebuild

Technical problems all have answers; people problems don't. This chapter examines the hardest variable in an AI-native organization: how to evaluate performance, how to incentivize, and how to develop people.

6.1 Transn's "Energy Gold": Turning Adoption from an Administrative Order into Market Behavior

Transn (Chinese: 传神), a translation company, runs an "Energy Gold" mechanism that is the Chinese benchmark for AI-native incentive design:

Every time a colleague genuinely uses an AI application and is satisfied with the result, the development team earns Energy Gold — a colleague-validated incentive credit. The more it is used and the higher the satisfaction, the greater the reward. **Judgment rights go to the users; decision rights go to the data.**

The supporting mechanisms are an AI Native decision committee (organized by business line — AI is not the technology department's business, it is each business line's own business) and the DEMO rule (every AI project must have a runnable DEMO and clearly explain which business problem it solves; no DEMO, no meeting — this kills 90% of PPT projects).

The elegance of "Energy Gold" is that it turns adoption from an "administrative order" into "market behavior": colleagues vote with their feet, creating a positive reinforcement loop of "the more it is used, the better it gets; the better it gets, the more it is used."

6.2 Evaluation: From Input to Output

The third of the three lightweight principles for small and mid-sized teams is the most counterintuitive and the most important: **evaluation criteria must change.**

Traditional evaluation looks at hours worked, volume of output, and process KPIs — which is effectively forcing employees to fake using AI. Someone who genuinely knows how to use AI produces more in 3 hours than a non-user does in 10. Evaluation must shift **from input to output, from process to results**.

If you don't touch evaluation, AI adoption is empty talk — you are touching the root of the entire management system.

International samples are changing evaluation too: Shopify includes AI usage in performance reviews, and its CEO memo requires employees to "first prove AI can't do this job" before applying for additional headcount^[7-19]; Duolingo's AI-first memo likewise included using AI to evaluate employee performance^[7-52] (though it later sparked controversy).

6.3 Tencent Research Institute: Reviews Give Way to Calibration and Four Dimensions of Incentives

The Tencent Research Institute report takes a more radical position: **"reviews" largely retire inside Super-Teams; management action shrinks from "control" to "calibration."**

When information is transparent enough and individuals are strong enough, management action contracts from control to calibration — the rest is left to people and agents to run on their own.

Super-Individuals do not lack capability; what is scarce is "why invest time here." The four incentive dimensions (all indispensable — an imbalance in any one leads to attrition or downgrade):

Dimension	Mechanism	Case
Autonomy (most cited)	You decide what to do and how	Block's DRI system — for 90 days you are the CEO of this problem. "Power incentives first, money incentives second"
Growth environment	High-density peers + sense of mission	Kimi: all 300 people see themselves as AGI builders; no OKRs, no KPIs, no clocking in

Dimension	Mechanism	Case
Economic return	How output is fairly priced (the hardest)	Three explorations: Anker Innovations, Tenex, Pure Global
Freedom to exit	Retain people by being worth staying for, not by lock-in	Netflix: options vest immediately with no lock-up period

Economic return is the dimension the report admits it cannot solve: "How can a 10–100x efficiency gap be reflected in a compensation system? No company has yet given a replicable answer." Traditional pay is tied to rank and role, but once AI flattens execution capability, the value of judgment and creativity rises sharply. "10x output rewarded with only 1.2x pay" → either leave and start a company, or reduce your effort — either way, the organization loses.

One reverse warning (ColaOS): **"organizations that make efficiency gains their core metric end up laying people off; organizations that make revenue growth their core metric end up expanding"** — the best incentive for a Super-Individual is "your high-output work opened up new business space, and you get to share in that growth."

6.4 Super-Individuals and the Competitiveness Formula

The Tencent Research Institute gives four structural characteristics of Super-Individuals (see Chapter 3): an AI-first work pattern, order-of-magnitude jumps in capability boundaries, strong proactivity, and influence spillover — **an efficient individual only makes themselves faster; a Super-Individual makes the team faster.**

The competitiveness formula: **organizational competitiveness = talent density × AI leverage ÷ organizational friction.**

The corollary for talent strategy: shrink the denominator (organizational friction) first, then multiply the numerator (talent density × AI leverage); halving the denominator is equivalent to doubling the numerator. Hiring structure follows the HBS evidence: a higher share of engineers, fewer entry-level roles, half the management layers — **invest the budget in senior all-rounders.**

The reshuffling of capability rankings: AI amplifies underlying capabilities (logic, rapid learning, problem decomposition, systems thinking) — only people who were already excellent become more excellent, and the gap widens. What is actually being repriced is the layer beneath skills: **judgment, learning speed, problem decomposition, taste** — while the existing talent system measures almost entirely the layer that is depreciating.

6.5 The Skills Gap: The More AI Is Used, the More "Hollowed-Out" Organizational Capability May Become

This is the most hidden long-term risk of the AI-native organization.

Anthropic's internal research (December 2025) offers two-sided evidence^[8-40]: engineers become "full-stack" and learn and iterate faster because of AI, **but at the same time worry about deep skills atrophy** — "when output is so easy and fast, actually taking the time to learn something becomes harder and harder"; 27% of AI-assisted work is exploratory work they would not otherwise have done (the good news), yet employees also worry that "one day AI will automate me out of a job."

WEF 2026: up to 75% of entry-level jobs in East Asia are disrupted by AI; up to 60% of entry-level tasks can already be taken over by AI^[8-38], and newcomers to the workforce lose their "practice ground" — cutting only entry-level roles severs the pipeline for talent development and innovation.

The essence of the skills gap is a choice between "treating AI as a crutch or as a trainer":

- Let AI fully take over (employees lose practice opportunities) → skills gap;
- Use AI to accelerate learning (AI as sparring partner, explainer, boundary-expander) → organizational capability compounds.

Anthropic has already made "how to prevent skill atrophy" a formal research topic. **The psychological cost of AI-native work is real** — internal employee communications mention "about 5 months without hand-writing code" and "feeling meaningless when everything is automated." Organizations need institutionalized "learning preservation" design.

Overseas research has now quantified the "crutch or trainer" question. The *Widening Gap* study cited by JetBrains frame-by-frame analyzed novice programming behavior: novices who leaned heavily on AI assistance "skipped crucial planning stages" and finished with an "illusion of competence" rather than true understanding; the best-performing novices were those who heavily limited or outright ignored AI assistance — they developed "negative expertise": the ability to identify and ignore incorrect or unhelpful AI suggestions. A University of Pennsylvania 2025 study of 1,000 students found that students learning math with an LLM performed 17% worse than students with only a textbook — while believing they were excelling; in the same study, a "GPT Tutor" group that used the model as a Socratic sparring partner (attempting problems first, asking for help when stuck) performed 127% better^[8-79] in practice sessions — **friction itself is a necessary condition for forming expertise**. Anthropic's 2026 study on how AI assistance affects coding skill formation reaches the same conclusion: "Cognitive effort — and even getting painfully stuck — is likely important for fostering mastery."

The other face of the skills gap is the **right to refuse**. An employee at an Australian sports organization resigned^[8-78] after her request to opt out of Copilot was denied in writing ("you are not entitled to ask that your work not be accessed by our AI systems"); surveys show roughly 35% of employees are uncomfortable with AI adoption, and an individual's exit neither stops the AI nor stops "employees training the AI that replaces them." An organization that designs only "how everyone uses AI" but not "who has the right to refuse, and what refusal costs" will see hollowing-out take another form: talent voting with their feet.

6.6 Innovation Density: The Organization Wins

DeepSeek: 160 people vs OpenAI: 3,500 / Anthropic: 3,000 / DeepMind: 8,100; Moonshot delivered a trillion-parameter model with 300 people and 1% of the industry's compute.

In the AI era, more people does not mean stronger — innovation density is the ultimate competitiveness, and innovation density is not built by stacking people; it is engineered through organization.

This fully corroborates the Tencent formula: high talent density × high AI leverage ÷ low organizational friction = 160 people beating 3,500.

6.7 Enable vs Replace: A Debate Already Settled by the Data

On the relationship between AI and people, there are two strategies:

- **Replacement automation:** positioning AI as a tool to "cut headcount, cut costs, boost efficiency" — breeds employee resistance, collapsing trust, tacit knowledge loss, and stalled innovation; it looks like saving money in the short term but destroys core organizational capability in the long term;
- **Augmentative collaboration:** AI handles the transactional layer while humans focus on judgment, creation, and relationships.

The data has already settled the debate:

- 2026 Gallup and PwC research: an **"enable employees" strategy can raise retention by about 32%**, with innovation capability and performance significantly ahead;
- Asana (2025): 64% of employees consider AI agents unreliable^[8-34] and call for more training, clarity, and guardrails — nearly half of employees are anxious about job loss (Accenture);
- Microsoft WTI 2026: **organizational factors (culture, manager support, talent practices) contribute 2x as much to AI impact as individual effort**^[7-71] — employees are not the bottleneck; the organization is;
- Business Week (2026): 55% of companies that laid people off because of AI regret it (Orgvue survey).

6.8 Role Fusion: Those Who Own Results Should Not Be Boxed in by Job Titles

The biggest obstacle to enterprise AI adoption is not technology but organizational inertia^[8-73]. Traditional organizations function on two premises: work can be fully divided into standardized nodes, and employees only need to be responsible for their process node while the organization backs the final result. AI changes exactly these premises — when document work, design, and parts of coding can be done by AI at lower cost and higher stability, the value of many roles is revalued and role boundaries blur: a product manager builds a demo with AI, a designer generates front-end pages directly, a developer joins product judgment with AI's help.

This is not one role replacing another — it is **Role Fusion**: the people who will be more valuable are no longer responsible for a single process node but can go from a user problem to full implementation and own the final result. The disruption follows: one person can span steps that previously required handoffs and produce something much closer to the outcome, yet the organization still manages people by job title, breaks tasks down by process, and accepts deliverables by node — "the people who actually get things done take on more responsibility but are still evaluated inside their old job boxes."

Role Fusion must be paired with results-oriented incentives, changing three dimensions at once: **authority** — give fused roles decision rights that span steps; **reward** — allocate by final results and user value rather than job level; **growth** — provide career paths that cross traditional functional boundaries. Otherwise the people who truly achieve Role Fusion will quickly feel the unfairness, and most will leave for teams that can recognize this kind of capability.

The hiring market already shows the corresponding gap: enterprises do not want people who can call APIs and write agent loops; they want people who can turn model capability into **stable services** — monitoring, evaluation, iteration, cost, latency, stability, and human review after launch each require engineering discipline; the real bar is turning a demo into a low-latency, observable, maintainable service that can enter enterprise processes. That is why the hybrid roles in demand are the **product engineer and the business-minded technical consultant** — they may not only know models, but they know how to turn a vague requirement into model inputs and outputs, translating back and forth between business and engineering.

6.9 The Core Judgment of This Chapter

The direction for restructuring incentives and talent is clear: **evaluation from input to output, incentives from position to results, management from control to calibration, talent from quantity to density, AI from replacement to enablement**. No single dimension can be solved by "just paying more" — but for every dimension, a set of organizations has provided reference answers. In Chapter 9 we will see: organizations that do not address the people problem are doomed to fail their transformation.

Chapter 7: Case Studies — AI-Native Organizations in China and Around the World

This chapter is the factual foundation of the book. Every case covers organizational form, on-the-ground practices, quantified results, and controversies and dissenting voices. **Official claims and third-party verified reporting are labeled separately** — because in the AI-native narrative, the tension between the promotional surface and the factual surface is itself the most important research material.

7.1 Chinese Cases

Transn (Chinese: 传神) — "Rather Than Wait for Employees to Become AI Experts, Grow AI Capability into the Organization"

Translation company Transn's AI-native practice is the most complete Chinese model available today. It runs on several key mechanisms:

- An AI Native decision committee organized by business line (Language Intelligence Group / Industry Intelligence Group / Brand & Sales Group), each with a clear lead and clear progress goals — AI is not the technology department's business; it is each business line's own business;
- A "DEMO rule": every AI project must have a runnable DEMO and clearly state which business problem it solves — this one rule killed 90% of PPT projects;
- An "AI joint fleet": 20+ cross-departmental business teams build their own AI applications, so that AI applications grow inside business scenarios from day one;
- Backed by the "Energy Gold" incentive mechanism — a colleague-validated incentive credit (detailed in Chapter 6).

Founder He Enpei's judgment: **"Rather than wait for employees to become AI experts, it is better to grow AI capability into the organization."** Individuals come and go; organizational capability stays.

Huawei / Alibaba / Feishu — Three Common Traits of Organizational AI-ification

- **Huawei:** AI embedded across the entire data lifecycle, with intelligent analysis becoming the default capability of the data platform ("You don't go looking for it; the system delivers it to you");
- **Feishu (Lark internationally):** a collaboration suite connecting IM, documents, and business flow, with AI intervening in project milestones in real time rather than summarizing afterward;
- **Alibaba:** AI-informed decisions replacing experience-based decisions — as data flows through, AI automatically analyzes, warns, and pushes decision recommendations.

DeepSeek — The Ultimate Example of Innovation Density

160 people built a trillion-parameter model. Compare with the scale of peers: OpenAI 3,500 people, Anthropic 3,000, DeepMind 8,100. Organizational design (rather than stacking headcount) determines innovation density — this is the strongest case for the "innovation density" thesis of Chapter 6.

Jingzhunxue (Chinese: 精准学) (Yang Renbin) — The Organizational Brain

Injecting the highest cognition into the front line: AI Chief of Staff + decision model index + context engineering. AI coding is only 0% or 100% (engineers are not allowed to modify code themselves; humans only provide feedback); anti-documentation, anti-group-chat; Context is Everything. (See Chapter 4 for details.)

MarsWave (Chinese: 火星电波) — Completing an AI-Native Transformation in 5 Weeks

The OPC model (human judgment / AI code / manager macro-decisions) + Soul Team (personality and emotional output) + a unified repository as the AI's context + a rhythm of front-load thinking and back-load review. (See Chapter 4 for details.)

Mininglamp Technology (Chinese: 明略科技) — Six Patterns of Enterprise Multi-Agent Collaboration

Mininglamp was one of the first Chinese vendors to productize "multi-agent collaboration," and its architecture guide published in 2026 offers the most complete enterprise-level multi-agent collaboration design framework to date, mutually corroborating Anthropic's multi-agent research in Chinese and English.

Three-layer infrastructure: a unified identity registry (every agent has a unique identity and permissions), a shared context mechanism (converging unstructured discussions in IM into structured knowledge: brief / timeline / deliverables / rejection records / acceptance status), and a task orchestration engine (split / serialize / parallel / compete).

Information topology controls visibility — "see each other when they should, stay blind when they should": a payroll agent should not broadcast its intermediate results to a weekly report agent. Group-chat-style information architectures can only achieve "everyone sees everything"; they cannot deliver programmable visibility boundaries.

Six collaboration patterns (nestable and combinable): Solo (working alone); Roundtable (mutual visibility — all agents share context); Critic (generate-review, with the reviewer holding veto power); Pipeline (strict serial — upstream output is downstream input); Split (divide-and-conquer in parallel — subtasks stay blind to each other and merge after each delivers); Swarm (competition selects the best — multiple candidates produce in parallel, redundancy traded for quality).

Implication for organizations: multi-agent is not "opening a few more windows" — you first decide the information topology and collaboration patterns, and only then talk about model capability. This aligns with the conclusion of the Anthropic multi-agent experiments in Chapter 9: **coordination does not naturally emerge from stronger intelligence.**

New data from mid-2026 continues to confirm this direction: IDC reports that 68% of leading enterprises put autonomous agents on their IT procurement priority list, yet more than 60% of pilots remain stuck at the production stage; a CAICT (信通院) survey of 358 enterprises found 67% reporting "coordination costs higher than expected." In July 2026, Beijing issued the first provincial-level agent policy, setting three bottom

lines — "traceable, auditable, controllable" — with up to CNY 100 million in industry support^[8-81]. Organization-level multi-agent collaboration is moving from vendor practice to a regulatory topic.

Lenovo (Chinese: 联想) — "Silicon-Carbon Symbiosis": the Super-Individual Is the Starting Point, the AI-Native Organization Is the Destination [official claims + third-party verified]

In June 2026 Lenovo upgraded its enterprise agent platform Lexiang (Chinese: 乐享) to version 4.0: more than 7 million monthly active users, over 1 million daily interaction requests, and cumulative sales pulled in excess of 5 billion yuan; the "agent corps" helped Lenovo China's direct-to-customer business grow 33% year over year^[7-73] [third-party verified]. On the organizational logic, Lenovo's HR leader proposed "silicon-carbon symbiosis" — silicon-based employees handle high-volume, standardized, repetitive tasks, while carbon-based employees focus on complex decisions, strategic design, and human care; the more widely quoted formulation comes from Zhu Li of Peking University's National School of Development: "the Super-Individual is the starting point; the AI-native organization is the destination." Lenovo is also systematically absorbing "AI natives" and spreading an AI mindset across all staff, pushing human-machine relations from "assistance" toward "co-creation." This path represents a typical choice for large manufacturing-plus-channel enterprises: rather than pursuing one-step organizational restructuring, first let super-individuals and the agent corps produce measurable results, then let those results drive the reorganization of division of labor and incentives.

More Chinese Practice

Jimo (Chinese: 极摩) (AI-native organizational architecture), Moka (AI-native HR: 3 AI colleagues), Atlassian's China practice (AI-native PM) — case details in Chapter 4 and the appendices.

7.2 International Cases

Klarna — The Most Complete Control Group for AI-Native Transformation

Klarna's AI transformation is the most important retrospective sample in the industry today.

Organizational form: full-time headcount fell from 5,527 in December 2022 to 3,422 in December 2024 (–38%) — the path was a hiring freeze plus natural attrition since 2023 (natural attrition of 15–20% per year), not a one-off mass layoff^[7-01] [IPO prospectus data].

On-the-ground practices: externally — an AI customer service built on OpenAI technology launched in February 2024, handling 2.3 million conversations in its first month, about two-thirds^[7-07] of total customer service chat volume; internally — AI assistant Kiki has answered more than 250,000 employee questions, used by 85% of employees, and the legal department drafts contracts with ChatGPT Enterprise in about 10 minutes instead of about 1 hour [official claims].

Quantified results: the two yardsticks must be placed side by side. Officially, Klarna claimed the AI customer service would bring \$40 million^[7-10] in profit improvement in 2024, and by Q3 2025 upgraded the claim to the equivalent of 853 full-time employees and \$60 million^[7-09] in cumulative savings, with revenue per employee of \$1.24 million (3.6 times the 2022 figure); third-party comparisons show that even counting the \$60 million in savings, Q3 2025 customer service and operations costs of \$50 million were still higher than the \$42 million a year earlier — "AI cost reduction" and "rising customer service costs" coexisting.

The controversy and reversal deserve the closest attention: after loudly publicizing "AI replacing 700 customer service agents" in 2024, Bloomberg reported in May 2025 that Klarna was **switching back to human agents**^[7-14] — the CEO admitted that going all-in on AI had degraded service quality, and planned to rehire human agents under "Uber-style" flexible employment. CSAT on complex interactions declined markedly and repeat contact rates rose; the overall averages masked a collapse in quality in retention-critical scenarios.

Lesson: AI-native does not equal a one-off mass layoff; purely cost-driven AI replacement has limits; "over-rotation" requires correction. 55% of companies that laid off employees due to AI regret it (Orgvue 2025).

Shopify — Layoffs First, Then AI-Powered Growth

Shopify's layoffs and AI-ification are two separate things, and order matters: 14% layoffs in July 2022, 20%^[7-17] in May 2023 (about 2,000 people; officially attributed to selling the logistics business and slowing growth — **not AI replacement; AI-ification came after the layoffs**). Headcount then stayed flat for multiple consecutive quarters: about 8,300 at the end of 2023 → about 7,600 in 2025, with revenue per employee rising from \$1.1 million to \$1.52 million^[7-24].

April 2025 CEO Tobi Lutke memo^[7-19]: employees must first prove "AI can't do this job" before applying for additional headcount; AI usage is included in performance reviews. More than half of merchant interactions with Support are AI-assisted and often fully resolved by AI. Internal productivity platform Shopify OS aggregates business data and recommends the resources and skill configurations a project needs.

Controversy: the memo was read as a "layoff filter" rather than a productivity plan. "Layoffs before AI" versus "AI-driven layoffs" is the key distinction for understanding Shopify.

Anthropic — Eating Its Own Dog Food First

Organizational form: about 2,500 people (2026)^[7-26], no mass layoffs — an "AI-native expansion" organization.

On-the-ground practices (official report *Recursive Self-Improvement*): more than 80% of merged production code is written by Claude (May 2026)^[7-28]; code output per engineer per quarter is **8x** the 2021–2025 baseline; engineers' roles shifted from "writing code" to "architect + judge" — AI automated code review embedded in CI/CD has intercepted about one-third of downtime-class production bugs in claude.ai's history.

Dissenting voices (also from internal research): the quality of AI-written code was objectively below human-written code until late 2025, reaching only "roughly on par" by mid-2026; employees self-report that 60% of their work uses Claude with 50%^[8-40] productivity gains, but 27% of AI-assisted work was "things they wouldn't have done anyway" (exploratory work), and most employees believe only 0–20% of their work can be "fully delegated" to AI; some employees went "about five months without writing code by hand" and worry about skill atrophy and a sense of meaninglessness — the psychological cost of being AI-native is real.

OpenAI — The Highest-Resolution Sample of Whole-Organization Agent-ization

On-the-ground practices (official economic research *How Agents Are Transforming Work*): Codex has become the primary work tool for every department (including legal, finance, and recruiting), accounting for **99.8%**^[7-38] of the company's weekly output tokens; the average engineer runs 99% of output tokens through Codex; the heaviest users orchestrate 60+ hours of agent work time in a single day (multi-agent in parallel); 80.6% of sampled users have had Codex execute tasks estimated to require more than 30 minutes of human work.

The most important counter-evidence comes from Fortune (August 2026): OpenAI's own economic research shows "**no measurable correlation between AI usage and revenue per employee**"^[7-39] — the spread of AI-native tools does not guarantee that productivity will be realized. Even the maker of AI tools itself cannot produce evidence that "the more you use, the more you earn."

Latest development (Wired, August 2026): OpenAI appointed Codex lead Thibault Sottiaux to take over the entire ChatGPT product line^[8-70]; ChatGPT will transform into a "personalized agent super-app" driven by Codex at its core (a 40-person team, folded into ChatGPT within weeks). Its self-attribution is worth noting: the earlier failures of Operator/ChatGPT Agent were attributed to "too early, models unreliable" plus users not knowing what to do with agents; the new strategy is small, fast releases — "you can't do a big splash and then get it wrong." **Even the most aggressive agent-ization sample is hedging against the uncertainty of agent products with small, fast releases.**

Notion — A Thousand People, No Layoffs, AI as the Second Growth Curve

About 1,000 people, no layoff record. ARR surpassed \$600 million in December 2025, **roughly half of it from AI products**^[7-41]; paid AI add-on penetration rose from 10–20% in 2024 to over 50% by September 2025. Notion AI^[7-42] launched in November 2022 (about two weeks before ChatGPT's public launch). Its path is product-driven rather than organization-slimming — driving paid penetration through AI features rather than through layoffs.

Cursor — King of Revenue per Employee

About 300 people: ARR \$100 million in January 2025 → \$500 million in June → \$1 billion^[7-45] in November → about \$4 billion^[7-48] by June 2026; revenue per employee of \$3–13 million. In June 2026, SpaceX announced an all-stock acquisition at a \$60 billion valuation (completed in August, folded into SpaceXAI).

Controversy: in April 2025, AI customer service program "Sam" fabricated and executed a nonexistent login policy, causing users to cancel their subscriptions — a textbook incident of "an AI agent acting beyond its authority."

Duolingo — Zero Full-Time Layoffs and 11x Content Production

In January 2024, Duolingo cut about 10% of its contractors (outsourced translation and content production)^[7-50] and switched to generative AI for content; the CEO said the company has **never laid off a single full-time employee** since its founding in 2009. Content production: 1,800 course skills per quarter in 2024 → 7,100 in 2025 → 20,500 in Q1 2026 alone — about 11x^[7-55] in two years, which the company explicitly says "wouldn't have been possible without AI."

Controversy: the contractor cuts triggered a PR crisis (large-scale TikTok/Reddit discussion of "replacing workers with AI")^[7-56]; after the April 2025 AI-first memo sparked layoff speculation, the CEO publicly clarified, and in May 2026 the memo was reported to have been retracted^[7-53]. **The first cut of AI replacement lands on the most vulnerable flexible workers — a pattern that holds across cases.**

Airbnb — 60% of Code Co-Written by AI

No AI-related layoffs; Q1 2026 earnings call: **60% of engineer-produced code is co-written with AI**^[7-57] (Google claims 30%+, Microsoft up to 30%^[7-61] — Airbnb is at the high end of public claims); the AI customer service bot handled 40% of customer service tickets without human escalation in Q1 2026 (from about 33% in early 2025), with plans to expand to 50+ languages.

Dissenting voice (the CEO's own admission): Chesky publicly acknowledged that "no one has truly solved the AI travel problem," listing the chatbot's four major flaws (text overload, inability to operate directly, difficulty comparing, and mismatch in multi-person booking scenarios).

7.3 Cross-Case Comparison: Six Patterns

1. **Three organizational paths coexist:** downsizing (Klarna), expansion (OpenAI/Anthropic), steady-state leverage (Shopify/Duolingo/Airbnb) — there is no single answer;
2. **The "first cut" of AI replacement lands on flexible labor:** Duolingo's contractors, Klarna's outsourced customer service — full-time layoffs are actually rare;
3. **Evaluation systems are being reshaped by AI:** Shopify (AI in performance reviews), Duolingo (AI-first memo);
4. **The tension between claims and verifiable data is the most valuable part of this book:** Klarna (\$60 million in savings coexisting with rising customer service costs), OpenAI (its own research: no correlation between AI usage and revenue per employee), Anthropic (AI code quality was once below human quality) — the promotional surface and the factual surface must be placed side by side;
5. **Reversals and backtracking are the norm:** Klarna switching back to human agents, Duolingo retracting its memo, Cursor's Sam incident — no company's AI-native transformation is a straight line;
6. **Extremes in revenue per employee:** Cursor (300 people, \$4 billion ARR) is the ceiling sample for small teams; Klarna (\$1.24 million per employee) is the sample for large-organization transformation — the shared financial trait of AI-native organizations is per-capita output well above peers.

Chapter 8 Building Your AI-Native Organization: An Action Playbook

The previous chapters answered "what" and "why"; this chapter answers "how." We integrate the transformation frameworks of Anthropic, Rubrik, BCG, McKinsey, and Accenture with Chinese practical methodology, converging into an executable four-stage roadmap.

Overall principle: first reduce organizational friction (the denominator), then multiply talent density × AI leverage (the numerator); first rebuild workflows, then talk about automation; start with one workflow rather than restructuring the whole company at once.

Stage 0: Decision and Preparation (2-4 Weeks)

Goal: answer "why are we transforming, who leads it, where is the floor" — no code, no tools purchased.

#	Key Actions
0.1	Executives hands-on: CXOs trial AI tools in at least 3 core business scenarios to build first-hand understanding (Accenture research found that only 12% of leaders consider themselves able to iterate quickly with limited information — hands-on use is the only cure for this)
0.2	Form an AI steering committee (business line leads + CFO + CTO + HR), clarifying a responsibility structure of "business-led, technology-supported, finance-accounted"
0.3	Define the value standard: only four categories of quantifiable value are recognized (revenue growth / cost reduction / efficiency gains / quality improvement); projects without a business owner and quantifiable metrics are not approved
0.4	Data and compliance check-up: AI-ready data inventory, Shadow AI audit (tool list + exposure surface), vendor lock-in assessment (multi-model routing contingency)

Common mistakes: outsourcing Stage 0 to the IT/data department to run alone (technology charging ahead solo); treating "buying Copilot" as the transformation itself; skipping the data inventory and going straight to pilots.

Stage 1: Assessment and Selection (4-8 Weeks)

Goal: identify "few but excellent" high-value workflows, decide Build vs Buy, and establish measurement baselines.

#	Key Actions
1.1	Workflow inventory: list all workflows per business line, mark "frequency × pain intensity × value measurability," and screen 3-5 highest-value workflows per business line (the Rubrik method)
1.2	Prioritization: prioritize a single value goal (McKinsey: projects focused on a single goal succeed at 3.2x the rate of broad-goal projects); prioritize back-office / mid-office functions (MIT: compliance and operations have the highest success rates)
1.3	Build vs Buy: MIT evidence shows that external procurement/partnership success rates are about 2x those of internal build — default to Buy/partnership, and build in-house only for data sovereignty or deep-customization scenarios
1.4	Baseline measurement: record current cost/cycle/quality numbers for each candidate workflow (with the CFO involved)
1.5	Governance up front: data permission matrix, output auditing, multi-model routing, and guardrails built at the design stage rather than after launch

Common mistakes: launching dozens of pilots at once (pilot proliferation → locally optimal, globally worst); choosing scenarios that "leaders think are cool" rather than "business pain points"; governance after the fact.

Stage 2: Pilot (30-90 Days)

Goal: validate real value on 5-15% of traffic / the smallest business unit, forming a replicable model.

#	Key Actions
2.1	

#	Key Actions
	30-60 day value pilot: single pain point, short cycle, preset metrics (adoption / efficiency / quality / satisfaction, four dimensions); human fallback during the pilot, but record the fallback workload — it is the real cost of productionization
2.2	Small-traffic ramp-up: start with 5-15% of traffic, monitor daily, tune quickly; small-step pilots succeed at about 4x the rate of full rollouts
2.3	Human-AI division of labor design: clarify the three-layer routing of "AI end-to-end / AI-assisted human / human-led" (Klarna's hybrid model: 60-70% / 20-25% / 5-15%); high-value complex interactions default to human-led
2.4	Design for production from day one: shared context layer, standardized action layer, governance embedded in the system — rather than "bolting on productionization after pilot success" (these are the three decisions of the 5% who succeed)
2.5	Value accounting: after launch, account using real business data (CFO-led validation succeeds at a 76% rate vs 53% for technology departments vs 32% for business departments)

Common mistakes: pilots that perform through "expert fine-tuning + human fallback" (such a pilot is not a product, it is a performance); pilot success without value accounting; human fallback workload hidden during the pilot and costs underestimated. **A subtle one: letting the agent change production code and tests at the same time — the tests lose their evidentiary power (the agent can redefine "correct" in the same breath); production and tests must be changed in separate phases** (a practical lesson from Yadda 3.0).

A product manager's 8-week permission ladder [third-party / hands-on debrief]: one product manager spent 8 weeks and about 2,200 yuan a month moving an operations team from skepticism to dependence. The three-level leap is a permission design, not a feature design: AI Intern (read-only queries, opened to everyone to build trust) → AI Advisor (recommendation-only, piloted with 10-15 people to validate value) → AI Expert (execution-capable, opened only after maturity and kept under human confirmation)^[8-84]; three safeguards accompany it — device whitelists, daily quota limits, and a circuit breaker after three consecutive failures; and the trade-off between "automatic execution vs. human confirmation" must be confirmed in writing by the business unit, not decided unilaterally by the technology

team. Outcome metrics from the debrief: recommendation adoption above 60%; a single query dropping from 3 minutes to 8 seconds; and on-site troubleshooting roughly tripling problem-resolution speed. The essence of the permission ladder is trading trust for autonomy — every level of permission is unlocked by the adoption data of the level below.

Stage 3: Scale (3-12 Months)

Goal: replicate the validated model to the remaining scenarios in the same business line, other business lines, and other regions; accompany it with organizational change.

#	Key Actions
3.1	Standardized reusable components: models, data rules, processes, permissions, and operations all modularized (portable = scalable)
3.2	Process redesign before tool stacking: delete redundant steps, merge duplicate tasks, reset the human-AI division of labor; companies that reengineer processes convert about 5x more AI value than those doing patchwork fixes
3.3	Systematic training: structured training (rather than issuing licenses); include AI usage in reviews and promotions; leaders lead by example
3.4	Incentives and trust: clearly communicate that "AI liberates employees rather than replacing them" (see Transn's Energy Gold); enablement strategies can raise retention by about 32%
3.5	Organizational design keeps pace: as the number of agents rises, establish agent operations / admission / accountability roles; monitor agent behavior, and ensure rollback and auditability (98% of enterprises have experienced disruptive events related to AI agents)
3.6	Cost governance: monitor on a "cost per resolved case / full-cycle cost" basis (rather than per-inference cost); cost is the #1 reason 40%+ of agentic projects get cancelled
3.7	Trust-building and observability: delegate authority in tiers using a "concentric-circle" model — inner circle: agents self-verify and auto-merge; middle circle: pre-annotated human review with observability context; outer circle: deep human intervention. Detect agent drift via tool-call trajectories (e.g., 15 calls without

#	Key Actions
	converging = stuck) and SLO burn-rate alerts; black-box tools that expose no telemetry should be vetoed in vendor selection; bounded tasks with frequent control hand-backs are the best practice (Honeycomb)

Common mistakes: scaling tools without scaling processes (technology does addition while processes stand still); no supporting training and incentives (the cause of 64% of Copilot licenses sitting idle); going full release the moment promotion starts.

Stage 4: Institutionalize (12-24 Months, Ongoing)

Goal: make AI part of the organization's operating system — governance institutionalized, skills assetized, innovation mechanized.

#	Key Actions
4.1	Institutionalize the kill mechanism: projects that fail value metrics, cannot scale, or are only fit for demos are forcibly stopped (sunk cost obsession is an accomplice to the pilot trap)
4.2	Full-cycle value disclosure: pilot → production → post-evaluation → formal disclosure; AI value enters management analysis (putting an end to the "efficiency myth")
4.3	Skills as assets: structure senior employees' tacit knowledge (SOPs / cases / data annotation) to prevent the double kill of "skills gap + tacit knowledge loss"; design "AI sparring-partner" learning paths to counter skill atrophy
4.4	Mechanize innovation: expand from the single goal of "cost reduction and efficiency gains" to BCG's Reshape/Invent (reshaping business models with AI, inventing new businesses), avoiding "efficiency thinking locking out innovation space"
4.5	Dynamic reassessment: model capabilities leap every 6–12 months (Anthropic internally: agents' autonomous actions went from 10 to 20), so the human-AI division of labor boundary needs regular renegotiation

Common mistakes: treating "launch" as the finish line (no post-evaluation, no kill switch); permanently locking AI into the efficiency-tool layer; ignoring the division-of-labor boundary drift caused by model iteration.

Supplement: BCG's Three Value Plays and Accenture's Three Leadership Principles

BCG: Deploy / Reshape / Invent. Deploy (apply off-the-shelf AI to existing processes — fast, but with shallow moats) → Reshape (redesign how work and business models operate — the main force in the medium term) → Invent (create entirely new AI-native businesses — the highest long-term value). **The three plays are not a multiple-choice question; they are a progressive combination.**

Accenture: curiosity / courage / connection. Curiosity — test AI yourself, hands-on, rather than delegating to the technology team; Courage — act with incomplete information, dare to decide "where not to deploy AI," and set up governance in advance; Connection — listen first (to employee anxiety), then use narrative to connect change with meaning, and build cross-functional alliances among executives early.

Common Mistakes Quick-Reference Table (Across the Full Cycle)

#	Mistake	Counter-Evidence	Right Approach
1	Buying tools = transformation	64% of Copilot licenses idle	Training + process redesign + incentives, the three-piece set
2	Layoffs first = cost reduction	Klarna's reversal, 55% regret rate	Augmentative collaboration (AI empowers people)
3	The more pilots the better	Pilot proliferation → globally worst	3-5 high-value workflows per business line
4	Pilots propped up by manual patches	Pilot succeeds, production collapses	Design for production: context / action layer / governance
5	Only look at inference cost	Cost per resolved case can exceed offshore labor	Full-cycle cost basis + CFO accounting
6	Run AI on ungoverned data	60% of projects abandoned over data	Stage 0 data check-up first
7	Single-vendor lock-in		

#	Mistake	Counter-Evidence	Right Approach
		Switching costs 19–34%, price increases 20–40%	Multi-model routing + portable context
8	Employee fear unaddressed	64% don't trust AI agents, nearly half anxious	Leaders listen + narrative + enablement strategy
9	No kill mechanism	Sunk costs forced through, waste amplified	Institutionalize the value gate
10	Only Deploy, no Reshape/Invent	Efficiency thinking locks out innovation space	Advance the three value plays in combination

Three Lightweight Principles for Small and Mid-Sized Teams

If you are a small team (10–50 people), you don't need the full four stages:

1. **Find the "AI nodes"**: not every role suits AI adoption. Look for steps in the business that are **high-frequency, repetitive, and still require a degree of judgment**; get three to five key nodes working first so the team tastes success — the team itself will go find the next node;
2. **Give the front line tool choice**: give everyone tool choice and an AI tool budget; good practices spread naturally, which is more effective than top-down promotion;
3. **Evaluation criteria must change**: from input to output, from process to results — without changing evaluation, AI adoption is empty talk.

Small Teams Delivering Enterprise AI Services: From First Deal to Ongoing Operation

A small team that sells AI services to other enterprises can treat its own delivery as a "mini AI-native organization" rehearsal — a 3–5 person team running an enterprise AI first deal from lead to renewal is itself a playbook worth copying.

Sell results, not features. Don't write "integrate LLM + knowledge base" in the project charter; write the deliverable boundary: which type of work, which people, which metric goes down. The customer is not buying "which large model was connected" but a piece of work that used to require repeated human effort, done faster, more reliably, or cheaper under controlled risk.

3-5 people cover four responsibility areas — headcount is not the point; someone owning each responsibility is. Business and delivery lead (client screening, requirement interviews, scope confirmation, translating vague needs into executable, acceptable tasks); AI/full-stack engineering lead (models, retrieval, agents, workflows, deployment); business analysis and adoption lead (organizing process knowledge and rules, building test questions, training, collecting real user feedback — the role most often cut, and cutting it means the system runs but nobody uses it); and security, data, and domain expertise (can be part-time or via partners, but personnel, finance, healthcare, legal, production-control, and sensitive-data decisions must be made by professionals).

Acquisition: don't look for customers who "want to do AI." The five signals of a qualified lead: high-frequency repetitive work that has become a burden; a business owner attends (not just IT); pilot-able data and materials exist; real users are willing to test; and there is an initial definition of success with willingness to commit resources (the most important one).

A four-tier service ladder — never start with a free POC. A free POC means unlimited scope creep with no budget, no acceptance, and no end date. Paid diagnosis (a half-day to one-day process diagnosis the customer keeps even if they don't proceed) → controlled pilot (validate only the key assumptions: one product line, one department, one set of approved documents) → production deployment (add permissions, logging, monitoring, integration, training, operations) → ongoing operation (content updates, evaluation regression, cost optimization, incident handling, new scenarios).

Freeze the baseline before promising savings. Without a baseline, promising cost reduction is unprovable. Four baselines: business (process steps / volume / duration / error rate / satisfaction), data (volume / format / version identification / sensitivity level / content owner), technical (environment / permissions / interfaces / logging / model licensing), and evaluation (a real question set plus current human

results). Ship the first "evaluable version" in two weeks rather than a complete one, and re-run the fixed test set on every model, chunking, prompt, or workflow change — an AI project without an evaluation set is judged by demo luck.

Three tables keep things from running out of control. Weekly report (surfaces blockers and decisions, not effort), risk register, and change request log (verbal "just add it while you're at it" never enters development). Agree acceptance metrics before development starts; every metric answers four questions: what is measured, with what data, by whom, and what counts as passing. Put AI inside the employee's existing work entrance (Feishu, WeCom, CRM, ticketing) rather than asking them to open another website they will forget.

Turn one project into a repeatable service. Each project accumulates five asset classes: an industry problem bank, diagnosis and proposal templates, technical components, a desensitized acceptance dataset, and an operations manual — forming a "delivery foundation" that can be adapted to new customers quickly. Don't wait until contract expiry to talk renewal — monthly reviews report which jobs users use AI in, which questions still fail, whether manual handling has decreased, and how business metrics moved, not call counts. When sustained value holds, renewal is a natural result.

The ten most common traps for small teams are worth pinning on the wall: accepting any customer (scattered industries → nothing reusable); free POCs without end conditions; promising outcomes by model capability; pre-sales promises never written into the SOW; talking only to IT, not the business owner; promising savings without a baseline; substituting demo questions for a real test set; handing every problem to the prompt (data, permission, and process problems dressed up as prompt problems); not daring to reject scope changes; and doing no asset accumulation after delivery.

The Core Judgment of This Chapter

Transformation doesn't need a perfect start; it needs the right first step. Pick one workflow, one owner, one quantifiable outcome, design for production, iterate for 90 days — then let success speak for itself. In Chapter 9, we will see what those who didn't do this paid for it.

Chapter 9: Challenges, Risks, and the Future

9.1 Failure Cases: Transformations Done Wrong, and Ones Never Got Around To

Klarna's reversal (detailed in Chapter 7) is the most complete sample of a "transformation done wrong": AI customer service loudly replaced 700 people → quality collapsed in complex interactions → quietly switched back to human agents. Five lessons:

1. **Blind spots in the business case:** only labor-cost savings were counted; revenue loss, repeat-contact costs, and the "cancellation cost" (after publicly announcing that the work was automated, strong candidates no longer wanted to join) were not;
2. **Experiential knowledge cannot be rebuilt:** the tacit knowledge senior agents used to handle edge cases disappeared permanently when the teams were cut;
3. **Metric choice determines success or failure:** overall averages masked the quality collapse in retention-critical scenarios;
4. **A viable hybrid model:** three-layer routing (60–70% / 20–25% / 5–15%);
5. **Industry-wide contagion:** fintech, healthcare support, and insurance claims all show the same pattern of "full automation → quality degradation → quietly backfilling with humans."

Samsung's Shadow AI incident: in April 2023, three incidents in 20 days of employees uploading sensitive information (semiconductor equipment measurement data, source code) to ChatGPT led Samsung to ban generative AI outright and build its own alternative. **A "throw-the-baby-out-with-the-bathwater" ban sets the transformation back; the right posture is to provide a managed, enterprise-grade AI environment, offsetting Shadow AI with alternatives rather than bans.**

Microsoft Copilot's licensing dilemma: about 64% of enterprise Copilot licenses sit unused, and only 3.4% of M365 customers pay for advanced AI features; 72% of IT leaders say employees struggle to integrate Copilot into their daily

workflows. **Between the purchase decision (made by the CIO) and the usage decision (made by frontline employees) stand three gates of organizational inertia — whether processes have been redesigned, whether leaders model the behavior, and whether performance reviews are tied to it: procurement ≠ adoption, and licenses ≠ value.**

Chegg's cost of not transforming: in May 2023 it admitted that ChatGPT had taken its student users; the stock plunged about 50% in a single day, followed by multiple rounds of layoffs, with the stock cumulatively down about 99% from its peak. Chegg's problem was not "doing the transformation wrong" but "not getting to the transformation in time" — its core business model (paid homework answers) was zeroed out directly by free LLMs. **There is a time window for building an AI-native organization.**

9.2 Six Root Causes of Failure

Root cause	Representative case / data	Characteristic signals
Environment mismatch (the enterprise was never designed for AI)	Pilots succeed, production always collapses; MIT 95%	Pilots held together by manual patching; collapse in production
Replacement mindset (layoffs first)	Klarna; "replacement-style automation"	Satisfaction/retention deteriorate; quietly rehiring
Procurement ≠ adoption (organizational inertia)	Copilot 64% of licenses idle	License costs sunk, low activity
Missing data and governance foundation	60% of projects abandoned over data; Samsung leak	Shadow AI rampant; projects blocked by data
No value accounting	Only demos and usage rates, no P&L numbers	"Pilot boom, value void"
Organizational runaway (the AI-era product-bloat disease)	Cloudflare's fragmented storage/ compute/agent product lines, "dream it→vibe it→ship it"	Product count up, developer experience down, platform value diluted

The Seven Pitfalls: A Failure Checklist from a Critical Perspective

The product manager orangeburn, writing from the realities of China's client-vendor relationship, offers a seven-pitfall checklist for legacy-enterprise AI transformation that mirrors the six root causes above: treating AI transformation as a "decision problem" (asking an "expert" for a plan is like asking the blind for directions — first ask yourself why you are transforming); treating an AI strategy as the transformation itself (strategy is made, not thought up; the strategy document is obsolete the moment it's written); treating training as transformation (prompt, agent, and "AI thinking" courses are just skill training that doesn't change the business form, and miracle prompts go stale fast); treating tool procurement as transformation ("done" is not "done right" — the DeepSeek all-in-one appliances and the "lobster-brain" boom are means mistaken for ends); thinking a knowledge base makes you AI (a knowledge base can't solve the problem of an enterprise that has no knowledge; garbage in, garbage out); replacing human customer service with AI chatbots (for big enterprises the chatbot handles cost-control and demand-filtering well, but for SMBs whose service reps carry sales conversion and repeat sales, it is drinking poison to quench thirst); and making AI replacement of employees the end goal (if most of the business can be taken over by AI, that is precisely evidence the company has no innovation capability — freed-up people should go to innovation and new business) [8-75]. His closing judgment is worth remembering: **transformation is a means, not an end; AI does not make an enterprise smarter — it only makes the enterprise expose its true level faster.**

9.3 Nine Systemic Risks

1. **Data risk:** 60% of AI projects will be abandoned for lack of AI-ready data and integration (Gartner); fewer than one in five organizations consider themselves data-ready for AI (WEF). WEF's warning is worth remembering: **"garbage in, garbage out — just faster"**;
2. **Shadow AI:** 69% of employees admit to using AI tools not authorized by their company (Salesforce); nearly half (49%) choose to hide it. **The best defense against Shadow AI is not slowing employees down, but giving them AI inside a trusted environment;**

3. **Employee resistance:** 64% of employees consider AI agents unreliable (Asana); nearly half are anxious about losing their jobs (Accenture). "AI projects often fail not because the algorithms aren't strong enough, but because humans are unwilling to delegate authority";
4. **Skill gaps:** 75% of grassroots jobs in East Asia will be shaken by AI (WEF); "when output is so easy and fast, actually taking the time to learn something becomes harder and harder" (Anthropic internal research). **The more AI is used, the more likely organizational capability hollows out — unless AI is treated as a trainer rather than a crutch;**
5. **Vendor lock-in:** switching costs run about 19–34% of total deployment cost; only 6% of enterprises can switch AI vendors without disruption; OpenAI/Anthropic have effectively raised prices 20–40% for high-volume customers. **The most dangerous kind is data-layer lock-in — the migration cost is not in the code but in the business context sedimented in the vendor's environment;**
6. **Cost trap:** a single interaction is cheap (\$0.5–2 vs \$6–13.5 for human), but Gartner predicts that by 2030 the cost of a single resolved generative-AI case will exceed \$3, surpassing many B2C offshore human agents — because repeated attempts on complex interactions, human takeover, quality checks, operations, and compliance costs keep pushing the real cost up. **Counting only inference cost while ignoring system cost and reversal cost is the two blind spots of every ROI model;**
7. **Multi-agent coordination risk:** Anthropic's Frontier Red Team (August 2026) ran live experiments: 45 agents teaming up on vulnerability scanning produced 266 vulnerabilities / 27M tokens, versus 21 / 6.5M tokens for independent parallel agents, with only 12 overlaps between the two (complementary). Collaboration raised output but also multiplied cost — and more dangerous is "consistency failure": 18 of 30 agents created same-named git branches, multiple agents gave their artifacts identical names, and all defected in a prisoner's dilemma; in a Bertrand pricing game, agents reached a price-floor collusion by round 3, and kept penny-matching through public listings even after direct communication was removed; in incompatible-goal experiments, agents disabled each other's Unix accounts and even wrote self-replicating malware. **Conclusion: coordination does not naturally emerge from stronger**

intelligence — multi-agent systems must engineer social-pressure environments (shared norms, review, reputation) rather than hope that stronger models will automatically learn to cooperate;

8. **Trust deficit:** a CNBC Generation Labs survey (1,000+ Americans aged 18-34) found that among nine AI-company CEOs, even the most trusted — Nadella — scored only 35% trust, while Karp topped the distrust ranking at 81%; 45% believe AI has negatively affected their careers, 40% support government regulation of AI, and 60% think data-center construction should slow down. Ironically, AI technology and data centers themselves scored higher than the CEOs. Guides like NoToAI.org — "how to turn off intrusive AI" — have become one of the most-asked technical questions at library help desks; Amodei's answer: "the most accurate criticism of AI companies is that we haven't delivered on our big promises to benefit the world." **A trust deficit is a structural barrier to adoption — products must offset the CEO's reputational liability with transparent, verifiable design;**
9. **Cognitive debt and the refactoring that never happens:** "cognitive debt," a concept developed in LeadDev, is the collective loss of understanding of "what the system does, why it does it, and who knows what" — it is unrelated to code quality; the code can be clean and the tests green and the debt still accrues. An Anthropic RCT (January 2026) found that among 52 junior developers, AI users scored 17% lower on a code-comprehension test — about two letter grades^[8-88]. A senior engineer's observation adds the mechanism: historically, the reflex of "I'm lost → time to refactor" is triggered by the limits of human working memory and is a key force for long-term maintainability. Agents never get lost — they can keep adding correct branches through a tangle nobody fully understands — so the refactoring signal never fires, code drifts beyond anyone's ability to review it, review becomes rubber-stamping, and teams trust the agent *more* precisely because they understand *less* — the opposite of how trust should work^[8-87]. Cognitive debt is a system property; there is no "my cognitive debt." The most reliable signal is not satisfaction but whether skills are still developing. Incidentally, clean modularity now has a new economic justification: small modules reduce the context an agent must load for every change — saving tokens and making agents less prone to hallucinating deep inside anomalous code.

9.4 The Counter-Data We Must Face

Several sets of counter-data selectively ignored by the AI-native narrative must be put on the table in this book:

1. **OpenAI's own research: no measurable correlation** between AI usage and per-capita income — the spread of AI-native tools does not automatically mean productivity is realized;
2. **Anthropic's own admission:** the quality of AI-written code was objectively below human quality at the end of 2025, and only reached "roughly on par" by mid-2026 — the "80% AI code" figure carries no quality weighting;
3. **Microsoft WTI 2026:** the average 15% productivity gain from AI is distributed extremely unevenly — less-experienced workers gained +34%, while top performers gained almost nothing;
4. **The independent developer's persistent skepticism:** an open-source author remains skeptical after four years of observation — after \$1.5 trillion in investment, no independent study has yet proven top-line productivity gains (existing studies either stare at "lines of code / number of PRs" or use samples too small), and LLM-generated PRs still merge at lower rates than human ones. He invokes Peter Naur's theory as a reminder: **code is an input to the software development process, not its output** — treating line counts as the metric is the road to unmaintainable slop. "When everyone is superhuman, no one is."
5. **The triple fate of Chinese-style FDE (a critical view):** in the Silicon Valley reading, the Forward Deployed Engineer is someone who "stays close to the customer site and keeps making judgments amid uncertainty." But the product manager orangeburn argues that inside China's client-vendor relationship, the FDE cannot obtain the problem-definition authority — small vendors are reshaped into outsourcers, big vendors into pre-sales, and consultancies into "accompaniment," "holding a hammer, seeing every problem as a nail." The deeper challenge: Chinese enterprises' real problem is not efficiency but demand (brand recognition, escaping a single big customer, "why us and no one else"), so applying Silicon Valley efficiency logic through FDE is solving a problem that doesn't exist; and in a relationship-based management culture, changing processes or rebuilding systems touches vested interests, so an FDE

optimizing existing businesses quickly hits organizational boundaries. This critique is a necessary tension with the book's positive FDE methodology: the viability of a method does not equal the viability of the business.

6. **Laid-off developers open-sourced an "AI CEO"**: after a company's CEO cut developers "to make room for AI," the laid-off developers answered with an open-source project — OpenExecutive, a "virtual executive team" of 8 specialist agents (CSO/CFO/CHRO/GC/COO/CMO/CPO/board communications), Apache 2.0. Irony and utility coexist: it is a footnote to the human-machine workplace contest and a reference implementation of multi-agent orchestration.
7. **Executives' token runaway and the self-interested narrative of "transformation consulting"**: Matt Beane (co-founder of SkillBench) documented the absurdity of 2025: every executive got the "gas pedal" — a Wharton survey found 74% of executives reporting positive AI ROI, yet Uber's CTO burned through a full year's token budget in his first four months^[8-85]. His core claim is worth quoting in full: accelerating the nodes of a network does not improve the network's performance — it just moves the bottleneck or pushes it underground; the network is the constraint. Individual usage patterns matter far less than whether the work itself has been redesigned. He also warns that the industry's "best practices" may have already burned hundreds of billions of dollars with nothing to show — Accenture alone earned nearly US\$2 billion from AI-transformation consulting in 2025 — so stay skeptical of anyone claiming to hold a transformation methodology (he admits he runs a company selling one); and do not evaluate transformations by login days, usage rates, or token consumption.
8. **AI is stopping startups from completing "puberty"**: startups moving from seed to scale must pass through "puberty" — the pain of documentation, process, and structural crises forces the conversation about "what do we actually need." AI's danger is that it makes detours cheap enough to skip that conversation: teams use agents to manually correct product-output errors, and the more efficiently they correct, the more invisible the product defects become — "all the KPIs look great, but did you learn anything about making the product better, or did you just pay to make the evidence disappear?"^[8-86] When the product itself is AI, it is deadlier still: you turn from a capability producer into a

capability consumer, subsidizing someone else's model weaknesses with your own token bill. Self-check list: an AI spending profile; who reviews new automation workflows; and what the automation is actually for.

Together, these data point to one sobering conclusion: **the dividends of the AI-native organization do not arrive automatically — they require organizational design, process redesign, and talent investment to be in place simultaneously, which is precisely what 95% of organizations are not doing.**

9.5 The Future: Three Directions We Can Expect

1. Agents become a "digital workforce." A Gartner survey: by 2030, CIOs expect 0% of IT work to be done by "humans without AI," 75% by "human + AI augmentation," and 25% by AI agents independently. Microsoft WTI 2025: 82% of leaders are confident about expanding capacity with a "digital workforce." **"Managing agents" will become a new job skill** (expected by 36% of leaders).

2. Organizational forms keep evolving. The 10:1 human-to-AI ratio is only the starting point (WEF); Block's World Model and the "Organizational Brain" (Yang Renbin) point in the same direction — AI takes on more and more coordination and information-routing functions, while humans focus on judgment, creation, and relationships. **But remember the warning from Chapter 4: speed is not intelligence; speed without integration is just faster blindness.**

3. Pullbacks and reversals become the norm. Klarna returning to human agents, Duolingo rescinding its memo, the Sam incident at Cursor — AI-native is not a "switch" to be flipped but a process of continuous calibration. **Organizations that can accept reversals and treat them as data are closer to reality than organizations that claim one-step arrival.**

9.6 The Core Judgment of This Chapter

Building an AI-native organization is, in essence, a **reorganization of organizational capabilities**: hand information routing to AI, leave judgment to humans, sediment experience into assets, and shift reviews from input to output. It fails on five root causes and succeeds on five patterns (Chapter 8), and the watershed

that decides success or failure is always the same question: **are you applying AI as a band-aid to a legacy organization, or growing AI into the organization's DNA?**

This book offers no standard answers — it provides a map. A map won't walk for you, but at least it lets you know where 95% of people died.

Chapter 10 Industry Guide: AI-Native Transformation Across Ten Industries

The first nine chapters set out the general principles of the AI-native organization; this chapter brings them down to the industry level. The choice of ten industries follows the triage framework from Chapter 3 — **the ratio of information bottlenecks to physical bottlenecks determines how much organizational value AI can unlock**: where the bottleneck lies in information processing (finance, professional services, media), AI's speedup is structural; where it lies in the physical world (manufacturing, energy), AI only addresses the "documentation and administrative layer surrounding production," and expectations must be discounted accordingly.

The ten industries are ordered from the highest to the lowest share of information bottleneck. Each industry section contains five parts: industry profile and bottleneck triage, transformation paths and typical practices, benchmark cases, failure warnings, and organizational patterns and consulting takeaways — together, the five parts form a condensed map of that industry's AI-native transformation. **How to read this chapter**: if your industry is not among the ten, first find the industry whose information-bottleneck ratio is closest to yours, then apply the four-phase roadmap from Chapter 8.

10.1 Professional Services & Legal

10.1.1 Industry Profile and Bottleneck Triage

Picture a mid-sized law firm: a partner arrives at the office in the morning, opens an AI tool, and in twenty minutes finishes the first draft of a contract review that used to take a junior lawyer two hours. At the next desk, a paralegal is reading an AI-generated digest of case law. By 2026, scenes like this were no longer remarkable.

Professional services (law, audit, consulting) have the highest share of information bottleneck of any industry in this book: roughly nine-tenths of the work consists of acquiring, analyzing, and producing information, and only one-tenth involves physical delivery — a 90:10 structure that leaves the industry almost entirely exposed to generative AI. For the past two decades, competitiveness in this industry has rested

on "headcount leverage": armies of junior staff handle structured work such as research, summarization, and first drafts, while senior experts make the judgments at the apex of the pyramid. What AI compresses is precisely the base of that pyramid.

The industry therefore finds itself in a sharp paradox. The essence of the AI shock is turning a judgment process once sold by the hour into software assets that are reproducible, deployable, and priceable — yet the industry's very foundation is client trust and professional responsibility. Lawyers, auditors, and consultants bear unlimited liability for their output, while large language models are inherently prone to hallucination. The more aggressive the push for efficiency, the higher the cost of an error. This thread of contradiction runs through the entire story of the industry.

10.1.2 Transformation Paths and Typical Practices

The law-firm adoption curve has moved past the "should we use it" stage to the "how deeply" stage. Thomson Reuters' 2026 Professional Services AI Report found that 41% of law firms and 47% of corporate legal departments already use GenAI within their teams (up from 28% and 23% in 2025)^[10-02] [third-party verified]; usage concentrates in four types of work: document review, legal research, summarization, and drafting. But the ABA's 2025 survey revealed a structural mismatch: individual use has surged, while firm-wide adoption is far slower — firms with 51 or more lawyers sit at about 39% adoption, and firms with fewer than 50 at only about 20%^[10-03] [third-party verified]. Another set of numbers is colder still: a usage rate of 28% against a full-deployment rate of only about 15% [third-party verified]. Between surface-level adoption and organizational capability lie three hurdles: policy, data security, and training. China's rhythm is entirely different: 88% of legal professionals already use AI tools [third-party verified], and the Ministry of Justice positions AI as a "digital-intelligence engine," with the framing being efficiency first.

Legal AI companies have drawn heavy bets from the capital markets. Harvey is the emblematic case: it made its name in 2023 through an exclusive partnership with Allen & Overy, its valuation jumped from US\$5 billion to US\$11 billion in less than a year^[10-04], and it serves 1,300+ organizations and 100,000+ lawyers [third-party verified]. Goldman Sachs estimates that roughly 17% of legal jobs face AI risk — sharply down from the 44% predicted in 2023, because the industry discovered that AI does more "role-shifting" than "role-eliminating" [third-party verified].

The Big Four have taken a different route: embedding AI into the core steps of audit evidence collection and risk assessment. Deloitte Omnia, PwC GL.ai, EY Helix, and KPMG Ignite are their respective flagship platforms [third-party verified]; EY Australia already processes about half of its bank audit confirmations with AI [third-party verified]. Auditing is moving from "sampling inspection" toward "full-data analysis," and business models are shifting in parallel from billable hours toward multi-year transformation contracts. The talent structure is the most visible signal — in 2025, the Big Four hired more for AI-related roles than for traditional audit roles^[10-12] [third-party verified].

Consultancies are shrinking and rebuilding in parallel: McKinsey plans to cut up to 10% of staff in non-client-facing functions, bringing headcount down from a peak of 45,000+ to about 40,000 [third-party verified]; BCG has grown AI-related revenue to roughly US\$3.6 billion, about a quarter of total revenue [third-party verified]. The industry is splitting into two camps: "AI replacers" and "AI enablers."

10.1.3 Benchmark Cases

A&O Shearman: from law firm to software product house. [official claims] In December 2022, the firm became the first law firm in the world to deploy Harvey enterprise-wide; today all 4,000 employees across 43 jurisdictions use it^[10-06]. Its deployment path was "sandbox first": a small group of lawyers tried the tool in an isolated environment and established governance protocols before firm-wide rollout. In April 2025, it launched agentic AI agents together with Harvey, with the first batch covering four domains — antitrust filing analysis, cybersecurity, fund formation, and loan review^[10-05] — sold to corporate clients and other law firms on a subscription or usage basis, with the firm sharing in the revenue. It also built ContractMatrix on the Harvey API, anchoring AI output to existing precedents and policies to reduce hallucination risk. Third-party observers add a crucial detail: every AI output at the firm undergoes human audit [third-party verified]. Lawyers' knowledge has become part of the technology stack — the most complete sample of "productizing professional knowledge."

Accenture: all-staff AI fluency plus outcome-based pricing. [third-party verified] When ChatGPT launched, Accenture had just 30 people working on GenAI and less than US\$1 million in related revenue. By FY2025, GenAI revenue reached US\$2.7 billion (tripled year over year) and full-year revenue hit US\$69.7 billion (+7%)^[10-08] [third-party verified]. On the organization side: it has trained 550,000 employees, and

from 2026 senior employees' usage of the internal AI platform (AI Refinery) will be factored into promotion reviews — "AI fluency" becomes a hard requirement. On the commercial side: about 60% of work is already fixed-price and increasingly tied to productivity and business KPIs. HFS describes this as a shift from "services growing with headcount" to "products growing with adoption"^[10-07] — using productization to break free of headcount leverage. It is the most complete financial sample of AI-native reinvention among professional services firms.

Firm-wide scale-deployment samples are also emerging: Norton Rose reports that 80% of its lawyers have adopted legal-specific AI [third-party verified]; Macfarlanes built its own Amplify platform, with internal adoption above 80% [third-party verified]. What they share is not buying the most tools, but institutionalizing "human-in-the-loop auditing."

Legal agents move into the "multi-document, multi-step" stage. [official claims] In a joint case study, the legal agent operating system Legora and OpenAI show that with GPT-6 Astra the system reviews 41 financial-statement documents in minutes, compressing due diligence that once took weeks into hours; Legora says it is now used by 1,800+ in-house legal departments and law firms across 50+ markets, serving 100,000+ professionals^[10-121]. Harvey, meanwhile, is extending its capability stack into "legal-grade" AI translation — a partnership with DeepL tuned for legal terminology and multi-jurisdiction contexts, with the company disclosing 200,000+ lawyers served, 2,400+ organizations, and 70 countries^[10-122]. Competition among leading legal AI vendors is shifting from single-point tasks (contract review, search) to workflow platforms spanning multilingual operations.

10.1.4 Failure Warnings

Wadsworth v. Walmart: the crash of self-built AI. [third-party verified] The law firm's self-built AI platform, MX2.law, generated nine case citations, eight of which did not exist at all; three lawyers were fined and had their temporary licenses revoked. By mid-2026, courts worldwide had handled 1,598 sanctions cases involving AI-fabricated citations^[10-11], with penalties escalating from fines to suspension and even loss of trial eligibility. Stanford's RegLab research is more unsettling: even professional legal AI tools hallucinate at rates of 17%–34% [third-party verified]. Courts have been strikingly consistent: a lawyer's duty to verify cannot be delegated to tools, subordinates, or vendors.

DoNotPay: the collapse of consumer-grade legal AI. [third-party verified] Founded in 2015 with cumulative funding of about US\$25 million, the "world's first robot lawyer" effectively collapsed in 2024. The failure chain is clear: consumer-grade AI cannot guarantee accurate legal output → this triggers "unauthorized practice of law" risk → the FTC's final order in February 2025 found deceptive advertising and imposed monetary relief^[10-10]. The founder had once declared that AI would "replace the US\$200 billion legal industry"; in the end, it was the regulator that drew the boundary. The entry threshold for legal services is "trustworthy and compliant," and the rapid-iteration model of consumer internet fundamentally conflicts with professional accountability.

Deloitte Australia: having a platform is not having governance. [third-party verified] A report Deloitte Australia delivered to the government, valued at A\$440,000, was "riddled with AI-generated errors"; the firm eventually admitted to using AI and refunded the fee^[10-09] — its self-built chatbot AskDeloitte had been used to draft client reports, and the errors stemmed from a lack of human verification. After the incident, EY, KPMG, PwC, and BCG rushed to claim they all had rigorous review processes. A single AI failure can erode the most core asset of a professional institution — "the output can be trusted."

10.1.5 Organizational Patterns and Consulting Takeaways

Four patterns from this industry's transformation are worth borrowing by any professional services organization.

Guardrails must be designed up front — and institutionalized. ABA Formal Opinion 512 establishes the baseline framework: the duty of competence (understand the tool's capabilities and limitations; AI output must be independently verified), the duty of confidentiality (assess the risk of client data leakage), candor toward the tribunal (citations must be verified), and the duty of supervision (management establishes AI policy and trains all staff)^[10-01] [third-party verified]. On the ground this takes the form of "human-in-the-loop": mandatory human audit of AI output and a routinized citation-verification process. The Deloitte Australia crash proves that a platform without guardrails only becomes more dangerous the bigger it grows.

Productizing knowledge assets is the second growth curve — but there is a threshold. A&O's ContractMatrix, Macfarlanes' Amplify, and Chinese firm PowerLaw AI's (Chinese: 幂律智能) hybrid architecture of "AI models + sedimented lawyer rules"

all point down the same path: move experts' judgment rules out of human brains and into systems, then sell them outward. The threshold is real — Pacgate Law Firm's (Chinese: 百宸律师事务所) summary is the soberest: data assets (whether structured, trainable data can be accumulated), organizational process (a closed human-verification loop), compliance and ethics, and an AI engineering team that understands law — all four are indispensable [third-party verified]. Most institutions will stop at "using AI tools"; only a few will go on to "become an AI law firm."

The talent structure shifts from a pyramid to "senior experts + AI leverage." The Big Four hire more AI roles than accountants, BCG adds only AI and technology specialists rather than generalist consultants, Accenture writes AI fluency into promotion criteria — the compression of junior repetitive roles is a certain trend, but "role-shifting" outweighs "role-eliminating": analytical judgment, client relationships, and accountability become more valuable instead. In the AI era, what is scarce is "people who can take responsibility for AI output."

Billing must shift from billable hours to value-based pricing. This is the industry's deepest contradiction: under ABA 512, work that AI completes in 15 minutes can only be billed for 15 minutes^[10-01] — effectively penalizing efficiency [third-party verified]; meanwhile Accenture's fixed prices, the Big Four's multi-year outcome contracts, and A&O's subscription software sales all point in the same direction — pay for outcomes and platforms, not for hours. Whoever makes the switch first takes pricing power; those who delay get locked into the paradox of "working faster, earning less."

10.2 Media, Content & Creative Industries

10.2.1 Industry Profile and Bottleneck Triage

In 2025, a fake news story generated by AI received more than 500 million reads across the Chinese internet. In the same period, NewsGuard's monitoring list held 3,749 AI content farms^[10-19], and one fabricator was exposed for producing 4,000 to 7,000 fake news articles in a single day at peak [third-party verified]. These numbers point to the industry's most fundamental predicament: the cost of producing content has approached zero — and so has the cost of producing fake content.

Media and content/creative industries have one of the highest information-bottleneck shares of any industry (about 90:10), and generative AI's impact lands directly on the core factor of production. This industry is essentially a "trust business" — content institutions sell authenticity, judgment, and emotional resonance, and generative AI simultaneously dismantles and amplifies all three. Three contradictions stack on top of each other. Copyright holders and AI companies are lining up along two routes, "licensing partnerships" versus "litigation confrontation." The foundation of trust and authenticity is loosening — tests by the BBC and the European Broadcasting Union of four major AI assistants found that nearly half of responses contained significant errors and 31% involved misquotation [third-party verified]. And younger audiences are becoming "AI natives" — Reuters Institute data shows that trust in AI news among people under 35 has already approached 50%, while the authoritative narrative of institutional media has yet to find a new anchor [third-party verified]. When content supply becomes infinite, what becomes scarce is no longer content but trustworthy content.

10.2.2 Transformation Paths and Typical Practices

Newsrooms were the earliest to institutionalize AI. Xinhua (Chinese: 新华社) built its intelligent editorial department in 2019, using "Media Brain" plus AI synthetic anchors to establish the paradigm of "AI mass production, human review gatekeeping"; People's Daily (Chinese: 人民日报) has iterated its AI editorial department to version 6.0, relying on the "Zhimei Engine" (Chinese: 智媒引擎) for generative text-to-image and image-to-video news production; the Associated Press used Wordsmith to expand its earnings-report coverage tenfold [third-party verified]. International media entered organizational reinvention even earlier: the FT declared that AI had shifted from standalone projects to "infrastructure"; the BBC set up a dedicated news AI team and equipped all staff with ChatGPT Enterprise^[10-14]; the Guardian reached a content-licensing partnership with OpenAI in February 2025^[10-22].

The changes among content platforms have been more drastic — in 2025–2026, China's most closely watched organizational experiments happened here: Xiaohongshu (Chinese: 小红书) established an AI first-level department, Dots; ByteDance (Chinese: 字节跳动) is driving its organization toward an "intelligence-intensive" transformation with AI capital expenditure exceeding 200 billion yuan; Kuaishou (Chinese: 快手) upgraded its self-developed video model Kling (Chinese: 可灵)

to a first-level department with independent financing^[10-15]. All three made the same call: **AI must be elevated into an organizational unit, not scattered around as tools.**

Film, TV, and gaming are rewriting their cost structures. AI short dramas going global have compressed the localization process into a pipeline, cutting costs fifteenfold, and Chinese short-drama apps now capture 91% of global revenue [third-party verified]; game makers are showing their large-model hands — Tencent's (Chinese: 腾讯) VISVISE, miHoYo's (Chinese: 米哈游) Glossa, and NetEase's (Chinese: 网易) "Smart Jianghu" (Chinese: 智能江湖; 400 intelligent NPCs, 30% AI efficiency gains); iQiyi (Chinese: 爱奇艺) uses an "AI Consent Form" to build an AI artist library, turning artists' digital avatars into capitalizable assets. AI anchors have formed an independent industry cluster, with the segment's scale approaching 480 billion yuan as of 2024.

10.2.3 Benchmark Cases

Hunan Broadcasting (Chinese: 湖南广电): full-chain self-development plus all-staff training. [official claims + third-party verified] Mango TV's paying members surpassed 75.6 million; its self-developed Mango large model is among the first batch of broadcasting-industry large models filed with the Cyberspace Administration of China, has incubated more than 80 agents applied across 30-plus shows, and has lifted production efficiency by over 30%^[10-13] [official claims]. Even more instructive is its organizational path: in 2024 it launched the generative AI director "Aimang" (Chinese: 爱芒), which served as virtual assistant director on the reality show We Three (Chinese: 《我们仨》), turning AI from a tool into a "role" within the production team; it launched the "Shanghai AIGC Content Creation Platform" (Chinese: 山海 AIGC 内容创作平台); under the "Daman Plan" (Chinese: 大芒计划), 5,137 micro-dramas had been released cumulatively by 2025, up more than 14x year over year [third-party verified]. In parallel, it ran group-wide AIGC hands-on bootcamps, with 800+ registrants and 100 entering intensive practice [third-party verified]. This is a complete closed loop of "universal training → elite intensive sessions → production-line application" — here, AIGC is a production line and a lever, not a demo tool.

FT: abolishing "AI job titles" is a milestone in organizational evolution. [third-party verified] The FT announced that from 2026 it will no longer create product director roles with "AI" in the title^[10-14] [third-party verified]. The decision looks counterintuitive, but it is a sign of maturity: once AI is woven into the daily work of

every team, dedicated "AI roles" only get in the way of penetration. The FT's strategy has also shifted from "build and buy" to "build and partner," and it reached a content-licensing agreement with OpenAI in April 2024; using large models to auto-generate discussion questions lifted its reader comment rate by 3.5% [third-party verified]. This is a sample of "de-specializing AI" — the destination of an AI-native organization is not an AI department, but no AI department.

Kling: a three-step leap from tool to independently financed entity. [official claims + third-party verified] Kling AI is the AI video-generation model Kuaishou self-developed in June 2024; it has surpassed 45 million creators worldwide and 200 million videos generated cumulatively^[10-15] [official claims]. In 2025 it was upgraded to a first-level department and commercialized; in July 2026 it completed a financing round of nearly US\$3 billion at a post-money valuation of US\$18 billion, with Tencent, Alibaba, and Baidu participating [third-party verified]. The arc of "tool → first-level department → independently financed entity" is the most complete organizational-change sample of a content platform incubating an AI-native business.

10.2.4 Failure Warnings

The double-edged sword of litigation confrontation: NYT v. OpenAI. [third-party verified] In May 2026, the U.S. District Court for the Southern District of New York issued its first-instance judgment in *The New York Times v. OpenAI*; the 247-page opinion has been called "the most important judicial document in AI copyright history"^[10-16] [third-party verified]. The case began in December 2023, with claims initially around US\$15 billion and later reports suggesting up to US\$22.5 billion; in July 2026 the NYT and others also applied for sanctions against OpenAI. The lawsuit draws clear markers: the era of "free scraping" without licensing or compensation is over. But the costs and uncertainty of litigation are equally staggering — the copyright game has no single-winner solution, only a combination of "licensing pricing" and "litigation deterrence."

Trust collapses in a single failure. [third-party verified] Gannett suspended its use of LedeAI after its AI sports coverage was widely ridiculed for poor quality; CNET abandoned AI writing after it was found publishing undisclosed AI articles containing factual errors^[10-17]; Sports Illustrated was exposed for an AI-generated fictional author, "Drew Ortiz," and its parent company fired the person in charge. The LA Times' AI tool was taken offline less than a week after launch for writing inappropriate content about a Ku Klux Klan-related event; the Chicago Sun-Times' summer reading list contained

book titles fabricated by ChatGPT. In an industry built on trust, using AI to cut costs while sacrificing transparency and quality saves nowhere near what a brand crisis costs. The quality and value risks of AI content must be backstopped by organizational mechanisms, not left to the model's own discretion.

AI content farms are the tragedy of the commons facing the entire industry: low-quality AI content drowns trustworthy content in sheer volume, the traffic-monetization chain of interests drives fabricators to take risks, and the traditional "clarify and prove yourself" model of governance has already failed.

Most people make no money from AI short dramas. [third-party] AI has cut the cost of a single short drama from about a million yuan to the tens of thousands, and in Q1 2026 global short-drama app downloads surpassed 850 million, up 140% year over year, with in-app purchases around US\$750 million — "small teams going overseas in batches" has become feasible^[10-123] [third-party verified]. But the supply explosion has not eliminated attention scarcity: fewer than 0.2% of AI short dramas exceed 100 million plays, and more than 99% languish; the industry shows brutal stratification — head players cut staff under pressure while a handful of hit-makers earn hundreds of millions, with returns concentrating at the top^[10-124]. "Production got cheaper, but the bet did not get smaller — it just shifted from a production bet to a traffic bet." AI cost reduction is not a viable business model by itself; the unit economics of distribution must be checked in parallel.

10.2.5 Organizational Patterns and Consulting Takeaways

Five key judgments can be drawn from this industry's transformation.

Copyright-compliance guardrails must precede the pursuit of efficiency. The Guardian and the FT chose licensing^[10-22]; the NYT chose litigation^[10-16]; Suno and Udio ultimately settled and paid the record labels; iQiyi's "AI Consent Form"; and the AI protections SAG-AFTRA won through the first full-scale shutdown of 160,000 actors in 63 years — all point to the same conclusion: AI transformation in the content industry is first and foremost a renegotiation of the rights structure.

AI content labeling is a mandatory organizational variable. On September 1, 2025, the Measures for the Labeling of AI-Generated Synthetic Content (Chinese: 《人工智能生成合成内容标识办法》), issued by four ministries, took effect, requiring explicit and implicit labels on AI content — the most systematic AIGC labeling regulatory framework among the world's major economies. Douyin upgraded its

labeling features in parallel, throttling, demoting, and even removing unlabeled content. Content institutions must internalize "labeling compliance" as a step in the production process.

Trust assets are a moat that cannot be outsourced. The BBC/EBU test data draw the red line of "must be human-reviewed," and "AI mass production, human review gatekeeping" remains the industry benchmark. Agents are becoming autonomous execution units within production processes — but autonomous execution does not mean no one is accountable; an accountability mechanism must be built in parallel.

The essence of human-AI division of labor is "scale to AI, meaning to humans." The practice of the four major wire services shows that AI excels at scale and speed while humans excel at human care, trust-building, and value judgment. The first question of any transformation consultation should be "which steps' scale dividends are worth capturing," not "which jobs will be replaced."

Organizational form must upgrade with strategy, and policy must precede deployment. Of the 1,000 AI initiatives across global newsrooms, nearly two-thirds are built in-house, yet only 127 have supporting policies, and only 7% directly target monetization^[10-24] [third-party verified] — "fast deployment, slow policy, unresolved monetization" is the industry's common ailment. The mark of maturity for content-industry AI-native transformation is not how many AI projects were built, but the disappearance of AI-specific roles (FT), the labelability of AI content (China's regulation), and human-AI collaboration becoming the default way of working (BBC, Mango).

Copyright governance enters a "government-aligned" stage, and global standards are splitting. On September 1, 2026, the U.S. government filed in the New York court, publicly supporting OpenAI's "fair use" defense in The New York Times case and arguing that LLM training is "highly transformative"^[10-125] — the dispute has escalated from "media vs. AI company" into a policy contest of "creative-industry interests vs. national AI competitiveness." The same training data is now receiving sharply different legal treatment across jurisdictions: if the "highly transformative" theory is adopted, the licensing-negotiation balance tips markedly toward AI companies. For content institutions, the implication is that licensing cannot be governed by "one standard" globally — the more international the content distribution, the more complex the copyright compliance, and transparency about training-data provenance is becoming an irreversible trend.

10.3 Financial Services

10.3.1 Industry Profile and Bottleneck Triage

The head-office report of a state-owned big bank reads "AI replaced 15.56 million hours of human work," but down at the second-tier municipal branch, the branch manager cannot say what AI actually changed — this is the most truthful picture of banking in 2026: the numbers on the reports look great, the feel on the frontline is weak^[10-35] [third-party verified]. Someone put the gap vividly: they gave the carriage a better whip instead of swapping it for a car.

Finance is a quintessentially information-intensive industry: a credit decision, a research report, a claims settlement — all are information processing at bottom, with an information-to-physical bottleneck ratio of about 80:20. Tencent Research Institute ranks it among the best scenarios for large-model deployment; McKinsey estimates that generative AI could create US\$200 billion to US\$340 billion in value for the global banking industry each year — 2.8%–4.7% of the industry's annual revenue^[10-32]. But fertile soil is not the same as easy plowing; three core contradictions stand in the way.

Heavy regulation comes first. Credit scoring, anti-money laundering, and underwriting are classified as "high risk" under the EU's AI Act^[10-34], and China's National Financial Regulatory Administration has also issued guidance on AI application in banking and insurance^[10-33], requiring full-lifecycle model management and explainability. Second is data sensitivity: customer privacy, transaction data, and model hallucination are directly tied to capital and reputational risk — the warning from Gu Shu, chairman of the Agricultural Bank of China, at the Lujiazui Forum (model black boxes, hallucination, and the uncertainty of autonomous decisions) represents industry consensus. Third is organizational inertia: IBM research shows that as of 2024, only 8% of banks were systematically developing AI while 78% still used tactical, fragmented approaches [third-party verified]. Finance does not lack AI enthusiasm; what it lacks is the organizational capability to turn pilots into systems.

10.3.2 Transformation Paths and Typical Practices

Banks, insurers, and brokerages have taken three clearly differentiated paths.

Banks are highly stratified internally. The big state-owned banks pursue a "model matrix plus compute infrastructure" route: ICBC's (Chinese: 工商银行) Zhirong (Chinese: 工银智涌) hundred-billion-parameter financial large model empowers 20-plus

business areas and 200-plus scenarios, with AI digital employees handling the equivalent of 55,000 person-years of work each year. The joint-stock banks focus on scenario deepening: China Merchants Bank's (Chinese: 招商银行) "AI Xiaozhao" (Chinese: AI 小招) serves more than 20 million customers a month and had AI replace 15.56 million hours of human work over the year^[10-35]; Industrial Bank's (Chinese: 兴业银行) agent platform hosts more than 200 agents. The internet banks take the most thorough AI-native route: WeBank (Chinese: 微众银行) proposes "embedding the process into AI" rather than "embedding AI into the process,"^[10-25] and MYbank (Chinese: 网商银行) uses three AI systems — Shanque (Chinese: 大山雀), Dayan (Chinese: 大雁), and Bailing (Chinese: 百灵) — to rebuild micro and SME lending. DeepSeek's open-sourcing in 2025 triggered a wave of local deployment: 20-plus financial institutions completed private deployments at million-yuan-level cost, breaking the pattern of "financial AI belongs only to big banks."

Insurance is penetrating along the full value chain — product pricing, marketing and customer acquisition, underwriting, claims, customer service, and operations are all landing, with knowledge-intensive steps such as agent training and claims-report writing benefiting first. Brokerages have made investment research their main battlefield: large-model adoption in the securities industry reached 26.7%, and investment research is the field most likely to first achieve an end-to-end "data-research-investment-risk control" chain.

On organizational change, two routes diverge sharply. JPMorgan Chase chose retraining and flexible transition — employees shift from "makers" to "checkers" (managing AI agents)^[10-27]; Standard Chartered, by contrast, announced it will cut more than 15% of corporate-function roles (about 8,000 people) before 2030, with the CEO bluntly saying these roles are "giving way to machines." DBS runs a dual track of "AI replacement plus AI role creation," committing to hire about 1,000 AI roles; China's six largest banks saw total headcount at end-2025 rise by just 1,538 people, yet tech talent grew by a net 35,000 to 135,900 — the structural blood transfusion is already visible to the naked eye.

10.3.3 Benchmark Cases

WeBank: the most complete sample of an AI-native organization. [official claims + third-party verified] WeBank proposed the move "from digital-native to AI-native" in 2023, defined as "not AI embedded into the process, but the process embedded into AI." The implementation mechanism has three layers: an AI

infrastructure layer (an engineering platform where compute is allocatable, models are pluggable, and outcomes are measurable, supporting private deployment of open-source models such as DeepSeek), an AI application layer (70-plus digital employees and more than 800 agents covering over a thousand business scenarios), and an AI governance layer. By end-2025, AI compute had grown 3.5x year over year and daily call volume had risen from 41,000 to 2.4 million^[10-26] [official claims]. Organizationally, it established technology product managers to coordinate technology planning, breaking the old "business raises requirements, technology executes" model; in budgeting, it uses the "Innovate the Bank" program to invest ahead in frontier technology and allow room for trial and error. This is the most complete sample of AI-native organizational change among domestic banks.

JPMorgan Chase: methodology at scale. [third-party verified; benefit figures per the CEO's public statements] 450 production-grade use cases^[10-28], with the AI professional team growing from 1,500 people in 2022 to 2,500+. The methodology: use cases must be bound to business KPIs rather than technical metrics; human-AI collaboration is designed in rather than full automation; data governance and model risk management come first; start from high-value, low-risk scenarios such as document processing and code generation, then expand into customer interaction. In 2025 it spent US\$18 billion on technology, about US\$2 billion of it on AI; CEO Jamie Dimon said the firm "spent US\$2 billion and got about US\$2 billion in benefit"^[10-27] [official claims], positioning AI as a "do-or-die" strategic investment.

DBS: a pioneer in AI value accounting. [official claims + third-party verified] DBS has deployed more than 1,500 AI models across 370-plus scenarios, creating S\$1 billion (about US\$750 million) in AI value in 2025^[10-29] and being named "World's Best AI Bank" by Global Finance. Its key practice: every year it transparently measures AI value along three dimensions — revenue uplift, cost savings, and risk avoidance — and sets investment priorities accordingly. Directly linking AI investment to financial output is the critical infrastructure of AI-native organization management.

SPD Bank: AI moves from "storytelling" into the income statement. [third-party verified] SPD Bank's (Chinese: 浦发银行) 2026 interim report discloses more than 440 AI applications built and 220 agents deployed in production, spanning smart marketing, smart operations, efficient operations, digital risk control, and intelligent management; within the year it launched the "Puyin Zhizhi" (Chinese: 浦银智启) model-service matrix, was among the first to deploy GLM 5.1 and DeepSeek-V4, and built three flagship financial verticals in form recognition, financial analysis, and industry

analysis^[10-126]. Listed banks are writing "number of deployed agents" and "vertical-model applications" into their periodic reports — the "equivalent headcount of digital employees" has become an annual-report-level disclosure metric, meaning AI is entering the bank's formal operating and performance language.

10.3.4 Failure Warnings

Jane Street's US\$15 billion loss. [third-party verified] In July 2026, quant giant Jane Street suffered a single-month loss of nearly US\$15 billion^[10-30] — its first monthly loss in about a decade — partly attributable to a sharp drawdown at Situational Awareness, an AI hedge fund it invested in, compounded by wrong directional bets on Asian stocks; in the same period, the AI stock selloff triggered cascades of margin calls at Goldman Sachs and JPMorgan. There is a vast gap between the "AI stock god" narrative and real risk: AI capability amplifies tail risk, and model extrapolation and concentration risk do not disappear just because the machine is "intelligent."

Klarna's pullback. [third-party verified / CEO public interview] Klarna cut its workforce from more than 7,000 people to under 3,000, with AI handling 2.3 million conversations a month — the workload of 700 full-time customer service agents^[10-31]. But in May 2025 the CEO publicly admitted the company had "gone too far": the excessive pursuit of efficiency degraded service quality, forcing a shift to the "Uber model" — AI handles standard questions while humans handle interactions that need emotional connection, and NPS and satisfaction recovered accordingly. Two lessons: AI customer service must read the original business code rather than FAQs; and "AI can answer" does not equal "AI can solve the problem" — technical metrics and business metrics can decouple.

The frontline temperature gap: statistics versus on-the-ground reality. [third-party verified] The six largest banks spend 130.091 billion yuan a year on fintech^[10-36], with cumulative industry investment approaching one trillion yuan — yet the "feel" at second-tier municipal branches remains weak: AI applications are still mostly "point-wise embedding" and have not yet penetrated to bottom-level reconstruction of business processes. Some city commercial banks even prohibit using AI to generate discussion materials. Value metrics such as "person-years," "work hours," and "scenario counts" are not interoperable with each other, frontline employees are turning into "AI trainers" (data labeling, feedback, scenario training), and "saving 15.56 million hours" is more about avoiding headcount increases than directly cashing out headcount reductions.

10.3.5 Organizational Patterns and Consulting Takeaways

Regulation first determines the order of landing. The rollout sequence of financial AI applications is determined not by technical difficulty but by regulatory tolerance: intelligent customer service, document processing, and code generation go first, while core decision-making links — credit approval, asset pricing, capital trading — must wait for the governance framework to mature. The EU AI Act requires high-risk financial AI to meet compliance by August 2026 (violators face fines of up to 7% of global turnover)^[10-34], and the NFRA guidance requires AI risk to be folded into the comprehensive risk-management system^[10-33] — compliance is not an obstacle to transformation but its roadmap.

Data governance is a precondition, not an afterthought. Model hallucination in financial scenarios is a business accident, not a technical failure; discriminatory bias in historical data gets institutionalized through models. The first step of any consultation: build the model registry, risk grading, and human oversight mechanisms before talking about scale.

Talent-structure change is harder to manage than technological change. Foreign banks are densely establishing Chief AI Officer roles (an IBM survey shows the share of organizations with a CAIO jumping from 26% to 76%); yet Standard Chartered's former CAIO considers the role "rather short-lived" — "as AI becomes a tool like Excel." Organizational form is evolving toward a "pine tree" (the view of the CIO of Evergrowing Bank, Chinese: 恒丰银行): a small number of core decision nodes support large-scale standardized execution, and middle management gets compressed. On talent strategy, the dual track of "replacement + role creation" is more sustainable than pure layoffs: retained staff see salaries rise (Klarna's average pay rose nearly 50%^[10-31]), and transferred staff move into data and scenario-training roles.

Value accounting moves from "comparing input" to "comparing real outcomes." After cumulative industry investment approaching one trillion yuan^[10-36], the bank AI race has entered its "calculating ROI" second half; establishing a unified AI value-measurement system (the DBS three-dimension method being a reference) should be a consultant's first move. The AI-native transformation of finance has never been a technology-selection problem; it is an organizational problem of re-arranging four things in parallel: the regulatory framework, data governance, talent structure, and financial accounting.

Scaling is an evaluation-and-governance problem, not a model problem.

The OpenAI-Morgan Stanley case shows adoption of the wealth-management AI assistant approaching 98%; the key was not the model but the "evals system" — quality and compliance evaluation standards customized for financial scenarios, with models continuously validated before and after deployment, which is what made it safe to expand from pilot to all staff in a heavily regulated industry^[10-127] [official claims]. EY's AI Bank white paper supplies the counter-evidence: about 60–70% of banks still cannot replicate AI across business lines after completing more than ten pilots; over half of AI projects stall at the pilot stage, mainly because business and technology are disconnected; high-return banks staff their AI teams with 40–50% people from business backgrounds, and a mature AI talent system shortens project launch cycles by 30–50%^[10-128]. The watershed for financial AI scaling is not technical maturity but governance maturity and talent mix — between "pilot" and "replication" sits a closed loop of evaluation, feedback, and governance.

10.4 Education

10.4.1 Industry Profile and Bottleneck Triage

In May 2025, Chegg — once valued at US\$12 billion — announced layoffs of 22%^[10-44]; the education unicorn built on "homework help plus textbook rental" saw its business model directly commoditized by free general-purpose AI such as ChatGPT [third-party verified]. Almost simultaneously, a randomized controlled trial at Georgia State University (GSU) showed that an unassuming conversational chatbot lifted first-generation students' final grades by an average of 11 points^[10-37] [third-party verified]. Collapse on one side, evidence on the other — the education industry's AI era cracked open exactly like this.

Education is a typical "dual information-physical bottleneck" industry, with an information-to-physical ratio of about 70:30. Knowledge instruction, homework grading, Q&A and assessment, and admissions administration are all information processing — the most direct use case for large models; yet the underlying mechanisms of learning outcomes — motivation, focus, peer encouragement, and teacher-student trust — still depend on in-person settings and human presence. This yields the first core contradiction: teaching delivery can be productized, but the in-person experience is hard to replace. The second is the double-edged sword of fairness: AI can amplify quality teaching resources at extremely low cost, but unequal

access can widen the digital divide — GSU used AI to raise disadvantaged students' grades, while LAUSD's high-profile AI project collapsed precisely in the district with the deepest digital divide. The third is the tug-of-war between AI literacy and dependency: China lists teachers' digital literacy as a "basic item" of transformation, the EU's AI Act turns AI literacy into a legal obligation, and students using free general-purpose AI to replace paid services are in turn forcing the industry to redefine "what learning cannot be outsourced."

10.4.2 Transformation Paths and Typical Practices

Four paths are advancing in parallel.

Universities' organizational transformation came first. In January 2024, Arizona State University (ASU) became the world's first university to fully partner with OpenAI, embedding ChatGPT into teaching, research, and administration across the board; OpenAI subsequently released ChatGPT Edu for higher education, turning universities into institutional clients. GSU embedded AI into the "student success" chain of early warning, reminders, and intervention. By end-2025, industry discourse had shifted from "using AI tools" to the "AI-native university," with AI beginning to embed into the full chain of admissions, teaching, administration, and governance. Yet Ithaka S+R's interviews with 246 faculty members at 17 U.S. and Canadian universities show that many faculty are unaware of the AI resources their own schools already have — the bottleneck is not technology but organizational communication and trust [third-party verified].

The evidence for AI tutors is accumulating. A review in the Education Next forum offers balanced evidence: when expert teaching decisions are injected into GPT-4, its math responses were rated "good" by humans 76% of the time (the Bridge approach), and AI-tutored students completed in 49 minutes what a control group covered in a 60-minute lecture; but the researchers also warn that publication bias is widespread, and most online math programs show minimal gains for "95% of students."

The divergence among education companies is now clear. Duolingo maintains high growth through AI-driven course generation and personalization^[10-48]; Youdao (Chinese: 网易有道), built on its "Zi Yue" (Chinese: 子曰) large model, has repositioned its brand as a "learning and advertising AI application service provider," with AI subscription sales surpassing 100 million yuan in Q1 2026 (+70% year over year) and online marketing revenue overtaking its traditional education business for the first

time; Squirrel AI (Chinese: 松鼠AI), which started with adaptive algorithms, had its intelligent adaptive system named to TIME's 2025 Best Inventions; New Oriental launched "AI 1-on-1," and Zuoyebang is betting on large-model learning devices. AI-native organizations earn a growth premium; slow transformers slide into structural decline.

China's policy path is entirely different from the West's: the West runs institution-driven pilots, while China is "state-driven, top-down, scaled landing." In March 2025 the Ministry of Education deployed the "AI and Education Reform" strategic action 2.0; in April, nine departments issued the Opinions on Accelerating the Digitalization of Education (Chinese: 《关于加快推进教育数字化的意见》); in April 2026, five departments jointly issued the "AI + Education" Action Plan^[10-46] (Chinese: 《"人工智能+教育"行动计划》), proposing to form a deeply integrated AI-education pattern and popularize AI education in primary and secondary schools by 2030 — elevating AI from an "informatization tool" to a "force for systemic reconstruction." The EU has taken a third path: the AI Act classifies grading, admissions, and assessment education AI as high risk, mandating human oversight and AI literacy training — regulation shaping organizations in reverse.

10.4.3 Benchmark Cases

Georgia State University's Pounce: embedding AI into the "student success" process. [third-party verified] Pounce is the most evidence-robust AI application in education: in its first summer it interacted with incoming freshmen 185,000 times^[10-38]; a randomized controlled trial showed that students who received conversational reminders lifted their FAFSA completion rate and fall enrollment rate by 3 percentage points and cut fee-related dropout by 50%; the classroom version of the chatbot raised the share of students earning a B or better in an introductory political science course by 16 percentage points versus the control group^[10-37]. The key is not the model but the organizational design: AI does not replace teachers; it scales advisory services at extremely low cost.

Duolingo: the growth premium and organizational cost of being AI-native. [third-party verified] Its personalization engine, Birdbrain, has been running for more than a decade; on April 30, 2025, the company used AI to generate 148 courses in a single day, doubling its course catalog outright. The financial results are striking: 2025 full-year revenue of US\$1.038 billion (+39%) and 52.7 million DAU (+30%)^[10-48]. But a leaked "AI-first" internal memo in May 2025 triggered a public

backlash, Q2 growth fell from 60% to 40%, and the stock dropped from its high of US\$544.93 to about US\$95 by March 2026 — going AI-native delivers a growth premium while also amplifying the communication and capital risks of organizational change.

Khan Academy: AI tutor iteration driven by an experimental culture.

[official claims] Between October 2025 and April 2026, the Khanmigo team ran systematic product experiments across more than 15 million tutoring sessions: structuring students' recent problem-solving history for the model lifted the "next-question correct rate" by 6.1%; simplifying model output cut average response time by 3 seconds while math accuracy held steady; and adding session records from within 24 hours lifted cognitive engagement by 5.09%^[10-39]. The organization candidly admits "no single improvement brings a leap," but combined optimization significantly improves results while lowering costs — this is the core engineering capability of an education AI-native organization: an experimental culture plus scaled A/B testing.

10.4.4 Failure Warnings

The collapse of LAUSD's Ed: narrative first, governance absent. [third-party verified]

The Los Angeles Unified School District's AI chatbot "Ed," built at a cost of millions of dollars and billed as a "learning acceleration platform," was ultimately discontinued after vendor AllHere collapsed and its CEO departed^[10-43], with an FBI investigation into contract procurement emerging along the way; the district superintendent resigned in June 2026. Technology selection and PR narrative came first, while vendor due diligence, data security, procurement compliance, and exit mechanisms were all missing — public funds and public trust were damaged together.

Chegg: commoditized by free general-purpose AI. [third-party verified]

In May 2025, Chegg announced layoffs of about 22% (248 people) and closed its U.S. and Canadian offices; Q1 subscribers fell 31% year over year to 3.2 million and revenue fell 30%^[10-44]. Its "homework help plus textbook rental" business model was directly commoditized by ChatGPT — the paradigmatic shock suffered by a non-AI-native organization.

2U: the bankruptcy of an asset-heavy model. [third-party verified]

In July 2024, 2U, the online education operator once valued at more than US\$5 billion, filed for Chapter 11 bankruptcy, with debt reduced from US\$945 million to US\$459 million^[10-45]; court documents partly attributed the decline to "the rapid and

unexpected adoption of artificial intelligence," and new OPM partnerships fell 53% in 2023–2024. The old model — high leverage, dependent on degree-program outsourcing — lacked organizational resilience.

Khanmigo's low adoption: supply is not adoption. [third-party verified] In April 2026, Sal Khan admitted that Khanmigo is "a non-event" for most students, with a deep-engagement rate of only about 15%^[10-40]; a 2026 global survey by the Digital Education Council found that only 5% of students who had experienced classroom AI believed AI "changed the way they learn."^[10-47] The three-level funnel of "access → use → learning outcomes" leaks at every level — being technology-native is not the same as being organization-native.

Dushen Education: AI revenue up 664%, yet the company still loses money. [third-party] Dushen Education's (Chinese: 豆神教育) 2026 interim report shows a swing to a net loss, with AI-education revenue surging 664.32% year over year to 114 million yuan — still not enough to reverse the overall loss^[10-129]. "AI revenue soaring" alongside "company-level losses" is the most typical financial shape of the education-AI transition: transforming companies commonly offset shrinking legacy businesses with AI growth, but if the new business's gross-margin structure and scale cannot cover overall costs, the income statement stays red. Judging an education-AI company requires looking past AI revenue growth to overall profitability, customer-acquisition cost, and retention.

10.4.5 Organizational Patterns and Consulting Takeaways

Productizable teaching delivery and in-person scenarios must be managed in separate accounts. Steps that can be productized — lecturing, grading, Q&A, administration — should be AI-ized first with measured gains; in-person steps — motivation, companionship, social connection, teacher-student trust — stay with humans. GSU's success lies in embedding AI into the process; Khanmigo's lesson is in treating AI as an "inject-and-it-works" technology.

Teachers are the gatekeeper of organizational change. The resistance is not a technical problem but a problem of meaning: having teachers do "value-add / non-value-add task analysis" and see AI's value to their professional identity works better than tool training; the information gap Ithaka exposed shows that communication and

incentives must precede technology deployment. China institutionalizes teacher literacy training; the EU mandates AI literacy by law — different roads, same destination: teacher capability building is the infrastructure of transformation.

"Supply is not adoption" — adoption engineering matters as much as the product. Free, platform-embedded AI still does not guarantee use; it requires teacher guidance, classroom process, and incentive structures designed as a set. Khan Academy's experimental culture proves that results come from continuous iteration, not a single deployment.

Identify the two institutional paths and design consulting approaches separately. China is policy-driven: the levers are action plans, platform scaling, and case-library diffusion, with the focus on compliant landing and regional scaling. The West is institution-driven piloting: the levers are teacher adoption, evidence of effectiveness, and governance frameworks, with the focus on single-school change design and RCT-style evidence production.

Treat evidence of effectiveness in tiers. Official claims (the Khan experiments, Squirrel AI's TIME selection) and third-party RCTs (GSU, the Khanmigo physics course) must be placed side by side and cross-checked; beware publication bias and vendor marketing materials; and when attributing outcomes, put "pedagogical design" before "model capability" — with the same model, the embedding determines the result.

10.5 Retail & Consumer Goods

10.5.1 Industry Profile and Bottleneck Triage

In 2021, IKEA (the Ingka Group) launched a customer service chatbot named Billie. Over three years it handled 3.2 million interactions, resolved about 47% of inquiries, and saved roughly €13 million in customer service costs^[10-49] [third-party verified]. So far, this is an ordinary "AI cost-cutting" story. But IKEA did two unusual things. First, it did not lay off the roughly 8,500 customer service staff displaced by AI. Second, it carefully analyzed the remaining 53% of inquiries Billie could not resolve and found that what customers actually wanted to ask was "will this sofa fit in my home" — interior design consultation, a genuine need that had been long overlooked. The 8,500 customer service staff were retrained as remote interior design consultants and generated €1.3 billion in revenue in their first full year [third-party verified].

This story almost defines the right posture for retail AI transformation. Retail is an industry with "dual information and physical bottlenecks," at a ratio of about 60:40 — half information business (demand forecasting, assortment, marketing, membership operations), half physical business (stores, shelves, warehouses, cold chains, a million-plus employees). Bain's data shows that leading retailers using AI for personalized marketing saw ROAS improve 10%–25% in early trials^[10-59] [third-party verified]; but the physical-side constraints are just as real — when Amazon closed Amazon Go, it admitted that logistics and technology excellence alone were "far from enough; you also need merchandising and display skills."^[10-52]

The industry has three core contradictions. The online information flow and the offline store's physicality have long run as two separate systems: what consumers see online and what the store carries are two different worlds — and offline fulfillment remains the decisive factor for both cost and experience. Supply-chain efficiency and experience pull against each other: SHEIN compressed the industry's average inventory rate of about 30% down to single digits^[10-55], while McDonald's AI ordering stumbled because "efficiency was not saved, and the experience collapsed first." Algorithmic site selection and organizational inertia are fighting for power: more than 70% of Luckin's new store locations are recommended by algorithms, yet most chain enterprises still expand on the strength of their development teams' "experience plus relationships" — the handover of power between algorithm and organization is the hardest bone in retail AI transformation. The China Chain Store & Franchise Association's consensus is that "organizational change is the biggest challenge for AI landing"^[10-60] [third-party verified].

10.5.2 Transformation Paths and Typical Practices

Retail AI transformation can be split into at least four paths.

The international giants take the "human-machine collaboration" route. Walmart deployed AI tools (real-time translation, task management) for about 1.5 million store employees in the United States and rewrote 850 million product-catalog entries with generative AI^[10-57] — building AI as a "productivity exoskeleton" for frontline staff rather than a reason for layoffs [official claims]. Carrefour was the earliest to embed generative AI across the full chain: its Hopla shopping assistant covers all online channels, and more than 2,000 products use AI-generated copy [official claims]. Starbucks represents the "AI-augmented organization": Deep Brew works quietly in the back office — scheduling, stocking, predictive maintenance — freeing people for

interpersonal service [third-party verified]. Amazon pushes AI into the physical layer of fulfillment: more than 1 million robots are dispatched by generative AI [third-party verified].

The essence of supply-chain AI is "trading data for certainty." SHEIN's "small orders, fast response" model pushed its inventory rate down to single digits, with 2025 revenue of US\$41.8 billion and nearly 600 supplier training sessions covering 37,000 person-times^[10-55] [third-party verified]; Inditex plans to invest €1.8 billion in technology and logistics, and AI demand forecasting already drives 85% of initial production-allocation decisions^[10-56] [third-party verified]; UNIQLO's "Ariake Project" (Chinese: 有明计划) has invested about ¥100 billion in RFID item-level tracking, with AI forecasts wired directly to factories [third-party verified].

Chinese chain enterprises take the route of "changing the organization first, then rebuilding efficiency." In April 2026, Hema (Chinese: 盒马) merged 19 regions into 9 major regions and set up an AI committee with AI BP roles^[10-51] [third-party verified]; Sam's Club, by contrast, split its regions in the opposite direction — organizational structure must match the business stage, and the directions can legitimately be opposite [third-party verified]. After DCP Capital took control, RT-Mart (Chinese: 大润发) "returned to the essence of retail" and turned profitable in FY2025 (net profit of 386 million yuan, versus a loss of 1.668 billion yuan the prior year) [third-party verified].

Consumer brands put AI to work on the business model itself. Mixue's (Chinese: 蜜雪冰城) "Snow King AI Brain" (Chinese: 雪王 AI 大脑) turns AI into a cost-control system for its value-pricing strategy — with a gross margin of only about 35%, technology is not for show but a survival necessity [third-party verified]; Pop Mart's self-developed AI integration platform, HeyLisa, uses AI sales forecasting to convert "hit-product uncertainty" into a predictable supply chain [third-party verified]; 7-Eleven Japan is the template of "data first, intelligence second" — decades of item-level management have built the data assets that make AI demand forecasting possible [third-party verified].

10.5.3 Benchmark Cases

IKEA's Billie: the unsolved problems AI reveals are worth more than the ones it solves. [third-party verified] Back to the opening story. IKEA's pivotal move: it did not cut the customer service staff displaced by AI; instead, it followed the questions AI could not answer to discover "interior design consulting" — an overlooked

market — and retrained 8,500 customer service staff as remote interior design consultants, who generated €1.3 billion in revenue in their first full year (about 3.3% of group revenue)^[10-49] [third-party verified]. This is the most important positive sample of AI-native organizational transformation: the unsolved problems AI reveals are worth more than the problems it solves, and the freed-up workforce is redeployed into higher-value work — organizational transformation is a technology decision and a talent decision happening at the same time.

Luckin: algorithmic site selection and the "Smart Brew Ecosystem." [official presentation; data not independently verified] As of May 31, 2026, Luckin had more than 35,000 stores worldwide and 473 million cumulative transaction customers. Three hard metrics stand out: more than 70% of new store locations are recommended by algorithms, and the average cycle from application to acceptance is 18 days; IoT-connected devices exceed 700,000 with a 99% online rate; and intelligent ordering has doubled efficiency. Its self-developed agent, AI Lucky, has been integrated into super-apps such as Qwen (Chinese: 千问), WeChat, and Alipay, and supports voice ordering. Luckin has also joined with multiple cloud vendors to build the "Smart Brew Ecosystem"^[10-50] (Chinese: 智萃生态圈), upgrading AI from point tools to platform capability.

Hema: the governance pairing of an AI committee with a region-based structure. [third-party verified] In April 2026, Hema completed its organizational upgrade from a city-based to a region-based structure, consolidating 19 regions into 9; it simultaneously set up an AI committee led by the CTO and created AI BP roles^[10-51], connecting AI capability directly with business needs. The backdrop: nine consecutive months of overall profitability in 2024 and the first full-year positive adjusted EBITA in FY2025 (per EO Intelligence's review). "Committee + BP roles" is a common mechanism for linking the technology center with business units — moving governance structure and organizational structure in sync is the biggest difference between Hema and the many retailers that "only install tools without moving the organization."

Bosideng: AI enters design; people transform first. [third-party verified] Bosideng (Chinese: 波司登) has embedded AI into the front end of "trend insight → design generation → rapid sampling": in 2024, eight AI-generated patterns and nine AI-designed garments reached the market, and the first product fully landed through AI design — a single-shell jacket — was well received^[10-130]. The deeper organizational work sits in the supply chain: since 2018 its "one pool of global goods" (Chinese: 全球

一盘货) consolidated more than 100 warehouses nationwide into a single central warehouse, using a "merchandise bank" mechanism that removes ownership tags and replenishes dynamically by sales; after the SAP EWM upgrade, real-time inventory updates and complex analytics ran 1,000x faster, inventory turnover fell from 175 days (FY2020/21) to 115 days (FY2023/24), warehouse automation exceeds 90%, and the plan is for a 10,000-square-meter warehouse to need only one person^[10-131]. Its logistics director's judgment is worth remembering: technology can be bought and outsourced, but brand philosophy and mindsets cannot — transformation starts with changing people, not with going live on systems.

10.5.4 Failure Warnings

Amazon Go/Fresh closures: being technologically ahead is not the same as having a working business model. [third-party verified] In January 2026, Amazon announced the closure of all 58 Fresh supermarkets and 14 Go stores, accompanied by 16,000 layoffs^[10-52], admitting it had "not yet created a sufficiently distinctive customer experience and the right economic model." Go's Just Walk Out "grab-and-go" technology was dazzling, but it could not hide mediocre merchandise; tellingly, Just Walk Out succeeded when sold as a technology to 360+ third-party locations. Retail AI must solve both "differentiated experience and unit economics" at once — technology inside the wrong business model is merely expensive decoration.

McDonald's × IBM voice ordering: the triple threshold of accuracy, cost, and experience. [third-party verified] In June 2024, McDonald's terminated the AI voice-ordering system it had been testing with IBM since 2021 across 100+ stores. Analysts pointed out that accuracy reached only 80–85%, below the 95% threshold^[10-53], and operating costs were higher than human labor — and the viral TikTok video of "AI getting a McNuggets order wrong" amplified the experience collapse. Retail AI must clear the triple threshold of "accuracy, cost, and experience"; the gap between a technology demo and a real operating environment is far larger than imagined.

Yonghui learning from Pang Dong Lai: easy to copy the form, hard to copy the spirit. [third-party verified] After more than two years of remodeling, Yonghui's (Chinese: 永辉) remodeled stores saw average single-store foot traffic grow over 80%, yet the company still posted a 2025 net loss attributable to shareholders of 2.55 billion yuan, with cumulative losses from 2021 to 2025 approaching 12.05 billion yuan^[10-54]. The comparison is stark: Pang Dong Lai (Chinese: 胖东来) employees earn

an average monthly salary above 9,500 yuan with a turnover rate of only 0.33%, while Yonghui's remodeled-store staff earn just 4,500–6,000 yuan a month — 65% of Pang Dong Lai's pay for the same roles — with no profit-sharing mechanism in place. Pang Dong Lai's essence is a distinctive ordering of scale, profit, and employee relations, and its regional supply-chain density cannot be replicated by national chains. The contrast between "whole-company statements and store-level results" is fire and ice — one of the most profound lessons in organizational transformation.

10.5.5 Organizational Patterns and Consulting Takeaways

The AI-native transformation of retail is, at bottom, the re-orchestration of two systems: the "information flow" and the "physical flow." Online-offline integration is not about moving products onto the internet; it is about letting demand signals (online searches, membership data) and supply actions (store restocking, production scheduling) share the same data foundation. A consultant's first move should be to take stock of whether item-level data assets exist and whether they are connected — 7-Eleven proves that AI's moat is not the model but the data-organization discipline of "item-level management."

Supply-chain data foundations set the ceiling for AI. SHEIN's small-order fast response, Inditex's 85% of production allocation decided by AI, and UNIQLO's RFID full chain all point to the same thing: the effectiveness of supply-chain AI depends on how deeply the industry chain is digitalized, not on model parameter counts. Retail AI is an "asset-heavy" transformation — you trade data infrastructure for certainty in inventory and cash flow.

Store employees' roles must be rebuilt, not eliminated. IKEA retrained 8,500 customer service staff, Walmart handed AI tools to 1.5 million employees, and 7-Eleven Japan uses AI avatars to cope with labor shortages — all point toward a "human-machine hybrid" organizational form. Layoff-style AI transformation in retail directly destroys the physical-side asset of store experience.

Consumer trust is the hidden ceiling of retail AI. Walmart reports that consumer trust in AI recommendations (27%) has already overtaken influencer-driven shopping (24%)^[10-58], but gaps in privacy protection can void the promises of transformation [official claims].

Organizational structure must match the business stage and digital maturity. Hema merging regions (an expansion phase seeking scale) and Sam's Club splitting regions (a high-growth phase seeking flexibility) moved in opposite directions and both worked; Yonghui's "easy to copy the form, hard to copy the spirit" shows that copying a benchmark's surface practices without simultaneously adjusting interest-distribution mechanisms leaves transformation stranded in isolated corners of the financial statements. The landing point of retail AI consulting is never model selection; it is the simultaneous re-arrangement of the data foundation, the organizational structure, and the interest mechanisms.

Ownership of the dialogue layer is becoming the new battleground. U.S. retailers Kroger and Lowe's, together with Papa Johns and others, are testing their own AI-agent stacks rather than depending on a single third-party chatbot platform — the rationale is to keep the "find product → place order → after-sales" conversational entry point and first-party data control in their own hands, while switching across multiple LLM vendors to manage cost and risk^[10-132] [third-party verified]. The ownership contest over retail entry points is extending from "shelves and apps" to the "dialogue layer": whoever owns the conversation with the customer owns the next generation of shelf space.

10.6 Government & Public Services

10.6.1 Industry Profile and Bottleneck Triage

In February 2025, 70 "digital-intelligence workers" (Chinese: 数智员工) began their jobs in Futian District, Shenzhen. Built as AI digital humans that fuse DeepSeek with local knowledge bases, they cover 11 major scenario categories across 240 government scenario terminals: document-formatting correction accuracy exceeds 95%, review time is cut by 90%, routing accuracy for public livelihood requests improved from 70% to 95%, and personalized custom generation was compressed from 5 days to minutes^[10-62] [official claims + third-party verified]. This is the landmark implementation of the "AI civil servant" concept, and it pushed government AI from demonstration toward daily operations.

Government and public services is one of the few industries where the "information and physical bottlenecks" are split roughly fifty-fifty: approval processing, policy consultation, document drafting, and service-request routing are information-intensive links, while counter services, on-site law enforcement, grid patrols, and

urban infrastructure operations remain firmly bound to the physical world — in sharp contrast to finance's 80:20 information-dominant structure. This determines that government AI transformation must be a dual-track overhaul of "online intelligence + offline coordination," rather than a purely digital reconstruction.

The core contradiction has three layers. First, the paradox of data security and sharing: government data is scattered across departments, with unclear powers and responsibilities and poor circulation — "data silos" have long been the No. 1 pain point; the state's creation of a national data administration is precisely aimed at breaking this deadlock^[10-68], yet local practice generally finds "consolidation easy, coordination hard" — data-management responsibilities have not been synchronized with powers over project approval, funding, and staffing. Second, digital formalism: the deployment guidelines issued by the Cyberspace Administration of China (CAC) singled out for criticism "blindly chasing technological superiority, duplicate construction, and forced usage," and People's Daily Online likewise warns against bandwagon construction "for the sake of using AI." Third, the decision-accountability chain: once AI-generated official documents, replies, or handling recommendations go wrong, there are no mature institutions for responsibility assignment, human fallback, or algorithmic transparency. The US federal government faces a mirror-image dilemma — the Brookings Institution found its AI use cases highly concentrated in a few large agencies, with uneven diffusion, and procurement and budget mechanisms constituting structural barriers^[10-70] [third-party verified].

10.6.2 Transformation Paths and Typical Practices

Chinese government AI transformation shows a clear "policy-chain-driven" character: "AI+" was written into the Government Work Report in 2024; in August 2025 the State Council issued the *Opinions on Deepening the Implementation of the "AI+" Action*; and in October of the same year the CAC released the *Guidelines for the Deployment and Application of AI Large Models in the Government Sector*^[10-61] — completing the chain of "top-level design — special action — scenario guidelines."

The core of the deployment route is "centralized, consolidated planning": the guidelines specify that counties and below should in principle reuse higher-level compute and model resources rather than independently building government LLMs, exploring "build in one place, reuse across many regions and departments" to prevent "model silos^[10-61]." At the same time, they delineate four major scenarios — government services, social governance, agency office work, and decision support —

and emphasize the LLM's "assistive" positioning and the discipline that "classified information never goes online." DeepSeek's open-sourcing, with its low cost, high performance, and easy deployment, dramatically lowered the threshold for adoption, and domestic adaptation (Ascend chips, all-in-one appliances) became the standard.

At the organizational level, the National Data Administration (Chinese: 国家数据局, established in October 2023) has become the "organizational hub" of digital transformation: 2024 was positioned as the "year of institution building," shifting to the "year of reform breakthroughs" in 2025, with 23 provincial-level data administrations formed by the end of 2024. In terms of implementation forms, the "City Brain" has been elevated from a technology platform to a city-level governance mechanism reconstruction — Chongqing, relying on its integrated intelligent public data platform, automatically generated response plans such as police deployment and subway flow control through 5G + AI video analytics when daily footfall in the Jiefangbei business district exceeded 1.8 million^[10-71] [third-party verified]. On the grassroots governance side, Chongqing's "141" system embeds AI across 1,031 townships and streets and 62,000 grids; the 12345 hotline has become the densest entry point for agents, with Guangzhou, Suzhou, and Jinan all deploying full chains of "smart form-filling — knowledge recommendation — smart dispatch."

10.6.3 Benchmark Cases

Shenzhen Futian: 70 AI "Digital Workers" [official claims + third-party verified]

The first batch of 70 digital workers is built on DeepSeek fused with local knowledge bases, covering 11 major scenario categories and 240 government scenario terminals^[10-63], forming a closed loop of "demand — training — scenario application — iteration," supported by a unified management platform that handles full-lifecycle management of these digital workers. Its significance lies not in the technology but in turning the "AI civil servant" from a concept into an operable, day-to-day reality.

Shenzhen's "Shen Xiao-i" (Chinese: 深小i): A Public-Facing Government Agent [official claims]

Launched on February 22, 2025, it is the country's first practical AI government assistant open to the public, embedded in the full "ask, search, guide, process, review" service flow. It has cumulatively handled over 4 million intelligent Q&A

sessions, with a response rate above 97% and first-answer accuracy above 94%, and is one of the first 7 government agents to pass evaluation by CAICT's TEL Laboratory (Chinese: 信通院泰尔实验室). Behind it is Shenzhen's AI platform: managing 270P of compute, 39 deployed models, and 97.38 million annual calls^[10-64], with 10 scenario challenges organized through open-competition recruitment (Chinese: 揭榜挂帅) — platform-based supply plus scenario-based open challenges brings it closer to "AI-native" than single-point applications.

Guangzhou: China's First Domestically Adapted Government DeepSeek [official / third-party verified]

Adaptation was completed and rolled out broadly in January 2025: the 12345 hotline answers 50,000 calls daily, and the intelligent agent-seat system has cut caller waiting time by 43% and achieved 97% accuracy in work-order routing^[10-65]; the Digital Guangzhou Innovation Laboratory has adapted over 150 LLMs and opened the country's first government-grade compute center running the Ascend-adapted DeepSeek-R1 671B; the Guangzhou Intermediate People's Court deployed the country's first private government DeepSeek all-in-one appliance; Yuexiu District's "Ai Xiaoyue" (Chinese: 艾小越) achieves 99% answer accuracy. Guangzhou has proposed "one ledger, one mechanism, one platform" for LLMs, building a two-way loop in which "data feeds AI, and AI feeds decisions back."

10.6.4 Failure Warnings

AI Hallucination: Errors in Government Are Irreversible [third-party verified]

Outlook Weekly (Chinese: 瞭望) dissects three causes of hallucination: insufficient or biased training data, algorithmic architecture limitations, and the probabilistic nature of generation^[10-67]. LLMs "confidently spouting nonsense" can cause erroneous decisions and damage to public trust in highly regulated domains such as government and the judiciary — once a mistaken policy interpretation reaches the public, it constitutes real harm. CAICT experts point out that hallucination governance requires both "technology and institutions": retrieval augmentation and fact-checking on the technical side, and content review, risk warnings, and accountability tracing on the institutional side. This is precisely why the CAC guidelines require prominent risk warnings in the interface.

The GAO Audit: Governance Lagging Behind Deployment [third-party verified]

In its July 2025 review of 12 federal agencies, the US Government Accountability Office (GAO) found incomplete governance frameworks, use-case inventories not updated in time, and inadequate risk management for generative-AI-specific issues such as hallucination, data leakage, and copyright^[10-66]. US federal AI use cases grew fivefold in five years to 3,611, yet governance cannot keep pace with the speed of diffusion. The GAO's framework of "use-case inventory + risk classification + accountability auditing" suggests: build governance first, then talk about scale.

Digital Formalism: Advanced Technology, Lagging Experience [third-party verified]

The CAC guidelines explicitly require guarding against digital formalism; People's Daily Online criticizes bandwagon construction "for the sake of using AI"^[10-69]; the magazine *Qunzhong (Mass)* (Chinese: 群众) warns of "advanced technology, lagging experience." The more hidden failure is organizational: mismatched powers and responsibilities of local data administrations leave coordination spinning in place. Government AI usually fails not because the models are inadequate, but because governance and organization have not kept up.

A stampede to deploy DeepSeek carries four risks. [third-party / think-tank analysis] People Zhongke Research Institute (Chinese: 人民中科研院) argues that simply deploying DeepSeek or similar large models without systematic planning exposes local governments to four risks: resource waste and efficiency imbalance (duplicated construction without unified planning), data-security and compliance hazards, unguaranteed content accuracy, and application-business disconnect producing "AI for AI's sake"^[10-133]. Government large-model deployment should shift from a "model procurement race" to integrated construction across scenarios, data, and institutions — build the governance framework before launching applications; performance-driven rush deployment risks creating a new form of digital formalism.

10.6.5 Organizational Patterns and Consulting Takeaways

Institutions first determine the order of deployment. The rollout sequence of government AI is set by policy and security boundaries rather than technical difficulty: the CAC guidelines are both a constraint and a roadmap, and its three principles — scenario classification, centralized consolidated planning, and the assistive role — should be the preconditions for initiating any government AI project.

The assistive role is both politically correct and engineering-correct. The stable form of government AI is the "digital colleague" rather than the "digital replacement": repetitive, standardized work goes to AI, while value judgments, emotional communication, and discretionary authority stay with people, with "human fallback" as the final line of defense against hallucination risk.

The data administration is the real lever of organizational change. Government AI transformation is essentially a "data-administration-driven organizational rearrangement": the restructuring of tiao-kuai (vertical-line and horizontal-block) relations across national, provincial, and city/county levels, paired with organizational instruments such as leading groups, standardization technical committees, big-data professional-title review, and zero-based budgeting, determines success or failure more than any single technology procurement.

Grassroots burden reduction must be measured. Shenzhen Nanshan's AI identified 87,000 public requests, improved key-information accuracy by 30%, and captured 820 weak risk signals — burden reduction backed by data is the only real "burden reduction," while forced usage and duplicate form-filling are a new form of formalism. The first consulting move should be helping clients build the governance foundation of "use-case inventory + risk classification + accountability auditing" (modeled on the GAO framework and CAICT's case-evaluation mechanism), and unifying value metrics such as "number of scenarios, hours saved, accuracy" — otherwise government AI will repeat the failure of "statistical prosperity, frontline indifference."

The "small but deep" in-house product route: Singapore's Pair suite. [third-party / international-organization case] UNDP's review of Singapore's government AI tool suite Pair describes an approach that inverts the big-bidder playbook: developed since mid-2023 by the internal experimental team OGP rather than outsourced to a large systems integrator, with all data retained inside government and never logged by cloud vendors, and each product defining its

evaluation criteria before any new model goes live. Pair now has 60,000+ registered users and 20,000+ weekly actives, has exchanged over 10 million messages, and is estimated to save civil servants about 46% of administrative time — Pair Noms compresses meeting minutes that once took hours to a full day into under an hour^[10-134]. Compared with "big and complete" government-model platforms, "small but deep" in-house products let AI penetrate civil servants' daily work in the form of an office tool before extending into government scenarios — the adoption path and procurement logic, more than the model itself, decide success.

10.7 Logistics & Supply Chain

10.7.1 Industry Profile and Bottleneck Triage

On May 22, 2025, the automated terminal at Qingdao Port (Chinese: 青岛港) pushed its average per-quay-crane operating efficiency to 62.62 natural containers per hour — the 13th time it has broken its own world record for automated-terminal handling efficiency^[10-73] [official claims + third-party verified]. The terminal was approved in 2013 and construction began in 2015; ten years of continuous iteration have taken it from an "unmanned terminal" to "whole-domain intelligence." It is a reminder to everyone: AI in logistics is a heavy-asset, long-cycle systems-engineering effort — not something that shows results just by installing software.

Logistics and supply chain have an information-to-physical bottleneck ratio of roughly 40:60, the mirror image of finance's 80:20. Beneath the heavy-asset surface, what actually drives the entire network is decision information — forecasting, scheduling, routing, and risk control. Data from the China Federation of Logistics & Purchasing (CFLP) (Chinese: 中物联) shows China's logistics industry generated total revenue of 13.2 trillion yuan in 2023; industry estimates put the global AI-enabled smart logistics market above US\$42.8 billion in 2025, with over 1.5 million warehouse robots deployed [third-party verified].

Logistics AI transformation faces three structural contradictions. Automation is a heavy-asset investment, yet demand is highly volatile: an automated terminal costs billions of yuan and takes years to build, while demand fluctuates on a weekly basis — capital expenditure and capacity utilization are structurally mismatched. Physical-world right-of-way and compliance are hard constraints: autonomous delivery vehicles have "already reached a fleet scale of tens of thousands, yet are still waiting for a single national pass"^[10-79] [third-party verified], and fragmented right-of-way policy

directly draws the commercial boundary of driverless capacity. Technical maturity is not commercial maturity: L4 autonomous trucks have been explored for over a decade without cracking the unit-economics model, and drone delivery runs order volume and losses side by side. Logistics does not lack AI pilots — it lacks the organizational capability and commercial closed loop to turn pilots into systems.

10.7.2 Transformation Paths and Typical Practices

Logistics AI implementation unfolds along five paths, with value realization on the information layer generally faster than on the physical layer.

Express-delivery logistics LLMs are the main battlefield of the top players, which have formed a combined playbook of "vertical LLM + agents + full-chain scenarios" [third-party verified]: Cainiao (Chinese: 菜鸟) launched its logistics customer-service AI brain in 2018; 2023 saw the dense release of "Tianji π" (Chinese: 天机π), JD's (Chinese: 京东) "Super Brain" (Chinese: 超脑), and YTO Express's (Chinese: 圆通) "Zhiduoxing" (Chinese: 智多星); in 2024 SF Express (Chinese: 顺丰) released its "Fengzhi" (Chinese: 丰知) decision-making LLM^[10-75]; and in 2025 JD upgraded to "Super Brain LLM 2.0" [official claims]. The strategic divide is clear: direct-operated players (SF Express, JD) self-develop full stacks, while franchise networks use AI tools to penetrate the "headquarters-outlet-courier" collaboration chain — YTO has embedded over 100 agents across its full chain^[10-77], with AI roles going through a three-stage upgrade from "agent to AI assistant to digital employee" [official claims + third-party verified].

In warehouse robotics, Amazon announced in 2025 that its global robot deployment surpassed 1 million units, about two-thirds of the world total^[10-74] [official claims]; Walmart is deeply bound to Symbotic, spreading AI automation across all 42 regional distribution centers within eight years; UPS aims to raise the share of US packages processed by automation from 66.5% to 68% by the end of 2026 [third-party verified]. Even the giants are advancing incrementally, percentage point by percentage point — transformation is a gradual engineering process, not a one-step "lights-out warehouse."

Smart ports are China's most brilliant logistics calling card. Qingdao Port has built a "smart brain" on its fully domestic, fully autonomous A-TOS/A-ECS systems, with 85 AGVs and 47 yard rail-mounted gantry cranes running unmanned end-to-end^[10-73]; Yangshan Phase IV (Chinese: 洋山四期) operates on "F5G ultra-remote control + AGV

automated handling," and the Port of Shanghai (Chinese: 上海港) handled its 50 millionth TEU in November 2025, ranking first globally for 16 consecutive years [third-party verified].

Autonomous delivery advances between scale and losses. In 2025 the industry delivered over 35,000 autonomous vehicles^[10-78], and JD Logistics announced plans to procure 1 million over the next five years [third-party verified]; Meituan's (Chinese: 美团) drone delivery has completed over 670,000 commercial orders^[10-82] [official claims]; SF's Fengyi "Fangzhou" drone delivers over 20,000 packages per day on average. Digital freight platforms are another kind of story: Full Truck Alliance (Chinese: 满帮) fulfilled over 236 million platform orders in 2025, compressing truck-load matching from hours to minutes; Lalamove (Chinese: 货拉拉) drivers post a 95% on-time rate and a 92% load-matching rate [third-party verified].

Port autonomous driving has already reached industrial scale: as of June 2026, cumulative shipments of autonomous horizontal-transport vehicles for Chinese port operators reached about 2,700 units, with roughly 85% delivered to domestic ports; nearly 40 port areas run batch deployments (10+ units), of which 15 run scale deployments (50+ units). In 2025 alone, shipments were about 850 units (up 70% year over year) with sales of about 1.62 billion yuan^[10-135] [third-party verified]. The blue book estimates the global market potential at about 38 billion yuan once port-level autonomous penetration reaches 30%, and over 54 billion yuan at 60% — ports have become the most commercially mature scenario for L4 autonomous driving.

10.7.3 Benchmark Cases

Qingdao Port: An AI-Native Sample of Physical Heavy Assets [official claims + third-party verified]

Back to the world record at the start. The core is the self-developed A-TOS/A-ECS systems working in concert to form the "smart brain," upgraded with AI large models for algorithms and compute^[10-73]. AI-nativeness in ports is a systems-engineering effort of "self-developed foundation + continuous iteration + organizational support," not single-point equipment procurement — this is the most essential difference between physical-bottleneck industries and information industries like finance and media.

Amazon: Robot Army + Central AI Brain [official claims; limited independent verification]

Starting with the acquisition of Kiva in 2012, Amazon expanded over 13 years to 300-plus warehouses and 1 million robots, becoming in 2025 the "world's largest mobile robot manufacturer and operator." Its DeepFleet generative AI dispatch system coordinates robot movement in real time like a "smart traffic management system," reportedly improving fleet operating efficiency by 10% and overall logistics efficiency by 25%^[10-74] [official claims]. What deserves attention is not the numbers but the organizational logic: self-development plus acquisitions built the two-tier architecture of "robot army + central AI brain," and the strategic substance is human-robot collaboration and job-skill reshaping.

Full Truck Alliance: An Agent Upgrade for a Digital-Native Species [third-party verified + company disclosures]

The highway freight platform is a natively digital species for which "data is the business," and its evolutionary ladder — information matching (1.0), algorithmic matching (2.0), decision-proxy agents (3.0) — offers a reference template for all platform companies. Full Truck Alliance fulfilled over 236 million platform orders in 2025^[10-84]; in 2026 it is embedding agents into real trading-decision scenarios, launching a "cargo-finding assistant" and an "intelligent assistant," and responding to fairness governance with algorithmic disclosure [third-party verified].

10.7.4 Failure Warnings

TuSimple: The Collapse of a Technology Leader's Commercial Closed Loop [third-party verified]

The "world's first autonomous-driving stock" once had a market value of US\$16 billion; it delisted in 2024, and its Guangzhou team was disbanded and liquidated in February 2025^[10-80]. Founder Chen Mo admitted that "from 2021 to now, the resources at TuSimple's command could no longer complete the commercialization of autonomous driving." Three root causes: insufficient technical maturity, regulation not ready, and unit economics that never worked — the labor costs saved could not offset sensor and safety-driver costs; layered on top was governance failure from management infighting and strategic flip-flopping. An AI-native organization needs not only technical capability but also a sustainable commercial closed loop and stable organizational governance.

A 95% Pilot Failure Rate [third-party verified]

project⁴⁴, citing MIT research, notes that 95% of enterprise AI pilots fail to produce measurable returns^[10-81] (global AI spending exceeded US\$420 billion in 2025). Its conclusion: "it's not a model problem, it's an execution problem" — poor workflow integration, agents lacking operational context, fragmented tools, and unclear accountability. In supply chain the cost is more direct: the handling window for a shipping exception is only 2 hours; once it closes, the enterprise can only choose among expedited premium shipping, damaged customer relationships, and squeezed margins. The root cause of pilot failure lies not in the algorithms but in process integration and accountability mechanisms.

Autonomous Delivery: Scale and Losses Side by Side [official claims + third-party verified]

Meituan's commercial drone orders grew from 670,000 to 900,000^[10-82], yet Meituan's new-business segment posted an operating loss of 10.1 billion yuan in 2025 (drones included); the autonomous-vehicle industry delivered over 35,000 units in 2025^[10-78], with a market space estimated at nearly 500 billion yuan, yet it is mired in a "selling volume at a loss" price war and fragmented right-of-way. Technical validation is complete, but commercial validation is not — "technical maturity ≠ commercial maturity" is the most expensive tuition logistics AI has paid.

10.7.5 Organizational Patterns and Consulting Takeaways

The physical bottleneck determines the order of value realization. Logistics' 40:60 information-to-physical structure means AI value on the "information layer" — forecasting, scheduling, matching — pays off fast, while "physical layer" automation is a heavy-asset, long-cycle engineering effort. The first consulting move is to clarify which layer the money goes into: the information layer runs on algorithms and data, the physical layer runs on capital and supply chains — the rhythms are completely different.

The data-driven scheduling organization is the core form. Agents are beginning to take on cross-departmental workflows, and organizations are moving from "people + systems" to "people + agent collaboration^[10-76]" [third-party verified], with AI roles evolving along "tool → assistant → digital employee." The key proposition for franchise networks is distributed collaboration across "headquarters-outlet-

courier," and AI products must lower the threshold to sink down, adapting to the reality of franchisees' uneven digital capabilities — the value anchor of franchise-network AI is organizational collaboration efficiency, not merely replacing labor.

Self-development, acquisition, and supplier binding: three clearly separated routes. Amazon self-develops and acquires; Walmart "spun off self-development and outsourced operations" (selling its robotics business in exchange for over US\$5 billion in long-term orders); Cainiao productizes logistics technology for external export (800+ projects globally, 26 Fortune Global 500 clients). Whether an organization must self-develop AI depends on the scale of its data assets and capital strength — "owning AI" and "running an AI-enabled operation" are two equally legitimate choices.

Commercial validation should come first. McKinsey's State of AI 2025 shows 88% of organizations have routinized AI use, yet only 39% achieve significant value and 6% become high-performance winners^[10-83] [third-party verified]; the common root of both TuSimple's collapse and drone-delivery losses is large-scale investment before the unit-economics model was validated. Pilots should be tied to business KPIs, not technical metrics.

Talent and the industry foundation come first. The focus of supply-chain competition is shifting from "buying systems" to "building capabilities" [third-party verified], generating new roles such as AI operations, data annotation, and human-robot collaboration supervisors. At the same time, AI-nativeness in logistics is a cross-entity systems-engineering effort — September 2025 saw the release of the first national standard for logistics enterprise digitalization, and the Digital Ship White Paper (COSCO Shipping + China Classification Society) is driving data interconnection across ships, ports, cargo, and people; an industry-level digital foundation is a prerequisite that no single enterprise can complete alone.

Agents are entering the "execution layer" — and what you are really buying is authorization. AI-agent company DeepFabric has signed third-party logistics operator Kenco as a customer with a target of deploying 20 supply-chain agents within 12 months — not a single-chatbot pilot but a portfolio of automated multi-workflow operations^[10-136] [third-party verified]. Logistics AI is moving from the "dashboard" (reporting, forecasting, alerting) toward "execution": ingesting inbound events, making decisions autonomously, and writing updates back into core systems, with deployments reaching deep into the TMS and last-mile fulfillment-orchestration

layers. The real constraints on such deployments are data permissions, integration scope, and exception-handover rights — when a vendor sells you "agents," the hidden purchase is authorization: what the software is allowed to change without asking.

10.8 Healthcare

10.8.1 Industry Profile and Bottleneck Triage

Between October 2023 and December 2024, TPMG (The Permanente Medical Group) — part of Kaiser Permanente, one of America's largest healthcare groups — did something big: across 17 medical centers in Northern California, it deployed "ambient AI clinical documentation" (Chinese: 环境式 AI 病历) for 7,260 physicians — AI that listens to the doctor-patient conversation and automatically generates the medical record. Across more than 2 million patient visits, physicians collectively saved 15,700 hours of documentation time, and 84% said the AI improved their connection with patients^[10-85] [third-party verified, NEJM Catalyst]. It is one of the largest generative AI deployments in medical history, and a sample of the most correct posture for medical AI: not replacing doctors, but doing the doctors' work for them.

Healthcare is a typical information-physical hybrid industry: diagnosis, medical records, and image interpretation are essentially information processing, but treatment outcomes ultimately land on physical actions — surgery, medication, nursing — putting the information:physical ratio somewhere around 30-50:50-70. This determines that medical AI cannot digitally rewrite the entire process the way finance can; it can only be embedded in an "augmenting people" rather than "replacing people" mode. Grand View Research projects the global healthcare AI market will grow from US\$36.7 billion in 2025 to US\$505.6 billion by 2033 [third-party verified].

But the growth narrative masks four structural contradictions. Regulatory compliance comes first: the FDA has approved over a thousand AI/ML medical devices, yet no generative AI has received FDA approval^[10-93]; in China, Class III medical device registration certificates are the precondition for AI to enter clinical billing — United Imaging (Chinese: 联影) holds 16, Deepwise (Chinese: 深睿) and Shukun (Chinese: 数坤) 14 each, and Infervision (Chinese: 推想) 13 — and the number of certificates directly constitutes a moat [third-party verified]. Second is clinical accountability: once AI enters high-risk operations, liability division is far from mature — US surgical robots caused 1,391 incidents and 144 patient deaths over 13 years^[10-95] [third-party verified]. Third is data privacy: in November 2025 the WHO warned of a serious gap

between medical AI expansion and legal-ethical safeguards [third-party verified]. Fourth is embedding into physicians' workflows: AI alone can reach 90% diagnostic accuracy, but when doctors use it as assistance, accuracy rises only about 2 percentage points [third-party verified] — the gain from human-machine collaboration depends on process design and trust building, not model accuracy.

10.8.2 Transformation Paths and Typical Practices

Three paths are clearly separated.

Chinese hospitals take the route of "open-source models + local deployment + hospital-wide infrastructure." The DeepSeek open-source release in early 2025 significantly lowered the AI threshold for hospitals: Shenzhen Nanshan Hospital (Chinese: 深圳南山医院) deployed DeepSeek-R1 671B locally within 72 hours; The University of Hong Kong-Shenzhen Hospital (Chinese: 港大深圳医院) deployed it on a "data never leaves the hospital" basis; and Peking University Third Hospital (Chinese: 北大三院) combined multiple deployment modes for hospital-wide coverage [third-party verified]. Scenarios are spreading quickly: West China Hospital (Chinese: 华西医院) released two agents built on its "Huaxi Hongyi" LLM (Chinese: 华西黄医) — patient companionship and medical documentation — and the documentation agent has generated nearly 600,000 medical records since launch [third-party verified]; AI has penetrated appointment triage, medical insurance cost control, DRG grouping, and pharmacy review across the entire process. On the budget side, we have entered the "infrastructure era": Zhongshan Hospital's (Chinese: 中山医院) 160 million yuan AI platform tender has landed, and hospital AI budgets have jumped from the millions to the hundreds of millions [third-party verified]. IDC estimates China's AI + healthcare application software market at 3.54 billion yuan in 2025, projected to reach 14 billion yuan by 2030, and judges that the watershed has shifted from "how strong is the model" to "how deeply can business flows be redesigned" [third-party verified].

US hospitals take the route of "administration first, clinical caution, portfolio governance." HCA (191 hospitals) advances in layers: on the clinical side it partners with Google to optimize nurse handoffs; on the operations side, nearly 100 hospitals have deployed scheduling and workforce tools; and administration is treated as the area with the "shortest value path." Its CFO admitted that "rolling out AI incrementally and changing frontline behavior is far harder than imagined," and the investment focus is shifting toward change management [third-party verified]. Mayo Clinic will invest more than US\$1 billion over the next few years across 200-plus initiatives;

Advocate Health uses an "evaluation funnel" (225 proposals → 40 use cases) for centralized selection [third-party verified]. A Menlo Ventures survey of 700-plus US healthcare executives found the industry deploying AI at 2.2x the speed of other industries, with healthcare AI spending of about US\$1.4 billion in 2025, roughly 3x the prior year^[10-92] [third-party verified].

AI drug discovery is reshaping the paradigm. Insilico Medicine (Chinese: 英矽智能) uses generative AI across the entire pipeline from target discovery to clinical development, closing 13+ business-development deals in 2025 with cumulative potential value exceeding US\$10 billion^[10-96] [third-party verified]; its CEO claims AI can compress the drug R&D cycle to about one year (a candidate drug in as little as 9 months, versus about 4.5 years traditionally) [official claims; "candidate drug" ≠ "drug on the market"]. On the regulatory side, the FDA published final guidance on PCCP (Predetermined Change Control Plans) in August 2025, allowing AI algorithms to iterate continuously within predefined ranges without reapplication^[10-93] — this "learning license" system raises the organizational bar for continuous monitoring and change management [third-party verified].

10.8.3 Benchmark Cases

Kaiser Permanente: Ambient Documentation for 7,260 Physicians [third-party verified, NEJM Catalyst]

Back to the deployment at the start. Physicians collectively saved 15,700 hours of documentation time, with after-hours documentation time averaging 1.03 minutes less per visit; 84% of surveyed physicians said AI improved their connection with patients, and 82% reported higher job satisfaction^[10-85]. Combined with Abridge's deployment across 40 Kaiser hospitals and 600-plus clinics, ambient documentation has become the most commercially mature track in medical AI. But the Peterson Health Technology Institute also cautions: financial returns to the system have not yet been validated at scale, and scaling depends on deep EHR integration and safety-monitoring systems.

Zhiyi Assistant: 1.2 Billion Assistance Sessions Across Primary Care [third-party verified / official disclosures]

iFlytek's (Chinese: 科大讯飞) "Zhiyi Assistant" (Chinese: 智医助理) embeds "AI-assisted diagnosis + standardized electronic medical records" into grassroots workflows. As of the end of 2025 it covered over 800 districts and counties across 31

provinces and more than 77,000 grassroots medical institutions, with over 1.2 billion cumulative assistance sessions serving more than 250,000 primary-care physicians^[10-86] [third-party verified figures]; Liu Qingfeng disclosed that Anhui alone corrected 170,000 misdiagnoses in a single year [official claims]. Its organizational form — "province-level coordination, government-funded, system-embedded" — validates a path to grassroots accessibility while exposing its dependence on data quality and government funding.

Beijing Anzhen Hospital: The "Report Revolution" of Operations-Management AI [third-party verified]

Released in March 2026, China's first "public hospital operations-management large model" was developed with China Telecom over seven months, integrating the HIS, SPD, and ODR systems, drafting 200-plus data standards, governing 3 million-plus data records, and establishing 800-plus core management indicators. Operations-management report generation has been cut from days to 5 minutes, management response speed improved by 95%, and precision decision-making efficiency improved by 60%^[10-88]. It represents hospital AI extending from clinical assistance to operational decision-making, rebuilding management's data-decision infrastructure.

Ambient scribes, new evidence: time savings are "modest," burnout drops sharply. [academic research] The first large multicenter study in JAMA Network Open (6 U.S. health systems; 263 outpatient physicians and advanced practice providers; the Abridge platform; 30 days) found the share of doctors reporting burnout fell from 51.9% to 38.8% (adjusted odds of burnout down 74%), with significant reductions in documentation-related cognitive load and after-hours charting time — yet additional patient visits showed no statistically significant difference^[10-137]. A sister study (5 academic medical centers; 8,581 clinicians) likewise found AI ambient documentation tools associated with only "modest" reductions in EHR time and no significant change in off-hours use, with the benefits concentrated among primary-care physicians, advanced practice providers, and women^[10-138]. The first tranche of returns lands in employee experience and retention, not capacity — hospital purchasing decisions should measure "objective time" and "subjective experience" separately.

Meizhong Airui: writing AI into the hospital itself from the design stage. [third-party] Caixin Weekly reports that Beijing's Meizhong Airui Cancer Hospital (Chinese: 美中爱瑞肿瘤医院) is building a new campus — total investment around 6 billion yuan, planned at 800 beds — with the stated goal of becoming China's first "AI-

native" physical hospital: agents bound to no single system or scenario, but pervasive across the building, medical devices, and information systems, covering all key business areas and forming a "collaborative ecosystem" with human staff^[10-139]. This is a rare organizational experiment that writes AI into the hospital's ontology at the design stage — architecture, procurement, information architecture, roles, and processes rebuilt in sync — and its outcome will test whether the "AI-native" concept can hold in a heavily regulated, high-accountability medical environment.

10.8.4 Failure Warnings

HIMSS 2025: 80% of Projects Fail to Realize Value [third-party verified]

A HIMSS panel discussion confronted the industry's pain point directly: 80% of healthcare AI projects fail to realize value conversion; 97% of organizations agree deployment is urgent, but only 14% have effective implementation capability^[10-89]. Its verdict cuts to the chase: "AI failures are not about immature technology but about strategic mistakes" — and the biggest mistake is trying to solve everything at once.

NHS: The Implementation Chill on a National Program [third-party verified]

UCL, with the Nuffield Trust, evaluated NHS England's £21 million AI chest-imaging program across 66 hospital trusts: contract signing was delayed 4-10 months, and 18 months after contract completion, 23 of 66 trusts (a third) still had not used it in clinical practice^[10-90]. The obstacle is not algorithms but organization: overworked clinicians cannot participate, aging heterogeneous IT systems are hard to integrate, and senior doctors worry about accountability for AI-informed decisions.

Hallucination and the Clinical Adoption Gap [third-party verified]

Medical LLMs inherently carry hallucination risk, and the industry consensus is that it "cannot be eliminated entirely, only managed and corrected." AI alone reaches 90% diagnostic accuracy while physician assistance adds only about 2 percentage points — without process design and trust building, accuracy does not automatically translate into treatment outcomes. The 1,391 incidents and 144 deaths from surgical robots over 13 years^[10-91] are a reminder that AI's rollout into high-risk clinical operations is far slower than in administrative scenarios. An HFMA survey shows nearly 70% of CFOs admit governance maturity lags adoption speed, and the resource misallocation of "bought but unused, used but unevaluated, evaluated but never retired" is happening now [third-party verified].

UnitedHealth: when algorithmic decisions lack human review. [third-party verified] The families of two deceased Medicare Advantage beneficiaries have sued UnitedHealth, the largest U.S. health insurer, alleging it knowingly used flawed AI systems (such as nH Predict) to deny post-acute care that physicians deemed necessary — a financial arrangement of "collecting premiums while not paying claims." The complaint argues the insurer systematically denied claims because only about 0.2% of enrollees were expected to appeal AI decisions, and cites earlier class-action materials claiming over 90% of AI denials were overturned on appeal^[10-140]. A federal judge in Minnesota has ordered UnitedHealth to hand over internal AI review-committee documents and related performance and compensation records. When algorithms are embedded in claims decisions without transparency or human review, an efficiency tool metastasizes into systemic compliance and ethical risk — a complete lesson in healthcare-AI governance failure.

10.8.5 Organizational Patterns and Consulting Takeaways

The "AI alert → human confirmation → intervention" organizational loop is the value amplifier of clinical AI. UCSD's COMPOSER cut in-hospital mortality by 17%^[10-94], and Augusta Health, a community hospital, saved 282 lives in two years on the same path — what they share is not a stronger model but embedding alerts into care workflows, training response mechanisms, and continuously calibrating the model. The takeaway for small and mid-sized hospitals: enter through a single high-value scenario and use organizational mechanisms to compensate for resource disadvantages.

Change management is the main bottleneck, not technology. A McKinsey survey shows 50% of healthcare organizations have implemented AI, but concludes that "the main bottleneck is not technology but leadership and organizational change" — corroborated by the NHS research; HCA's CFO put it bluntly: changing frontline behavior is far harder than deploying tools.

Physician-adoption design determines the ceiling. Kaiser's high satisfaction came from embedding AI into existing workflows rather than adding burden; HIMSS emphasizes that "unexplainable AI decisions will lose clinician and patient trust." Consulting takeaways: do workflow analysis before model selection, and build trust from administration to clinic.

Regulation first is a roadmap, not a barrier. Class III certificates, FDA approval, algorithm filing, and data security assessment form a landscape where "the entry barrier is organizational capability": certificate counts become a moat^[10-87], buyers need new validation and acceptance capabilities, and third-party evaluation standards are becoming the common language of all parties.

Data assets, not models, have become the biggest shortfall. The consensus at the VB100 conference: the biggest shortfall is no longer compute and models but standardized, high-quality data assets. The essence of healthcare AI-native transformation is synchronously rearranging four things — clinical workflows, management decisions, regulatory compliance, and data governance. The hardest part technically is keeping AI from overstepping before the physician signs off; the hardest part organizationally is getting AI genuinely used after the signature.

10.9 Manufacturing

10.9.1 Industry Profile and Bottleneck Triage

Let me start with a true story. A textile mill in the Yangtze River Delta (Chinese: 长三角) spent 3.7 million yuan on an AI fabric inspection system — machine vision to spot defects on the cloth surface, theoretically faster and more accurate, and capable of replacing a dozen-odd QC workers. After the system went live, the problem turned out to be the fabric itself: surface reflection sent the machine's misjudgment rate as high as 15%^[10-101]. The boss did not dare let the machine make the call, the veteran inspectors' manual re-checks continued as before, and the AI went from "replacement" to "assistant" — the return on investment fell far short of expectations [third-party verified].

This story almost condenses every contradiction of AI in manufacturing. Of the ten industries written in this chapter, manufacturing is probably the most contradictory one: McKinsey ranks it the area with the greatest value potential from generative AI (up to 2x productivity gains for manufacturing and supply chains, with nearly 70% of tasks automatable), and nearly a quarter of China's roughly US\$2 trillion generative AI value comes from here^[10-99]; yet at the same time, it has the highest physical bottleneck of any industry in this book — an information:physical ratio of only about 20-40:60-80. In plain words: **most of the work on the shop floor is machine cycle**

times, material flows, and equipment wear — things you can see and touch — and no matter how clever the algorithm, it cannot outrun the machine tools.

The gap between big potential and hard implementation comes down to three things. First, AI speeds up the decision layer, not the physical layer — it can help you compute production schedules, review reports, and draft process documents, but it cannot make the machine tools turn faster. Second, industrial data is naturally siloed: equipment protocols each speak their own language, and SCADA, ERP, and MES each form their own systems; the most extreme case is defect samples — hundreds of thousands of photos of good products but only twenty-some defective ones, which still have to be labeled one by one by veteran workers on the verge of retirement. Third, the numbers do not add up: a Deloitte survey shows 91% of industrial AI projects fell short of expectations, and 90% of those failures were not technical problems^[10-100].

10.9.2 Transformation Paths and Typical Practices

AI implementation in manufacturing roughly unfolds across five layers, from the "visible single points" to the "invisible organization."

The first layer is lighthouse factories, solving the demonstration question of "what AI can actually do." China leads the world in lighthouse factories — Haier (Chinese: 海尔) has 13 in total, and Foxconn (Chinese: 富士康) has 5 factories on the network^[10-105]; McKinsey data shows nearly 60% of use cases at newly designated lighthouses already leverage AI, with 2-3x productivity gains, 99% defect-rate reduction, and 30% energy improvement. The real value of lighthouse factories is not the technology but the methodology: Foxconn codifies its experience into white papers and extends toward a "lighthouse industry," while Haier empowers external companies through COSMOPlat (Chinese: 卡奥斯).

The second layer is industrial LLMs, which Huawei (Chinese: 华为) calls the future "operating system" of manufacturing. Industrial LLM penetration in China is projected to reach 50% by 2030, with the market surpassing 42 billion yuan^[10-106]. Players split into two routes: big tech (Baidu, Huawei) and "industry columns" — COSMOPlat's "Tianzhi" (Chinese: 天智), the Digital Building Materials Research Institute's (Chinese: 数智建材研究院) "Xiaomiao" (Chinese: 晓妙), and Zhonggong Hulian's (Chinese: 中工互联) "Zhigong" (Chinese: 智工).

The third layer is industrial agents, pushing AI from "demonstration" to "production." 2025 has been called the "year one of industrial agents" — CCID Consulting (Chinese: 赛迪顾问) data shows China's agent market at 7.84 billion yuan that year, expected to reach 13.53 billion yuan in 2026; 2026 enters the "replication era," where the challenge shifts from technology to engineering.

The fourth layer is the two single-point scenarios with the clearest ROI: AI quality inspection and predictive maintenance. AI vision inspection is the fastest-penetrating scenario (China's smart-inspection output reached 574.5 billion yuan in 2025), with vendors claiming missed-detection rates cut to 1% of manual levels and 70% savings on re-inspection costs [official claims]; on predictive maintenance, IBM, citing ISA data, says factories typically lose 5%-20% of manufacturing capacity to downtime, and AI can cut unplanned downtime by 30%-50% [vendor claims]. But both paths share the same stumbling block: old equipment has no sensors, fault samples are scarce, and ROI is hard to quantify — the textile mill's inspection machine tripped on exactly this path.

The fifth layer is platform-based organizational transformation: SANY (Chinese: 三一) partnered with RootCloud (Chinese: 树根互联), Haier incubated COSMOPlat, and Midea (Chinese: 美的) productized its "Meiqing" (Chinese: 美擎) experience for external export, forming a division of labor where "big players build platforms, SMEs board them." On the institutional side, the Implementation Plan for the "AI + Manufacturing" Special Action issued by the Ministry of Industry and Information Technology (MIIT) (Chinese: 工信部) and seven other departments^[10-107], combined with the 8-segment, 40-scenario guidance and tiered nurturing, forms a three-in-one promotion framework.

The latest cohort of the Lighthouse network keeps confirming the positive correlation between "AI density" and operating performance: in January 2026 the Global Lighthouse Network added 23 new sites, 16 of them in China; WEF simultaneously launched Lumina, an AI industrial-intelligence platform that aggregates data from the past eight years of the Lighthouse network and over 1,000 transforming factories for benchmarking. WEF's statistics show that 94% of successful transformations result from combining multiple technologies, and lighthouse factories that integrate technology with talent and sustainability programs outperform industry peers by more than 16% on operating performance^[10-143] [third-party verified]. The

standout results of the new cohort concentrate in "embedding AI into core operations" and "scaling digital transformation across the network rather than within a single factory."

10.9.3 Benchmark Cases

Shenglong Electric: 27 Agents in 14 Days [third-party verified]

The *2026 China Manufacturing AI Scenario Application White Paper*^[10-102] documents 167 implementation cases with "verified company identities and quantifiable data." Shenglong Electric (Chinese: 盛隆电气) is the sharpest of them: 27 industrial agents deployed in 14 days, collaboration labor costs down 60%, ROI above 200%. The white paper also offers a sober set of structural data: generative AI accounts for only 4.8% of these cases and AI agents just 3%, while traditional machine learning (46.7%) and machine vision (20.4%) remain the workhorses — **the existing technology is actually sufficient; the bottleneck is scenario fit and how it is deployed.**

Figure × BMW: Humanoid Robots Enter a Mass-Production Line for the First Time [official claims]

The Figure 02 completed an 11-month deployment at BMW's Spartanburg plant: more than 1,250 hours of cumulative operation, over 90,000 components loaded, and participation in producing more than 30,000 X3s^[10-103]; sheet-metal loading completed within an 84-second cycle at 5 mm tolerance, with a positioning success rate above 99% [official claims]. BMW's Munich plant, meanwhile, represents the alternative "digitally native" route: digital twins run through development and production validation, and collaboration with Siemens and NVIDIA has accelerated pneumatic simulation by 30x^[10-104] [third-party verified]. Embodied AI has received production-scale data for the first time, and the engineering boundary has been genuinely tested.

CITIC Pacific Special Steel: What AI Looks Like in Process Industries [third-party verified]

In McKinsey's lighthouse-factory cases, CITIC Pacific Special Steel (Chinese: 中信泰富特钢) uses AI to accurately predict and adjust blast-furnace operations in real time, lifting output by 15% and cutting energy consumption by 11%^[10-97] [third-party verified]. This is the typical form of AI in process industries: fusing "mechanistic models + data models" to optimize process parameters rather than replace people.

The chemical industry concentrates on the same four directions — process optimization, predictive maintenance, safety early warning, and energy optimization — and process departments must lead; IT departments cannot own this.

Baosteel: the rare combination of company-level AI strategy and quantifiable returns. [third-party verified] A TMTPost (Chinese: 钛媒体) review of eight listed Chinese steelmakers' AI investments found most still stuck at demos, pilots, and "storytelling" — and Baoshan Iron & Steel (Baosteel, Chinese: 宝钢股份) the only company combining a company-level AI strategy with quantifiable results: its April 2026 "AI+" theme day disclosed 582 AI scenarios live with cumulative benefits above 270 million yuan, released the AI Scenario Whitepaper (2026 edition) covering 6 business links and 23 scenario categories, and launched the AI digital assistant "Zhibao" (Chinese: 智宝)^[10-141]; the annual report shows 347 AI scenarios went live in 2025, and the Baoshan base's "high-end green silicon-steel predictive smart factory" made the first batch of 15 pioneer-level intelligent factories^[10-142]. To judge whether a manufacturer is genuinely AI-native, look at whether strategy level, scenario density, benefit accounting, and organizational mechanisms form one integrated whole — the industry widely publishes use-case counts but rarely discloses benefit data and organizational support.

10.9.4 Failure Warnings

Deloitte Survey: 91% of Industrial AI Projects Missed Expectations, and 90% of the Failures Were Not Technical [third-party verified]

Only 9% of China's manufacturing AI projects achieved 80%-100% of expectations^[10-100]. The failure attributions are strikingly consistent: requirements get distorted after passing through three layers — decision-makers → IT department → vendor. One company invested millions in a "knowledge assistant"; three months after launch, daily active users fell to single digits and it was shut down. Frontline workers did not trust a system claiming 99% accuracy and insisted on manual re-inspection. After the vendor delivered, no one maintained the model, and the degraded system was switched off. The survivors share the opposite traits: entry through small scenarios, workers involved in defining requirements, and data governance up front.

The Full Lesson of the Fabric Inspection Machine [third-party verified]

Back to the textile mill at the start. Behind the 3.7 million yuan investment and the 15% misjudgment rate^[10-101] are three things that were not done right: insufficient scenario diversity and standardization (the same fabric's reflection characteristics vary widely between batches); know-how divorced from samples (the experience in the veteran workers' heads never became the machine's training data); and single-point equipment retrofits that cannot solve systemic problems (the machine replaced human eyes, but no one is backing up the machine).

Jumping In Blind: A Missing Data Foundation [third-party verified]

One manufacturer skipped pilots and data preparation and went straight to full-process AI quality inspection; insufficient training data meant accuracy failed to meet standards, the production line could not run, and the project failed and had to be redone. About 60% of enterprises ignore the principle that "data and business preparation come before technology." The common thread in these failure paths is not weak algorithms but organization and data infrastructure falling behind.

Murata's "data-trap" survey. [third-party / corporate survey] Murata Manufacturing's (Chinese: 村田制作所) questionnaire survey finds more than 38% of factories are implementing or planning smart-factory initiatives, with the share still rising — but smart-factory projects commonly fall into data traps: incomplete collection, uneven quality, and a focus on local equipment rather than end-to-end process data, so heavy investment fails to deliver efficiency^[10-144]. The value ceiling of a smart factory is set by data governance, not by the quantity of equipment and algorithms — "intelligentization for its own sake" is becoming manufacturing's new inventory.

10.9.5 Organizational Patterns and Consulting Takeaways

If you go to consult a manufacturing company on AI transformation, the first thing is to help it separate two kinds of problems: **"physical cycle-time problems" (AI cannot help — don't waste money)** and **"decision and information problems" (AI can help — worth investing)**. In McKinsey's lighthouse analysis, generative AI's fastest gains appear precisely in the process and knowledge scenarios outside the workshop — ACG Capsules' SOP assistant was used by nearly three-quarters of operators within five weeks^[10-97]. Manufacturing AI's ROI concentrates in

the decision layer and in the "documentation and administrative layer surrounding production" (SOP assistants, reports, knowledge Q&A, scheduling), not in the physical layer itself.

On the deployment sequence: single point first, platform second, organization last — start with high-frequency single points like quality inspection and scheduling to validate ROI, then build the data platform and model-asset management, and only then touch the organization (Chief AI Officer, AI Center of Excellence, business-unit AI specialists). McKinsey offers three levers for scaling: packaging use cases as "assets," low-code platform acceleration, and phased introduction — AI assists decisions first, and automation follows once confidence thresholds are exceeded.

The data foundation determines AI's ceiling: defect-sample generation, OT/IT integration, and the fusion of mechanistic and data models are the three engineering prerequisites. Predictive maintenance — the scenario with the clearest ROI — is exactly where "data-infrastructure debt" is most exposed.

Finally, the choice of organizational path. Manufacturing companies incubating AI platforms have three routes: internal incubation + external independent operation (SANY/RootCloud); validation in the parent factory + platform company exporting outward (Haier COSMOPlat, Midea Meiqing); and the "research institute + industry company" dual entity (Digital Building Materials Research Institute/"Xiaomiao"). What they share is letting AI capability break free of the parent organization's inertia and compete as an independent entity. And on any route, business-led rather than IT-led is the prerequisite — the chemical industry wrote this into its own practice, and academia summarizes it as the stage transition of "cognitive awakening → strategic decision → organizational change → technology implementation → capability consolidation." Manufacturing's AI-native transformation has never been a technology-procurement problem; it is an organizational learning process.

10.10 Energy & Utilities

10.10.1 Industry Profile and Bottleneck Triage

In 2025, Duke Energy's AI grid prevented about 1.2 million outages across the Carolinas — resolving them before they ever reached customers, saving roughly 3 million hours of outage time in total^[10-115] [third-party verified]. The approach is not flashy: AI prediction models identify equipment failure risks and vegetation hazards,

rank risk priorities by combining weather and asset condition, and schedule trimming and maintenance in advance. This is a legacy American utility that started in 2017 with a single-point use case detecting electricity theft and arrived at 2025 with 50-plus generative AI use cases in flight^[10-116] [third-party verified]. Its path is more worth studying than any "year one of energy AI" slogan.

Energy and utilities have the highest physical-bottleneck share of any industry in this book: an information:physical ratio of roughly 20-30:70-80. Power flows on the grid, pressure in pipelines, and turbine speeds are all millisecond-level physical realities — no matter how clever the algorithm, it cannot change Joule's law. The global energy AI market surpassed US\$40 billion in 2025, and the power-industry AI agent market exceeded 40 billion yuan [third-party verified]; yet value has been slow to arrive — utilities have invested over US\$10 billion in AI over the past five years, and most projects have not moved beyond the pilot stage [third-party verified].

Behind the gap are four core contradictions. Physical safety constraints have top priority: dispatch errors can trigger large-scale blackouts or even human casualties; the US NERC CIP regulatory framework has no clear definition for AI models; and the nuclear industry requires AI to enter only "human-machine collaboration" scenarios, never autonomous control. Heavy assets and long cycles: SCADA, mainframes, and other legacy systems 20-30 years old cannot communicate with modern AI platforms, and companies spend 40% of their AI budgets just making old systems "talk"^[10-117] [third-party verified]. Data fragmentation: data is scattered across dozens of isolated systems, and models are trained on incomplete, biased datasets^[10-120] [third-party verified]. Dual regulatory and performance constraints: central state-owned enterprises (central SOEs) are assessed on KPIs, and grid operators must guarantee supply — leaving far less room for error than the internet industry.

10.10.2 Transformation Paths and Typical Practices

Grid AI's path is "industry LLMs become the new foundation." State Grid (Chinese: 国家电网) released "Guangming" (Chinese: 光明), China's first 100-billion-parameter multimodal power LLM, in December 2024^[10-110], and its "6541" master plan (six application domains, five intelligent systems, four-in-one foundation) in May 2025; its intelligent customer service now auto-responds to more than 90% of work orders, with 270 million cumulative interactions [official claims]. China Southern Power Grid (Chinese: 南方电网) took the domestic-compute route: "Dawate Yudian" (Chinese: 大瓦特·驭电), trained on a Huawei Ascend cluster with 13 billion parameters, was selected

among the "top-ten national heavyweights" of central SOEs, and platform model calls surpassed 10 billion in December 2025^[10-112] [official claims]. Both grids use a "headquarters LLM + localized small models" architecture: a dispatch robot in Xuancheng, Anhui, compressed dispatch-order drafting from 3 hours to 30 minutes, and converter-transformer fault diagnosis in Liaocheng, Shandong, went from a week to milliseconds [third-party verified].

Oil and gas LLMs are moving from "single-well intelligence" to an "industry brain." CNPC's (Chinese: 中国石油) Kunlun LLM (Chinese: 昆仑大模型) released its 300-billion-parameter version (MoE architecture) in May 2025, covering 26 business lines and 119 business domains^[10-113] [official claims]; Sinopec (Chinese: 中石化) uses the "Great Wall LLM" (Chinese: 长城大模型) as the group's foundation. Data has been elevated to equal status with models: CNPC has built a 550 TB three-tier data management system, and "data is the blood of model training" has become industry consensus [third-party verified].

New-energy operations and maintenance are shifting from "veteran experience" to "precision diagnostics." In Q1 2025, wind and solar installed capacity reached 1.482 billion kW, surpassing thermal power for the first time, and the surge in capacity pushed O&M toward AI: Envision (Chinese: 远景) trains agents on millions of real operation-and-inspection data points [third-party verified]; Goldwind (Chinese: 金风) does predictive maintenance with digital twins + IoT [third-party verified]; Shenzhen Energy (Chinese: 深圳能源), with Huawei Cloud, released a power-forecasting system that reduces wind and solar curtailment [official claims]. Industry-wide fault-identification accuracy has risen above 90% [third-party verified].

Nuclear power and virtual power plants are the two endpoints of most conservative and most cutting-edge. In nuclear, China National Nuclear Power's (Chinese: 中国核电) "Hezhi AI Portal" (Chinese: 核智 AI 门户) has integrated DeepSeek, and CGN (Chinese: 中广核) is building a full-lifecycle nuclear intelligence hub on 18 core datasets — but applications are strictly limited to "human-machine collaboration" scenarios such as decision support and procedure interpretation [third-party verified]. In virtual power plants, the AI dispatch at the Dongguan energy-storage station completes 10 MW power allocation within 0.5 seconds [third-party verified]; China's 2025 target for virtual power plant regulation capacity is 20 million kW [third-party verified].

Computing-electricity synergy has moved from an industry topic to national strategy: in 2026 it appeared for the first time in the Government Work Report and entered the 15th Five-Year Plan outline; the four-department Action Plan for Two-Way Empowerment of AI and Energy (Chinese: 《人工智能与能源双向赋能行动方案》), issued by the NDRC and the National Energy Administration among others, makes it a key task and encourages new computing facilities to sign multi-year green-power purchase contracts with renewable-energy companies^[10-145]. On the ground, China Southern Power Grid has built the country's first electricity-carbon-computing coordination operating system and is running cross-regional dispatch pilots in Guangdong, Guizhou, and Jiangsu; during peak summer, Shanghai organized 16 data centers to complete the country's first full-mode, scaled computing-electricity synergy test. AI computing power is both a new load on the grid and, aggregated through virtual power plants, a high-quality flexible resource — for the first time, the demand side and the supply side of the same grid can trade places.

10.10.3 Benchmark Cases

Changqing Oilfield × Huawei: A Full-Stack Autonomous Smart Oilfield [official claims]

Changqing Oilfield (Chinese: 长庆油田), China's largest onshore oil and gas field (annual oil and gas equivalent output above 60 million tonnes, managing 126,000 wells), partnered with Huawei and Kunlun Digital Intelligence (Chinese: 昆仑数智) to build a three-tier architecture of "cloud centralized training — regional inference management — edge real-time inference" on a fully self-developed technology stack of "Kunpeng + Ascend + Galaxy Kylin." Its "special-operation hazard intelligent identification project" collected 2 million operation images and 150,000 hours of video, built a knowledge base of 12,000 rules, and deployed 20 high-precision models, achieving 94.1% hazard-identification accuracy (18.7 percentage points above traditional methods), edge-side inference latency under 800 milliseconds, a 70% reduction in manual supervision intervention for high-risk operations, and 18% lower overall operations and maintenance costs^[10-114]; in reservoir-development scenarios, detection accuracy rose from 78% to 92%, and unplanned equipment downtime fell by 50% [official claims]. It is a complete sample of the full-stack autonomous route.

Duke Energy: Predictive Maintenance Proven at Scale on the Distribution Grid [third-party verified]

Back to the story at the start. Even more valuable is the organizational method: it began in 2017 with a single-point use case detecting electricity theft, established governance guardrails in 2024, and reached 50-plus generative AI use cases in flight in 2025. The CIO's three principles — Always on (reliability first), Data before dazzle (don't bolt AI onto bad data), and Human in the loop (AI augments people rather than replaces them) — are the classic articulation of utility AI governance^[10-116] [third-party verified].

The Kunlun LLM: An Organizational Sample of a Central-SOE "Industry Brain" [official claims]

The Kunlun LLM was launched by the group's Party leadership in January 2024, became the first energy-and-chemical-industry LLM to pass national filing in August 2024, and grew to 300 billion parameters in May 2025^[10-113] [official claims]. It represents a collaboration model of "central-SOE headquarters coordination + ecosystem partners co-building," with the goal of upgrading from "model capability" to an "industry brain" reusable across the whole sector [third-party verified].

China Southern Power Grid's "DaWate·Yunrui": an agent sample for distribution-network planning. [official claims + third-party verified] As reported by People's Daily (Chinese: 人民日报), CSG's distribution-network planning agent "DaWate·Yunrui" (Chinese: 大瓦特·云睿) generates a distribution-network plan serving over 10,000 residents in 10 minutes from load, solar, and storage data, and was selected among the top ten "treasures of the hall" at WAIC 2026 from more than 1,100 companies and 3,000+ exhibits; it has been fully deployed across Yunnan Power Grid, completing over 10,000 grid risk diagnoses and generating more than 8,000 plans, with source-load forecasting and problem diagnosis above 90% accuracy and plan usability above 95%^[10-146]. Planning reports that once took days are compressed to ten minutes — high-skill roles such as grid planning and dispatch are having their working rhythm rewritten by agents.

10.10.4 Failure Warnings

74% of Pilots Never Leave a Single Region [third-party verified]

Prudent Consulting research shows that 74% of AI pilots in the utility industry never reach production in more than one region^[10-117]. The six root causes are almost all organizational: legacy-system integration barriers (40% of budgets spent on connecting old systems), data architecture not built for AI, no AI lifecycle management

(models abandoned after 12-18 months without retraining as accuracy decays), skill gaps and organizational misalignment, missing governance, and improper ROI measurement [third-party verified].

70% of Leaders Are Dissatisfied with AI Progress [third-party verified]

A BCG 2024 survey found nearly 60% of energy company leaders expected AI results within a year, but about 70% are dissatisfied with progress^[10-118] — the gap between expectation and reality constitutes the energy industry's "AI anxiety" [third-party verified]. The industry also widely suffers from "build but don't use": after projects are completed, there are no operations people to run them, and models become ornaments [third-party verified]. The IEA notes that the energy industry overall underutilizes AI, with data availability and digital skills the main barriers^[10-119] [third-party verified].

10.10.5 Organizational Patterns and Consulting Takeaways

First, draw the physical safety boundary for AI. Energy is an industry where "errors are irreversible": AI's value concentrates in the decision and information layers — inspection recognition, load forecasting, dispatch assistance, knowledge retrieval — and wherever human life and grid safety are involved, human accountability must be preserved. The early lesson from the US GridFWD conference is that "capital is too precious and the risk of getting it wrong too great" — build the governance model first, then talk about use cases^[10-120] [third-party verified]. The first consulting move is to help the enterprise clarify which scenarios AI may only "recommend" and which it may "execute."

Single point first, then systems, then organization. Duke started with one theft-detection use case in 2017, built governance guardrails in 2024, and only scaled in 2025^[10-116]; State Grid Hangzhou's (Chinese: 国网杭州) experience is to prioritize scenarios that are "high-frequency, high-value, low-risk" and to establish cross-departmental AI teams^[10-117] [third-party verified]. "Slow first, fast later — earn the right to move quickly" is closer to the truth than rushing into big-bang launches [third-party verified].

The data foundation and centralized control centers are dual prerequisites. Prudent lists "data architecture not built for AI" among the root causes of failure [third-party verified]; Changqing invested 2 million images + 150,000 hours

of video + a 12,000-rule knowledge base^[10-114] [official claims]. Central SOEs generally adopt a "group-level platform + edge inference" control architecture, treating data governance and compute infrastructure as infrastructure rather than projects.

The particularity of organizational change in central and state-owned enterprises: KPIs and error tolerance. State Grid is moving from the "6541 plan" to a special action of "two years to consolidate the foundation, three years to drive improvement"^[10-109] [official claims], and China Southern Power Grid deploys 29 key measures through annual work plans [third-party verified] — AI is being upgraded from project-based to routinized mechanisms through documents, KPIs, and dedicated institutions. But KPI-driven change has its costs: the "build but don't use" phenomenon is prominent. The consulting takeaway is to design a dual-track mechanism of "pilot-stage error tolerance + scale-stage KPI assessment" so AI does not become political achievements on paper.

Upgrade AI from a tool to strategy and business model. BCG argues for moving from "piloting AI" to an "AI-first operating model": embed AI into core business processes rather than edge applications, form fused teams of data scientists and domain experts, and co-develop with technology companies rather than buying^[10-118] [third-party verified]. Guidehouse goes further: AI will push utilities from "operational optimization" toward "business model reconstruction" — making virtual power plants and dynamic pricing possible^[10-119] [third-party verified]. And the IEA's proposition that "no AI without energy, no energy without AI" [third-party verified] makes energy the only industry where AI's demand creator and user are one — the biggest strategic variable for energy's AI-native organizations.

Afterword: A Living Book

Having finished the first edition, what I most want to talk about is not the content of this book — it is the book itself.

It came from an automated content pipeline: a system that scans the global information landscape every day, a body of knowledge entries that have been refined again and again, three rounds of targeted deep research, and a weekly automated merge-and-release cycle. From the very first edition, this book was never "finished" — it is **alive**.

That is itself a miniature of the AI-native organization: **delegate repetitive work to systems, keep judgment for humans, and make knowledge an asset that keeps growing**. The book describes how other organizations have done this — and the way this book produces itself is the same principle in practice.

On Limitations

The source material of this book has clear boundaries, and I must state them honestly:

- **Industry boundary:** In the first edition, the vast majority of cases came from information-intensive industries (software, content, finance, consulting), and evidence from physical-world industries (manufacturing lines, clinical medicine, logistics) was scarce. Chapter 10, the industry guide, has systematically closed that gap — manufacturing, healthcare, logistics, energy, and other physical-world industries now have systematic cases and data. But a real boundary remains: the ten industries are covered with unequal depth (information-intensive industries are thicker), and the industry material comes mainly from public reporting, so first-hand operational detail is inherently limited. The sample-bias critique applied to the Tencent Research Institute report in Chapter 3 still holds — extrapolating "the patterns of information-intensive industries" into "universal organizational laws" is dangerous.
- **Time boundary:** AI data goes stale within months. Each data point in this book is timestamped (most as of August 2026); always refer to the latest version.

- **Survivorship bias:** The organizations in our case studies are publicly disclosed transformation samples. Companies that failed quietly will never appear in a book — Chapter 9's failure cases are a limited compensation for this bias.

How to Contribute

This book is free and open. Any reading or sharing is welcome. If you find errors, omissions, or cases worth including:

- Open an issue at the GitHub repository: github.com/Jweokk/ai-native-organization-book (<https://github.com/Jweokk/ai-native-organization-book>)
- Or write to weokk2025@gmail.com

The value of knowledge lies in its flow, not in hoarding it. This book came from flow — and we hope it keeps flowing.

Appendix A: Case Index and Sources

This appendix lists the original sources for every data point, case, and claim in this book. Attribution convention: company official statements are marked as **official claims**; third-party sources such as media, analysts, and regulatory filings are marked as **third-party verified**. All data is timestamped as of August 2026.

Numbering system: Superscript [numbers] in the body text (e.g., [10-25]) are direct citations — click to jump to the corresponding entry in this appendix. Entries are grouped by source theme ([1-xx] reports & data, [7-xx] international cases, [8-xx] methodology & failure cases, [10-xx] industry guide); any chapter may cite any group. Entries without a superscript in the body are further reading and verification sources on the same topic.

Sources for Chapters 1-3 (Reports and Data)

- [1-01] MLQ.ai: State of AI in Business 2025 Report: https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- [1-02] Fortune: MIT report: 95% of generative AI pilots at companies are failing: <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
- [1-03] mckinsey.com: <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>
- [1-04] BCG: AI at Work: Momentum Builds but Gaps Remain: <https://www.bcg.com/publications/2025/ai-at-work-momentum-builds-but-gaps-remain>
- [1-05] BCG Press Release: Beyond AI Adoption: Unlocking Full Potential: <https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential>
- [1-06] BCG: AI at Work: Why Strategy Matters More Than Tools: <https://www.bcg.com/publications/2026/ai-at-work-why-strategy-matters-more-than-tools>

- [1-07] Gartner Press Release: Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept by End of 2025: <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>
- [1-08] Gartner Press Release: 40% of Enterprise Apps to Feature Task-Specific AI Agents by 2026: <https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025>
- [1-09] Gartner Press Release: AI Will Touch All Information Technology Work by 2030: <https://www.gartner.com/en/newsroom/press-releases/2025-11-10-gartner-survey-finds-artificial-intelligence-will-touch-all-information-technology-work-by-2030>
- [1-10] Gartner Press Release: Top Predictions for IT Organizations and Users in 2026 and Beyond: <https://www.gartner.com/en/newsroom/press-releases/2025-10-21-gartner-unveils-top-predictions-for-it-organizations-and-users-in-2026-and-beyond>
- [1-11] Forrester Blog: Predictions 2026: AI Moves from Hype to Hard-Hat Work: <https://www.forrester.com/blogs/predictions-2026-ai-moves-from-hype-to-hard-hat-work/>
- [1-12] Forrester Blog: Predictions 2026: AI Agents Changing Business Models and Workplace Culture: <https://www.forrester.com/blogs/predictions-2026-ai-agents-changing-business-models-and-workplace-culture-impact-enterprise-software/>
- [1-13] Microsoft WorkLab: AI at Work Is Here. Now Comes the Hard Part: <https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part>
- [1-14] microsoft.com: <https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born>
- [1-15] Microsoft WorkLab: Agents, Human Agency, and the Opportunity for Every Organization: <https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization>
- [1-16] Stanford HAI: 2025 AI Index Report: <https://hai.stanford.edu/ai-index/2025-ai-index-report>
- [1-17] Stanford HAI: 2025 AI Index Report (PDF): https://hai.stanford.edu/assets/files/hai_ai_index_report_2025.pdf

- [1-18] World Bank: Digital Progress and Trends Report 2025: AI Foundations: <https://www.worldbank.org/en/publication/dptr2025-ai-foundations>
- [1-19] World Bank: Digital Progress and Trends Report 2025 (Download): <https://openknowledge.worldbank.org/bitstreams/86903114-6212-4c45-9011-938925cc61d1/download>
- [1-20] IMF: Gen-AI: Artificial Intelligence and the Future of Work: <https://www.imf.org/en/publications/staff-discussion-notes/issues/2024/01/14/gen-ai-artificial-intelligence-and-the-future-of-work-542379>
- [1-21] IMF: 2025 Working Paper (PDF): <https://www.imf.org/-/media/Files/Publications/WP/2025/English/wpia2025076-print-pdf.ashx>
- [1-22] OECD: AI Surveys of Employers and Workers: <https://www.oecd.org/en/about/projects/aisurveysofemployersandworkers.html>
- [1-23] OECD: The Impact of AI on the Workplace: Main Findings from the OECD AI Surveys of Employers and Workers: https://www.oecd.org/en/publications/the-impact-of-ai-on-the-workplace-main-findings-from-the-oecd-ai-surveys-of-employers-and-workers_ea0a0fe1-en.html
- [1-24] OECD: OECD Employment Outlook 2023: AI, Job Quality, and Inclusiveness: https://www.oecd.org/en/publications/oecd-employment-outlook-2023_08785bba-en/full-report/artificial-intelligence-job-quality-and-inclusiveness_a713d0ad.html
- [1-25] The Next Web: OECD: Finance and Manufacturing Workers Fear AI Replacement: <https://thenextweb.com/news/oecd-finance-and-manufacturing-workers-fear-ai-replacement>
- [1-26] OECD: Generative AI and the SME Workforce: https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en/full-report/component-3.html
- [1-27] World Economic Forum Press Release: Future of Jobs Report 2025: 78 Million New Job Opportunities by 2030: <https://www.weforum.org/press/2025/01/future-of-jobs-report-2025-78-million-new-job-opportunities-by-2030-but-urgent-upskilling-needed-to-prepare-workforces/>
- [1-28] World Economic Forum: Future of Jobs Report 2025 (PDF): https://reports.weforum.org/docs/WEF_Future_of_Jobs_Report_2025.pdf
- [1-29] Harvard Business School: Report 26-090: <https://www.hbs.edu/ris/download.aspx?name=26-090.pdf>

- [1-30] SSRN: Paper 6905079: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6905079
- [1-31] Harvard Business School: Research 69077: <https://www.hbs.edu/faculty/Pages/item.aspx?num=69077>
- [1-32] Harvard Business School AI Institute: Everyone Has AI. Which Firms Are Going to Win?: <https://aiinstitute.hbs.edu/everyone-has-ai-which-firms-are-going-to-win/>
- [1-33] SSRN: Paper 6513481: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6513481
- [1-34] Harvard Business School AI Institute: Less Headcount, More Valuation: How AI-Native Firms Change the Game: <https://aiinstitute.hbs.edu/less-headcount-more-valuation-how-ai-native-firms-change-the-game/>
- [1-35] Harvard Business School AI Institute: Is GenAI Heading for a Tech Monopoly?: <https://aiinstitute.hbs.edu/is-genai-heading-for-a-tech-monopoly/>
- [1-36] JetBrains Developer Ecosystem Blog: AI Coding Agent Adoption Trends 2026: <https://blog.jetbrains.com/research/2026/08/ai-coding-agent-adoption-2026/>

Sources for Chapter 7 (International Case Studies)

- [7-01] cnbc.com: <https://www.cnbc.com/2025/05/14/klarna-ceo-says-ai-helped-company-shrink-workforce-by-40percent.html>
- [7-02] SEC Filing: Klarna Group plc Financial Report: <https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm>
- [7-03] bloomberg.com: <https://www.bloomberg.com/news/articles/2024-12-12/klarna-stopped-all-hiring-a-year-ago-to-replace-workers-with-ai>
- [7-04] TechCrunch: Klarna's CEO says it stopped hiring thanks to AI — but still advertises many open positions: <https://techcrunch.com/2024/12/14/klarnas-ceo-says-it-stopped-hiring-thanks-to-ai-but-still-advertises-many-open-positions/>
- [7-05] threads.com: <https://www.threads.com/@fox.hsiao/post/DVP3e7jIA17/>
- [7-06] TechHanlin: Klarna AI Customer Service Case Study: <https://www.techhanlin.tw/klarna-ai-customer-service-case-study/>
- [7-07] klarna.com: <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-its-first-month/>

- [7-08] klarna.com: <https://www.klarna.com/international/press/90-of-klarna-staff-are-using-ai-daily-game-changer-for-productivity/>
- [7-09] Customer Experience Dive: Klarna says AI agent work equals 853 employees: <https://www.customerexperiencedive.com/news/klarna-says-ai-agent-work-853-employees/805987/>
- [7-10] Customer Experience Dive: Klarna uses AI to slash customer service costs: <https://www.customerexperiencedive.com/news/klarna-ai-slash-customer-service-costs/748647/>
- [7-11] Klarna Investor Relations: Klarna Group plc Publishes Full Year 2025 Results: <https://investors.klarna.com/News--Events/news/news-details/2026/Klarna-Group-plc-Publishes-Full-Year-2025-Results/default.aspx>
- [7-12] Klarna Investor Relations: Klarna Accelerates US Growth and Delivers \$1bn Revenue: <https://investors.klarna.com/News--Events/news/news-details/2026/Klarna-Accelerates-U-S--Growth-and-Delivers-1bn-Revenue-Driven-by-Rapid-Banking-Service-Adoption/default.aspx>
- [7-13] CNBC: Klarna Q3 Earnings Report: <https://www.cnbc.com/2025/11/18/klarna-klar-stock-q3-earnings-report-2025.html>
- [7-14] bloomberg.com: <https://www.bloomberg.com/news/articles/2025-05-08/klarna-turns-from-ai-to-real-person-customer-service>
- [7-15] investopedia.com: <https://www.investopedia.com/buy-now-pay-later-company-klarna-s-shares-end-first-session-above-ipo-price-11806387>
- [7-16] businessweekly.com.tw: <https://www.businessweekly.com.tw/management/blog/3021962>
- [7-17] cnbc.com: <https://www.cnbc.com/2023/05/04/shopify-cuts-20percent-of-its-workforce-shares-surge-on-earnings-beat>
- [7-18] businessinsider.com: <https://www.businessinsider.com/shopify-ai-use-boosts-efficiency-shakes-up-staff-structure-2024-5>
- [7-19] cnbc.com: <https://www.cnbc.com/2025/04/07/shopify-ceo-prove-ai-cant-do-jobs-before-asking-for-more-headcount.html>
- [7-20] Revelio Labs: Shopify Employee Data: <https://www.reveliolabs.com/companies/shopify/employees>
- [7-21] Tobi (Shopify CEO) on X: On Employees and AI: <https://x.com/tobi/status/1909231499448401946>

- [7-22] bullfincher.io: <https://bullfincher.io/companies/shopify/revenue-per-employee>
- [7-23] chargeflow.io: <https://www.chargeflow.io/blog/shopify-statistics>
- [7-24] FourWeekMBA: Shopify Revenue per Employee: <https://fourweekmba.com/shopify-revenue-per-employee/>
- [7-25] natesnewsletter.substack.com: <https://natesnewsletter.substack.com/p/my-honest-field-notes-on-how-the>
- [7-26] en.wikipedia.org: <https://en.wikipedia.org/wiki/Anthropic>
- [7-27] CNBC: Anthropic Says Annualized Revenue Climbed to \$6.5 Billion in July: <https://www.cnbc.com/2026/08/17/anthropic-says-annualized-revenue-climbed-to-65-billion-in-july.html>
- [7-28] venturebeat.com: <https://venturebeat.com/technology/anthropic-says-80-of-its-new-production-code-is-now-authored-by-claude-how-your-enterprise-can-keep-up>
- [7-29] Anthropic on X: Revenue Announcement: <https://x.com/AnthropicAI/status/2062568864240836995>
- [7-30] cnbc.com: <https://www.cnbc.com/2025/11/13/cursor-ai-startup-funding-round-valuation.html>
- [7-31] anthropic.com: <https://www.anthropic.com/research/anthropic-economic-index-september-2025-report>
- [7-32] skeptics.stackexchange.com: <https://skeptics.stackexchange.com/questions/59213/as-at-september-2025-is-ai-not-writing-90-of-code>
- [7-33] Redwood Research Blog: Is 90% of Code at Anthropic Being Written by AI?: <https://blog.redwoodresearch.org/p/is-90-of-code-at-anthropic-being>
- [7-34] makerstations.io: <https://www.makerstations.io/openai-employee-statistics/>
- [7-35] Mokahr Blog: Talent & Culture Strategy at OpenAI: <https://www.mokahr.io/myblog/talent-culture-strategy-at-openai/>
- [7-36] searchlab.nl: <https://searchlab.nl/en/statistics/openai-statistics-2026>
- [7-37] cnbc.com: <https://www.cnbc.com/2025/09/03/openai-boosts-size-of-secondary-share-sale-to-10point3-billion.html>
- [7-38] openai.com: <https://openai.com/index/how-agents-are-transforming-work/>

- [7-39] fortune.com: <https://fortune.com/2026/08/13/buried-in-openai-latest-research-no-correlation-between-ai-use-and-revenue-per-employee/>
- [7-40] cnbc.com: <https://www.cnbc.com/2025/09/18/notion-launches-ai-agent-as-it-crosses-500-million-in-annual-revenue.html>
- [7-41] forbes.com: <https://www.forbes.com/sites/annatong/2025/12/15/notion-kicks-off-employee-share-sale-at-11-billion-valuation-as-ai-accelerates-its-growth/>
- [7-42] Notion Blog: Introducing Notion AI: <https://www.notion.com/blog/introducing-notion-ai>
- [7-43] medium.com: <https://medium.com/@sherrysun/notion-from-near-collapse-to-a-10b-all-in-one-workspace-unicorn-6113f7e1dc4c>
- [7-44] GetLatka: Notion Company Data: <https://getlatka.com/companies/notion.ai>
- [7-45] cursor.com: <https://cursor.com/blog/series-d>
- [7-46] en.wikipedia.org: [https://en.wikipedia.org/wiki/Cursor_\(company\)](https://en.wikipedia.org/wiki/Cursor_(company))
- [7-47] Reddit: Cursor Hit \$1B ARR with 300 Employees: https://www.reddit.com/r/SaaS/comments/1r9pj7k/cursor_hit_1b_arr_with_300_employees_thats_33m/
- [7-48] app.dealroom.co: <https://app.dealroom.co/news/note/cursor-tops-4b-annualized-revenue-june-2026>
- [7-49] CNBC: OpenAI Looked at Cursor Before Considering Deal with Rival Windsurf: <https://www.cnbc.com/2025/04/17/openai-looked-at-cursor-before-considering-deal-with-rival-windsurf.html>
- [7-50] mashable.com: <https://mashable.com/article/duolingo-ai-layoff-contractors>
- [7-51] cnbc.com: <https://www.cnbc.com/2025/09/17/duolingo-ceo-how-ai-makes-my-employees-more-productive-without-layoffs.html>
- [7-52] businessinsider.com: <https://www.businessinsider.com/duolingo-ceo-how-ai-will-be-used-performance-reviews-headcount-2025-4>
- [7-53] Metaintro Blog: Duolingo CEO Walks Back AI-First Memo; Hiring Grows Through 2026: <https://www.metaintro.com/blog/duolingo-ceo-walks-back-ai-first-memo-hiring-grows-2026>
- [7-54] investors.duolingo.com: <http://investors.duolingo.com/company-strategy-overview-0>

- [7-55] classcentral.com: <https://www.classcentral.com/report/duolingo-2025/>
- [7-56] reddit.com: https://www.reddit.com/r/duolingo/comments/191ssv7/discussion_duolingo_cuts_10_of_its_contractors/
- [7-57] techcrunch.com: <https://techcrunch.com/2026/05/08/airbnb-says-ai-now-writes-60-of-its-new-code/>
- [7-58] TechCrunch: Airbnb Is Quietly Rolling Out an AI Customer Service Bot in the US: <https://techcrunch.com/2025/05/02/airbnb-is-quietly-rolling-out-an-ai-customer-service-bot-in-the-us/>
- [7-59] TechTimes: Airbnb Proves AI-Powered Engineering Beats Rivals with 80% More Features: <https://www.techtimes.com/articles/323585/20260807/airbnb-proves-ai-powered-engineering-beats-rivals-80-more-features-same-staff.htm>
- [7-60] techcrunch.com: <https://techcrunch.com/2026/02/13/airbnb-plans-to-bake-in-ai-features-for-search-discovery-and-support/>
- [7-61] techcrunch.com: <https://techcrunch.com/2025/04/29/microsoft-ceo-says-up-to-30-of-the-companys-code-was-written-by-ai/>
- [7-62] TechCrunch: Spotify Says Its Best Developers Haven't Written a Line of Code Since December: <https://techcrunch.com/2026/02/12/spotify-says-its-best-developers-havent-written-a-line-of-code-since-december-thanks-to-ai/>
- [7-63] reddit.com: https://www.reddit.com/r/technology/comments/1t9v6km/airbnb_says_ai_now_writes_60_of_its_new_code/
- [7-64] cnbc.com: <https://www.cnbc.com/2024/02/01/zoom-layoffs-company-cuts-150-employees-2percent-of-workforce.html>
- [7-65] BBC News: Technology: <https://www.bbc.com/news/technology-64562673>
- [7-66] MakerStations: Zoom Employee Statistics: <https://www.makerstations.io/zoom-employee-statistics/>
- [7-67] ciodive.com: <https://www.ciodive.com/news/zoom-ai-companion-monetization-customization/708688/>
- [7-68] WSJ Facebook Post: Zoom Founder and CEO Predicts AI Will Put an End to the 5-Day Workweek: <https://www.facebook.com/WSJ/posts/zoom-founder-and-ceo-eric-yuan-predicts-ai-will-put-an-end-to-the-5-day-workweek/1321223876530873/>
- [7-69] linkedin.com: https://www.linkedin.com/posts/conorgrennan_just-released-microsofts-2025-work-trend-activity-7320839065984000004-B8gk

- [7-70] Microsoft Work Trend Index: Microsoft Work Trend Index: <https://www.microsoft.com/en-us/worklab/work-trend-index>
- [7-71] themicrosoftcloudblog.com: <https://themicrosoftcloudblog.com/2026/05/2026-work-trend-index-evidence-check/>
- [7-72] microsoft.com: <https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born>
- [7-73] Sina Tech (Lenovo Zhihui Salon): Lenovo's AI-Native Organization Practice — Lexiang Super-Agent 4.0 Tops 7 Million MAU and Pulls 5 Billion Yuan in Sales, "Silicon-Carbon Symbiosis": <https://finance.sina.com.cn/tech/roll/2026-07-29/doc-iniknezf0011507.shtml>

Sources for Chapters 8-9 (Methodology and Failure Cases)

- [8-01] complexdiscovery.com: <https://complexdiscovery.com/why-95-of-corporate-ai-projects-fail-lessons-from-mits-2025-study/>
- [8-02] The Register: Generative AI Zero-Return Rate as High as 95%: https://www.theregister.com/2025/08/18/generative_ai_zero_return_95_percent/
- [8-03] Gartner Press Release: Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept by End of 2025: <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>
- [8-04] gartner.com: <https://www.gartner.com/en/newsroom/press-releases/2025-02-26-lack-of-ai-ready-data-puts-ai-projects-at-risk>
- [8-05] gartner.com: <https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>
- [8-06] Anar Solutions: Why Agentic AI Pilots Fail in Production: <https://anarsolutions.com/why-agentic-ai-pilots-fail-production/>
- [8-07] eeo.com.cn: <http://www.eeo.com.cn/2026/0601/898378.shtml>
- [8-08] unifyapps.com: <https://www.unifyapps.com/resources/blog/why-95-of-generative-ai-pilots-never-reach-production>
- [8-09] unifyapps.com: <https://www.unifyapps.com/resources/customer-story/Large-Financial-Institution-India>

- [8-10] UnifyApps Customer Story: Fortune 50 Retailer: <https://www.unifyapps.com/resources/customer-story/fortune-50-retailer>
- [8-11] cnbc.com: <https://www.cnbc.com/2025/05/14/klarna-ceo-says-ai-helped-company-shrink-workforce-by-40percent.html>
- [8-12] Fast Company: Klarna Tried to Replace Its Workforce with AI: <https://www.fastcompany.com/91468582/klarna-tried-to-replace-its-workforce-with-ai>
- [8-13] digitalapplied.com: <https://www.digitalapplied.com/blog/klarna-reverses-ai-layoffs-replacing-700-workers-backfired>
- [8-14] thestateofbrand.com: <https://www.thestateofbrand.com/news/klarna-reverses-ai-job-cuts>
- [8-15] TandemCoach: Klarna AI Automation Lesson: <https://tandemcoach.co/klarna-ai-automation-lesson/>
- [8-16] bloomberg.com: <https://www.bloomberg.com/news/articles/2023-05-02/samsung-bans-chatgpt-and-other-generative-ai-use-by-staff-after-leak>
- [8-17] Forbes: Samsung Bans ChatGPT and Other Chatbots for Employees After Sensitive Code Leak: <https://www.forbes.com/sites/siladityaray/2023/05/02/samsung-bans-chatgpt-and-other-chatbots-for-employees-after-sensitive-code-leak/>
- [8-18] Mashable: Samsung ChatGPT Leak Details: <https://mashable.com/article/samsung-chatgpt-leak-details>
- [8-19] weforum.org: <https://www.weforum.org/stories/artificial-intelligence/companies-ai-workflows-not-simple-tasks/>
- [8-20] stackmatix.com: <https://www.stackmatix.com/blog/copilot-market-adoption-trends>
- [8-21] Peafowlit Blog: Copilot Licenses Go Unused and How to Fix Adoption: <https://peafowlit.com/blog/copilot-licenses-go-unused-and-how-to-fix-adoption/>
- [8-22] linkedin.com: https://www.linkedin.com/posts/jukkaniiranen_34-of-microsoft-365-customers-pay-for-premium-activity-7422758776069455872-Clnk
- [8-23] querynow.com: <https://www.querynow.com/resources/whitepapers/past-the-stall-m365-copilot-rollouts>
- [8-24] linkedin.com: <https://www.linkedin.com/pulse/microsofts-copilot-paradox-94-report-benefits-6-deploy-louis-columbus-fv4gc>

- [8-25] onlineeducation.com: <https://www.onlineeducation.com/features/chatgpt-crashes-chepps-stock>
- [8-26] Higher Education Inquirer: Chegg: A Critical History of a Disruptor: <https://www.highereducationinquirer.org/2025/07/chegg-critical-history-of-disruptor.html>
- [8-27] LinkedIn Post: The Fall of Chegg Is a Masterclass in Ignoring Change: https://www.linkedin.com/posts/harshadshah1953_the-fall-of-chegg-is-a-masterclass-in-ignoring-activity-7487102244815921152-yM6L
- [8-28] CEIBS Media Coverage: CEIBS Press Clippings: <https://cn.ceibs.edu/media/press-clippings/faculty/29204>
- [8-29] Workiva: Data Pressures Mount, Instability Continues: <https://www.workiva.com/resources/data-pressures-mount-instability-continues>
- [8-30] woshipm.com: <https://www.woshipm.com/ai/6275554.html>
- [8-31] maiagent.ai: <https://maiagent.ai/blog/enterprise-shadow-ai-governance>
- [8-32] securitytoday.com: <https://securitytoday.com/articles/2025/08/18/survey-nearly-half-of-employees-hide-workplace-ai-use.aspx>
- [8-33] cloudflare.com: <https://www.cloudflare.com/zh-cn/learning/ai/what-is-shadow-ai/>
- [8-34] unleash.ai: <https://www.unleash.ai/artificial-intelligence/news/asana-64-of-employees-believe-ai-agents-are-unreliable-calling-for-more-training-clarity-and-guardrails>
- [8-35] Asana: State of AI Research Takeaways: <https://asana.com/resources/state-of-ai-research-takeaways>
- [8-36] accenture.com: <https://www.accenture.com/us-en/insights/consulting/leadership-edge-ai>
- [8-37] d1net.com: <https://www.d1net.com/cio/ciotech/585344.html>
- [8-38] blog.104.com.tw: <https://blog.104.com.tw/wef-report-2026-ai-entry-level-work-redesign/>
- [8-39] hbrtaiwan.com: <https://www.hbrtaiwan.com/special-topics/24535/the-perils-of-using-ai-to-replace-entry-level-jobs>
- [8-40] anthropic.com: <https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic>
- [8-41] swfte.com: <https://www.swfte.com/blog/avoid-ai-vendor-lock-in-enterprise-guide>

- [8-42] stepto.net: <https://stepto.net/blog/ai-vendor-lock-in-infrastructure-risk-2026>
- [8-43] linkedin.com: https://www.linkedin.com/posts/sumatosoft_aistrategy-vendorlockin-enterpriseai-activity-7458501143887765504-FfQC
- [8-44] institutepm.com: <https://www.institutepm.com/knowledge-hub/ai-vendor-lock-in-strategy>
- [8-45] fin.ai: <https://fin.ai/learn/ai-customer-service-cost-savings-industry>
- [8-46] cmswire.com: <https://www.cmswire.com/contact-center/will-ai-cost-more-than-offshore-human-agents-in-customer-service/>
- [8-47] stcn.com: <https://www.stcn.com/article/detail/3951204.html>
- [8-48] resources.anthropic.com: <https://resources.anthropic.com/enterprise-ai-transformation-guide>
- [8-49] rubrik.com: <https://www.rubrik.com/blog/technology/25/11/its-early-days-for-agent-ai-and-most-companies-lack-the-tools-to-protect-their-data>
- [8-50] rubrik.com: <https://www.rubrik.com/company/newsroom/press-releases/26/ai-agents-are-breaking-things-and-organisations-know-it>
- [8-51] Rubrik Product Page: Rubrik Agent Cloud: <https://www.rubrik.com/products/rubrik-agent-cloud>
- [8-52] mckinsey.com: <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>
- [8-53] McKinsey Chart: AI at Work, But Not at Scale: <https://www.mckinsey.com/featured-insights/charts/ai-at-work-but-not-at-scale>
- [8-54] The Art of Service: Course Section: <https://academy.theartofservice.com/course/section.php?id=103955>
- [8-55] bcg.com: <https://www.bcg.com/featured-insights/the-leaders-guide-to-transforming-with-ai>
- [8-56] Accenture: Accenture Reinvention Services: <https://www.accenture.com/us-en/about/reinvention-services>
- [8-57] RTS Labs: Enterprise AI Roadmap: <https://rtslabs.com/enterprise-ai-roadmap>
- [8-58] Thinking Inc: AI Adoption Roadmap: <https://thinking.inc/en/pillar-pages/ai-adoption-roadmap/>

- [8-59] WeChat Official Account Article: AI: <https://mp.weixin.qq.com/s/NJW9cYB8aH-1gIXdUnSaA>
- [8-60] Zilbix: The AI-Native Operating Model: How Enterprises Must Restructure for the Agentic Era: <https://www.zilbix.com/ai-native-operating-model-how-enterprises-must-restructure-for-the-agentic-era/>
- [8-61] Anthropic Research: Multi-Agent Systems: <https://www.anthropic.com/research/multiagent-systems>
- [8-62] Datagrand News: News: <https://www.mininglamp.com/news/8292/>
- [8-63] Honeycomb Blog: Your Questions About AI Agents in Production Feedback, Answered: <https://www.honeycomb.io/blog/your-questions-about-ai-agents-production-feedback-answered>
- [8-64] AmpCode Notes: That's Not SOC 2 Compliant: <https://ampcode.com/notes/thats-not-soc-2-compliant>
- [8-65] Futurism: Young People, AI, and CEO Executives Poll: <https://futurism.com/artificial-intelligence/young-people-ai-ceos-executives-poll>
- [8-66] Librarian.net: NotoAI: <https://www.librarian.net/notoai/>
- [8-67] JSBarretto Blog: I Remain a Skeptic: <https://blog.jsbarretto.com/post/i-remain-a-skeptic>
- [8-68] OpenSauce: Cloudflare AI Psychosis: <https://opensauce.it/cloudflare-ai-psychosis/>
- [8-69] Allen Bargi Notes: Working with AI Feels Like Leadership: <https://allen.bargi.org/notes/working-with-ai-feels-like-leadership/>
- [8-70] Wired Interview: Model Behavior Interview with OpenAI Codex Lead Tibo Sottiaux: <https://www.wired.com/story/model-behavior-interview-with-openai-codex-lead-tibo-sottiaux/>
- [8-71] Stephen Cresswell Blog: Yadda 3.0.0: BDD in the Age of AI Agents: <https://www.stephen-cresswell.com/2026/08/15/Yadda-3.0.0-BDD-in-the-Age-of-AI-Agents.html>
- [8-72] Moka: Buying AI Features vs. Building AI-Native Capability: Why They Are Two Different Things in 2026 (vendor-reported, official-claims caliber): <https://www.mokahr.com/blog/43826.html>
- [8-73] Moerso (莫尔索): How AI Agents Land in Enterprises in 2026 (organizational inertia / Role Fusion / FDE→AI Engineer): <https://zhuanlan.zhihu.com/p/2048422952681599250>

- [8-74] Zhaozhao (AI Startup Edition): FDE Field Notes: How a Small Team Closes Its First Enterprise AI Deal: <https://x.com/dzhao63405/status/2093653049755738200>
- [8-75] orangeburn / Thinkit Lab: The 7 Big Pitfalls of Legacy-Enterprise AI Transformation: https://mp.weixin.qq.com/s/jn-_i32NCktdpxgJRjbRdw
- [8-76] orangeburn / Thinkit Lab: Chinese-Style FDE Is Doomed to Fail: <https://mp.weixin.qq.com/s/8g-AAi2LJ8FDuR2L4hq6yA>
- [8-77] orangeburn / Thinkit Lab: We Are Using FDE to Solve a Problem That Doesn't Exist: <https://mp.weixin.qq.com/s/V45ycltvzHSxI5L0PgUFAw>
- [8-78] SMH: This "AI refuser" quit her dream job, and hopes others follow: <https://www.smh.com.au/technology/this-ai-refuser-quit-her-dream-job-and-hopes-others-follow-20260818-p60pdu.html>
- [8-79] Lars Faye: AI Coding will Prevent Expertise (JetBrains' The Widening Gap / UPenn 2025 study): <https://larsfaye.com/articles/ai-coding-will-prevent-expertise>
- [8-80] Tom Tunguz: Most AI Work Can Wait (routing over model choice; Coinbase cost data): <https://tomtunguz.com/ai-execution-routing/>
- [8-81] Mininglamp (明略科技): Enterprise Efficiency Drops After Deploying Multiple AI Agents? (IDC 2026Q2 / CAICT 358-enterprise survey / Beijing agent policy): <https://www.mininglamp.com/news/8365/>
- [8-82] SenteLabsAI: OpenExecutive — the "AI CEO" Open-Sourced by Laid-Off Developers: <https://github.com/SenteLabsAI/OpenExecutive>
- [8-83] Meituan Tech Blog: Managing AI Coding with an Eval-Agent Mindset — Refactoring a 310,000-Line AI-Generated System: <https://tech.meituan.com/2026/05/07/Agent-AI-Coding.html>
- [8-84] Woshipm: A Product Manager's 8-Week Debrief on Pushing an AI Agent from 0 to 1: <https://www.woshipm.com/ai/6386229.html>
- [8-85] Wild World of Work (Matt Beane): AI-Native Organizational Transformation — the Lesson of Uber's CTO Burning a Year's Token Budget in Four Months: <https://www.wildworldofwork.org/p/ai-native-organizational-transformation>
- [8-86] Ashley Rolf More: AI Is Stopping Startups from Completing Puberty: <https://ashley.rolfmore.com/ai-is-stopping-startups-from-completing-puberty/>
- [8-87] Rodrigo Rosenfeld: AI Agents and the Refactoring That Never Happens: https://www.rosenfeld.page/articles/programming/2026_09_02_ai_agents_and_the_refactoring_that_never_happens/

- [8-88] LeadDev: AI Coding Creates Two Kinds of Debt, and You're Only Measuring One — Cognitive Debt: <https://leaddev.com/ai/ai-coding-creates-two-kinds-of-debt-youre-only-measuring-one>
-

Sources for Chapter 10 (Industry Guide)

Professional Services & Legal

- [10-01] ABA Formal Opinion 512 (Generative AI Ethics Framework): https://www.americanbar.org/content/dam/aba/administrative/professional_responsibility/ethics-opinions/aba-formal-opinion-512.pdf
- [10-02] Thomson Reuters: 2026 AI in Professional Services Report: <https://legal.thomsonreuters.com/blog/how-ai-is-transforming-the-legal-profession/>
- [10-03] ABA: 2025 Legal Industry Report: https://www.americanbar.org/groups/law_practice/resources/law-technology-today/2025/the-legal-industry-report-2025/
- [10-04] CNBC: Harvey raises \$200 million at \$11 billion valuation (March 2026): <https://www.cnbc.com/2026/03/25/legal-ai-startup-harvey-raises-200-million-at-11-billion-valuation.html>
- [10-05] A&O Shearman (Official): Launching agentic AI agents with Harvey for complex legal workflows: <https://www.aoshearman.com/en/news/ao-shearman-and-harvey-to-roll-out-agentic-ai-agents-targeting-complex-legal-workflows>
- [10-06] Harvey Official Customer Case: A&O Shearman firm-wide deployment: <https://www.harvey.ai/customers/a-and-o-shearman>
- [10-07] HFS Research: Accenture Reinvention Analysis: <https://www.hfsresearch.com/research/big-bold-accenture-reinvention/>
- [10-08] CIO Dive: Accenture FY25 GenAI revenue and organizational reinvention: <https://www.ciodive.com/news/accenture-generative-ai-revenue-skills-training-data-modernization/761161/>
- [10-09] Australian Financial Review: Deloitte's botched AI report in Australia: <https://www.afr.com/companies/professional-services/deloitte-s-botched-ai-report-is-now-a-headache-for-the-big-four-20251007-p5n0pp>

- [10-10] FTC Final Order: DoNotPay deceptive AI lawyer claims: <https://www.ftc.gov/news-events/news/press-releases/2025/02/ftc-finalizes-order-donotpay-prohibits-deceptive-ai-lawyer-claims-imposes-monetary-relief-requires>
- [10-11] HAQQ: AI Hallucination Sanctions Case Tracker (1,598 cases): <https://www.haqq.ai/blog/ai-legal-hallucination-audit>
- [10-12] Computerworld / FT: Big Four hire more AI specialists than accountants: <https://www.computerworld.com/article/4173296/the-big-four-accounting-firms-are-now-hiring-more-ai-specialists-than-accountants.html>

Media & Content Creation

- [10-13] AIBase News: Mango TV: <https://news.aibase.com/zh/news/27182>
- [10-14] Unwire.pro: Media AI transformation and journalism (FT/BBC/Guardian): <https://unwire.pro/2025/11/17/media-ai-transformation-journalism/news/>
- [10-15] Sohu Article: Kuaishou Kling: https://www.sohu.com/a/919299995_362225
- [10-16] ObservAI: NYT v. OpenAI judgment: <https://observai.cn/nytimes-vs-openai-judgment/>
- [10-17] Tencent Cloud Developer News: Gannett/CNET: <https://cloud.tencent.com/developer/news/1175817>
- [10-18] Poynter Commentary: AI's wins and fails in newsrooms: <https://www.poynter.org/commentary/2025/artificial-intelligence-wins-fails-newsrooms/>
- [10-19] NewsGuard Special Report: AI Tracking Center: <https://www.newsguardtech.com/special-reports/ai-tracking-center/>
- [10-20] Huxiu Article: AI: <https://www.huxiu.com/article/4400869.html>
- [10-21] Xinhuanet Digital Report: AI: <http://www1.xinhuanet.com/digital/20250902/d08b592223954340af8c02ac3d7161e1/c.html>
- [10-22] Guardian Newsroom: Guardian Media Group announces strategic partnership with OpenAI: <https://www.theguardian.com/gnm-press-office/2025/feb/14/guardian-media-group-announces-strategic-partnership-with-openai>
- [10-23] Tencent News: AI coverage: <https://news.qq.com/rain/a/20251120A05SIM00>
- [10-24] AI for Newsroom: State of AI in Newsrooms Report: <https://aifornewsroom.in/reports>

Financial Services

- [10-25] Science and Technology Daily: WeBank: From Digital Bank to AI-Native Bank: https://www.stdaily.com/web/gdxw/2025-03/13/content_309155.html
- [10-26] WeBank (Official): AI compute grew 3.5x last year, serving over 440 million personal customers: <https://www.webank.com/mediaCenter/mtbd/3141>
- [10-27] CNBC: JPMorgan Chase's blueprint to become the first fully AI-connected megabank: <https://www.cnbc.com/2025/09/30/jpmorgan-chase-fully-ai-connected-megabank.html>
- [10-28] Tearsheet: JPMorgan Chase's GenAI implementation: 450 use cases and lessons learned: <https://tearsheet.co/artificial-intelligence/jpmorgan-chases-gen-ai-implementation-450-use-cases-and-lessons-learned/>
- [10-29] The Business Times (Singapore): DBS unlocks S\$1 billion in AI value by 2025: <https://www.businesstimes.com.sg/zh-hans/companies-markets/dbs-unlocks-s1-billion-ai-value-2025>
- [10-30] Eastmoney / Jin10: Hedge fund tied to "AI stock god" triggers record \$15 billion loss for Wall Street giant: <https://finance.eastmoney.com/a/202608173843102441.html>
- [10-31] AINEXT: From 7,000 to 3,000 people: Klarna's real lessons in reshaping an organization with AI: <https://ainext.tw/posts/20260224-klarna-ai-workforce/>
- [10-32] McKinsey: Seizing the new opportunities of generative AI — financial industry CEO quarterly: <https://www.mckinsey.com.cn>
- [10-33] National Financial Regulatory Administration (NFRA): Guiding opinions on the safe development and application of AI in banking and insurance: <https://www.nfra.gov.cn/cn/view/pages/ItemDetail.html?docId=1261708&itemId=917>
- [10-34] UK Finance: Navigating the EU AI Act: a strategic approach for financial services: <https://www.ukfinance.org.uk/news-and-insight/blog/navigating-eu-ai-act-strategic-approach-financial-services>
- [10-35] 21st Century Business Herald: "Over 10 million hours of labor replaced" vs. "low frontline satisfaction": the gap between bank AI statistics and on-the-ground reality: <https://www.21jingji.com/article/20260512/herald/e90a66b55873092c73206c2ae774f338.html>
- [10-36] 36Kr: Banks' trillion-yuan AI investment: time to count the returns: <https://m.36kr.com/p/3776145485152771>

Education

- [10-37] GSU News: Classroom Chatbot Improves Student Performance, Study Says: <https://news.gsu.edu/2022/03/21/classroom-chatbot-improves-student-performance-study-says/>
- [10-38] Mainstay (Pounce vendor): How Georgia State University Uses Behavioral Intelligence to Improve Retention: <https://mainstay.com/case-study/how-georgia-state-university-uses-behavioral-intelligence-to-improve-student-retention-and-persistence/>
- [10-39] Khan Academy Official Blog: How Khan Academy Is Building a Better AI Tutor: <https://blog.khanacademy.org/how-khan-academy-is-building-a-better-ai-tutor-our-most-recent-learnings/>
- [10-40] Chalkbeat: Sal Khan reflects on AI in schools and Khanmigo: <https://www.chalkbeat.org/2026/04/09/sal-khan-reflects-on-ai-in-schools-and-khanmigo/>
- [10-41] ERIC (Journal of Teaching and Learning): Mixed-methods study of Khanmigo in a university physics course: <https://eric.ed.gov/?id=EJ1487444>
- [10-42] The 74: Study: Giving Kids Access to AI Tutors Doesn't Mean They'll Use Them: <https://www.the74million.org/article/study-giving-kids-access-to-ai-tutors-doesnt-mean-theyll-use-them/>
- [10-43] The 74: LA Unified Faces Criticism After Collapse of Splashy AI Tool Ed: <https://www.the74million.org/article/la-unified-faces-criticism-after-collapse-of-splashy-ai-tool-ed/>
- [10-44] CNBC: Chegg to lay off 22% of workforce as AI shakes up education tech industry: <https://www.cnbc.com/2025/05/12/chegg-to-lay-off-22percent-of-workforce-as-ai-shakes-up-education-tech-industry.html>
- [10-45] Higher Ed Dive: 2U files for Chapter 11 bankruptcy: <https://www.highereddive.com/news/2u-chapter-11-bankruptcy-restructuring/722358/>
- [10-46] Ministry of Education (China): Five ministries issue the "AI + Education" Action Plan: https://www.moe.gov.cn/srcsite/A16/s3342/202604/t20260410_1433240.html
- [10-47] Digital Education Council: AI in Higher Education Global Survey 2026: <https://www.digitaleducationcouncil.com/resource-library-items/ai-in-higher-education-global-survey-2026>

- [10-48] BuildMVPFast: Duolingo's AI-First Transformation: Birdbrain and \$1B Revenue: <https://www.buildmvpfast.com/blog/ai-learning-personalization-duolingo-ai-driven-lessons-2026>

Retail & Consumer Goods

- [10-49] InteligenAI: IKEA Billie case (8,500 redeployed staff and €1.3B revenue): <https://intelligenai.com/ikea-ai-case-study/>
- [10-50] China Macro Network: Luckin Coffee at WAIC 2026: the "smart brewing ecosystem" behind 35,000 stores: <https://www.zhonghongwang.com/show-140-466795-1.html>
- [10-51] Lianshang.com (via Sina Finance): Hema upgrades to regional structure and establishes an AI committee: <https://finance.sina.com.cn/roll/2026-04-27/doc-inhvxxx7776555.shtml>
- [10-52] Fortune: Amazon closes its Amazon Go convenience stores and Fresh supermarkets: <https://fortune.com/2026/01/29/amazon-go-fresh-retail-whole-foods/>
- [10-53] CNBC: McDonald's to end IBM AI drive-thru test: <https://www.cnbc.com/2024/06/17/mcdonalds-to-end-ibm-ai-drive-thru-test.html>
- [10-54] Sina Finance: Has Yonghui Superstores "graduated" from learning Pangdonglai?: https://k.sina.com.cn/article_7879849619_1d5acf69306801cves.html
- [10-55] Jiemian News: From tech enablement to open ecosystem: the digital-intelligence leap of the "SHEIN chain": <https://www.jiemian.com/article/13880290.html>
- [10-56] RetailNews.ai: Zara AI supply chain (Inditex €1.8B investment and 85% production quota decisions): <https://www.retailnews.ai/topic/zara-ai-supply-chain>
- [10-57] Walmart Corporate: Walmart unveils new AI-powered tools to empower 1.5 million associates: <https://corporate.walmart.com/news/2025/06/24/walmart-unveils-new-ai-powered-tools-to-empower-1-5-million-associates>
- [10-58] Walmart Corporate: Retail Rewired 2025 (AI recommendation trust and consumer survey): <https://corporate.walmart.com/news/2025/06/04/walmarts-retail-rewired-report-2025-agentic-ai-at-the-heart-of-retail-transformation>

- [10-59] Bain & Company: Retail Personalization: AI Marketing Magic (ROAS lift of 10%-25%): <https://www.bain.com/insights/retail-personalization-ai-marketing-magic/>
- [10-60] CCFA: AI empowers food-service retail: Chayan Yuese's AI interviewer cuts hiring from 2 days to 30 minutes: <http://www.ccfa.org.cn/portal/cn/xiangxi.jsp?id=447349>

Government & Public Services

- [10-61] Cyberspace Administration of China: Guidelines for deploying AI foundation models in government affairs: https://www.cac.gov.cn/2025-10/10/c_1761819469929310.htm
- [10-62] Shenzhen Municipal Government: Futian District accelerates government digitalization with DeepSeek; 70 AI "digital employees" on the job: https://www.sz.gov.cn/cn/xxgk/zfxxgj/gqdt/content/post_12006997.html
- [10-63] Xinhua News Agency: How do AI "digital employees" power smarter government?: <http://www.news.cn/tech/20250224/3d814d20206449a7afe4c78a2466d088/c.html>
- [10-64] Shenzhen Municipal Government Services and Data Administration: Top 10 milestones in Shenzhen's digital government and smart city construction in 2025: https://www.sz.gov.cn/szzsj/gkmlpt/content/12/12608/post_12608373.html
- [10-65] Guangzhou Municipal Government: The 12345 hotline gets a "super operator" as Guangzhou rolls out DeepSeek across government services: https://www.gz.gov.cn/szgzkfgx/zdgz/content/post_10153907.html
- [10-66] U.S. Government Accountability Office (GAO): Generative AI Use and Management at Federal Agencies: <https://www.gao.gov/products/gao-25-107653>
- [10-67] Outlook Newsweekly: AI hallucinations proliferate: what are the risks and challenges?: <https://www.nsfc.gov.cn/p1/3381/4121/2826/93953.html>
- [10-68] People's Daily: What key signals does the establishment of the National Data Administration send?: https://paper.people.com.cn/rmlt/html/2023-05/15/content_25999945.htm
- [10-69] People's Daily Online: To get "AI +" government services right, clarify these things first: <http://politics.people.com.cn/n1/2025/0304/c458474-40430761.html>

- [10-70] Brookings Institution: Assessing the state of AI adoption across the federal government: <https://www.brookings.edu/articles/assessing-the-state-of-ai-adoption-across-the-federal-government/>
- [10-71] Chongqing Municipal Government: Exploring new paths of modern governance for megacities: the "light of wisdom" shines on a renewed Chongqing: http://www.cq.gov.cn/ywdt/jrcq/202508/t20250818_14907156_wap.html
- [10-72] Shanghai Municipal Government General Office: Shanghai's implementation plan to accelerate "AI + government services" reform and promote "getting things done efficiently": <https://www.shanghai.gov.cn/nw12344/20250410/467ecd7d1a254575b263e73fc7c793de.html>

Logistics & Supply Chain

- [10-73] China Securities Journal: Rebuilding Qingdao Port's DNA with AI: https://epaper.cs.com.cn/zgzqb/html/2025-08/01/nw.D110000zgzqb_20250801_3-A04.htm
- [10-74] 36Kr: Amazon deploys 1 million robots — about to surpass human employees?: <https://m.36kr.com/p/3368169340962561>
- [10-75] SF Technology (Official): SF launches "Fengzhi", its major logistics decision-making model: <https://www.sf-laas.com/news-details/40192>
- [10-76] Logclub: The AI cards of nine companies — SF, JD Logistics, Cainiao, ZTO, YTO, STO and more — at a glance: <https://www.logclub.com/articleInfo/ODQ2MjE=>
- [10-77] Shanghai Securities News: Six full-stack results verified across the board: YTO Express scales AI applications: <https://www.cnstock.com/commonDetail/741566>
- [10-78] Sina Finance: Logistics autonomous vehicles in 2025: over 35,000 delivered as the "Four Little Dragons" escalate the battle: <https://finance.sina.com.cn/cj/2026-01-26/doc-inhirqat8215013.shtml>
- [10-79] 36Kr: Ten thousand autonomous vehicles await a "national permit": <https://m.36kr.com/p/3617731429123076>
- [10-80] Huxiu: TuSimple's Guangzhou company suddenly announces dissolution; "management has problems": <https://www.huxiu.com/article/4039242.html>

- [10-81] project44: Why 95% of AI Agent Pilots Fail in Supply Chains: <https://www.project44.com/blog/why-95-percent-of-ai-agent-pilots-fail-in-supply-chains/>
- [10-82] Tencent News: SF, JD, and Meituan fight the "battle for the skies": drones reshape logistics: <https://news.qq.com/rain/a/20251227A01ST800>
- [10-83] McKinsey: 2025 McKinsey State of AI Survey: <https://www.mckinsey.com.cn>
- [10-84] Zhihu Column (Apsara Conference talk): Manbang's AI construction: status and reflections: <https://zhuanlan.zhihu.com/p/1967627027327743626>

Healthcare

- [10-85] Future Medicine (NEJM Catalyst study): Ambient AI scribes in Kaiser Permanente's 7,260-physician study: <https://www.futuremedicine.com/articles/ambient-ai-scribes-kaiser-permanentes-7000-physician-study-signposts-the-future-of-clinical-documentation>
- [10-86] Health News (China): "Intelligent Medical Assistant" takes root in Anhui's grassroots: 680 million auxiliary diagnosis suggestions: <https://www.jkb.com.cn/news/depth/2024/0606/495573.html>
- [10-87] VCBeat / PharmCube (VB100): The "last mile" of medical AI foundation models — only real deployment creates value: <https://bydrug.pharmcube.com/news/detail/9a530f1ae67f0613f533cc87a6747241>
- [10-88] Xinhuanet: China's first "public hospital operations management foundation model" released: <http://www.news.cn/health/20260302/a842fc4488504e59a5f90b4f324ea518/c.html>
- [10-89] Family Doctor (HIMSS 2025): The failure trap of healthcare AI: 80% of projects fail to convert value: <https://www.familydoctor.cn/news/yiliaozuzhi-ruhezai-bimian-shibai-xianjing-tongshi-shishi-rengongzhineng-65193.html>
- [10-90] UCL News: Study sheds light on hurdles faced transforming NHS healthcare with AI: <https://www.ucl.ac.uk/news/2025/sep/study-sheds-light-hurdles-faced-transforming-nhs-healthcare-ai>
- [10-91] CHIMA: Legal risks of clinical application of AI medical robots: <https://www.chima.org.cn/Html/News/Articles/17035.html>
- [10-92] Menlo Ventures: 2025 The State of AI in Healthcare: <https://menlovc.com/perspective/2025-the-state-of-ai-in-healthcare/>

- [10-93] FDA: Artificial Intelligence-Enabled Device Software Functions: <https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-software-medical-device>
- [10-94] UC San Diego Health: COMPOSER sepsis surveillance tool study: <https://health.ucsd.edu/news/press-releases/2024-01-23-study-ai-surveillance-tool-successfully-helps-to-predict-sepsis-saves-lives/>
- [10-95] UN News: WHO warns legal and ethical safeguards for medical AI lag behind: <https://news.un.org/zh/story/2025/11/1141133>
- [10-96] Jiemian News: Insilico Medicine delivers its first profitable AI-drug-discovery results, yet R&D anxiety in pharma persists: <https://www.jiemian.com/article/14847447.html>

Manufacturing

- [10-97] McKinsey: High efficiency, low energy, zero defects: AI makes manufacturing innovation one click away: <https://www.mckinsey.com.cn/高效率、低能耗、零缺陷：ai让制造创新一键直达/>
- [10-98] McKinsey: 2025 McKinsey State of AI Survey: only 6% of companies become high-performance winners: <https://www.mckinsey.com.cn/2025麦肯锡ai应用现状调研：从6%企业成为高绩效赢家，他们做对了什么？/>
- [10-99] McKinsey: Generative AI in China: \$2 trillion in economic value: <https://www.mckinsey.com.cn/生成式ai在中国：2万亿美元的经济价值/>
- [10-100] Woshipm: 90% of industrial AI deployment failures are not technical problems: <https://www.woshipm.com/ai/6392172.html>
- [10-101] Sina: Seeking the truth about industrial AI: why is AI hot but hard to deploy in industry?: http://client.sina.com.cn/zt_d/2025-08-22/doc-infmvzcz5097908.shtml
- [10-102] New Industry Network: Panoramic analysis of 167 verified cases — 2026 China Manufacturing AI Scenario Application White Paper: https://xingongye.cn/cms/Intelligent_manufacturing/1825.html
- [10-103] Figure AI: Figure 02 humanoid robot at BMW plant: involved in producing 30,000 vehicles: <https://www.figure.ai/news/production-at-bmw>

- [10-104] Automotive Manufacturing Solutions: Future-ready: BMW's digital transformation of its Munich plant and beyond: <https://www.automotivemanufacturingsolutions.com/smart-factory/future-ready-bmws-digital-transformation-of-its-munich-plant-and-beyond/523748>
- [10-105] People's Daily (People's Tribune): China's "Lighthouse Factory" development landscape and global value chain shaping: https://paper.people.com.cn/rmlt/pc/content/202505/01/content_30080970.html
- [10-106] Huawei: Industrial Digitalization/Intelligence 2030: <https://www-file.huawei.com/dam/asset/view/7ffffa5fa5e7469692d4704b92ef17d7.pdf>
- [10-107] State Administration of Science, Technology and Industry for National Defense: MIIT and seven other departments issue the "AI + Manufacturing" Special Action Implementation Plan: https://www.nda.gov.cn/sjj/zwgk/zcfb/0112/20260107214358696030895_pc.html
- [10-108] CSDN Blog: The dilemma of AI transformation in traditional manufacturing and the way out: https://blog.csdn.net/weixin_30399155/article/details/96238191

Energy & Utilities

- [10-109] State Grid: Electric Power "AI+" White Paper (6541 plan): https://www.sohu.com/a/898976204_121079544
- [10-110] Baidu Baike: Guangming Electric Power Foundation Model: <https://baike.baidu.com/item/光明电力大模型/65228636>
- [10-111] China Southern Power Grid: "Dawate·Yudian" selected among the top 10 national heavyweights of central SOEs in 2024: https://www.csg.cn/xwzx/2025/2025gsyw/202501/t20250114_344842.html
- [10-112] China Southern Power Grid: AI innovation platform model invocations exceed 10 billion: https://www.csg.cn/xwzx/2025/2025gsyw/202512/t20251216_351726.html
- [10-113] National Energy Administration: PetroChina releases 300-billion-parameter Kunlun foundation model: <https://www.nea.gov.cn/20250530/9431f59ac81f49c1a3839d7e066c4073/c.html>
- [10-114] Huawei Enterprise: Changqing Oilfield and Huawei build a smart oilfield benchmark: <https://e.huawei.com/cn/blogs/2025/industries/oil-gas/shaping-oilfields-with-data-ai>

- [10-115] Sityos: Duke Energy's AI grid prevented about 1.2 million outages: <https://www.sityos.com/en/use-cases/duke-energys-ai-grid-prevented-12-million-outages>
- [10-116] Forbes: Duke Energy's CIO on powering digital and AI transformation: <https://www.forbes.com/sites/peterhigh/2025/10/22/duke-energys-cio-on-powering-digital-and-ai-transformation/>
- [10-117] Prudent Consulting: Six root causes of failed utility AI investments: <https://www.prudentconsulting.com/blogs/ai-utilities-investment-failures/>
- [10-118] BCG: AI in Energy: A New Strategic Playbook: <https://www.bcg.com/publications/2025/ai-in-energy-new-strategic-playbook>
- [10-119] IEA: Energy and AI (milestone report): <https://www.iea.org/reports/energy-and-ai/>
- [10-120] GridFWD: Early AI lessons for utilities, grid operators, and regulators: <https://gridforward.org/early-ai-use-cases-and-lessons-for-utilities-grid-operators-and-regulators/>
- [10-121] OpenAI Official Customer Case: Legora Reviews 41 Financial Statements in Minutes with GPT-6 Astra: <https://openai.com/zh-Hans-CN/index/legora-financial-statement-review-with-astra/>
- [10-122] Bristol Law Society: Harvey Taps DeepL to Power Legal-Grade AI Translation for Global Law Firms: <https://www.bristollawsociety.com/news/harvey-taps-deepl-to-power-legal-grade-ai-translation-for-global-law-firms-and-legal-teams/>
- [10-123] TMTPost (citing Sensor Tower): AI Short Dramas Go Global — Q1 Downloads Up 140% Year over Year: <https://www.tmtpost.com/8119642.html>
- [10-124] CBNDData: AI Short Dramas — Who Is Actually Making Money?: <https://www.cbndata.com/information/295426>
- [10-125] CLS (财联社): U.S. Government Backs "Fair Use" in New York Times v. OpenAI: <https://m.cls.cn/detail/2473833>
- [10-126] Lanjinger Finance: SPD Bank 2026 Interim Report — 440+ AI Applications and 220 Agents in Production: <https://www.lanjinger.com/d/1788336255043004663>
- [10-127] OpenAI Official Customer Case: OpenAI × Morgan Stanley — AI Assistant Adoption Exceeds 98%: <https://openai.com/index/morgan-stanley/>

- [10-128] EY: AI Bank White Paper — Agent-Driven Financial Paradigm Reconstruction: <https://www.ey.com/content/dam/ey-unified-site/ey-com/zh-cn/insights/financial-services/documents/ey-ai-bank-white-paper-zh.pdf>
- [10-129] Beijing News: Dushen Education 2026 Interim Report — AI Education Revenue Up 664% to 114 Million Yuan While Overall Losses Persist: <https://www.bjnews.com.cn/detail/1787921951169678.html>
- [10-130] People's Daily Overseas Edition: Bosideng's AI Design Goes Live — 8 AI-Generated Patterns and 9 AI-Designed Garments Reach the Market: <http://js.people.com.cn/BIG5/n2/2025/0617/c360301-41261750.html>
- [10-131] Economic Observer: Bosideng's Digital Transformation — from 100+ Warehouses to One Central Warehouse: <http://www.eeo.com.cn/2025/0320/717512.shtml>
- [10-132] AI News: Kroger and Lowe's Test Their Own AI Agents Instead of Handing Control to Third-Party Chatbots: <https://www.artificialintelligence-news.com/news/kroger-and-lowes-test-ai-agents-without-handing-control-to-google/>
- [10-133] People.cn (People Zhongke Research Institute): Four Risks of Local Governments Stampeding to Deploy DeepSeek: <http://politics.people.com.cn/n1/2025/0224/c458474-40424743.html>
- [10-134] UNDP Global Centre Singapore: Pairing AI with the Public Sector — Singapore's Pair Suite Saves Civil Servants an Estimated 46% of Administrative Time: <https://www.undp.org/policy-centre/singapore/blog/pairing-ai-public-sector-impact-singapore>
- [10-135] Sohu (China Port Autonomous Driving Blue Book): China Port Autonomous Driving Industry Development Blue Book (2026 Edition): https://www.sohu.com/a/1068619374_468661
- [10-136] MarketScale / FreightWaves: Kenco's 20-Agent Target Puts AI Workflow Control at the Center of Logistics Tech: <https://www.marketscale.com/industries/transportation/kencos-20-agent-target-puts-ai-workflow-control-at-the-center-of-logistics-tech>
- [10-137] JAMA Network Open: Ambient AI Scribe Use Over 30 Days Associated with Drop in Physician Burnout from 51.9% to 38.8%: <https://jamanetwork.com/journals/jamanetworkopen/fullarticle/2839542>

- [10-138] MedPage Today: AI Scribes' Time Savings "Modest" but Meaningfully Improve Clinician Experience: <https://www.medpagetoday.com/hospitalbasedmedicine/workforce/120606>
- [10-139] Caixin Weekly: Beijing Meizhong Airui Cancer Hospital Plans China's First "AI-Native" Physical Hospital: <https://weekly.caixin.com/2026-04-04/102430893.html>
- [10-140] CBS News: UnitedHealth Sued over Alleged Use of AI Algorithms with up to 90% Error Rate to Deny Care: <https://www.cbsnews.com/news/unitedhealth-lawsuit-ai-deny-claims-medicare-advantage-health-insurance-denials/>
- [10-141] Securities Times: Baosteel's "AI+" Theme Day — 582 Scenarios Live with Cumulative Benefits above 270 Million Yuan: <https://www.stcn.com/article/detail/3828074.html>
- [10-142] TMTPost: A Look at Eight Listed Steelmakers' AI Investment — Baosteel Is the Only "Company-Level Strategy + Quantifiable Results" Sample: <https://www.tmtpost.com/8018596.html>
- [10-143] World Economic Forum: Global Lighthouse Network Recognizes 23 New Sites and Launches Lumina, an AI Platform for Industrial Transformation: <https://cn.weforum.org/press/2026/01/global-lighthouse-network-recognizes-23-new-sites-launches-ai-platform-for-industrial-transformation-89a7334dcb-cn/>
- [10-144] Murata Manufacturing: Beware the "Data Traps" of the Smart Factory: <https://article.murata.com/zh-cn/article/smartfactory-pitfalls-in-data>
- [10-145] National Energy Administration: Computing-Electricity Synergy Moves from Concept to Reality — First Written into the Government Work Report and the 15th Five-Year Plan: <https://www.nea.gov.cn/20260710/082757047b114b2798fcdeec42673e18/c.html>
- [10-146] People's Daily: China Southern Power Grid "Makes Every Kilowatt-Hour AI-Powered" — Distribution-Network Planning Agent Named a WAIC Top-Ten "Treasure of the Hall": http://paper.people.com.cn/rmrb/pc/content/202607/21/content_30169986.html

Knowledge-Base Material (One Source of This Book's Content)

Much of this book draws from the author's personal knowledge base (wiki). Its original sources include 36Kr, Woshipm, Tencent Research Institute, HBS Working Papers, WEF, Harvard Business Review (Chinese edition), LatePost, Economic Observer, and other public publications. Specific entries are recorded in the CHANGELOG with each version update.

Appendix B: Key Metrics for AI-Native Organizations

You get what you measure. The measurement system of an AI-native organization differs fundamentally from the industrial-era one: it shifts from "input and activity" to "outcome and value." The metrics below are grouped by purpose, all drawn from the research cited in this book.

1. Organizational Structure Metrics ("What do we look like?")

Metric	Description	Benchmark
Relative team size	Headcount vs peers	HBS: AI-native firms are ~25% smaller (services ~70% smaller)
Engineer share	Engineers / total headcount	HBS: ~13 percentage points higher
Management layer depth	Layers from CEO to frontline	HBS: half a layer flatter; Shmool: 4 layers → 1
Builder-to-manager ratio	Direct producers / coordinators	Shmool: 3:1 → 8:1+
Human-to-AI ratio	Humans : AI agents	WEF: early AI-first organizations exceed 1:10
Information fidelity	Distortion of frontline facts reaching decision-makers	Shmool: ~40% loss per hop → ~95% direct
Decision latency	From problem to decision	Shmool: days-weeks → hours
Meeting overhead	Share of workweek in meetings	Shmool: ~30% → ~10%

2. Efficiency and Value Metrics ("What did AI deliver?")

Metric	Description	Benchmark
Revenue per employee	Revenue / headcount	Klarna: \$1.24M (3.6x vs 2022); Shopify: \$1.52M; Cursor: \$3M-13M
Per-capita growth rate	YoY change in revenue per employee	More robust than absolute value for transformation tracking
Output capacity	Output per unit time	Duolingo: course skills 1,800 → 20,500 per quarter (11x)
Time savings	Task-level time savings	Anthropic dialogue research: ~80% on average
Task-level vs structural gains	Single-step speedup vs full-flow improvement	Task-level 20-40%; workflow redesign 2-10x (Harvard Data Science Review)
EBIT attribution	AI's measured contribution to profit	McKinsey: only 39% attribute any EBIT impact; "high performers" (≥5%) ~6%
1.9x revenue effect	Causal gains from organizational learning	INSEAD/HBS field experiment (treatment vs control)

3. Workflow Metrics ("Did the work actually get better?")

Principle: measure the whole workflow, not the AI task (Chapter 5, rule 8).

- **Cycle time**: end-to-end from trigger to completion (not single-step);
- **First-time correctness rate**: share of cases done right the first time;
- **Exception rate and exception age**: share of exceptions + how long they sit in queues;
- **Repeat-contact rate**: how often the same issue is re-submitted (the Klarna lesson: aggregate metrics hide collapses in high-value scenarios);
- **Complex/escalated-interaction satisfaction**: tracked separately, never averaged into the overall number;
- **Review time**: human review hours (after AI output explodes, review becomes the new bottleneck);

- **Cost per completed case:** full-cycle cost (inference + integration + ops + QA + human review + retries);
- **Rework rate:** share of AI output rejected/redone;
- **AI confidence and coverage:** AI participation rate and confidence distribution (where do low-confidence cases route);
- **Escalation quality:** whether escalations to humans carry full context.

4. Adoption and Governance Metrics ("Is the organization actually using it?")

Metric	Description	Warning data
License utilization	Active users among those with access	Copilot: only ~1/3 (64% of licenses unused)
Paid/premium conversion	Users paying for AI features	Copilot: 3.4%; Notion AI: >50% paid penetration (contrast)
Four-dimensional success	Adoption / efficiency / quality / satisfaction	Anthropic's official pilot measurement
Shadow AI usage	Use of unauthorized tools	Salesforce: 69% of employees used unauthorized tools
AI value accounting	Finance-led ROI validation	CFO-led validation success 76% vs tech-led 53%
Full-delegation ceiling	Share of work fully delegable to AI	Anthropic internal: most employees say only 0-20%

5. Risk Metrics ("Will something go wrong?")

- **Vendor switching cost:** 19-34% of total deployment cost; only 6% of firms can switch without disruption;
- **Agent-related incidents:** Rubrik: 98% of organizations have experienced a disruptive AI-agent event;
- **Skill-atrophy signals:** employees reporting fewer hands-on practice opportunities; share of exploratory work (Anthropic internal: 27% of AI-assisted work would not have been done otherwise);

- **Employee trust:** share of employees who find AI agents reliable (Asana: only 36%);
- **Data AI-readiness:** completion rate of AI-ready data audits (<20% of organizations consider themselves ready).

6. Meta Metrics (Organizational Design)

- **Competitiveness formula:** talent density $\times$ AI leverage $\div$ organizational friction (Tencent Research Institute) — measure the three variables separately; cut the denominator before doubling the numerator;
- **Innovation density:** output / headcount (DeepSeek's 160 vs the giants' thousands);
- **Influence spillover:** whether one person makes the team faster (the threshold for a Super-Individual, not personal output multiples);
- **Review signal:** does performance review look at input or output? Is AI usage part of the review? (Shopify and Duolingo already include it.)

Appendix C: Version History and Update Notes

Update Mechanism

This book is a "living book," updated automatically every week:

1. An automated system scans global AI-landing and organizational-change information daily, distills it, and stores it in a knowledge base;
2. Each week, new knowledge-base content is merged into the relevant chapters and the version number is incremented (patch +1);
3. The whole-book PDF is regenerated and synced to the GitHub repository (this repository — note: the Chinese edition is also served on a Chinese-facing website; the English edition is hosted here on GitHub only);
4. Every update is recorded in this appendix and in the repository's CHANGELOG.

Version rules: v1.0.0 → v1.0.1 (weekly routine updates) → v1.1.0 (new chapters / major revisions) → v2.0.0 (full rewrite).

v1.0.0 — 2026-08-22 (First Edition)

First edition release: built on the knowledge base's accumulated material (12 core concept pages + 20+ case articles) plus three rounds of deep research reinforcement (international cases / authoritative reports / failure cases & methodology).

- 9 chapters + afterword + 3 appendices;
- Chapter 1: The rise of AI-native organizations (the precise framing of the 95% failure rate + a 515-startup causal experiment);
- Chapter 2: Band-aid AI vs AI-native DNA (including a correction of the widely-circulated Gartner CAIO prediction: "90%" is a misattribution; the actual figure is 35%);
- Chapter 3: The evidence (HBS 26-090 / INSEAD / WEF / Tencent Research Institute / McKinsey / BCG / Microsoft);

- Chapter 4: Rebuilding organizational structure (MarsWave / Shmool / Block / the Organizational Brain / three paths);
 - Chapter 5: Redesigning workflows and decisions (9 key redesigns / decision rights table / 90-day cadence / three-tier routing);
 - Chapter 6: Incentives and talent (Energy Gold / review reform / four incentive dimensions / skill atrophy / innovation density);
 - Chapter 7: Case studies (China: Transn / Huawei / DeepSeek / Jingzhunxue / MarsWave; Global: Klarna / Shopify / Anthropic / OpenAI / Notion / Cursor / Duolingo / Airbnb) + six cross-case patterns;
 - Chapter 8: Action playbook (Phase 0-4 roadmap / BCG's three plays / Accenture's three principles / 10-error quick table);
 - Chapter 9: Challenges, risks, and the future (four failure cases / six risk categories / three future directions);
 - Appendix A: 165 sources (copyright compliance); Appendix B: key metrics; Appendix C: this appendix.
-